

Fractured Trust: The Spillover Effects of Police Violence on Satisfaction with Democracy and Political Trust*

Ignacio Urbina[†]

Vitória Sgorlon[‡]

Working paper — this version: September 2026.

The Supplementary Information for the three studies follows the manuscript in this file.

Abstract

How does police violence impact citizens' satisfaction with democracy? This paper explores how police violence against civilians influences citizens' satisfaction with democracy through its effect on trust in institutions. To examine the effects of direct and indirect police violence, we conducted two studies: an experimental study in Brazil and an observational analysis using LAPOP data across various Latin American countries. Our results reveal that both direct experiences of police violence and indirect exposure significantly erode trust in key democratic institutions, including the judiciary, legislature, and executive. This erosion of political trust is mediated through increased distrust in the police and negative perceptions of procedural justice. Additionally, the effects of indirect exposure are moderated by factors such as race, partisan identity, and trait aggression personality predisposition. Overall, this research underscores the critical impact of police violence on political trust and emphasizes that addressing such violence is crucial for sustaining democratic legitimacy, particularly in regions facing democratic backsliding. By highlighting the broader political consequences of police violence, our study contributes to the literature on political behavior and public opinion while suggesting law enforcement reforms to strengthen democracy's resilience.

*Special mention to EK and MC for providing insightful feedback, which has significantly improved our paper. We extend our thanks to the discussants and participants from the 2023 PolMeth Latam Conference, 2024 APSA, 2024 MPSA, 2024 URI Human Rights Workshop, 2024 PolMeth Conference, and the internal seminars at Stony Brook University for their detailed and constructive critiques. These interactions have been crucial in refining our arguments and enhancing the clarity and rigor of our work.

[†]PhD Candidate, Department of Political Science, Stony Brook University, 11790 NY. Corresponding author (ignacio.urbina@stonybrook.edu).

[‡]PhD Candidate, Department of Political Science, Stony Brook University, 11790 NY (vitoria.sgorlon@stonybrook.edu).

1 Introduction

Police violence is a persistent issue in Latin America, where law enforcement has been involved in activities such as unlawful killings, torture, monetary extortion, and collaboration with criminal organizations (Brinks, 2007). Reports from Bolivia, Honduras, Colombia, Argentina, El Salvador, and Peru, among other countries, highlight widespread misconduct (Human Rights Watch, 2020). In Chile, police killed at least 31 civilians and injured approximately 3,000 during the 2019 protests (Dammert and Sazo, 2020). In Brazil, police killings have exceeded 6,000 per year since 2018 (Bueno and de Lima, 2019, 2023). Despite democratization, police and military forces retain significant power and influence. In this context, it is important to learn whether exposure to police violence can influence citizens' evaluation of democracy and their trust in its institutions.

Here, we ask whether citizens' experiences with police violence influence their satisfaction with democracy and trust in its institutions. Systematic police violence undermines democratic principles by violating equal citizenship, breaking the rule of law, and eroding perceptions of procedural justice (Bonner, 2021; González, 2020; Wiatrowski, 2016; Manning, 2015). These violations weaken trust in institutions, shaping citizens' evaluation of the democratic system's effectiveness and fairness. Research shows that police violence reduces cooperation with law enforcement (McLean et al., 2024; Droic and Shoub, 2024; Ang et al., 2021; Carr, Napolitano, and Keating, 2007). It also influences political behavior by mobilizing some citizens, particularly in marginalized communities (Reid, 2024; Morris and Shoub, 2023; Ang and Tebes, 2021), shifting policy preferences (Mullinix, Bolsen, and Norris, 2021; Culhane, Boman, and Schweitzer, 2016), and shaping perceptions of political efficacy (Branton, Carey, and Martinez-Ebers, 2023), highlighting its consequences on democracy.

Satisfaction with democracy is a widely studied measure of system legitimacy, mostly explained by economic expectations, socioeconomic status, education, interpersonal trust, crime victimization, government efficacy, and perceived income fairness (Monsiváis-Carrillo and Cantu Ramos, 2022; Kriekhaus et al., 2014; Evans and Rose, 2012; Fernandez and Kuenzi, 2010; Mishler and Rose, 1996). However, existing research has largely overlooked whether police violence, an institution of the state, affects democratic satisfaction.

Studying the impact of police violence is particularly important in Latin America, where democracies are often described as 'uncivil democracies' (Holston and Caldeira, 1998). Although political rights such as voting are established, pervasive violence, impunity, and corrupt judicial systems hinder the full realization of civil rights (Pereira and Davis, 2000). Despite the central role of the police in society, analyses of politics and opinion about democracy in the region frequently overlook their influence. Instead, scholarly attention tends to focus on state, military, and criminal violence when examining political behavior and attitudes (e.g., Chouhy et al., 2022; Gomes and Alves de Aquino, 2018; Trelles and Carreras, 2012).

We test our theory in three studies. In Study 1, an original two-wave vignette experiment in Brazil isolates indirect exposure to police violence, the kind that reaches most citizens through media coverage. In Study 2, we turn to the 2019 Chilean *estallido social*, where police repression fell on some census zones and not others, and we ask whether residents exposed to that shock diverged from residents of comparable zones. In Study 3, we use LAPOP survey data across sixteen Latin American countries to examine direct victimization by the police, and to test whether the pathway we theorize appears in the cross-national record.

In Study 1, participants who read a vignette depicting a violent police encounter reported lower trust in the police and lower satisfaction with democracy than those who read a non-violent scenario, while the effect on political trust emerged mainly among non-whites, respondents who like the PT, and individuals high in trait aggression, the groups in which every effect concentrated. In Study 2, living in an exposed zone eroded trust in political institutions and satisfaction with democracy, and it did so independently of the nationwide collapse in police trust that followed the repression. In Study 3, respondents who report police abuse show lower trust in the police, lower political trust, and lower satisfaction with democracy across all sixteen countries, and the association between police abuse and dissatisfaction runs mainly through political trust rather than through trust in the police directly.

To our knowledge, this is the first study to theorize and examine the effect of exposure to police violence on trust in democratic institutions, under one framework. In a period of concern over democratic backsliding, our study argues for treating police violence as more than a human rights violation. It also threatens democratic legitimacy, through its negative spillover onto citizens' trust in the institutions democracy rests on and system satisfaction.

2 The Spillover Effects of Police Violence

The police are a physical embodiment of state authority and, as such, they are probably the most visible state institution to the majority of ordinary citizens. As a street-level bureaucratic organization, the police deliver policy in the everyday lives of the citizens they are meant to serve (Brodin, 2011). In many places they are the most proximate face of the state, especially in small towns and villages where institutions such as the Judiciary and the Legislature are absent. This visibility is what leads us to ask whether police violence negatively affects broader public opinion about democracy and its institutions.

One crucial distinction the police have over other street-level bureaucrats is that they are *armed*. They carry the legal right to use coercive power to fulfill their mandate, and citizens expect them to apply force proportionately, practically, and appropriately (Kleinig, 2014). Following Milani et al. (2018), we define police violence as the excessive and needless use of physical force in police-citizen encounters, where “excessive” refers to force disproportionate to what a given situation requires. This definition also extends more broadly to include psychological, sexual, and neglectful harm, consistent with the World Health Organization’s classification of violence (Cooper et al., 2004). Incidents of such violence against citizens have provoked contentious public debate and a growing literature on policing. These experiences weaken citizens’ assessment of procedural justice—the fairness of process and treatment, rather than outcomes (Tyler and Allan Lind, 2001), and foster more negative attitudes toward law enforcement, including distrust and perceptions of inefficacy (e.g., Branton et al., 2023; Jackson et al., 2021; Boudreau et al., 2019). That skepticism, in turn, reduces civilian cooperation with local police and local government, lowering, for example, willingness to call the police (McLean et al., 2024; Drolc and Shoub, 2024; Carr et al., 2007).

This dynamic should be particularly pronounced in Latin America, where the police have run drug operations, committed killings, kidnappings, torture, extorted money, supplied intelligence to criminal organizations, and mistreated detainees (Brinks, 2007). Despite democratization, the

military and the police retain significant power in many countries, while public confidence in law enforcement's ability to fight crime and uphold human rights has declined (Sung et al., 2022; Pion-Berlin and Carreras, 2017). Studying police violence becomes especially important here, in democracies often described as “uncivil” (Holston and Caldeira, 1998), where political rights such as voting are established, but pervasive violence, impunity, and corrupt judicial systems block the full realization of civil rights (Pereira and Davis, 2000).

A large body of work connects contact with police violence to public opinion and political behavior. Mullinix et al. (2021) and Culhane et al. (2016) show that exposure to police violence through text, audio, or video can increase support for policy changes such as the mandatory use of body cameras. On the perceived value of democratic participation, Branton et al. (2023) finds that exposure to police violence lowers beliefs about the external efficacy of voting and protesting. Attitudes toward police violence, however, can be polarized along ideological and partisan lines. In Latin America, high criminal violence pressures governments to act, and because public security remains a core state responsibility, many citizens, particularly on the right, embrace *mano dura* (iron-fist) policing, framing aggressive enforcement as a function of a strong state and tolerating or endorsing violence against perceived criminals. The pattern holds globally. In the United States, policing attitudes split sharply by ideology, with conservatives holding stable or positive views of law enforcement while liberals show greater concern about misconduct (Reny and Newman, 2021), and support for harsh policing tracks authoritarianism among populist bases of leaders such as Donald Trump and Jair Bolsonaro (Vidigal, 2022; Knuckey and Hassan, 2022; Womick et al., 2018), sometimes surfacing only among citizens on the ideological right (Cohen and Smith, 2016).

Effects on behavior are more mixed and often racialized. Morris and Shoub (2024) and Ang and Tebes (2024) find that police killings mobilize voter registration and turnout in Black and Hispanic communities, while Markarian (2023) observes a localized drop in Black turnout within a mile of an incident, and Laniyonu (2018) finds lower participation in more-exposed communities in the United Kingdom, again moderated by race. On vote choice, in the United States, Reid (2024) finds that voters in areas with recent police killings favored the Democratic candidate in 2020, net of the usual predictors. Race structures these effects because policing is itself racialized in most contexts, where marginalized racial groups are disproportionately profiled (e.g., French, 2013; Greenwald et al., 2003; Correll et al., 2002) and victimized (e.g., Bueno and de Lima, 2023; Buehler, 2017). Non-white individuals encounter police violence directly or through others in their community, so repeated exposure activates negative perceptions that shape their civic and political action more than among white individuals, who are both less victimized and less behaviorally responsive to these incidents (Morris and Shoub, 2024; Crabtree and Yadon, 2022; Ang and Tebes, 2024).

Importantly, systemic and excessive police violence is in tension with democratic principles, so it can degrade citizens' perceptions of how well democracy serves its citizens. We distinguish democratic evaluations as either normative or instrumental. Normative judgments stem from an intrinsic valuation of democracy and its core principles, as when someone cherishes free speech or equality before the law. Instrumental evaluations treat democracy as a tool for delivering outcomes such as economic growth or social inclusion. Police violence cuts against both. It violates core democratic principles directly, transforming institutions meant for protection into sources of fear and departing from the state's obligation to safeguard its citizens. When violence

is produced, or not stopped, by a street-level bureaucratic organization, this signals that the system fails to deliver the outcomes citizens expect of it.

Here, we argue that police violence also shapes public opinion about democracy, as it can foster distrust in democratic institutions, and procedural justice underpins that link because it turns on four features of citizens' interactions with authorities: being treated with dignity and respect, having a voice, facing neutral and transparent decisions, and encountering trustworthy motives (Hough et al., 2017; Tyler, 2003). More broadly, procedural justice is crucial for maintaining legitimacy, trust, and public confidence in democracy (Saxton, 2021; Magalhães and Aguiar-Conraria, 2019; Grimes, 2017; Magalhães, 2014). Violence by law enforcement violates procedural justice, and thereby breeds distrust. It weakens confidence in the rule of law and, by extension, the Judiciary, as citizens conclude that officials are not held accountable or bound by the same constraints as ordinary people.

As a spillover effect, then, we expect police performance to influence the evaluations of the state as a whole. Silva et al. (2022), for instance, find that exposure to police violence lowers trust in *local* government. But the spillover should not stop at local government. The police are the everyday embodiment of state authority, but the police sometimes *are* the state, especially where other democratic institutions are absent from daily life, and a citizen has no more tangible institution against which to judge the wider system. A violent encounter therefore updates beliefs not about one agency but about the fairness and trustworthiness of the state that arms and authorizes that agency. And because marginalized groups are disproportionately targeted, it is their belief in a biased system that attaches most strongly to the justice and legislative institutions charged with equality before the law. Police violence thus travels from the most visible arm of the state to the abstract national institutions.

Police violence can also degrade evaluations of democracy by threatening citizens' psychological and basic needs (DeVylder et al., 2017; Geller et al., 2014). Abusive practices conflict with the fundamental need for respect, inclusion, and dignity, undermining a sense of belonging and self-worth. Because this violence falls disproportionately on communities marginalized on racial, ethnic, or socioeconomic grounds, it intensifies inequality and pushes those communities out of economic and social life (DeVylder et al., 2017). Such violence leaves these communities materially and psychologically worse off, delivering the opposite of what a state is meant to provide for its citizens.

Despite these consequences, the role of police violence in shaping public opinion about democracy remains underexplored. The literature has extensively examined how economic expectations, socioeconomic status, education, interpersonal trust, government efficacy, and perceived income fairness, among many others, shape opinions about democracy and its institutions (e.g., Krieckhaus et al., 2014; Fernandez and Kuenzi, 2010; Mishler and Rose, 1996). Yet, when studying the effects of violence in Latin America, researchers tend to focus on state, military, and criminal violence while overlooking the police (e.g., Chouhy et al., 2022; Gomes and Alves de Aquino, 2018; Trelles and Carreras, 2012).¹

In this paper, we argue that assessments of public opinion about democracy and its institutions should also account for experiences of police violence. These experiences erode citizens' perceptions of procedural justice, which in turn breeds distrust in democratic institutions and

¹ See Cruz (2015) for an exception.

shapes their broader evaluations of democracy.

3 Satisfaction with Democracy and Trust in its Institutions

Citizens' assessment of and satisfaction with democracy depend on how they evaluate the system's effectiveness in upholding core democratic norms and its performance in meeting their needs (Canache et al., 2001). Conceptually, satisfaction with democracy indicates a mid-level support for the democratic system. Easton (1965, 1975) originally distinguished diffuse support, directed at the political community and the regime, from specific support, targeted at political authorities. Norris (1999) refined this scheme into evaluations of political actors, regime institutions, regime performance, regime principles, and the political community. Satisfaction with democracy sits at the level of regime performance, a middle tier of support between specific support for actors and institutions and support for the regime's principles, such as democracy itself (Norris, 1999).

As an instrumental evaluation, satisfaction with democracy is influenced by many factors related to system performance (Singh and Mayne, 2023). For instance, citizens who hold positive economic expectations, enjoy higher status, or see the income distribution as fair report more satisfaction (Loveless and Binelli, 2020; Krieckhaus et al., 2014), and it also rises when citizens perceive the government as efficient and delivering on its promises, since this signals that the system is working and meeting their needs (Linde and Dahlberg, 2021; Magalhães, 2014).

Satisfaction with democracy has long served as an indicator for many political behaviors and attitudes across countries. Although evidence on its correlation with self-reported political participation is mixed (Singh and Mayne, 2023), recent research in Mexico finds a positive correlation between satisfaction with democracy and administratively measured behaviors—turnout, training as a poll-worker, and staffing a polling station on Election Day (Vargas et al., 2026). Dissatisfaction can also shape vote choice, making citizens more likely to support populist parties of both the radical right and left (Ramiro, 2016; Arzheimer, 2009; Lubbers et al., 2002). In Latin America, satisfaction with democracy is often found to be positively correlated with support for democracy (Walker and Kehoe, 2013; Sarsfield and Echegaray, 2006). Claassen (2020) further finds that satisfaction is positively associated with subsequent change in democracy within democracies, where a one-standard-deviation rise in satisfaction predicts roughly one additional point in a country's V-Dem score the following year².

Trust in democratic institutions is the other face of these evaluations, conceptually sitting in the specific system support for actors and institutions (Easton, 1965, 1975; Norris, 1999). Where satisfaction with democracy gauges how well the system works in practice, trust registers whether the institutions that run it are seen as legitimate and fair, and we argue the two move together. Citizens who judge institutional processes as fair extend both specific support to particular authorities and diffuse support to the system itself (Hetherington, 1998). Procedural justice is what maintains legitimacy and public confidence in democracy (Saxton, 2021; Magalhães and Aguiar-Conraria, 2019; Grimes, 2017; Magalhães, 2014), and it is precisely the channel through

²Once support for democracy enters the model, satisfaction's statistical significance vanishes while support retains its own, suggesting that the standalone association reflects satisfaction proxying for principled support rather than any independent effect of performance evaluations.

which experiences with the state's most present institution, the police, reach citizens' evaluations of broader institutions.

Overall, because police violence can diminish both the normative and the instrumental evaluations of the democratic system, we hypothesize the following:

- **Main Hypothesis:** *Police violence will negatively affect satisfaction with democracy and trust in its institutions.*

4 Study 1: Brazil's Experiment

To test our theory, we conducted an original two-wave online survey vignette experiment in Brazil³. Fieldwork was done throughout October 2023, and 1,004 respondents completed the first wave, and 624 completed both waves. We contracted Opinion Box, an online sample recruiter, to recruit subjects and conduct the study using the Qualtrics platform. We chose Brazil as a key case for studying the effects of police violence because of the high number of incidents and of its complex social and political dynamics.

The number of police killings in the country has been over six thousand a year since 2018, with 6,429 deaths in 2022 (Bueno and de Lima, 2019, 2023). Individuals belonging to non-white racial groups disproportionately experience the impacts of police misconduct. The average percentage of fatal victims who were identified as white was approximately 19.5 from 2018 to 2022 (Bueno and de Lima, 2019, 2020, 2021, 2022, 2023)⁴, which makes police violence a racialized issue in the country. For this reason, we pre-registered race as one of the possible moderators of the effects of exposure to police violence on satisfaction with democracy and political trust.

Brazil's political dynamics around police violence are also polarized. In July 2023, São Paulo, the most densely populated state in the country, saw a police operation launched in response to the killing of an officer in the region. Within a month, it had escalated into an extreme case of police violence, with 28 killings (Human Rights Watch, 2023). Tarcísio de Freitas, the right-wing State Governor at the time and an ally of far-right former President Jair Bolsonaro, described himself as "extremely satisfied" with the operation, even as residents of the region alleged summary executions and torture of non-criminals. The governor's reaction was not an isolated stance but part of a broader right-wing trend that took hold with Jair Bolsonaro's rise to the presidency in 2019, which brought a surge in right-wing populism alongside growing public support for the police and military. Bolsonaro, a former army captain, folded military personnel into his government, including a high-ranking officer as vice-president, and consistently promoted the armed forces' perceived morality and efficiency. That endorsement extended to the police as well, both through their military branch and through a persistent narrative that a strong police presence was essential to ending criminal violence (Watmough, 2021). Freitas's continued support for police killings after Bolsonaro left office shows the durability of the defense of *mano dura*. Given

³Our experimental hypotheses were pre-registered, anonymized materials are in the Supplementary Information (SI).

⁴White Brazilians represent 43.5% of the population, according to the 2022 Census.

this rhetoric, we also exploratorily examine Bolsonarismo (support for Bolsonaro) and dislike of Bolsonaro as possible moderators.

Additionally, we pre-registered support for (*Petismo*) and opposition to (*antipetismo*) the left-wing Workers' Party (*Partido dos Trabalhadores*, PT) as a moderator, since these two identities anchor Brazil's political polarization. Identification with the PT has been the most—and arguably the only—stable partisan identity in the country for nearly three decades (Carreirão and Rennó, 2018; Baker et al., 2016; Samuels, 2006), and the party's perceived failures, including corruption scandals and economic crises, have also driven widespread disaffection among voters (Samuels and Zucco, 2018).

Drawing from political psychology, among other possible traits, we chose to pre-register trait aggression personality as a possible moderator of the reactions to these violent events, particularly due to its effects on eliciting strong emotional responses to violence. The trait aggression personality battery measures people's aggressive tendencies in everyday life (Buss and Perry, 1992). Individuals with high levels of aggression exhibit a pronounced sensitivity to injustices, responding not with mere disapproval but with intense moral outrage. This reaction is deeply rooted in the aggressive person's inherent predisposition towards hostility in the face of perceived threats or violations of justice (Kalmoe, 2013). When exposed to incidents of police brutality, aggressive individuals are likely to experience a stronger reaction, perceiving these events as not only violations of individual rights but also as emblematic of broader systemic failures.

Having laid out the theoretical rationale for our moderators, we now turn to how the study was designed to capture their effects. In the first wave of Study 1, we measured baseline levels of our outcome variables and our moderators—race, political attachments, and trait aggression⁵. One week after completing the first wave, respondents were re-contacted to complete wave 2, in which they were randomly assigned to read one of the two vignettes. They later answered our outcome variable questions, constituting items about perceptions of procedural justice in the context of the situation described in the vignette (respect, voice, transparency, trustworthiness), one item related to trust in the police, three items related to trust in democratic institutions or “political trust” (trust in Congress, the Presidency, and the Judiciary), and one item tapping on subjects satisfaction with democracy.

There were two conditions within our manipulation: excessive police violence ($\text{violence}_i = 1$) and a non-violent policing scenario ($\text{violence}_i = 0$). The vignettes were summarized news articles describing an interaction between police officers and a middle-aged man during a traffic stop. The treatment with excessive police violence presents a scenario that escalates to violence when the man has a hard time finding his identification, resulting in an officer shooting the civilian in the leg. Conversely, the non-violent condition describes an appropriate interaction between the person and the police officers, where the person, though nervous, eventually finds his identification with no intervention from the officers.

We employ the Two-Way Fixed Effects (TWFE) estimator to assess the impact of this exposure to police violence. We define the treatment variable $D_{it} = \text{violence}_i \times \text{Post}_t$, where $\text{Post}_t = 1$ for observations from wave 2 and $\text{Post}_t = 0$ for wave 1 data. Consequently, $D_{it} = 0$ for all participants in wave 1, and $D_{it} = 1$ for those randomly assigned to the police violence vignette

⁵We also measured other demographics, political interest, turnout, and vote choice in the last presidential election.

condition during wave 2. Through randomization in the second wave, we ensure that individuals are comparable across both waves in terms of their pre-treatment characteristics and attitudes. Therefore, our design corresponds to a pre-post between-subjects experiment. The Average Treatment Effect (ATE) of interest can be conceptually illustrated as follows:

$$ATE = (\bar{Y}_{\text{wave}=2, \text{violence}=1} - \bar{Y}_{\text{wave}=1, \text{violence}=1}) - (\bar{Y}_{\text{wave}=2, \text{violence}=0} - \bar{Y}_{\text{wave}=1, \text{violence}=0}) .$$

To further investigate heterogeneity in treatment effects, following our theoretical predictions, we calculate Conditional Average Treatment Effects (CATE) by interacting D with one moderator at a time: feelings about the PT (pro-PT for a 4 or 5 on the five-point party item, neutral for a 3, anti-PT for a 1 or 2, with nobody dropped), a non-white dummy, and a high trait aggression disposition (trait aggression scale > median).⁶

The full regressions, including the overall ATE in the full sample and the CATEs, are reported in Table A1 in the SI. The ATE results indicate a significant negative impact of indirect exposure to police violence on both satisfaction with democracy, at -0.194 (SE 0.061 , $p < .001$), and trust in the police, at -0.405 (SE 0.089 , $p < .001$). However, we do not find substantial evidence that the treatment directly and unconditionally influenced political trust (-0.099 , SE 0.064 , $p = .062$).


As shown in Figure 1, the treatment increased distrust in the police in every subgroup. Because our hypotheses are directional, we report one-sided tests in the text and two-tailed tests in the tables. The contrast of interest is between pro-PT and anti-PT respondents. The effect is nearly twice as large in the first group, -0.535 (SE 0.138) against -0.285 (SE 0.112), with $p = .052$ for the difference, and the differences by trait aggression ($p = .017$) and by race ($p = .049$) run the same way. For political trust, the treatment left whites, anti-PT respondents, and individuals with low trait aggression unaffected. It was negative for pro-PT respondents (-0.180 , SE 0.100 , $p = .036$), non-whites (-0.201 , SE 0.077 , $p = .005$), and aggressive individuals (-0.204 , SE 0.081 , $p = .006$). For satisfaction with democracy, the effect is negative in every subgroup and largest among pro-PT respondents (-0.284 , SE 0.095 , $p = .001$). Its between-group differences are smaller, at $p = .067$ for feelings about the PT, $p = .029$ for trait aggression, and $p = .079$ for race.

Finally, in the SI, we include a table with estimates for the treatment groups on the items measuring perceptions of procedural justice, which we treat as a manipulation check. We find that perceptions of procedural justice regarding the behavior of the police officers in the violent vignette strongly erode perceptions of procedural justice in contrast to the non-violent vignette.

The results provide substantial evidence for our theoretical framework and previous research. Exposure to a police violence scenario sharply reduces trust in the police and lowers satisfaction with democracy. Its effect on political trust points the same way but is indistinguishable from zero overall, surfacing only among non-whites, pro-PT respondents, and individuals high in trait aggression, the groups in which every effect concentrates. Perceptions of police violence thus erode democratic attitudes, most visibly where citizens are primed to register it, though the conditional effects across race, partisanship, and aggression rest on weaker evidence.

⁶Please refer to the SI for the construction of these scales. Feelings about Bolsonaro, measured the same way on a matching item, appear as an exploratory fifth panel of Figure 1; that moderator was not pre-registered.

Figure 1: Treatment Effects



Note: A five-point item asks about the Workers' Party, running from 1, I hate the party, to 5, I am a petista. The pro-PT group answered 4 or 5, the neutral group answered 3, and the anti-PT group answered 1 or 2; every respondent is sorted and nobody is dropped. The Bolsonaro panel sorts the same way on a matching item about Bolsonaro. Each panel reads its own item alone; nothing is combined across items. High trait aggression is a score above the sample median on the twelve-item Buss and Perry battery, with each item rescaled 0 to 1 and summed. Non-white respondents did not report their race as white (*branca*). The anti-PT group is the omitted one in the PT panel and the pro-Bolsonaro group is the omitted one in the Bolsonaro panel. All outcomes are standardised, so estimates are in standard deviations.

5 Study 2: Chilean Shocks

Study 1 showed that exposure to police violence, manipulated within a controlled vignette experiment, reduces trust in the police and satisfaction with democracy. Study 2 asks whether the same mechanism holds when exposure is not a hypothetical scenario but actual, geographically bounded episodes of police violence, in a country where such abuse is rare. We turn to Chile, where the 2019 *estallido social* produced a natural experiment: government repression of protesters generated documented human rights violations concentrated in specific census zones, exposing some residents directly while leaving residents of otherwise similar zones untouched. We test whether this real-world shock produced the same erosion in political trust and democratic satisfaction found in Study 1, using panel survey data linked to geocoded violence records. First, Chile’s low pre-2019 baseline of police violence—unlike Brazil’s chronic pattern—lets us observe how respondents react when police abuse arrives as a rupture rather than a routine feature of daily life. Second, because exposure here is observed rather than assigned, the design captures residents’ actual experience of an unanticipated shock, an ecological validity no survey vignette can replicate.

That shock began on October 18, 2019, when a student-led campaign against a metro fare increase ignited the largest wave of protests in Chile in 30 years. High school students organized mass fare evasions, sit-ins, and demonstrations across Santiago. Within days, the movement expanded nationwide, culminating in over a million Chileans marching in Santiago and hundreds of thousands more joining protests across the country. The at the time government of right-wing President Sebastián Piñera responded with unprecedented police and military repression, deploying armed soldiers and tanks to the streets, resembling the repression of Pinochet’s dictatorship (1973-1990). State forces used tear gas, rubber bullets, and live ammunition to disperse crowds, using force against civilians regardless of their involvement in the protests. Within weeks, reports of severe human rights violations emerged, including arbitrary arrests, torture, sexual violence, and mutilations caused by rubber bullets, with hundreds of protesters suffering permanent eye injuries (Bonner and Dammert, 2022). Social media and independent press captured harrowing images of law enforcement officers violently confronting unarmed civilians. Between October and December 2019, at the height of the *estallido social*, violence from the police and the military resulted in at least 31 deaths, roughly 3,000 injuries, and roughly 20,000 arrests (Dammert and Sazo, 2020).

The year before the *estallido social*, there were no reported cases of arbitrary or unlawful killings by government agents (United States Department of State, 2019). Reports of torture and other cruel, inhuman, or degrading treatment were also rare. This pattern aligns with the historical Chilean trend after the end of the dictatorship in 1990, in which police violence has been less common than in other Latin American countries. Against this baseline, the *estallido social* marks a sharp break from Chile’s pre-2019 pattern, producing the zone-level police violence events this study treats as shocks.

We estimate the impact of these events on political trust and democratic satisfaction using panel data from the Longitudinal Social Study of Chile (ELSOC), an ongoing survey fielded annually since 2016. Our analytical sample focuses on the first four waves (2016–2019), capturing respondent attitudes before and immediately following the *estallido social*. Each respondent is geocoded to their residential census zone (*zona censal*), and zone identifiers are merged with

geo-referenced police violence records from the National Institute of Human Rights (INDH). The INDH registry documents 2,832 victims logged between October 2019 and March 2020, 88.1% of them reporting physical-injury consequences, including 153 cases of ocular trauma, 32 of them resulting in irreversible vision loss (Instituto Nacional de Derechos Humanos, 2020).⁷ A respondent is treated ($D_z = 1$) if INDH records a documented case of police violence in their own census zone on or before their Wave 4 interview date, and untreated ($D_z = 0$) otherwise. Standard errors are clustered at the census zone level to reflect the spatial unit of treatment assignment.

Police violence during the *estallido* did not fall randomly across Chile. The municipalities most likely to register a documented violation were often the sites of the heaviest protest. Comparing exposed and unexposed zones within them would confound the effect of repression with the effect of the unrest that drew it. We therefore separate municipalities by their 2019 protest volume: those where protest did not surge above their own historical range, and those where it did. A third difference then asks whether the response to exposure differs between them:

$$y_{it} = \sum_{\tau \in \{2016, 2017, 2019\}} [\beta_\tau + \delta_\tau \mathbf{1}[S_c \geq 1]] (D_z \times \mathbf{1}[t = \tau]) + \alpha_i + \gamma_{ct} + \varepsilon_{it}$$

Here D_z marks a respondent's census zone as exposed, and S_c measures a municipality's 2019 protest volume against its own 2009–2018 average in standard deviations. The term α_i absorbs everything fixed about a respondent, and γ_{ct} absorbs everything that moved a whole municipality in a given year. Three quantities follow. The first, β_{2019} , is the effect of exposure where protest did not surge. The second, δ_{2019} , is how much that effect differs where it did. Their sum is the implied effect in the municipalities that surged. We take 2018 as the reference year, which leaves 2016 and 2017 as placebos. No zone is exposed before 2019, so those coefficients ask whether zones that would later be exposed were already drifting apart from the rest—not whether an existing treatment had an effect. Where that joint test raises concern, we adjust for differential trends using each respondent's own baseline value of the outcome. We keep the adjusted version only when it both raises the joint p -value and shrinks the pre-period coefficients.

Table 1 reports the results; p -values in Studies 2 and 3 are two-tailed. Where protest did not surge, exposure raises the share of respondents who report no trust in political institutions by 0.187 (SE 0.062, $p = .003$) and the share dissatisfied with democracy by 0.201 (SE 0.066, $p = .002$). Figure 2 plots the year-by-year coefficients behind those estimates. Neither measure of police trust moves detectably, a null we interpret below. The high-minus-low difference runs the other way on both outcomes, -0.171 (SE 0.067, $p = .011$) for political distrust and -0.230 (SE 0.073, $p = .002$) for dissatisfaction, very nearly cancelling the first. Adding the two leaves an implied effect in the municipalities that surged of 0.016 ($p = .508$) and -0.029 ($p = .340$), neither distinguishable from zero. The pre-trend column of Table 1 tests each row's own pre-period coefficients. The protocol described above tests all of them jointly. Dissatisfaction is the one outcome to cross its trigger, at a joint p of .179. Its reported estimate therefore uses the baseline-adjusted specification, which raises that value to .659, while political trust clears the unadjusted joint test on its own at .799.

The census-zone boundary that defines exposure is one defensible line among several, and if these estimates turned on where exactly it falls, redrawing it should move them. We refit the


⁷The SI (Study 2) provides the full breakdown by violation type, instrument, and consequence.

Table 1: Full-sample triple-difference estimates of own-zone police-violence exposure.

Outcome	Effect in 2019 (SE)	p	Pre-trend p	Obs.	Respondents	Zones (treated)
<i>A. Municipalities where protest did not surge: treatment DiD</i>						
Political trust: low share	0.187 (0.062)	0.003	0.584	9,962	3,169	748 (123)
Satisfaction with democracy: not satisfied share	0.201 (0.066)	0.002	0.619	9,676	3,114	747 (123)
Police trust: level	-0.130 (0.191)	0.495	0.535	10,170	3,212	749 (123)
Police trust: low share	0.041 (0.087)	0.641	0.450	10,170	3,212	749 (123)
<i>B. High-minus-low protest: DDD contrast</i>						
Political trust: low share	-0.171 (0.067)	0.011	0.455	9,962	3,169	748 (123)
Satisfaction with democracy: not satisfied share	-0.230 (0.073)	0.002	0.946	9,676	3,114	747 (123)
Police trust: level	0.260 (0.206)	0.207	0.544	10,170	3,212	749 (123)
Police trust: low share	-0.079 (0.091)	0.384	0.795	10,170	3,212	749 (123)
<i>C. Municipalities that surged: implied treatment DiD</i>						
Political trust: low share	0.016 (0.025)	0.508	0.749	9,962	3,169	748 (123)
Satisfaction with democracy: not satisfied share	-0.029 (0.031)	0.340	0.482	9,676	3,114	747 (123)
Police trust: level	0.130 (0.077)	0.094	0.859	10,170	3,212	749 (123)
Police trust: low share	-0.039 (0.027)	0.157	0.667	10,170	3,212	749 (123)

Note: $S_c < 1$ is the no-surge group, $S_c \geq 1$ the surge group; Panel B is their DiD contrast, Panel C the implied surge effect (delta-method SE). Democratic dissatisfaction uses the protocol-selected baseline-outcome specification; other outcomes use the unadjusted DDD. Respondent and municipality-by-wave fixed effects, 2018 omitted, p two-sided, SEs clustered by zone.

Figure 2: Event-Study Estimates, No-Surge Municipalities



Note: Treatment-year coefficients for no-surge municipalities ($S < 1$), estimated as the reference group in the full-sample DDD. Red line: estallido onset (2019-10-18). The three share outcomes share one vertical scale (± 0.40) and can be compared directly; police trust: level is a coefficient on the raw 1-5 item, a different unit, and carries its own scale.

model with four alternative definitions, built from the same geocoded events and the same timing rule. Three of them count an event within 500, 1,000, or 1,500 meters of the respondent's zone centroid. The fourth counts adjacency instead of distance. Both effects keep their sign and rough size throughout. The institutional-trust effect runs from 0.187 in the own zone to 0.118 under the widest line, and dissatisfaction from 0.201 to 0.136; every high-minus-low difference stays negative. What they lose is precision. Under the 1,500-meter definition, 1,970 of 3,169 respondents count as exposed. The comparison group then consists largely of people who merely live further from an event, and both estimates fall just short of conventional levels ($p = .051$ and $p = .056$). We read the pattern as the boundary doing what a wider boundary should.

Study 1 found the effects concentrated among the groups most primed to register police violence, so we ask the same question here with the political identity this design observes: whether a respondent voted for the government whose police carried out the repression. Who a respondent voted for in 2017 conditions how much of this they register. Among those who did not vote for the incumbent President, right-wing Piñera, exposure moves police trust by -0.544 (SE 0.327), and, among those who did, it moves it by $+0.526$ (SE 0.354). The difference between those two responses is 1.070 (SE 0.433, $p = .016$), the one subgroup contrast in this study that separates cleanly. Splitting instead by left-right self-placement points the same way, but reaches conventional levels in neither arm—which is what a smaller and noisier partition of the same sample would be expected to do. Supporting the government in charge appears to buffer people against revising their view of the police, even when the violence occurred where they live.

That buffering points to why police trust does not move in the main comparison. Trust in the Carabineros collapsed everywhere, not only where violations were documented. Among respondents observed in both 2018 and 2019, the share placing themselves at the floor of the police trust item rose from 11.9% to 28.4% in exposed zones. In unexposed zones it rose from 9.8% to 24.0%. The raw change among unexposed respondents was -0.501 scale points, and a municipality-by-year term absorbs exactly that kind of common movement. A nationally televised institutional failure leaves little room for local variation to show up on the institution at its centre. Democratic satisfaction and trust in political institutions had no comparable national collapse, so proximity to a violation still registers there.

Cluster-robust inference depends on how much independent variation the design actually has, so we test what our own procedure does under the null by permuting treatment across zones within municipalities, and report the results in the SI⁸. Read with those diagnostics, Study 2 shows that living where the state used violence erodes confidence in democratic institutions. It does so independently of the national reassessment of the police that ran alongside it.

⁸SI, *Statistical Power and Minimum Detectable Effects*. The exercise re-randomizes exposure within municipalities and reports the rejection rate and the minimum detectable effect for every quantity in Table 1. A treated-cluster count is not a count of the information the design carries. Residual intra-zone correlation in the estimating equation is near zero once the fixed effects are in, so the design effect is essentially one and the effective sample is the 9,962 respondent-waves themselves; clustering on the zone in fact yields a smaller standard error than its heteroskedasticity-robust alternative on three of the four outcomes, this one included. The triple difference is estimated across 748 census zones, 123 of them exposed, in 89 municipalities, while the no-surge stratum's own difference-in-differences rests on 13 treated zones.

6 Study 3: Cross-National Evidence

Studies 1 and 2 estimate the effect of exposure to police violence within single countries. In Study 3, we address whether this pattern generalizes across Latin America. To assess the extent of the association, we use the 2008 wave of LAPOP. It is the only cross-national wave ever fielded in the region that asks respondents both our dependent variables of interest and whether the police mistreated them: “In the past 12 months, has any police officer mistreated you verbally, physically, or assaulted you? How many times?”⁹. The item records direct exposure reported by the respondent, which differs from the exposures we examined in the previous studies. That difference, and the cross-sectional design, are why we read what follows as corroboration of the pattern rather than as a replication of the effect.

Exposure here is neither assigned nor quasi-random, and the people the police mistreat are not a random draw of respondents. We therefore control for two available rival exposures, as well as multiple other covariates. Crime victimization is included so that the estimate does not simply track who encounters crime, and police bribe solicitation so that it does not track who encounters corrupt officials. Demographic and socioeconomic covariates enter alongside them. We estimate a survey-weighted linear multilevel regression, one per outcome. Each country receives its own intercept and its own slope on victimization; the controls enter with one slope shared across countries.


The three outcomes, again, are trust in the police, political trust, and satisfaction with democracy. Trust in police is a single seven-point item. Political trust is a three-item index combining confidence in the legislature, the executive, and the courts; we rescale each item to run from 0 to 1, add them, and standardize the sum. Satisfaction with democracy is a four-point item. We standardize every outcome within the dataset, so a coefficient gives the change in outcome standard deviations for a respondent who reports abuse relative to one who does not.

Respondents who report police abuse show lower trust in police, lower political trust, and lower satisfaction with democracy (Figure 3, Panel A). The drop is largest for trust in the police, at -0.402 (SE 0.042 , $p < .001$), but it also appears for the legislature, the executive, and the courts taken together, at -0.205 (SE 0.032 , $p < .001$), and for satisfaction with democracy, at -0.121 (SE 0.036 , $p = .006$). The association is negative in all sixteen countries on both trust measures. In the units respondents answered in, reporting abuse corresponds to about three-quarters of a point of trust in the police on its seven-point scale, and to under a tenth of a point of satisfaction with democracy on its four-point scale; 5.8 percent of respondents report abuse.

All evidence found so far implies the spillover effect we theorize. A violent encounter is information not only about this one agency, but about the fairness and trustworthiness of the state that arms and authorizes it, and because the police are the institution of that state most present in daily life, the judgment travels from them to the abstract national institutions charged with equality before the law. Trust is where that judgment registers, and satisfaction with democracy is where it surfaces, as a verdict on how the system works in practice. Read this way, the pathway makes two claims that can be checked separately: that abuse depresses trust in the police and, with it, trust in broader institutions, and that it is this broader institutional distrust that carries

⁹The item was asked of all respondents, with “don’t-know” responses coded to missing. The item records a count; we code victimization as one or more incidents.

Figure 3: Police Abuse, Trust, and Satisfaction with Democracy. Panel A: victimization and each outcome, LAPOP 2008. Panel B: trust and satisfaction across three surveys.



Note: All estimates use survey weights, and outcomes are standardised within dataset, so coefficients are in outcome standard deviations. Every row comes from a multilevel model that gives each country its own intercept and its own slope on the predictor of interest. Bars are 95% confidence intervals. Stars mark * $p < .05$, ** $p < .01$, *** $p < .001$. All three datasets cover the same sixteen countries. The appendix reports estimation details and full coefficients.

through to dissatisfaction with democracy.

Therefore, next, we ask whether the steps from trust to satisfaction with democracy that this mechanism requires are present in the available datasets. In Figure 3, Panel B, we report these steps. With the 2008 LAPOP, fitting the trust-to-satisfaction specification, with controls, political trust predicts satisfaction with democracy at 0.282 (SE 0.020, $p < .001$) per standard deviation of trust, roughly twice the size of the victimization coefficient on the same outcome. We explore two 2023 surveys that reproduce this fifteen years later and find that political trust enters at 0.269 (SE 0.017, $p < .001$) in the Latinobarómetro and 0.334 (SE 0.014, $p < .001$) in LAPOP, trust in police at 0.040 (SE 0.009, $p < .001$) and 0.039 (SE 0.015, $p = .019$). Political trust is the far stronger predictor, and its slope is positive in all sixteen countries of both 2023 surveys. The police trust slope is positive in all sixteen Latinobarómetro countries and in fourteen of the sixteen LAPOP countries, and the two exceptions sit near zero.

The two measures are recorded at the same moment, so rather than assume an order, we fit both in Table 2. Treating trust in police as the predictor, most of its association with satisfaction runs through political trust, 63.6% in LAPOP 2008, 79.2% in LAPOP 2023, and 67.2% in the Latinobarómetro. Reversed, almost none of political trust's association runs through trust in police, 51.4% in LAPOP 2008 but only 5.8% in both 2023 datasets. Two cautions belong with that decomposition. The two trust measures are not independent. Corrected for measurement error they correlate at .52 to .55, and a single component accounts for 52 percent of their variance, past the level at which constructs are conventionally treated as separable. And a decomposition of associations measured in the same cross-section reports which measure carries the association, not what would follow from changing trust. In the recent data, political trust carries the direct association, and trust in police reaches satisfaction with democracy mainly through it.

Three checks suggest this pattern is not an artifact of how the models are built. First, adding the two trust measures raises the share of variation the model accounts for from 0.190 to 0.269 in the Latinobarómetro and from 0.084 to 0.208 in LAPOP. Second, political trust's coefficient falls only from 0.268 to 0.219 when winner–loser status, perceptions of corruption, and presidential approval are added, which suggests that it is not standing in for other judgments about the regime. Third, any unmeasured variable that explained the association away would have to be a strong one. We can ask how closely such a variable would need to be tied to both trust and satisfaction to account for the entire coefficient, its E-value (VanderWeele and Ding, 2017). For political trust, that threshold is 2.04 in one survey and 2.22 in the other, while the strongest variable we actually measured reaches only 1.35. The same threshold is lower for trust in police, at 1.27 and 1.25. This is what the decomposition anticipates: a measure whose association with satisfaction runs almost entirely through another variable should be more sensitive to confounding than one acting directly, and the two findings point the same way. Taken as a whole, the results point to political trust as the more immediate carrier of the spillover from police violence to satisfaction with democracy, with trust in the police reaching that judgment largely at one remove.

We also probe whether these associations vary by ideology and by race, mirroring Study 1's moderators as far as these surveys allow. The victimization associations are stronger on the left: among left-identified respondents, reporting abuse corresponds to -0.304 (SE 0.054) on political trust and -0.243 (SE 0.052) on satisfaction with democracy, against -0.135 (SE 0.049) and -0.047 (SE 0.047) among right-identified respondents, with $p = .013$ and $p = .004$ for the differences; the police-trust association does not differ by ideology ($p = .517$). In the 2023

Table 2: The mediation decomposition, both orderings, all three datasets.

Ordering	Direct	Indirect	Total	% mediated	E-value	Benchmark RR
<i>LAPOP 2008</i>						
A: abuse → police → institutional → satisfaction	-0.044	-0.077**	-0.121**	63.6%	1.292	1.061
B: abuse → institutional → police → satisfaction	-0.059	-0.062**	-0.121**	51.4%	1.292	1.061
<i>LAPOP 2023</i>						
A: police → institutional → satisfaction	0.039**	0.148**	0.187**	79.2%	2.216	1.134
B: institutional → police → satisfaction	0.334**	0.020**	0.354**	5.8%	1.254	1.134
<i>Latinobarómetro 2023</i>						
A: police → institutional → satisfaction	0.040**	0.083**	0.123**	67.2%	2.037	1.346
B: institutional → police → satisfaction	0.269**	0.017**	0.285**	5.8%	1.266	1.346

Note: 95% CIs from a 1,000-resample bootstrap over countries. Stars: *** $p < .001$, ** $p < .01$, * $p < .05$. Each country's own indirect effect is negative – the direction the pathway predicts, since LAPOP 2008's exposure is police abuse rather than trust – in 16/16 (A) and 16/16 (B) in LAPOP 2008; positive in 16/16 (A) and 14/16 (B) in LAPOP 2023, and 16/16 (A) and 15/16 (B) in Latinobarómetro. E-value is the smaller of the two links', since a mediated effect resists confounding only as well as its weaker link. Benchmark RR is the strongest association any measured covariate has with the outcome, as $\max(RR, 1/RR)$; an E-value above it means the confounder would have to beat every covariate we measured.

surveys, the police-trust slope on satisfaction is somewhat larger on the right (difference $p = .023$ in LAPOP, $p = .056$ in the Latinobarómetro), while the political-trust slope does not differ. None of the seven white versus non-white differences is distinguishable from zero (p between .22 and .83). The full moderation tables are in the SI.

We run the same nine robustness checks on all three datasets. They include alternative definitions of the trust index, ordinal rather than linear outcome models, and E-value sensitivity. Refitting every model with country dummies and standard errors clustered on country—which assumes nothing about how countries are distributed—leaves every coefficient within a hundredth of its reported value, changes no sign and no standing against zero, and returns errors slightly smaller than the multilevel ones, so the intervals we report are the wider of the two. For the 2008 analysis we additionally re-estimate all three associations by weighting rather than regression adjustment—propensity, entropy-balancing and doubly robust schemes that equalize the covariate distributions of victimized and non-victimized respondents, each screened on overlap and balance before its estimate is read—and every scheme that clears that screen reproduces the reported coefficients without assuming that the controls enter linearly. The SI reports all of them, with country-level diagnostics and complete estimates for every equation.

Table 3: How the three studies triangulate the same conclusion.

	Exposure	Population	Quantity estimated	Contributes
Study 1	Written scenario of a violent police stop	Brazilian adults, on-line panel	Randomized effect of indirect exposure	Internal validity
Study 2	Documented repression in the respondent's census zone	Chilean panel respondents, 2016–2019	Difference-in-differences effect of local exposure	Ecological validity
Study 3	Personal victimization reported by the respondent	LAPOP respondents in sixteen countries	Covariate-adjusted association	External validity

Note: Every estimated effect or association on police trust, political trust, and satisfaction with democracy is negative in all three studies. Specifications, samples, and estimates appear in the study sections and the SI.

7 Discussion

The three studies return the same answer to the paper's question: citizens in our settings did not treat their evaluation of the police as separate from their evaluation of the regime. When the police turned violent, respondents marked down institutions that had no part in the encounter, along with their satisfaction with democracy itself. This runs against a common expectation, in which diffuse support is buffered against complaints about the day-to-day performance of individual agencies, and that buffering helps explain why democracies remain stable when particular institutions disappoint. Whatever protection the buffer affords elsewhere, we find none of it where the police are concerned; in all three studies, evaluations of the police entered judgements about the system itself. Table 3 summarizes how the three designs reach that conclusion.

Each study, taken alone, is open to an objection, and the objections do not overlap. The Brazilian experiment assigns exposure at random, but because the exposure is a written scenario it cannot tell us what happens when a person is actually harmed. The Chilean panel captures real repression, but where protest surged we cannot separate a genuinely absent effect from exposure so widespread that no untreated neighbor remains. The cross-national study reaches sixteen countries and respondents' own victimization, but it is observational, and no set of controls can reconstruct the counterfactual that a causal claim requires. Had the three designs shared a weakness, their agreement would mean little, since the same flaw could have produced the same wrong answer three times. Because the weaknesses differ, a skeptical reading of the pattern requires three unrelated failures to have pointed the same way at once. Nor did we expect the designs to produce comparable magnitudes, since they measure different exposures in different populations with different instruments. The prediction common to all three was directional, and that prediction held in each case.

The moderation results bear on the process behind these effects, and they distinguish exposure at a distance from exposure in person. Trust in the police, the judgement closest to the event, is the clearest case. Where exposure was indirect, the response depended on identity and prior politics: in Brazil the vignette moved pro-PT, non-white, and high-aggression respondents furthest, while

in Chile the one clean subgroup contrast divides residents by their vote for the president whose government commanded the repression. Direct victimization follows a different pattern, in which victims on the left and on the right lost trust in the police at similar rates; because the cross-national sample is by far our largest, that null cannot be attributed to thin data. We read the contrast as evidence about interpretation. A person the police mistreated has little need of prior beliefs to understand what happened, whereas a person hearing about violence at a distance must decide what to make of it. Prior beliefs reappear one step further on, since among direct victims ideology does moderate how far the damage extends to political trust and to satisfaction with democracy. Victims' politics shaped how far the experience generalized to the regime rather than their view of the police who harmed them.

The partisan pattern also has an implication for measurement. In both countries where a political moderator was available, the respondents who buffered their view of the police were those aligned with the side that champions them. In Chile, Piñera voters' trust in the police survived exposure that lowered it among everyone else; in Brazil, the vignette moved anti-PT respondents about half as far as their pro-PT counterparts. Trust in the police is usually read as an evaluation of performance, but on this evidence part of what it records is political loyalty. Because the item is a standard fixture of comparative surveys, analyses that treat it purely as a performance measure will absorb some partisan sorting into what they report as assessments of policing. Two mechanisms could produce the alignment, and our cases cannot separate them. Concern over police violence has long belonged to the Latin American left, so the pattern may follow the left whoever governs. Alternatively, it may reflect opposition to the government that commands the police, and in that case it should reverse when the left holds that command. In the periods we study, the salient champions of the police sat on the right, and the two mechanisms therefore predict the same observations.

These results support a bounded claim: exposure to police violence worsens citizens' evaluations of democratic institutions, in the same direction in every setting we examined. They do not support comparisons of magnitude across the three studies, whose different exposures and instruments make the estimates non-commensurable. Nor do they support conclusions about policy, since no design here estimates the consequences of any particular reform, or a prediction about the partisan moderator under a government of the left, for the reason given above. What remains inside these boundaries rests on three designs rather than one.

The paper's contribution is to move police violence from the study of policing into the study of democracy. Measured only by attitudes toward the police, misconduct appears to be a problem for one agency and its reformers. In our results, the damage extends to the institutions of the regime itself, under three kinds of exposure and in settings that range from a single national panel to a sixteen-country survey. A question follows that our designs were not built to pursue: which institutional arrangements weaken the connection between police conduct and system legitimacy. Political systems plausibly differ in how well they absorb shocks of this kind, and identifying the sources of that difference is the natural extension of these results.

8 Conclusion

This paper asked whether police violence damages the attitudes that sustain democracy. Three studies varied the exposure: a randomized vignette in Brazil, geocoded repression during Chile's *estallido social*, and self-reported victimization across sixteen Latin American countries. Each found that exposure lowers evaluations of the political system, not of the police alone, and that the political damage varies with identity and partisanship.

Latin American electorates are offered *mano dura* policing, election after election, as the remedy for crime. Our results record a cost that rarely enters that debate: police violence weakens public attachment to democratic institutions in the region where that attachment is already strained. The cost is also unlikely to fall evenly. In our studies the political damage was concentrated among citizens on the left, among non-whites, and among those primed to register abuse, while supporters of the side that champions the police were buffered against it. Hard policing may therefore do more than lower average support for the system: it can sort disaffection along political cleavages the region already contends with, concentrating doubts about the regime in the groups least persuaded that the state's coercion is used fairly. Debates over iron-fist policing usually turn on whether it reduces crime. On our evidence, what it does to citizens' confidence in democratic institutions belongs in that debate as well.

References

- Ang, DESMOND and JONATHAN Tebes. 2024. "Civic responses to police violence". *American Political Science Review* 118 (2): 972–987.
- Arzheimer, Kai. 2009. "Contextual factors and the extreme right vote in western europe, 1980–2002". *American Journal of Political Science* 53 (2): 259–275.
- Baker, A., B. Ames, A. E. Sokhey, and L. R. Rennó. 2016. "The dynamics of partisan identification when party brands change: The case of the workers party in brazil". *The Journal of Politics* 78 (1): 197–213.
- Bonner, Michelle D. and Lucía Dammert. 2022. "Constructing police legitimacy during protests: frames and consequences for human rights". *Policing and Society* 32 (5): 629–645.
- Boudreau, C., S. A. MacKenzie, and D. J. Simmons. 2019. "Police violence and public perceptions: An experimental study of how information and endorsements affect support for law enforcement". *The Journal of Politics* 81 (3): 1101–1110.
- Branton, R. P., T. E. Carey, and V. Martinez-Ebers. 2023. "Lethal engagement: The relationship between contextual exposure to police killings and external political efficacy". *Polit Behav* 45 : 1263–1283.
- Brinks, Daniel M. 2007. *The Judicial Response to Police Killings in Latin America: Inequality and the Rule of Law*. Cambridge University Press.

- Brodkin, E. Z. 2011. "Policy work: Street-level organizations under new managerialism". *Journal of Public Administration Research and Theory* 21 : i253–i277.
- Buehler, James W. 2017. "Racial/ethnic disparities in the use of lethal force by us police, 2010–2014". *American Journal of Public Health* 107 (2): 295–297.
- Bueno, S. and R. de Lima. 2019. "13º anuário brasileiro de segurança pública". Fórum Brasileiro de Segurança Pública.
- Bueno, S. and R. de Lima. 2020. "14º anuário brasileiro de segurança pública". Fórum Brasileiro de Segurança Pública.
- Bueno, S. and R. de Lima. 2021. "15º anuário brasileiro de segurança pública". Fórum Brasileiro de Segurança Pública.
- Bueno, S. and R. de Lima. 2022. "16º anuário brasileiro de segurança pública". Fórum Brasileiro de Segurança Pública.
- Bueno, S. and R. de Lima. 2023. "17º anuário brasileiro de segurança pública". Fórum Brasileiro de Segurança Pública.
- Buss, Arnold H. and Mark Perry. 1992. "The aggression questionnaire". *Journal of Personality and Social Psychology* 63 (3): 452–459.
- Canache, Damarys, Jeffery J Mondak, and Mitchell A Seligson. 2001. "Meaning and measurement in cross-national research on satisfaction with democracy". *Public opinion quarterly* 65 (4): 506–528.
- Carr, PATRICK J., LAURA Napolitano, and JESSICA Keating. 2007. "We never call the cops and here is why: A qualitative examination of legal cynicism in three philadelphia neighborhoods". *Criminology* 45 (2): 445–480.
- Carreirão, Y. S. and L. R. Rennó. 2018. "Presidential voting: partisanship, economy, ideology". In *Routledge Handbook of Brazilian Politics*, ed. B. Ames. New York: Routledge.
- Chouhy, C., A. Singer, and P. Lehmann. 2022. "From punitiveness to authoritarianism: Examining the sources of public support for the use of state violence in latin america". *International Criminology* 2 .
- Claassen, Christopher. 2020. "Does public support help democracy survive?". *American Journal of Political Science* 64 (1): 118–134.
- Cohen, Mollie J. and Amy Erica Smith. 2016. "Do authoritarians vote for authoritarians? evidence from latin america". *Research & Politics* 3 (4): 2053168016684066.
- Cooper, H., L. Moore, S. Gruskin, and N. Krieger. 2004. "Characterizing perceived police violence: implications for public health". *American Journal of Public Health* 94 (7): 1109–1118.

- Correll, J., B. Park, C. M. Judd, and B. Wittenbrink. 2002. "The police officer's dilemma: Using ethnicity to disambiguate potentially threatening individuals". *Journal of Personality and Social Psychology* 83 (6): 1314.
- Crabtree, K. and N. Yadon. 2022. "Remaining neutral?: White americans' reactions to police violence and policing". *Political Behavior* 1 : 355–375.
- Cruz, José Miguel. 2015. "Police misconduct and political legitimacy in central america". *Journal of Latin American Studies* 47 (2): 251–283.
- Culhane, Scott E., John H. Boman IV, and Kimberly Schweitzer. 2016. "Public perceptions of the justifiability of police shootings: The role of body cameras in a pre- and post-ferguson experiment". *Police Quarterly* 19 (3): 251–274.
- Dammert, Lucia and Diego Sazo. 2020. "Scapegoats of the 2019 chilean riots: From foreign intervention to riff-raff involvement". *The SAIS Review of International Affairs* 40 (2): 121–135.
- DeVylder, J. E., H. Y. Oh, B. Nam, T. L. Sharpe, M. Lehmann, and B. G. Link. 2017. "Prevalence, demographic variation and psychological correlates of exposure to police victimisation in four u.s. cities". *Epidemiology and Psychiatric Sciences* 26 (5): 466–477.
- Drolc, Cody A. and Kelsey Shoub. 2024. "The particular and diffuse effects of negative interactions on participation: Evidence from responses to police killings". *Policy Studies Journal* 52 (3): 623–646.
- Easton, D. 1965. *A Systems Analysis of Political Life*. New York: Wiley.
- Easton, David. 1975. "A re-assessment of the concept of political support". *British journal of political science* 5 (4): 435–457.
- Fernandez, K. E. and M. Kuenzi. 2010. "Crime and support for democracy in africa and latin america". *Political Studies* 58 (3): 450–471.
- French, J. H. 2013. "Rethinking police violence in brazil: Unmasking the public secret of race". *Latin American Politics and Society* 55 (4): 161–181.
- Geller, A., J. Fagan, T. Tyler, and B. G. Link. 2014. "Aggressive policing and the mental health of young urban men". *American Journal of Public Health* 104 : 2321–2327.
- Gomes, M. D. and J. Alves de Aquino. 2018. "Violência e satisfação com a democracia no brasil". *Opinião Pública* 24 (1): 209–238.
- Greenwald, A. G., M. A. Oakes, and H. G. Hoffman. 2003. "Targets of discrimination: Effects of race on responses to weapons holders". *Journal of Experimental Social Psychology* 39 : 399–405.
- Grimes, Marcia. 2017. "Procedural fairness and political trust". In *Handbook on political trust*, 256–269. Edward Elgar Publishing.

- Hetherington, Marc J. 1998. "The political relevance of political trust". *American political science review* 92 (4): 791–808.
- Holston, J. and T. Caldeira. 1998. "Democracy, law, and violence: disjunctions of brazilian citizenship". In *Fault Lines of Democracy in Post-Transition Latin America*, eds. Felipe Aguero and Jeffrey Stark. Miami: North-South Center Press.
- Hough, Mike, Jon Jackson, and Ben Bradford. 2017. "Policing, procedural justice and prevention". In *Handbook of crime prevention and community safety*, 274–293. Routledge.
- Human Rights Watch. 2023, November. "'they promised to kill 30': Police killings in Baixada Santista, São Paulo state, Brazil". Technical report, Human Rights Watch.
- Instituto Nacional de Derechos Humanos. 2020. "Mapa de la memoria: Violaciones a los derechos humanos en el contexto de la crisis social, octubre 2019–marzo 2020". INDH. <https://mapamemoria.indh.cl/>.
- Jackson, A. N., L. Fedina, J. DeVlyder, and R. P. Barth. 2021. "Police violence and associations with public perceptions of the police". *Journal of the Society for Social Work and Research* 12 (2): 303–326.
- Kalmoe, Nathan P. 2013. "From fistfights to firefights: Trait aggression and support for state violence". *Political Behavior* 35 : 311–330.
- Kleinig, J. 2014. "Legitimate and illegitimate uses of police force". *Criminal Justice Ethics* 33 (2): 83–103.
- Knuckey, J. and K. Hassan. 2022. "Authoritarianism and support for trump in the 2016 presidential election". *The Social Science Journal* 59 (1): 47–60.
- Kriekhaus, J., B. Son, N. M. Bellinger, and J. M. Wells. 2014. "Economic inequality and democratic support". *The Journal of Politics* 76 (1): 139–151.
- Laniyonu, Ayobami. 2018. "Police, politics and participation: The effect of police exposure on political participation in the united kingdom". *The British Journal of Criminology* 58 (5): 1232–1253.
- Linde, Jonas and Stefan Dahlberg. 2021. "Quality of government and political support". *The Oxford Handbook of the Quality of Government*: 228.
- Loveless, Matthew and Chiara Binelli. 2020. "Economic expectations and satisfaction with democracy: Evidence from italy". *Government and Opposition* 55 (3): 413–429.
- Lubbers, Marcel, Mérove Gijsberts, and Peer Scheepers. 2002. "Extreme right-wing voting in western europe". *European Journal of Political Research* 41 (3): 345–378.
- Magalhães, Pedro C and Luís Aguiar-Conraria. 2019. "Procedural fairness, the economy, and support for political authorities". *Political Psychology* 40 (1): 165–181.

- Magalhães, P. C. 2014. "Government effectiveness and support for democracy". *European Journal of Political Research* 53 (1): 77–97.
- Markarian, G. Agustin. 2023. "The impact of police killings on proximal voter turnout". *American Politics Research* 51 (3): 414–430.
- McLean, Kyle, Bryan Lee Miller, Andrew Pyle, and Olivia Bauwens. 2024. "Democratic policing, building trust, and willingness to call 911: Examining the relationship between law enforcement legitimacy and calling the police". *American Journal of Criminal Justice* 49 (2): 183–200.
- Milani, J., B. Bradford, and J. Jackson. 2018. "Police violence". In *Oxford Research Encyclopedia of Criminology and Criminal Justice*. Oxford University Press.
- Mishler, W. and R. Rose. 1996. "Trajectories of fear and hope: Support for democracy in post-communist europe". *Comparative political studies* 28 (4): 553–581.
- Morris, KEVIN T. and KELSEY Shoub. 2024. "Contested killings: The mobilizing effects of community contact with police violence". *American Political Science Review* 118 (1): 458–474.
- Mullinix, Kevin J., Toby Bolsen, and Robert J. Norris. 2021. "The feedback effects of controversial police use of force". *Global Crime* 43 (2): 881–898.
- Norris, Pippa. 1999. *Critical Citizens: Global Support for Democratic Government*. Oxford University Press.
- Pereira, A. W. and D. E. Davis. 2000. "New patterns of militarized violence and coercion in the americas". *Latin American Perspectives* 111 (27): 3–17.
- Pion-Berlin, D. and M. Carreras. 2017. "Armed forces, police and crime-fighting in latin america". *Journal of Politics in Latin America* 9 (3): 3–26.
- Ramiro, Luis. 2016. "Support for radical left parties in western europe: social background, ideology and political orientations". *European Political Science Review* 8 (1): 1–23.
- Reid, Jonathan C. 2024. "Vote choice in the context of lethal police violence: A research note". *Journal of Criminal Justice* 95 : 102319.
- Reny, Tyler T. and Benjamin J. Newman. 2021. "The opinion-mobilizing effect of social protest against police violence: Evidence from the 2020 george floyd protests". *American Political Science Review* 115 (4): 1499–1507.
- Samuels, D. 2006. "Sources of mass partisanship in brazil". *Latin American Politics and Society* 48 (2): 1–27.
- Samuels, D. and C. Zucco. 2018. *Partisans, Antipartisans, and Nonpartisans: Voting Behavior in Brazil*. New York: Cambridge University Press.

- Sarsfield, Rodolfo and Fabián Echegaray. 2006. "Opening the black box: How satisfaction with democracy and its perceived efficacy affect regime preference in latin america". *International Journal of Public Opinion Research* 18 (2): 153–173.
- Saxton, Gregory W. 2021. "Governance quality, fairness perceptions, and satisfaction with democracy in latin america". *Latin American Politics and Society* 63 (2): 122–145.
- Silva, Andrea, Diego Esparza, Valerie Martinez-Ebers, and Regina Branton. 2022. "Perceived police performance, racial experiences, and trust in local government". *Politics, Groups, and Identities* 10 (3): 343–366.
- Singh, Shane P and Quinton Mayne. 2023. "Satisfaction with democracy: A review of a major public opinion indicator". *Public Opinion Quarterly* 87 (1): 187–218.
- Sung, H., J. Capellan, and B. Barthuly. 2022. "Trust in the police and the militarisation of law enforcement in latin america". *Policing and Society* 32 (3): 311–340.
- Trelles, Alejandro and Miguel Carreras. 2012. "Bullets and votes: Violence and electoral participation in mexico". *Journal of Politics in Latin America* 4 (2): 89–123.
- Tyler, Tom R. 2003. "Procedural justice, legitimacy, and the effective rule of law". *Crime and justice* 30 : 283–357.
- Tyler, Tom R and E Allan Lind. 2001. "Procedural justice". *Handbook of justice research in law*: 65–92.
- United States Department of State. 2019, March. "2018 country reports on human rights practices: Chile". Technical report, Bureau of Democracy, Human Rights, and Labor.
- VanderWeele, Tyler J. and Peng Ding. 2017. "Sensitivity analysis in observational research: Introducing the e-value". *Annals of Internal Medicine* 167 (4): 268–274.
- Vargas, Andrés Santos, Enrique Seira, and Alberto Simpser. 2026. "Satisfaction with democracy predicts democratic behaviors". *Proceedings of the National Academy of Sciences* 123 (14): e2517671123.
- Vidigal, R. 2022. "Authoritarianism and right-wing voting in brazil". *Latin American Research Review* 57 (3): 554–572.
- Walker, Lee Demetrius and Genevieve Kehoe. 2013. "Regime transition and attitude toward regime: The latin american gender gap in support for democracy.". *Comparative Politics* 45 (2): 187–205.
- Watmough, Simon P. 2021. "Jair bolsonaro: Far-right firebrand and cheerleader for dictatorship". *Populism Studies*.
- Womick, J., T. Rothmund, F. Azevedo, L. A. King, and J. T. Jost. 2018. "Group-based dominance and authoritarian aggression predict support for donald trump in the 2016 u.s. presidential election". *Social Psychological and Personality Science* 10 (5): 643–652.

Supplementary Information A: Study 1: Brazil's Experiment

Ignacio Urbina and Vitória Sgorlon

Stony Brook University

September 2026

Contents

Data collection and experiment procedures	2
Scale validation	8
Other ways of grouping people	11
Feelings about the PT and about Bolsonaro	14
Main results	21
Who left the panel between waves	22
Ethics and human subjects	23
Manipulation check and randomization balance	24
Summary statistics	26
Pre-registration materials	26

Table A1: Fixed-effects ATE and CATE estimates. Panel A: average treatment effect of the violent vignette (vs. non-violent). Panels B–D: conditional treatment effects by feelings about the Workers’ Party, by race, and by trait aggression.

Term	Satisf. w. Democracy	Political Trust	Trust in Police
<i>Panel A: Average Treatment Effect (ATE)</i>			
D = 1{Violence × Post}	-0.194** (0.061)	-0.099 (0.064)	-0.405*** (0.089)
<i>Panel B: CATE by feelings about the PT</i>			
D (anti-PT)	-0.127† (0.076)	-0.026 (0.081)	-0.285* (0.112)
D × Pro-PT	-0.157 (0.105)	-0.154 (0.112)	-0.250 (0.153)
D × Neutral	-0.093 (0.102)	-0.116 (0.108)	-0.195 (0.148)
<i>Panel C: CATE by Non-White</i>			
D = 1{Violence × Post}	-0.131† (0.075)	0.010 (0.079)	-0.298** (0.109)
D × Non-White	-0.121 (0.085)	-0.211* (0.090)	-0.206† (0.125)
<i>Panel D: CATE by High Trait Aggression</i>			
D = 1{Violence × Post}	-0.120† (0.072)	-0.012 (0.076)	-0.285** (0.105)
D × High Trait Aggr.	-0.163† (0.086)	-0.192* (0.091)	-0.264* (0.125)
<i>Model parameters</i>			
Respondent fixed effects	Yes	Yes	Yes
Wave fixed effects	Yes	Yes	Yes
Observations	1,248	1,248	1,248
Individuals	624	624	624

Outcomes are z-standardized. Individual-clustered robust standard errors in parentheses. † p<0.10, * p<0.05, ** p<0.01, *** p<0.001. D = 1 for the violent vignette; Post = 1 in wave 2. Panel B sorts people by the PT item alone: 4 or 5 pro, 3 neutral, 1 or 2 anti; nobody is dropped. The anti-PT group is the omitted one, so the D row is its effect and each interaction below it is a difference from it.

Data collection and experiment procedures

Opinion Box recruits participants for online surveys. Its pool in Brazil exceeds a million people who have said they are willing to take surveys. We recruited a convenience sample directly through the platform, restricted to participants over 18 living in Brazil who speak Portuguese, and directed them to a survey via Qualtrics.

In wave 1, 1,004 participants completed a survey. It covered demographics (age, gender, education, race, residency); the 12-item Buss–Perry trait-aggression measure; voter turnout and vote choice in the most recent presidential election; political interest; feelings towards the PT, Lula, the PL, and Bolsonaro; trust in the police and institutions (President, Congress, Judiciary);

satisfaction with democracy; and crime victimization. In wave 2 (about a week later), 624 participants were randomized into one of two vignette conditions, then answered the same trust and satisfaction items again plus four procedural-justice items.

Table A2: Experimental Vignettes (in English)

Violent Policing Vignette	Non-Violent Policing Vignette
<i>You will read a summarized news article about a recent interaction involving the police. For anonymity purposes, the source, location, and the police department details are omitted, and people's names are altered.</i> According to the news report, a middle-aged man named Carlos was driving home from work when two police officers pulled him over. Carlos was slow in responding to the officers' commands, and when asked for his identification, he could not find it. The officers asked Carlos to step out of the car and kneel on the ground. Despite appearing nervous, Carlos complied with their orders. Carlos then began to argue with the officers, asking them to explain the reason for the traffic stop. In response, one of the officers drew his weapon and aimed it at Carlos. Carlos panicked and started running away from the officers. One of the officers chased after him and fired his gun, striking Carlos in the leg. Carlos fell to the ground, was gravely injured, and was taken to the hospital. He underwent surgery and remained in the hospital for several weeks.	<i>You will read a summarized news article about a recent interaction involving the police. For anonymity purposes, the source, the location, and the police department details are omitted, and people's names are altered.</i> According to the news report, a middle-aged man named Carlos was driving home from work when two police officers pulled him over. Carlos was slow in responding to the officers' commands, and when asked for his identification, he could not find it. The officers asked Carlos to step out of the car. Despite appearing nervous, Carlos complied with their orders. Carlos then began to argue with the officers, asking them to explain the reason for the traffic stop. As Carlos continued to discuss with the officers, one of them gently asked him to look for his identification in his car. After ten minutes of searching, Carlos was able to find it. According to Carlos's testimony, he had a rough week, and that is why he was agitated.

Table A3: Variables' Dictionary - Brazil's Survey Experiment

Variable Name	Question Wording	Response Options
TEXT0	consent_text	Yes, I agree to participate
ID	ID_respond	
Tipo de coleta	Collection type	c("Answered part 1 and 2", "Answered only part 1")
age	How old are you?	1=18-24 years, 2=25-34 years, 3=35-44 years, 4=45-54 years, 5=55-64 years, 6=65-74 years, 7=75 years or more
gender	What is your gender: male, female or other?	1=Male/male, 2=Female/female

Table A3: Variables' Dictionary - Brazil's Survey Experiment (continued)

Variable Name	Question Wording	Response Options
educ	What is the highest level of education you have achieved?	1=I did not attend school, 2=Incomplete primary education, 3=Complete primary education, 4=Incomplete secondary education, 5=Complete secondary education, 6=Incomplete university education, 7=Complete university education
race	Do you consider yourself a white, black, brown, indigenous or yellow person?	1=Asian, 2=White, 3=Indigenous, 4=Other, 5=Brown, 6=Black
state	Which state do you live in in Brazil? Write the name of the state and select from the list.	
city	Now, write the name of the city where you live and select it from the list.	
knowszip	Do you know your zip code?	c("Yes, I remember", "I prefer not to say", "I don't remember")
zipcode	What is your zip code? Add numbers only.	
traitaggr_q1	For each of the following statements, indicate whether the statement is true or false for you. - "I have difficulty controlling my temper."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q2	For each of the following statements, indicate whether the statement is true or false for you. - "Sometimes I lose control for no good reason."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q3	For each of the following statements, indicate whether the statement is true or false for you. - "I get irritated quickly, but I get over it quickly."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q4	For each of the following statements, indicate whether the statement is true or false for you. - "Sometimes I feel like life is unfair to me."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q5	For each of the following statements, indicate whether the statement is true or false for you. - "It seems like other people are always luckier."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true

Table A3: Variables' Dictionary - Brazil's Survey Experiment (continued)

Variable Name	Question Wording	Response Options
traitaggr_q6	For each of the following statements, indicate whether the statement is true or false for you. - "I wonder why I sometimes feel so bitter about things."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q7	For each of the following statements, indicate whether the statement is true or false for you. - "There are people who pressured me so much that we even physically attacked each other."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q8	For each of the following statements, indicate whether the statement is true or false for you. - "If there is enough provocation, I am capable of hitting a person."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q9	For each of the following statements, indicate whether the statement is true or false for you. - "I threatened people I know."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q10	For each of the following statements, indicate whether the statement is true or false for you. - "Often, when I realize it, I'm disagreeing with people."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q11	For each of the following statements, indicate whether the statement is true or false for you. - "I can't avoid arguments when people disagree with me."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
traitaggr_q12	For each of the following statements, indicate whether the statement is true or false for you. - "My friends say I'm a bit confrontational."	1=Completely false, 2=Mostly false, 3=Slightly false, 4=Slightly true, 5=Mostly true, 6=Completely true
turnout22	Did you vote in the second round of the last presidential election, in 2022?	1=No, 2=I don't remember, 3=Yes
votechoice22	Who did you vote for?	1=Jair Messias Bolsonaro, 2=Luiz Inácio Lula da Silva, 3=Null/Blank
political_interest	What is your level of interest in politics? Would you say you are very interested, interested, somewhat interested, or not at all interested?	1=Not at all interested, 2=A little interested, 3=Interested, 4=Very interested

Table A3: Variables' Dictionary - Brazil's Survey Experiment (continued)

Variable Name	Question Wording	Response Options
pl_feeling	Which of the following sentences best describes your feelings towards the PL (Liberal Party)?	1=I hate the PL, 2=I don't like the PL, but I don't hate it, 3=I neither like nor dislike the PL, 4=I like it, but I don't feel like a PL supporter, 5=I'm a PL supporter
pt_feeling	Which of the following sentences best describes your feelings towards the PT (Workers' Party)?	1=I hate the PT, 2=I don't like the PT, but I don't hate it, 3=I neither like nor dislike the PT, 4=I like it, but I don't feel like a petista (PT supporter), 5=I'm a petista (PT supporter)
lula_feeling	Which of the following sentences best describes your feelings towards Luiz Inácio Lula de Silva?	1=I hate Lula, 2=I don't like Lula, but I don't hate him, 3=I neither like nor dislike Lula, 4=I like Lula, but I don't feel like a Lulista, 5=I like Lula a lot, I am a Lulista
bolso_feeling	Which of the following sentences best describes your feelings towards Jair Messias Bolsonaro?	1=I hate Bolsonaro, 2=I don't like Bolsonaro, but I don't hate him, 3=I neither like nor dislike Bolsonaro, 4=I like Bolsonaro, but I don't feel like a Bolsonarista, 5=I really like Bolsonaro, I am a Bolsonarista
trust_police_w1	For each of the following institutions, indicate whether you believe them to be very trustworthy, moderately trustworthy, somewhat trustworthy, or not at all trustworthy. - Police	1=Not at all trustworthy, 2=Slightly trustworthy, 3=Moderately trustworthy, 4=Very trustworthy
trust_president_w1	For each of the following institutions, indicate whether you believe them to be very trustworthy, moderately trustworthy, somewhat trustworthy, or not at all trustworthy. - President	1=Not at all trustworthy, 2=Slightly trustworthy, 3=Moderately trustworthy, 4=Very trustworthy
trust_congress_w1	For each of the following institutions, indicate whether you believe them to be very trustworthy, moderately trustworthy, somewhat trustworthy, or not at all trustworthy. - National Congress	1=Not at all trustworthy, 2=Slightly trustworthy, 3=Moderately trustworthy, 4=Very trustworthy

Table A3: Variables' Dictionary - Brazil's Survey Experiment (continued)

Variable Name	Question Wording	Response Options
trust_judiciary_w1	For each of the following institutions, indicate whether you believe them to be very trustworthy, moderately trustworthy, somewhat trustworthy, or not at all trustworthy. - The Judiciary	1=Not at all trustworthy, 2=Slightly trustworthy, 3=Moderately trustworthy, 4=Very trustworthy
satisf_democ_w1	In general, are you very satisfied, a little satisfied, Slightly satisfied or not at all satisfied with the functioning of democracy in Brazil?	1=Not at all satisfied, 2=Slightly satisfied, 3=Somewhat satisfied, 4=Very satisfied
crime_vic	Now, changing the subject, have you been the victim of any type of crime in the last 12 months? In other words, have you been a victim of aggression, robbery, kidnapping, fraud, blackmail, extortion, violent threats or any other type of crime in the last 12 months?	1=No, 2=Yes
procjust_dignity_C	To what extent do you believe Carlos was treated with dignity and respect by the officers during the interaction?	1=Not at all, 2=A little, 3=Moderately, 4=Very much, 5=Completely
procjust_voice_C	In your opinion, to what extent do you believe Carlos had the opportunity to express his views and concerns during his interaction with the police officers?	1=Not at all, 2=A little, 3=Moderately, 4=Very much, 5=Completely
procjust_impartial_C	Based on the information provided, to what extent do you believe the decision-making process of the officers involved in the interaction with Carlos was impartial and transparent?	1=Not at all, 2=A little, 3=Moderately, 4=Very much, 5=Completely
procjust_goodintent	How well-intentioned do you believe the officers were during their interaction with Carlos?	1=Not at all, 2=A little, 3=Moderately, 4=Very much, 5=Completely
procjust_dignity_T	To what extent do you believe Carlos was treated with dignity and respect by the officers during the interaction?	1=Not at all, 2=A little, 3=Moderately, 4=Very much, 5=Completely
procjust_voice_T	In your opinion, to what extent do you believe Carlos had the opportunity to express his views and concerns during his interaction with the police officers?	1=Not at all, 2=A little, 3=Moderately, 4=Very much, 5=Completely

Table A3: Variables' Dictionary - Brazil's Survey Experiment (continued)

Variable Name	Question Wording	Response Options
procjust_impartial_T	Based on the information provided, to what extent do you believe the decision-making process of the officers involved in the interaction with Carlos was impartial and transparent?	1=Not at all, 2=A little, 3=Moderately, 4=Very much, 5=Completely
procjust_goodintent_T	How well-intentioned do you believe the officers were during their interaction with Carlos?	1=Not at all, 2=A little, 3=Moderately, 4=Very much, 5=Completely
trust_police_w2	For each of the following institutions, indicate whether you believe them to be very trustworthy, moderately trustworthy, somewhat trustworthy, or not at all trustworthy. - Police	1=Not at all trustworthy, 2=Slightly trustworthy, 3=Moderately trustworthy, 4=Very trustworthy
trust_president_w2	For each of the following institutions, indicate whether you believe them to be very trustworthy, moderately trustworthy, somewhat trustworthy, or not at all trustworthy. - President	1=Not at all trustworthy, 2=Slightly trustworthy, 3=Moderately trustworthy, 4=Very trustworthy
trust_congress_w2	For each of the following institutions, indicate whether you believe them to be very trustworthy, moderately trustworthy, somewhat trustworthy, or not at all trustworthy. - National Congress	1=Not at all trustworthy, 2=Slightly trustworthy, 3=Moderately trustworthy, 4=Very trustworthy
trust_judiciary_w2	For each of the following institutions, indicate whether you believe them to be very trustworthy, moderately trustworthy, somewhat trustworthy, or not at all trustworthy. - The Judiciary	1=Not at all trustworthy, 2=Slightly trustworthy, 3=Moderately trustworthy, 4=Very trustworthy
satisf_democ_w2	In general, are you very satisfied, Slightly satisfied or not at all satisfied with the way democracy works in Brazil?	1=Not at all satisfied, 2=Slightly satisfied, 3=Somewhat satisfied, 4=Very satisfied

Scale validation

The *Attachment to the PT scale* adds the feeling-thermometer items toward the PT and Lula, each rescaled to run from 0 to 1. The *Political Trust scale* adds trust in Congress, the Presidency, and the Judiciary, each rescaled the same way. The *Trait Aggression scale* is the twelve-item Buss–Perry Aggression Questionnaire (Buss and Perry, 1992). Cronbach's alpha for each scale

is reported in the note under its polychoric-correlation table below.

Table A4: Polychoric correlations, wave-1 partisanship and institutional-trust items (scale-validation subset).

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)
(1) rev_pl_feeling	1	0.55	0.28	0.31	0.28	0.07	0.18	-0.1	0.28
(2) rev_bolso_feeling	0.55	1	0.57	0.58	0.43	0.18	0.33	-0.12	0.5
(3) pt_feeling	0.28	0.57	1	0.84	0.61	0.32	0.45	-0.03	0.54
(4) lula_feeling	0.31	0.58	0.84	1	0.62	0.29	0.43	-0.03	0.55
(5) trust_president_w1	0.28	0.43	0.61	0.62	1	0.56	0.64	0.15	0.59
(6) trust_congress_w1	0.07	0.18	0.32	0.29	0.56	1	0.63	0.33	0.4
(7) trust_judiciary_w1	0.18	0.33	0.45	0.43	0.64	0.63	1	0.31	0.56
(8) trust_police_w1	-0.1	-0.12	-0.03	-0.03	0.15	0.33	0.31	1	0.09
(9) satisf_democ_w1	0.28	0.5	0.54	0.55	0.59	0.4	0.56	0.09	1

Cronbach's alpha: Political trust (3 items) = 0.83; Attachment to the PT (2 items) = 0.92; full partisanship set (4 items) = 0.81.

Table A5: Polychoric correlations, the 12 Buss–Perry trait-aggression items.

	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)	(9)	(10)	(11)	(12)
(1) traitaggr_q1	1	0.73	0.51	0.39	0.36	0.4	0.39	0.41	0.44	0.38	0.49	0.33
(2) traitaggr_q2	0.73	1	0.5	0.42	0.4	0.45	0.47	0.43	0.52	0.4	0.51	0.33
(3) traitaggr_q3	0.51	0.5	1	0.46	0.37	0.38	0.28	0.3	0.28	0.41	0.39	0.32
(4) traitaggr_q4	0.39	0.42	0.46	1	0.77	0.73	0.36	0.36	0.33	0.38	0.33	0.26
(5) traitaggr_q5	0.36	0.4	0.37	0.77	1	0.71	0.34	0.34	0.31	0.32	0.31	0.22
(6) traitaggr_q6	0.4	0.45	0.38	0.73	0.71	1	0.38	0.37	0.36	0.36	0.34	0.28
(7) traitaggr_q7	0.39	0.47	0.28	0.36	0.34	0.38	1	0.57	0.61	0.37	0.43	0.38
(8) traitaggr_q8	0.41	0.43	0.3	0.36	0.34	0.37	0.57	1	0.51	0.38	0.45	0.32
(9) traitaggr_q9	0.44	0.52	0.28	0.33	0.31	0.36	0.61	0.51	1	0.38	0.46	0.38
(10) traitaggr_q10	0.38	0.4	0.41	0.38	0.32	0.36	0.37	0.38	0.38	1	0.51	0.49
(11) traitaggr_q11	0.49	0.51	0.39	0.33	0.31	0.34	0.43	0.45	0.46	0.51	1	0.47
(12) traitaggr_q12	0.33	0.33	0.32	0.26	0.22	0.28	0.38	0.32	0.38	0.49	0.47	1

Cronbach's alpha (12 items) = 0.89.

How we sort people into party groups

We sort people into three groups on each side of Brazilian politics, using a five-point question about each party or leader. The question asks which sentence best describes how the respondent feels. Its answers run from “I hate the PT” through “I neither like nor dislike the PT” to “I am a petista,” and the same for the PL, for Lula and for Bolsonaro. People who choose one of the top two answers count as supporters. People who choose one of the bottom two count as opponents. Everyone at the midpoint counts as neutral. Each side reads one question alone — the party question on the left, the Bolsonaro question on the right — so every respondent is sorted and nobody is dropped. The questions come from Batista Pereira et al. (2022), whose own grouping instead reads the party and Lula questions together on the left; that alternative, and a median split on the combined scale, are scored against the rule we use below. Every answer comes from the first survey, before anyone saw a vignette.

Tables A6 and A7 check the measure against information it was not built from. We compare it with the two other ways of sorting the same answers, and score all three on the answers about the other side, on who people said they voted for in 2022, and on how satisfied they were with democracy and how much they trusted the president before the experiment. The rule we use matches the answers about the other side at -0.540 for the PT groups and -0.548 for the Bolsonaro groups, and explains 0.485 and 0.495 of how people voted. The two-question alternative scores somewhat higher on several checks, at -0.565 against the other side and 0.574 of vote explained on the PT side, but part of that edge reads the Lula question, one of its own inputs, and it sorts 22 respondents into no group where the rule we use sorts everyone. The tables report every number for all three rules.

Table A6: Three ways of sorting people into party groups, and how well each one matches things it was not built from.

Check	Attachment to the PT			Attachment to Bolsonaro		
	Party answer only (used)	Party and Lula together	Median split on the scale	Bolsonaro answer only (used)	Bolsonaro and party, strict neutrals	Median split on the scale
Opponents (percent)	47.3	52.6	—	—	—	—
Matches own side	0.790	0.889	circular	0.493	0.513	circular
Matches other side	-0.540	-0.565	-0.513	-0.548	-0.578	-0.430
Ranks voters	0.903	0.921	0.896	0.907	0.915	0.793
Predicts vote	0.868	0.901	0.897	0.863	0.880	0.776
Vote explained	0.485	0.574	0.518	0.495	0.546	0.299
People	744	728	744	744	661	744

Read each column as one way of sorting people, and each row as a separate check on it. All three ways use the same five-point ratings, measured before the vignette. The first way is the one we use: it reads one rating alone — the party rating on the left, the Bolsonaro rating on the right — making someone a supporter on one of the top two answers, an opponent on one of the bottom two, and neutral at the midpoint, so every respondent is sorted. The second way reads both ratings for a side with the same cutoffs, making someone a supporter on either top-two answer and an opponent on either bottom-two, and sorting nobody who is warm about one and cold about the other. The third averages the two ratings for a side and splits the sample at the middle value, which produces no opponent group, so its share is left blank. Opponents (percent) is the share of the sample placed in the opponent group. Matches own side correlates the sorting with the remaining rating for that same side — Lula for the PT groups, the PL for the Bolsonaro groups. For the way we use this is a clean out-of-sample check. Read that row with care in the second column: the two-question way reads the Lula rating too, so its number is in part a comparison with one of its own inputs. The third way is built entirely from that rating, so its number would compare it with itself and is left out. Matches other side correlates the sorting with the rival camp's leader, which none of the three ways uses, and should run negative. The last three rows come from predicting who someone said they voted for in 2022, using the sorting alone. Ranks voters is the chance the sorting puts a voter for that side above a voter for the other side. Predicts vote is the share of people it gets right. Vote explained is how much of the difference in vote it accounts for. People who spoiled their ballot or did not answer are left out. None of the three ways was built using vote. On the PT side the two-question way scores above the way we use on every check, and places more people in the opponent group, 52.6 per cent against 47.3. The first of those two facts is the one to discount: three of the four checks it wins on read the Lula rating, which is an input to that way and not to the one we use. Its opponent share is larger because coldness towards Lula counts as opposition there, where the way we use requires coldness towards the party itself. What the two-question way costs is coverage: it sorts 22 people into no group at all.

Table A7: The same three ways of sorting people, compared against two attitudes measured before the vignette.

Side	How people are sorted	Satisfaction with democracy	Trust in the president
Attachment to the PT	Party answer only (used)	0.502	0.581
Attachment to the PT	Party and Lula together	0.525	0.601
Attachment to the PT	Median split on the scale	0.496	0.569
Attachment to Bolsonaro	Bolsonaro answer only (used)	-0.481	-0.418
Attachment to Bolsonaro	Bolsonaro and party, strict neutrals	-0.509	-0.447
Attachment to Bolsonaro	Median split on the scale	-0.374	-0.349

Each number is a correlation, on all 1,004 people who answered the first survey. We ran the survey in November 2023, when Lula was president, so people sorted towards the PT should be more satisfied and more trusting of the president, and people sorted towards Bolsonaro less so. Every one of the twelve numbers runs in that direction. On all four comparisons the way we use sits below the two-question alternative. One caution. These two attitudes are also the outcomes the experiment measures. People's party loyalties should shape how they rate a president they voted for or against, so the correlations mean something. But a sorting that scores higher here is being judged on the same kind of attitude the experiment is about, which the vote comparison in the previous table is not.

Other ways of grouping people

The main text builds the groups from the party question alone. The natural question is what that single question misses about Lula. The two answers correlate at .89. Of the 280 people who gave the middle answer about the party, 110 held a clear view about Lula: 55 said they like or love him and 55 said they dislike or hate him. The party answer alone leaves all 110 among the neutrals; the two-question rule below sorts them.

Table A8 sets out how each way of sorting people builds its three groups, cell by cell. Panel A covers the Workers' Party and Lula. Panel B covers Bolsonaro. Each row states what a respondent must answer to be placed in each group under that rule.

First, the rules read answers, not scores, except for the two median splits. Those two divide a scale that adds the answer about a party to the answer about its leader, each rescaled to run from 0 to 1. Splitting a scale gives two groups rather than three. Second, order matters where a rule reads two questions. Someone who answers 1 or 2 about the party counts as an opponent even if they answer 4 or 5 about Lula. Thirteen people answer that way. Third, only one rule leaves anyone out. The strict-neutral Bolsonaro rule requires a respondent at the Bolsonaro midpoint to be at the party midpoint as well; 76 of the 624 people who answered both surveys sit at the Bolsonaro midpoint but not at the party midpoint.

Table A8: How each way of sorting people builds its three groups.

Sorting rule	Supporter	Nonpartisan	Opponent
<i>Panel A: the Workers' Party and Lula</i>			
Party only (used)	4 or 5 about the party	3 about the party	1 or 2 about the party
Party and Lula	4 or 5 about the party or Lula	the rest, and nobody warm about one and cold about the other	1 or 2 about the party or Lula
Median split	above the median	—	at or below the median
<i>Panel B: Bolsonaro</i>			
Bolsonaro only (used)	4 or 5 about Bolsonaro	3 about Bolsonaro	1 or 2 about Bolsonaro
Bolsonaro, strict	4 or 5 about Bolsonaro	3 about Bolsonaro and 3 about the party	1 or 2 about Bolsonaro
Median split	above the median	—	at or below the median
Every rule reads one or two five-point questions. Each question asks which sentence best describes how the respondent feels, and the answers run from 1, I hate the party, to 5, I am a petista, with the same wording for the PL, for Lula and for Bolsonaro. The rule we use reads one question per side and sorts everyone: nobody is dropped and no answers are combined across questions. The two-question alternative reads both questions on a side, so warmth about either makes a supporter and coldness about either an opponent; the 22 people who are warm about one and cold about the other are sorted into no group at all under it. The two median splits divide a scale rather than reading answers. The party and Lula scale adds the answer about the party to the answer about Lula, each rescaled to run from 0 to 1. The Bolsonaro and PL scale does the same with the answers about Bolsonaro and the PL. Splitting a scale produces two groups rather than three, so those rows have no nonpartisan group. One further rule leaves people out. Sorting by the Bolsonaro answer with strict neutrals requires a respondent at the Bolsonaro midpoint to be at the party midpoint too, and 76 of the 624 people who answered both surveys are not.			

Figure A1 re-estimates the same models under five ways of grouping people. Every block reports the same three outcomes, and in each block the omitted group is the bottom row. All p -values below are two-tailed.

The figure and the numbers below it answer different questions. Each bar in the figure is one group's own effect against zero, so a bar clear of zero means the vignette moved that group. Every number below is a gap between two groups. A gap is a separate quantity, and its uncertainty is larger, so the two readings can disagree. Under the party answer alone, the pro-PT group moves -0.535 and the anti-PT group moves -0.285 , and both intervals clear zero. The gap between them, -0.250 , has an interval from -0.551 to 0.051 and does not clear zero.

The first block is the two-question grouping of Batista Pereira et al. (2022), which reads the party and the Lula answers together. Among the 624 people who answered both surveys it gives 181 pro-PT or pro-Lula respondents, 98 nonpartisans and 331 anti-PT or anti-Lula. The remaining 14 answer in the top two about one of the two questions and in the bottom two about the other, and that rule does not sort them.

The second block is the grouping used in the main text: the same three groups built from the party answer alone. This rule sorts everyone, including the 14 left unsorted by the two-question rule. Among the 624 it gives 152 pro-PT respondents, 169 neutrals and 303 anti-PT. Under the party-only grouping the supporter-opponent gap is smaller on all three outcomes: 0.250 against 0.325 on trust in the police, 0.157 against 0.186 on satisfaction with democracy, and 0.154 against 0.167 on political trust. The first number in each pair comes from the grouping we use. None of the three gaps is distinguishable from zero two-tailed under the grouping we use ($p=0.103$, 0.135 and 0.168; one-tailed 0.052, 0.067 and 0.084). Under the two-question rule

the trust-in-police gap is distinguishable from zero ($p=0.024$) and the satisfaction gap is close ($p=0.062$).

The third and fourth blocks test one attachment at a time. Both have three groups, and both use the same omitted base: the 303 people who dislike the Workers' Party. These two blocks read the party answer alone for the base: a base built partly from the Lula answer would depend on what the third block is testing. The third asks about Lula: 164 respondents rate him 4 or 5. The fourth asks about the party: 152 rate it 4 or 5. The middle row of each is the same 98 people, who answer 3 to both questions. Neutral here means the midpoint on both questions, not a residual category.

Everyone else is left out of the block — 59 people in the third and 71 in the fourth. They report neither the attachment being tested nor a settled neutrality. Assigning them to either group would impose a label they did not choose. Groupings that classify less than the whole sample appear beside the others rather than instead of them. The baseline is defined by the party answer alone, so it includes the eight respondents who dislike the party while liking Lula. The main text's grouping reads the same answer and sorts them the same way. The two-question rule instead leaves those eight unsorted: they are cold about the party and warm about Lula.

The two blocks together show which item produces the moderation: Lula. Attachment to Lula separates people from the anti-PT group on two of the three outcomes: 0.205 on satisfaction with democracy ($p=0.048$) and 0.398 on trust in the police ($p=0.008$). The third, political trust, sits at 0.193 ($p=0.077$). Attachment to the party separates them on none: 0.157 ($p=0.130$), 0.250 ($p=0.108$) and 0.154 ($p=0.170$). The two baselines are identical people and the two attached groups differ by only twelve respondents, so the contrast is not a matter of sample size.

The neutral row is flat in both blocks, 0.028 on satisfaction and 0.061 on trust in the police, neither distinguishable from the anti-PT row beneath it. People with no view of either the party or its leader look like people who dislike the party.

Read together, the two blocks show that the Lula answer adds information the party answer does not. The party question alone does not separate the groups; Lula does. The two-question measure counts a respondent as attached if either answer is warm; it separates the groups further than either question on its own.

The last block groups people by support for Bolsonaro. We did not pre-register this comparison, since H5 names Lula and the Workers' Party. It separates pro-Bolsonaro from anti-Bolsonaro respondents on trust in the police (0.270, $p=0.062$) but not on the other two outcomes (0.124, $p=0.205$ for satisfaction with democracy, and 0.006, $p=0.956$ for political trust). It also leaves 76 of the 624 people unclassified, because they give the middle answer about Bolsonaro while holding a view about the party.




Figure A1: The same models under five ways of grouping people.

Note: Each block groups people a different way and re-estimates the same models. The first block is the grouping used in the main text, built from the answers about the Workers' Party and about Lula. The second uses the answer about the party alone. The third and fourth test one attachment at a time against the same baseline, the 303 people who dislike the party. In the third that attachment is Lula: 164 respondents rate him 4 or 5. In the fourth it is the party: 152 rate it 4 or 5. The middle row of both is the same 98 people, who answer 3 to both questions, so neutral here means the midpoint on both rather than whatever is left over. Everyone else is left out of the block, 59 people in the third and 71 in the fourth, because they hold neither the attachment being tested nor a settled neutrality. The last block uses the answer about Bolsonaro. Each point is one group's own effect and each bar is its 95 percent confidence interval, so a bar away from zero means the vignette moved that group. It does not mean the groups differ from each other. The gap between two groups is a separate quantity with its own and larger uncertainty, and those are the p-values reported in the text and in the contrasts table. All outcomes are standardised, so estimates are in standard deviations.

Feelings about the PT and about Bolsonaro

Figure A2 reads each question on its own. People who answer 4 or 5 are counted as pro, people who answer 3 as neutral, and people who answer 1 or 2 as anti. Nothing is combined across questions and nobody is left out, so each block covers all 624 people who answered both surveys.

Both blocks order the groups as the theory expects, and the two orderings are mirror images. On trust in the police, the effect is largest among pro-PT respondents and smallest among anti-PT ones, at -0.535 against -0.285 . On the same outcome it is largest among anti-Bolsonaro respondents and smallest among pro-Bolsonaro ones, at -0.526 against -0.256 . Pro-Bolsonaro

respondents are the least affected group in either block. That is the expected pattern if tolerance for police violence weakens the story's effect.




Figure A2: Effect of the violent story on all three outcomes, by feelings about the Workers' Party and about Bolsonaro.

Note: Each block reads one question on its own. People who answer 4 or 5 like that party or figure, people who answer 3 are neutral, and people who answer 1 or 2 are against it. Nothing is combined across questions and nobody is left out, so each block covers all 624 respondents who answered both surveys. The omitted group is the one we expect to move least, which is people who dislike the PT in the first block and people who support Bolsonaro in the second. Each point is that group's own effect against zero and each bar is its 95 percent confidence interval. Differences between the groups are reported separately, in the contrasts table.

The same data in the vocabulary of the literature. Work on Brazilian partisanship usually sorts people into named groups rather than reading each question separately. Figure A3 follows that practice here, after Samuels and Belarmino (2025). The pro-PT group chooses the answer reading "I am a petista" and the anti-PT group chooses "I hate the PT". Among everyone else, the pro-Lula group rates Lula 4 or 5 and the pro-Bolsonaro group rates Bolsonaro 4 or 5. Eight respondents rate both leaders highly and are left out, since the scheme sorts people by which one leader they favor, and these eight favor both. We depart from the source twice. We use the endpoints rather than the top and bottom two answers, and we use the leader questions rather than the 2022 vote, since our outcomes are attitudes measured in the same sitting.

This grouping shows two things the previous figure cannot show. The pro-Lula group loses more trust in the police than the pro-PT group does. The anti-PT group, defined strictly as people who say they hate the PT, barely moves, while people who merely dislike the party do. Reading the question in three parts pools those two groups together.




Figure A3: Effect of the violent story on all three outcomes, by five groups after Samuels and Belarmino (2025).

Note: Groups follow Samuels and Belarmino (2025), with two changes. The pro-PT group chooses the answer reading "I am a petista" and the anti-PT group chooses "I hate the PT", rather than the top two and bottom two answers, because those two state an identity rather than a degree of warmth. Among everyone else, the pro-Lula group rates Lula 4 or 5 and the pro-Bolsonaro group rates Bolsonaro 4 or 5, where the source sorts these groups by the 2022 vote. Respondents who rate both leaders at 4 or 5 are left out, since the scheme sorts people by which leader draws them and those respondents do not answer that question. Eight of the 624 fall into that case. The anti-PT group is the omitted one. Each point is that group's own effect against zero and each bar is its 95 percent confidence interval.

Differences under the measure the main text uses. Table A9 reports the differences between groups for the two-question grouping. That rule counts a respondent as attached to the left if they rate the Workers' Party *or* Lula at 4 or 5, and as opposed to it if they rate either at 1 or 2. It covers 610 of the 624 respondents; the other 14 rate one at 4 or 5 and the other at 1 or 2, and the rule does not sort them. This is the pre-registration's own reading of attachment: H5 names Lula and the Workers' Party together. We report it beside the single-question grouping the main text uses.

The table gives both p-values. H5 predicts a stronger effect among people attached to Lula or the Workers' Party. The prediction was registered before the data were collected, so the one-tailed test is warranted for these comparisons. The two-tailed value is reported beside it as the conservative reading.

Under this grouping, attached respondents differ from anti-PT or anti-Lula respondents by 0.325 on trust in the police, 0.186 on satisfaction with democracy, and 0.167 on political trust. Trust in the police and satisfaction with democracy clear conventional levels one-tailed ($p=0.012$ and 0.031); political trust does not ($p=0.056$). Two-tailed, only trust in the police clears ($p=0.024$, 0.062 and 0.111). The middle group falls between the two ends on all three outcomes and is not distinguishable from either.

Table A9: Differences between groups under the measure the main text uses, which counts a respondent as attached to the left if they rate the Workers' Party OR Lula at 4 or 5.

Outcome	Estimate	SE	p two-tailed	p one-tailed
<i>Pro-PT or pro-Lula minus anti-PT or anti-Lula</i>				
Satisf. w. Democracy	-0.186	(0.100)	0.062	0.031
Political Trust	-0.167	(0.105)	0.111	0.056
Trust in Police	-0.325	(0.144)	0.024	0.012
<i>Pro-PT or pro-Lula minus neither</i>				
Satisf. w. Democracy	-0.161	(0.128)	0.208	0.104
Political Trust	-0.033	(0.134)	0.803	0.402
Trust in Police	-0.283	(0.184)	0.125	0.063
<i>Neither minus anti-PT or anti-Lula</i>				
Satisf. w. Democracy	-0.025	(0.117)	0.827	0.414
Political Trust	-0.133	(0.123)	0.277	0.139
Trust in Police	-0.043	(0.169)	0.800	0.400

Each cell is one group's effect minus another's, with its standard error in parentheses. The three groups are people who rate the Workers' Party or Lula at 4 or 5, people who rate the party at 1 or 2, and everyone else. Note the or. A respondent who is warm towards Lula while neutral on the party is counted as attached here, which is what separates this measure from the one that requires both answers to agree. Every one of the 624 respondents falls into one of the three groups. Standard errors come from the covariance matrix of the fit, so they account for the two estimates being correlated. Both p-values are given. H5 predicts that the effect is stronger among people attached to Lula or the Workers' Party, and that prediction was registered before the data were collected, so a one-tailed test is the licensed one for these comparisons. The two-tailed value is reported beside it as the conservative reading.

When the two answers have to agree. The blocks above read one question at a time. Figure A4 instead requires the two answers to agree. In the first block a respondent counts as attached to the left only if they rate both the Workers' Party and Lula at 4 or 5, and as opposed to it only if they rate both at 1 or 2. The second block pairs Bolsonaro against Lula. This is the crossed comparison in Samuels and Belarmino. The middle row is a residual in both: people with no attachment, together with people whose two answers disagree.

Requiring agreement sharpens the comparison. On trust in the police the gap between the two ends is 0.365 in the first block ($p=0.026$) and 0.413 in the second ($p=0.017$), against 0.250 when the party question is read on its own ($p=0.103$). Both are among the few differences in this appendix that reach conventional levels. Neither block shows a comparable gap on the other two outcomes.




Figure A4: Effect of the violent story on all three outcomes, where two answers must agree before a respondent is counted.

Note: Both blocks require two answers to agree, so a respondent who likes the PT but is cool on Lula is not counted as attached to the left. The first block reads the two left-wing items together. The second pairs Bolsonaro against Lula, which is the crossed comparison Samuels and Belarmino draw. The middle row is a residual in both and therefore holds two kinds of person, those with no attachment and those whose two answers disagree. The omitted group is the one we expect to move least. Each point is that group's own effect against zero and each bar is its 95 percent confidence interval.

Differences between the groups. Both figures plot each group against zero, which shows whether the vignette moved that group. The figures do not show whether the groups differ from each other. A difference between two estimates has larger uncertainty of its own, so the two readings can disagree. Table A10 gives those differences.

Most are not distinguishable from zero. The exception is the contrast the second figure was built to show. The pro-Lula group loses 0.605 more trust in the police than the anti-PT group ($p=0.002$) and 0.313 more satisfaction with democracy ($p=0.034$). Those are the only two differences in the table that reach conventional levels. When the questions are read in three parts, as in the first figure, the largest difference is between anti-Bolsonaro and pro-Bolsonaro respondents: 0.270 on trust in the police ($p=0.062$).

Table A10: Differences between groups, for the two figures above.

Contrast	Outcome	Estimate (SE)	p two / one-tailed
<i>Panel A: feelings about the PT</i>			
Pro-PT minus anti-PT	Satisf. w. Democracy	-0.157 (0.105)	0.135 / 0.067
Pro-PT minus anti-PT	Political Trust	-0.154 (0.112)	0.168 / 0.084
Pro-PT minus anti-PT	Trust in Police	-0.250 (0.153)	0.103 / 0.052
Pro-PT minus neutral	Satisf. w. Democracy	-0.065 (0.116)	0.576 / 0.288
Pro-PT minus neutral	Political Trust	-0.038 (0.123)	0.759 / 0.380
Pro-PT minus neutral	Trust in Police	-0.055 (0.169)	0.744 / 0.372
Neutral minus anti-PT	Satisf. w. Democracy	-0.093 (0.102)	0.363 / 0.182
Neutral minus anti-PT	Political Trust	-0.116 (0.108)	0.282 / 0.141
Neutral minus anti-PT	Trust in Police	-0.195 (0.148)	0.190 / 0.095
<i>Panel B: feelings about Bolsonaro</i>			
Anti-Bolsonaro minus pro-Bolsonaro	Satisf. w. Democracy	-0.124 (0.099)	0.213 / —
Anti-Bolsonaro minus pro-Bolsonaro	Political Trust	0.006 (0.105)	0.957 / —
Anti-Bolsonaro minus pro-Bolsonaro	Trust in Police	-0.270† (0.145)	0.062 / —
Anti-Bolsonaro minus neutral	Satisf. w. Democracy	-0.021 (0.108)	0.848 / —
Anti-Bolsonaro minus neutral	Political Trust	0.024 (0.115)	0.837 / —
Anti-Bolsonaro minus neutral	Trust in Police	-0.142 (0.158)	0.367 / —
Neutral minus pro-Bolsonaro	Satisf. w. Democracy	-0.103 (0.115)	0.371 / —
Neutral minus pro-Bolsonaro	Political Trust	-0.018 (0.122)	0.884 / —
Neutral minus pro-Bolsonaro	Trust in Police	-0.128 (0.168)	0.446 / —
<i>Panel C: pro or anti on BOTH the PT and Lula</i>			
Pro both minus anti both	Satisf. w. Democracy	-0.181 (0.112)	0.107 / 0.053
Pro both minus anti both	Political Trust	-0.164 (0.119)	0.168 / 0.084
Pro both minus anti both	Trust in Police	-0.365* (0.164)	0.026 / 0.013
Pro both minus neither	Satisf. w. Democracy	-0.085 (0.114)	0.455 / 0.227
Pro both minus neither	Political Trust	-0.087 (0.121)	0.472 / 0.236
Pro both minus neither	Trust in Police	-0.155 (0.166)	0.348 / 0.174
Neither minus anti both	Satisf. w. Democracy	-0.096 (0.097)	0.321 / 0.160
Neither minus anti both	Political Trust	-0.078 (0.103)	0.451 / 0.226
Neither minus anti both	Trust in Police	-0.210 (0.141)	0.138 / 0.069
<i>Panel D: Bolsonaro against Lula</i>			
Anti-Bolsonaro pro-Lula minus opposite	Satisf. w. Democracy	-0.180 (0.118)	0.129 / —
Anti-Bolsonaro pro-Lula minus opposite	Political Trust	-0.016 (0.126)	0.899 / —
Anti-Bolsonaro pro-Lula minus opposite	Trust in Police	-0.413* (0.173)	0.017 / —
Anti-Bolsonaro pro-Lula minus neither	Satisf. w. Democracy	-0.189† (0.106)	0.075 / —
Anti-Bolsonaro pro-Lula minus neither	Political Trust	-0.031 (0.113)	0.781 / —
Anti-Bolsonaro pro-Lula minus neither	Trust in Police	-0.288† (0.154)	0.062 / —
Neither minus pro-Bolsonaro anti-Lula	Satisf. w. Democracy	0.009 (0.102)	0.926 / —
Neither minus pro-Bolsonaro anti-Lula	Political Trust	0.015 (0.108)	0.887 / —
Neither minus pro-Bolsonaro anti-Lula	Trust in Police	-0.125 (0.148)	0.400 / —
<i>Panel E: Lula, against people who dislike the PT</i>			
Likes Lula minus dislikes PT	Satisf. w. Democracy	-0.205* (0.103)	0.048 / 0.024
Likes Lula minus dislikes PT	Political Trust	-0.193† (0.109)	0.077 / 0.039

Continued on next page.

Table A10 continued from previous page.

Contrast	Outcome	Estimate (SE)	p two / one-tailed
Likes Lula minus dislikes PT	Trust in Police	-0.398** (0.149)	0.008 / 0.004
Neutral minus dislikes PT	Satisf. w. Democracy	-0.028 (0.118)	0.815 / 0.408
Neutral minus dislikes PT	Political Trust	-0.131 (0.125)	0.293 / 0.147
Neutral minus dislikes PT	Trust in Police	-0.061 (0.170)	0.721 / 0.361
Likes Lula minus neutral	Satisf. w. Democracy	-0.177 (0.129)	0.170 / 0.085
Likes Lula minus neutral	Political Trust	-0.062 (0.136)	0.649 / 0.324
Likes Lula minus neutral	Trust in Police	-0.337† (0.186)	0.070 / 0.035

Panel F: the PT, against people who dislike the PT

Likes PT minus dislikes PT	Satisf. w. Democracy	-0.157 (0.104)	0.130 / 0.065
Likes PT minus dislikes PT	Political Trust	-0.154 (0.112)	0.170 / 0.085
Likes PT minus dislikes PT	Trust in Police	-0.250 (0.155)	0.108 / 0.054
Neutral minus dislikes PT	Satisf. w. Democracy	-0.028 (0.116)	0.813 / 0.406
Neutral minus dislikes PT	Political Trust	-0.131 (0.125)	0.297 / 0.148
Neutral minus dislikes PT	Trust in Police	-0.061 (0.174)	0.727 / 0.364
Likes PT minus neutral	Satisf. w. Democracy	-0.130 (0.129)	0.313 / 0.157
Likes PT minus neutral	Political Trust	-0.023 (0.139)	0.869 / 0.434
Likes PT minus neutral	Trust in Police	-0.189 (0.192)	0.325 / 0.163

Panel G: five groups

Pro-PT minus anti-PT	Satisf. w. Democracy	-0.227 (0.180)	0.206 / 0.103
Pro-PT minus anti-PT	Political Trust	-0.178 (0.193)	0.355 / 0.177
Pro-PT minus anti-PT	Trust in Police	-0.481† (0.266)	0.071 / 0.035
Pro-Lula minus anti-PT	Satisf. w. Democracy	-0.303* (0.136)	0.026 / 0.013
Pro-Lula minus anti-PT	Political Trust	-0.145 (0.146)	0.318 / 0.159
Pro-Lula minus anti-PT	Trust in Police	-0.609** (0.201)	0.002 / 0.001
Pro-Lula minus pro-PT	Satisf. w. Democracy	-0.076 (0.180)	0.673 / —
Pro-Lula minus pro-PT	Political Trust	0.033 (0.193)	0.864 / —
Pro-Lula minus pro-PT	Trust in Police	-0.128 (0.266)	0.632 / —
Pro-Bolsonaro minus anti-PT	Satisf. w. Democracy	-0.053 (0.138)	0.702 / —
Pro-Bolsonaro minus anti-PT	Political Trust	-0.124 (0.148)	0.403 / —
Pro-Bolsonaro minus anti-PT	Trust in Police	-0.285 (0.205)	0.165 / —
Pro-Bolsonaro minus pro-PT	Satisf. w. Democracy	0.174 (0.182)	0.337 / —
Pro-Bolsonaro minus pro-PT	Political Trust	0.054 (0.195)	0.780 / —
Pro-Bolsonaro minus pro-PT	Trust in Police	0.196 (0.269)	0.465 / —

Each cell is one group's effect minus another's, with its standard error in parentheses. The figures above plot every group against zero, which answers a different question. Two groups can both move clearly while the difference between them stays undetermined, because that difference carries its own and larger uncertainty. Panels A and B take the groups from the attitude figure, Panels C and D from the reinforcing-identities figure, Panels E and F from the third and fourth blocks of the groupings figure, and Panel G from the five-group figure. Panels E and F share a baseline and a middle row, the same people in both, and differ only in which attachment they test. Standard errors come from the covariance matrix of the fit each contrast is drawn from, so they account for the two estimates being correlated. Both p-values appear where a direction was registered in advance. H5 predicts a stronger effect among people attached to Lula or the Workers' Party, so the one-tailed test is licensed for those rows. A dash marks the comparisons where no direction was registered and halving would not be legitimate. Stars follow the two-tailed value throughout. † p<0.10, * p<0.05, ** p<0.01, *** p<0.001.

Main results

Table A1 reports the ATE and CATEs for all three outcomes; see the main text (or study1_paper_section.tex if reading this appendix standalone) for the full results and figure.

The political trust index, item by item

Political trust is an index: the mean of trust in the President, Congress and the Judiciary. An average effect on that index can hide two patterns. All three institutions might move together, or one might move while another stays flat or moves in the opposite direction. The index alone cannot tell them apart. The distinction is central here, because the argument is specifically that police violence spills over to the institutions charged with equality before the law.

Figure A5 repeats the models of the main figure with the index replaced by each of its three items. All three items move the same way, and none moves much on its own: -0.050 for the President, -0.120 for Congress and -0.085 for the Judiciary, no interval clear of zero. The index is not concealing an offset.

The differences between groups are larger, and they are not spread evenly. Non-white respondents lose a further 0.226 standard deviations of trust in Congress relative to white respondents ($p=0.034$), and respondents high in trait aggression lose a further 0.283 in the Judiciary ($p=0.007$). On the President, no difference reaches conventional levels; the largest is 0.176 by race ($p=0.060$). The two institutions that move are the legislative and judicial ones; the spillover is faintest for the executive, the most partisan of the three and the least associated with equal treatment under law.

We report this as a decomposition, not as a fourth test. These items were not separately hypothesized, and six additional comparisons per moderator raise a multiplicity concern.




Figure A5: Average and conditional effects on each component of the political trust index.

Note: The political trust index is the mean of these three items, so this figure reports the same models the main figure does with the index replaced by each of its parts. Every item is standardised on its own scale, which is why the three panels share one horizontal range. The average effect is negative on all three and reaches conventional levels on none: -0.052 for the President, -0.120 for Congress and -0.085 for the Judiciary. The index result is therefore three institutions moving a little in the same direction, not one moving while another cancels it. The differences BETWEEN groups are larger than the averages and are not spread evenly across the three: relative to white respondents, non-white respondents lose a further 0.226 of a standard deviation of trust in Congress ($p=0.034$), and respondents high in trait aggression

a further 0.283 in the Judiciary ($p=0.007$), while no difference on the President reaches conventional levels. Each point is a group's own effect against zero and each bar its 95 percent confidence interval, which is a different quantity from those differences and carries its own uncertainty. The omitted group in the attachment panel is the anti-PT or anti-Lula one. All outcomes are standardised, so estimates are in standard deviations.

Who left the panel between waves

Of the 1,004 respondents who completed wave 1, 624 returned for wave 2. The pre-registration promised a supplementary analysis of the 380 who did not; this section reports it.

Treatment was assigned in wave 2, after the 380 had already left the panel. Randomization therefore ran within the 624 who remained, and no wave-1 difference between leavers and stayers can bias the estimated effect. Such a difference can change who the estimate describes.

Table A11 shows that four of twelve wave-1 characteristics differ: respondents who stayed were slightly older ($p=0.005$), better educated ($p=0.045$), more interested in politics ($p=0.037$), and lower in trait aggression ($p=0.004$). Baseline trust in the police, in each institution, and satisfaction with democracy do not differ, nor do race, gender, or partisanship. Attrition was not selective on the attitudes the study measures.

Table A11: Wave-1 characteristics of respondents who completed both waves against those who did not.

Variable	N (left)	N (stayed)	Mean (left)	Mean (stayed)	SMD	p
Age	380	624	2.634	2.88	0.178	0.005
Education	380	624	5.889	6.053	0.132	0.045
Non-white	380	624	0.526	0.506	-0.04	0.541
Woman	380	624	0.518	0.503	-0.03	0.64
Political interest	380	624	2.429	2.561	0.135	0.037
Attachment to the PT (composite scale)	380	624	0.421	0.409	-0.042	0.515
Trait aggression scale (composite)	380	624	0.348	0.306	-0.189	0.004
Trust in police (wave 1)	380	624	2.611	2.617	0.007	0.912
Trust in president (wave 1)	380	624	2.224	2.224	0.001	0.992
Trust in congress (wave 1)	380	624	2.221	2.189	-0.035	0.595
Trust in judiciary (wave 1)	380	624	2.429	2.385	-0.045	0.487
Satisfaction with democracy (wave 1)	380	624	2.182	2.143	-0.04	0.54

380 of the 1,004 wave-1 respondents did not return for wave 2. SMD is stayed minus left, in pooled standard deviations. Treatment was assigned in wave 2, after these respondents had already left, so differences here cannot bias the estimated effect; they bear on who the estimates describe. Trait aggression is the difference that matters most, because it is also one of the moderators.

A wave-1 variable is a threat only if it predicts both leaving and the wave-2 outcomes; one that predicts only leaving shifts the sample's composition without disturbing a comparison made within the sample. Table A12 reports that age, race and political interest predict attrition and predict none of the three outcomes. Trait aggression, education and feelings toward Bolsonaro predict both.

Table A12: Wave-1 variables that predict leaving the panel, and whether they also predict the wave-2 outcomes.

Variable	p (attrition)	Smallest p across the three outcomes	Threat to the outcomes?
Trait aggression scale (composite)	0.004	0.001	Yes
high_traitaggr	0.006	0.010	Yes
Age	0.007	0.217	No
race	0.034	0.232	No
Political interest	0.038	0.505	No
Education	0.042	0.021	Yes
Feeling thermometer: Bolsonaro	0.046	0.000	Yes
rev_bolso_feeling	0.046	0.000	Yes

Restricted to variables that predict attrition at $p < 0.05$. A variable threatens the analysis only if it predicts attrition AND the outcomes; predicting attrition alone changes who the sample represents, which is a question of generalisation rather than bias. Smallest p is taken across the three wave-2 outcomes. Yes means the variable predicts at least one outcome at $p < 0.10$.

Respondents who stayed score 0.306 on the trait aggression scale against the leavers' 0.348, so the panel that reached wave 2 is less aggressive than the one recruited. The average effect is unaffected for the reason given above. The conditional effects by trait aggression, however, describe a sample from which the most aggressive respondents were disproportionately lost. The high- aggression group is therefore less extreme than it would have been in the wave-1 sample. These data cannot settle whether that attenuates or exaggerates the moderation. We report it so the estimates are read as applying to the completers rather than to the population the panel started from.

Ethics and human subjects

This research was reviewed and approved by the Stony Brook University Institutional Review Board (FWA #00000125), protocol IRB2023-00445, approved September 1, 2023 under expedited review category 7. The IRB determined that the study presents no more than minimal risk and granted a waiver of documentation of consent under 45 CFR 46.117(c). Participants were recruited through Opinion Box and gave informed consent before beginning wave 1; the consent text is reproduced in the variable dictionary above. Subject-facing materials were administered in Portuguese.

Both vignettes describe a routine traffic stop. In the treatment condition the encounter escalates and the driver is shot in the leg; no image is shown and the description is textual. The IRB-approved materials included the recruitment advertisement, the consent form, and a debriefing form, and respondents were debriefed after the survey.

The study shows a depiction of police violence to a Brazilian sample that is roughly half non-white, and the paper's own framing identifies non-white Brazilians as disproportionately subject to such violence. The depiction is textual only, and the debriefing followed every session. The research conforms to the APSA Principles and Guidance for Human Subjects Research.

Manipulation check and randomization balance

Table A13: Manipulation check: wave-2 (post-treatment) means by vignette condition.

Variable	N (control)	N (treated)	Overall mean	Mean (control)	Mean (treated)	Mean diff.
Trust in police (wave 2)	309	315	2.651	2.764	2.54	-0.224
Trust in president (wave 2)	309	315	2.09	2.12	2.06	-0.06
Trust in congress (wave 2)	309	315	2.149	2.191	2.108	-0.083
Trust in judiciary (wave 2)	309	315	2.381	2.417	2.346	-0.071
Satisfaction with democracy (wave 2)	309	315	2.13	2.207	2.054	-0.153
Procedural justice: dignity	309	315	2.546	3.207	1.898	-1.309
Procedural justice: good intent	309	315	2.737	3.217	2.267	-0.95
Procedural justice: impartiality	309	315	2.681	3.168	2.203	-0.965
Procedural justice: voice	309	315	2.596	2.955	2.244	-0.711

Mean diff. = treated (violent vignette) minus control (non-violent vignette). Procedural-justice items are the strongest manipulation check, since they describe the police interaction respondents were shown directly.

Table A14: Randomization check: wave-1 (pre-treatment) baseline means by vignette condition.

Variable	N (control)	N (treated)	Overall mean	Mean (control)	Mean (treated)	Mean diff.
Age	309	315	2.88	2.858	2.902	0.044
Education	309	315	6.053	6.136	5.971	-0.165
Non-white	309	315	0.506	0.495	0.517	0.022
Woman	309	315	0.503	0.502	0.505	0.003
Political interest	309	315	2.561	2.595	2.527	-0.068
Attachment to the PT (composite scale)	309	315	0.409	0.395	0.423	0.028
Trait aggression scale (composite)	309	315	0.306	0.309	0.303	-0.006
Feeling thermometer: PL	309	315	2.91	2.883	2.937	0.054
Feeling thermometer: PT	309	315	2.614	2.54	2.686	0.146
Feeling thermometer: Lula	309	315	2.659	2.618	2.698	0.080
Feeling thermometer: Bolsonaro	309	315	2.79	2.809	2.771	-0.038
Trait aggression item 1	309	315	2.627	2.592	2.66	0.068
Trait aggression item 2	309	315	2.24	2.246	2.235	-0.011
Trait aggression item 3	309	315	3.232	3.175	3.289	0.114
Trait aggression item 4	309	315	2.739	2.764	2.714	-0.050
Trait aggression item 5	309	315	2.787	2.803	2.771	-0.032
Trait aggression item 6	309	315	2.607	2.657	2.559	-0.098
Trait aggression item 7	309	315	1.804	1.851	1.759	-0.092
Trait aggression item 8	309	315	2.279	2.379	2.181	-0.198
Trait aggression item 9	309	315	1.564	1.634	1.495	-0.139
Trait aggression item 10	309	315	3.223	3.194	3.251	0.057
Trait aggression item 11	309	315	2.428	2.427	2.429	0.002
Trait aggression item 12	309	315	2.808	2.79	2.825	0.035
Trust in police (wave 1)	309	315	2.617	2.557	2.676	0.119
Trust in president (wave 1)	309	315	2.224	2.227	2.222	-0.005
Trust in congress (wave 1)	309	315	2.189	2.175	2.203	0.028
Trust in judiciary (wave 1)	309	315	2.385	2.379	2.39	0.011
Satisfaction with democracy (wave 1)	309	315	2.143	2.123	2.162	0.039

Mean diff. = treated (violent vignette) minus control (non-violent vignette). Includes both raw items (partisan-feeling thermometers, individual trait-aggression items) and their scale composites (petismo_sc, traitaggr_sc).

Summary statistics

Table A15: Sample summary statistics (wave-1 N=1,004; wave-2 completers N=624).

Variable	N	Mean	SD	Min	Max
Feeling thermometer: PL	1004	2.924	1.05	1	5
Feeling thermometer: PT	1004	2.632	1.19	1	5
Feeling thermometer: Lula	1004	2.677	1.23	1	5
Feeling thermometer: Bolsonaro	1004	2.853	1.273	1	5
Age	1004	2.787	1.407	1	7
Education	1004	5.991	1.228	1	7
Non-white	1004	0.514	0.5	0	1
Woman	1004	0.509	0.5	0	1
Political interest	1004	2.511	0.978	1	4
treatment_violent	624	0.505	0.5	0	1
Trust in police (wave 1)	1004	2.615	0.89	1	4
Trust in president (wave 1)	1004	2.224	1.045	1	4
Trust in congress (wave 1)	1004	2.201	0.927	1	4
Trust in judiciary (wave 1)	1004	2.401	0.987	1	4
Satisfaction with democracy (wave 1)	1004	2.157	0.98	1	4
Trust in police (wave 2)	624	2.651	0.823	1	4
Trust in president (wave 2)	624	2.09	1.047	1	4
Trust in congress (wave 2)	624	2.149	0.916	1	4
Trust in judiciary (wave 2)	624	2.381	0.959	1	4
Satisfaction with democracy (wave 2)	624	2.13	0.995	1	4
Procedural justice: dignity	624	2.546	1.308	1	5
Procedural justice: good intent	624	2.737	1.245	1	5
Procedural justice: impartiality	624	2.681	1.275	1	5
Procedural justice: voice	624	2.596	1.258	1	5

Pre-registration materials

Below is the complete content of our pre-registration materials, copied directly from the website (OSF registration [zykfm](#)).

Where the analysis departs from the registered plan

The registered analysis plan, reproduced verbatim below, specifies t-tests, Mann-Whitney tests and linear regression, evaluated two-tailed at the 0.05 and 0.01 levels. The paper reports something else: a panel model with respondent and wave fixed effects. We set out the difference here so a reader comparing the two does not have to reconstruct it.

We changed estimators because the registered tests compare group means and discard information the design collected. Every respondent answered the same outcome questions before and after the vignette, so each person can be their own control. Respondent fixed effects use that second measurement and absorb anything stable about a person; wave fixed effects absorb anything that moved all respondents between the two sittings. A between-group comparison of wave-2 answers alone would leave both out. The change is one of estimator, not of ques-

tion: the hypotheses tested, the direction each predicts, and the outcomes are the ones that were registered.

Two smaller departures remain. First, we report one-tailed p -values beside two-tailed ones, but only for contrasts whose direction the registration states in advance; every other comparison reports the two-tailed value alone, and the tables mark which is which. Second, the registration names three moderators, and the main text reports those three. The additional groupings were not registered and are exploratory: attachment to Bolsonaro, which the main text plots but does not tabulate, and this appendix's alternative partisanship codings and reinforcing-identity comparisons. They are here because readers of this literature will ask how sensitive the partisanship result is to how partisanship is coded, not because we predicted them.

The registered data-exclusion rule is followed as written: the main analysis uses respondents who answered both waves.

Study Information

Hypotheses

Based on the experimental design, we will test the following initial set of hypotheses:

- H1: Participants in the treatment condition will be more likely to perceive a lack of procedural fairness in the police encounter compared to those in the control group.
- H2: Participants in the treatment condition will be more likely to say they distrust the police compared to those in the control group.
- H3: Participants in the treatment condition will exhibit lower levels of political trust compared to those in the control group.
- H4: Participants in the treatment condition will exhibit increased dissatisfaction with democracy compared to those in the control group.

Additionally, drawing from theories of social identity, partisanship, and personality, we argue that the effect of making police violence salient on political trust and satisfaction with democracy should be moderated by partisan orientations, racial identity, and trait aggression. We propose the following hypotheses regarding these moderating effects:

- (Moderating effect of partisan orientations) H5: The negative effects of police violence on political trust and satisfaction with democracy will be stronger among participants identifying with Lula or the Workers' Party (PT).
- (Moderating effect of race) H6: The negative effects of police violence on political trust and satisfaction with democracy will be stronger among non-whites.
- (Moderating effect of trait aggression) H7: The negative effects of police violence on political trust and satisfaction with democracy will be stronger among individuals high in trait aggression personality.

Design Plan

Study type. Experiment — a researcher randomly assigns treatments to study subjects; this includes field or lab experiments, also known as an intervention experiment, and includes randomized controlled trials.

Blinding. For studies that involve human subjects, they will not know the treatment group to which they have been assigned.

Study design. The study is a between-subjects experiment with a completely randomized design with one factor having two mutually exclusive conditions.

Randomization. We will use block randomization, where each participant will be randomly assigned to one of the two equally sized, predetermined blocks. We will run the study on Qualtrics software that will perform the randomization as individuals enter the online survey.

Sampling Plan

Existing data. Registration prior to creation of data.

Data collection procedures. Our study design is a between-subjects randomized survey experiment. Subjects will participate by responding to an online survey. Our population of interest is adult Brazilians with access to the internet. The study will be conducted using Opinion Box, an online survey company, entirely online; researchers will not have access to identifying information, and because the study is conducted externally through a Qualtrics link, even anonymous Opinion Box worker IDs cannot be linked to survey responses.

There will be two waves. In wave 1, participants answer demographics, opinions about institutions, satisfaction with democracy, and political positions. In wave 2, the same participants are invited again; those who agree are given a fictional summarized news article about a police interaction. In the treatment, the police act inappropriately, relying disproportionately on physical violence in a routine traffic stop; in the control, the police act appropriately. Every participant is then asked the same questions about the article and about institutions/satisfaction with democracy.

Sample size. Our target sample size is 1,000 participants (500 per treatment condition).

Sample size rationale. The sample size was chosen to match standard sample sizes in comparable political-behavior experiments, generally sufficient for a two-condition design.

Stopping rule. We will stop data collection when we have reached 1,000 total participants.

Variables

Manipulated variables. We manipulate whether the participant is exposed to a vignette describing a traffic stop that turns into excessive police violence (treatment) or one that goes normally without violence (control).

Measured variables. Subjects' responses to survey questions gauging attitudes about the police, political institutions, and satisfaction with democracy, plus demographics and other political attitudes/behaviors measured pretreatment in wave 1.

Indices. The procedural-justice index is an additive index of the four procedural-justice items, normalized to 0–1, validated by a confirmatory factor analysis. Political trust is likewise an additive index of trust in Congress, the executive, and the judiciary, normalized to 0–1 and validated by a confirmatory factor analysis.

Analysis Plan

Statistical models. We will use simple t-tests and non-parametric Mann-Whitney tests to analyze differences in group means, and linear regression when including additional demographic covariates and when studying the conditional effects of the three moderators.

Inference criteria. Two-tailed tests, with significance levels of both 0.05 and 0.01.

Data exclusion. For the main analysis, we include data from subjects who answered both waves. In a supplementary analysis, we analyze wave-1-only respondents.

Missing data. List-wise deletion of missing cases in the relevant analysis.

References

Boudreau, C., MacKenzie, S. A., & Simmons, D. J. (2019). Police violence and public perceptions: an experimental study of how information and endorsements affect support for law enforcement. *The Journal of Politics*, 81(3), 1101-1110.

Buss, A. H., & Perry, M. (1992). The Aggression Questionnaire. *Journal of Personality and Social Psychology*, 63(3), 452-459.

Cohen, E., Gunderson, A., Jackson, K., McLachlan, P., Clark, T. S., Glynn, A. N., & Owens, M. L. (2019). Do officer-involved shootings reduce citizen contact with government? *The Journal of Politics*, 81(3), 1111-1123.

Cooper, H., Moore, L., Gruskin, S., & Krieger, N. (2004). Characterizing perceived police violence: implications for public health. *American Journal of Public Health*, 94(7), 1109-1118.

DeVylder, J. E., Oh, H. Y., Nam, B., Sharpe, T. L., Lehmann, M., & Link, B. G. (2017). Prevalence, demographic variation and psychological correlates of exposure to police victimisation in four US cities. *Epidemiology and Psychiatric Sciences*, 26(5), 466-477.

Erlingsson, G. Ó., Linde, J., & Öhrvall, R. (2014). Not so fair after all? Perceptions of procedural fairness and satisfaction with democracy in the nordic welfare states. *International Journal of Public Administration*, 37(2), 106-119.

Jackson, A. N., Fedina, L., DeVylder, J., & Barth, R. P. (2021). Police violence and associations with public perceptions of the police. *Journal of the Society for Social Work and Research*, 12(2), 303-326.

Magalhães, P. C. (2014). Government effectiveness and support for democracy. *European Journal of Political Research*, 53(1), 77-97.

Milani, J., Bradford, B., & Jackson, J. (2018). Police violence. In *Oxford Research Encyclopedia of Criminology and Criminal Justice*.

Mullinix, K. J., Bolsen, T., & Norris, R. J. (2021). The feedback effects of controversial police use of force. *Political Behavior*, 43, 881-898.

Mummolo, J. (2018). Militarization fails to enhance police safety or reduce crime but may harm police reputation. *Proceedings of the National Academy of Sciences*, 115(37), 9181-9186.

Nägel, C., & Nivette, A. E. (2022). Unexpected events during survey design and trust in the police: A systematic review. *Journal of Experimental Criminology*, 1-27.

Reny, T. T., & Newman, B. J. (2021). The opinion-mobilizing effect of social protest against police violence: evidence from the 2020 George Floyd protests. *American Political Science Review*, 115(4), 1499-1507.

Rhodes-Purdy, M. (2021). Procedures matter: strong voice, evaluations of policy performance, and regime support. *Political Studies*, 69(2), 412-433.

Supplementary Information B: Study 2: Chilean Shocks

Ignacio Urbina and Vitória Sgorlon

Stony Brook University

September 2026

Contents

Data Layer	1
Measurement Diagnostics	9
Specification and Pre-Trend Protocol Sensitivity	14
Sub-group Moderation and Political Heterogeneity	31
Police Dynamics, Cutpoints, and Conduct	37
Alternative Definitions of Exposure	44
Statistical Power and Minimum Detectable Effects	48
Methodological Standards Comparison	52

Data Layer

This section documents the panel sources, the screens applied in building the panel, the cleaning applied to each item, the neighborhood controls, the distribution of every outcome, and the cluster structure.

Table B1: Study 2 data layer: sources, stages, and what each produces.

Stage	Source	Produces	Role
01	ELSOC panel 2016-2022 (COES)	Respondent-wave item panel	Core
02	COES Protest Events 2009-2019	Municipality protest-shock statistic S_c	Core
03	INDH event geocoding (build chain)	Respondent exposure at interview date	Core
04	ELSOC-CIT 2019 / Censo 2017	Six neighbourhood controls, z-scored	Core
05	stages 01-04	Analytical panel + screen ladder	Core
06	stage 02 + SARIMA partition	Protest partitions (long)	Sample screen

The panel carries measured quantities and applies no screen the design owns, so the design's sample screen is a filter at estimation. Stage 06 supplies that screen: the design uses emp_1p0 ($S_c \leq 1$). Treatment is own-zone exposure from stage 03.

Table B2: Sample screen ladder for the shared analytical panel.

Step	Observations	Respondents	Municipalities
01_elsoc_panel (waves 1-4, municipality known)	12,565	4,447	105
drop movers (2016-2019 rule)	12,283	4,373	94
drop respondents with no INE zone match	10,591	3,343	94
05_analysis_panel (core)	10,591	3,343	94
(design a/c would also drop: shock_sc missing)	10,204	3,218	90

Only two screens are applied: the 2016-2019 mover rule, and respondents the INE crosswalk could not place (all exposure measures missing). The final row reports what the zone x bin2 design would additionally drop; that screen is NOT applied to the core, so each design's N is a difference against one common baseline. The mover rate reproduces the 1.7% documented for the source package.

Table B3: Treatment definition and the protest partitions used.

Measure	Treated	Percent
<i>A. Treatment (one row per respondent)</i>		
Own-zone event by interview date	450	13.5
<i>B. Protest partitions used (respondents per cell)</i>		
emp_1p0	3,343	100.0

Panel A: the treatment – an INDH-registered event in the respondent's own INE census zone by their interview date. Panel B: respondents assigned by the protest partition; unassigned respondents are those whose municipality has no usable pre-period protest series. emp_1p0 cuts S_c at 1.0 and defines the estimation sample.

Treatment is an indicator, but the underlying INDH registry records each victim's violence type, instrument, and consequence individually. Table B4 profiles the registry itself, independent of which respondents are geocoded into an exposed zone. 88.1 percent of victims report physical injuries, with a further 153 cases of ocular trauma, 32 of them resulting in irreversible vision loss. The most frequently logged instrument is the anti-riot shotgun, and the most frequently logged act is a gunshot wound.

Table B4: Profile of the INDH violation registry: 2,832 victim records, October 17, 2019 through March 14, 2020.

Category	N	Percent
<i>A. Consequence (one per victim)</i>		
Lesiones físicas	2,496	88.1
Otras lesiones	107	3.8
Lesión causada por trauma ocular	101	3.6
Sin determinar	55	1.9
Pérdida de visión por trauma ocular irreversible	32	1.1
Estallido de globo ocular	20	0.7
Fallecido	7	0.2
Quemaduras	5	0.2
Other (n<5 each)	9	0.3
<i>B. Violence type recorded (not mutually exclusive)</i>		
Violencia física	2,659	93.9
Otro tipo de violencia	1,184	41.8
Violencia sexual	390	13.8
Violencia psicológica	300	10.6
<i>C. Primary instrument (n=1,484 logged)</i>		
Escopeta antidisturbios	996	67.1
Carabina lanza gases	258	17.4
Gas pimienta	55	3.7
Carro lanza aguas	48	3.2
Other	127	8.6

Panel A: the single recorded consequence per victim. Panel B: violence type, logged across up to four acts per victim, so a victim can appear in more than one row and the column does not sum to the panel total. Panel C: the instrument logged for the first recorded act; an act such as an arbitrary detention logs no instrument, so the panel denominator is the non-missing count. Source: Instituto Nacional de Derechos Humanos, Mapa de la Memoria registry (Instituto Nacional de Derechos Humanos 2020).

Protest-Intensity Statistic (S_c) Construction

The design's third difference depends on S_c , each municipality's 2019 protest intensity expressed in its own pre-period standard-deviation units:

$$S_c = \frac{n_{2019} - \bar{n}_{\text{pre}}}{sd_{\text{pre}}},$$

where n is the asinh-transformed sum of COES-recorded protest-event sizes in a municipality-year, and the pre-period is 2009–2018 excluding 2011, the student-movement cycle. Including 2011 would inflate every municipality's baseline mean and shrink its protest-surge score. Category-7 (“unknown size”) events are imputed at the median of the winsorized known distribution rather than dropped. Table B5 reports the event-size construction and the resulting distribution across the 90 of the panel's 94 municipalities with a usable pre-period series; Figure B1 plots it, with the design's $S_c = 1$ partition marked.

Table B5: Construction of the protest-intensity statistic S_c .

Quantity	Value
<i>A. COES event-size construction, 2009-2019</i>	
Events (comuna/year screen)	21,151
Known attendance	13,738 (65.0%)
Imputed attendance (unknown-size events)	7,413 (35.0%)
Winsorize cap (99th pctl of known)	20,000 attendees
Imputation value (median of winsorized known)	75 attendees
<i>B. S_c across the panel's 94 municipalities</i>	
Municipalities with a usable pre-period series	90
...below historical norm ($S_c < 1$)	15 (16.7%)
...elevated ($S_c \geq 1$)	75 (83.3%)
No usable pre-period series (S_c undefined)	4
S_c : 25th / 50th / 75th percentile	1.45 / 3.24 / 6.69

S_c is 2019 protest activity minus a municipality's own 2009-2018 baseline mean, divided by its baseline standard deviation. Protest activity is the asinh-transformed sum of COES event sizes in a municipality-year; the baseline excludes 2011 (the student-movement cycle, which would inflate every municipality's baseline mean and shrink its own shock score). Panel A: category-7 (unknown size) events are imputed at the median of the winsorized known distribution, never dropped. Panel B: S_c is undefined where a municipality's baseline standard deviation is 0, a non-random exclusion selecting on low-variance, typically small, municipalities.




Figure B1: Distribution of S_c across the panel's municipalities, dashed line at the design's partition ($S_c = 1$).

Table B6: Item screening: sentinel codes removed and remaining valid responses.

Variable	ELSOC item	Sentinel (-999/-888)	Missing in source	Valid
pinera_vote_raw	c39	249	15,526	2,260
Aggression (BPAQ item)	f04_01	32	9,861	8,142
trust_pc_militants_raw	c06_03	446	8,630	8,959
c35_04	c35_04	176	8,251	9,608
c35_02	c35_02	159	8,251	9,625
c35_03	c35_03	135	8,251	9,649
c35_01	c35_01	47	8,251	9,737
Agreement with class placement	c34	18	8,251	9,766
Social-class placement	c33	158	8,093	9,784
m16	m16	51	6,957	11,027
t04_03	t04_03	395	5,470	12,170
t04_04	t04_04	155	5,470	12,410
t04_02	t04_02	136	5,470	12,429
trust_private_companies	c05_06	294	3,971	13,770
f05_02	f05_02	50	2,898	15,087
f05_01	f05_01	47	2,898	15,090
t02_03	t02_03	24	2,898	15,113
t02_02	t02_02	18	2,898	15,119
t02_04	t02_04	18	2,898	15,119
t02_01	t02_01	13	2,898	15,124
t09_03	t09_03	940	1,159	15,936
trust_neighbors	t01	103	1,159	16,773
t09_02	t09_02	99	1,159	16,777
t09_01	t09_01	58	1,159	16,818
t10	t10	32	1,159	16,844
Satisfaction with democracy (1-5)	c01	650	0	17,385
Left-right self-placement	c15	297	0	17,738
f05_04	f05_04	99	158	17,778
f05_03	f05_03	76	158	17,801
f05_06	f05_06	60	158	17,817
Trust in congress (1-5)	c05_07	214	0	17,821
f05_07	f05_07	43	158	17,834
Trust in the judiciary (1-5)	c05_05	163	0	17,872
s11_08	s11_08	126	0	17,909
generalized_trust_raw	c02	119	0	17,916
s02	s02	114	0	17,921
s11_09	s11_09	96	0	17,939
Trust in the government (1-5)	c05_01	83	0	17,952
s11_07	s11_07	83	0	17,952
s11_06	s11_06	67	0	17,968
s11_05	s11_05	64	0	17,971
s11_02	s11_02	49	0	17,986

Continued on next page.

Table B6 continued from previous page.

Variable	ELSOC item	Sentinel (-999/-888)	Missing in source	Valid
s11_01	s11_01	47	0	17,988
Political interest	c13	46	0	17,989
m02	m02	44	0	17,991
Police trust: level	c05_03	39	0	17,996
s11_03	s11_03	38	0	17,997
s01	s01	34	0	18,001
s11_04	s11_04	27	0	18,008
s03	s03	20	0	18,015
March attendance	c08_02	18	0	18,017
News consumption	c14_02	17	0	18,018
Education level	m01	15	0	18,020

ELSOC codes missing as -999 (no response) and -888 (no sabe), both negative. The source package binarized two secondary batteries off the unscreened columns, so those sentinels entered as substantive responses – for the class-respect battery in the wrong direction, affecting 1.9-9.1 percent of the positive class and unevenly across years. Every item is screened here, so the binary outcome forms are built from already-missing values.

Table B7: CIT-Censo neighbourhood controls: imputation and standardization.

Variable	Missing	Percent	Median imputed	Mean	SD
Population density	5	0.15	6736.8066	8709.3705	7673.2175
Mean schooling	5	0.15	9.0401	9.3078	1.3642
Unemployment rate	5	0.15	0.0365	0.0366	0.0087
Immigrant share	5	0.15	0.0152	0.0318	0.0524
Low-education household heads	5	0.15	0.7963	0.7602	0.1318
Overcrowding rate	5	0.15	0.0686	0.0741	0.0401

Moments are computed over the full CIT file, before any join and before any sample screen, so they are invariant to the design chosen downstream. The file is keyed on respondent, not on zone, so the standardization is respondent-weighted: a zone holding more respondents pulls the mean and standard deviation more. That is a measurement decision, preserved from the source. Respondents with no CIT record at all are set to the standardized mean of zero in stage 05, which is a different imputation from this one.

Table B8: Outcome distributions in the shared analytical panel.

Outcome	Non-missing	Missing	Mean	SD	At floor	At ceiling
Satisfaction with democracy: not satisfied share	10,127	464	0.4152	0.4928	58.5	41.5
Police trust: low share	10,569	22	0.1945	0.3959	80.5	19.5
Political trust: low share	10,385	206	0.4908	0.3969	30.6	27.8


Floor and ceiling shares are percentages of non-missing observations. All three outcomes are extreme-response indicators, so a high floor share is expected and is what motivates the linear-probability specification. The trust share is all-or-nothing across its three items: any missing item makes the index missing. It deliberately excludes police, which is its own outcome.

Table B9: Cluster structure of the analytical panel.

Quantity	Value
Respondents	3,343
Respondent-waves	10,591
Municipalities	94
Zones (cluster units)	762
Median respondents per zone	4
Max respondents per zone	38
Singleton zones	89

Treatment is constant within zone, so the zone is the cluster the design assigns on (Abadie, Athey, Imbens and Wooldridge 2023). Counts are for the full analytical panel; the estimation sample is its emp_1p0 subset.

A



B




Figure B2: Sample construction and exposure.

Note: Panel A shows respondent-waves in the shared analytical panel, by ELSOC wave, split on whether an INDH event had occurred in the respondent's own census zone by their interview date. Exposure is

structurally zero before 2019, so the pre-period bars are entirely unexposed. The pre-period coefficients of any event study therefore test parallel trends of the eventually exposed group rather than a time-varying treatment effect. Panel B shows the prevalence of each recorded exposure definition, one row per respondent. Which definition serves as the treatment is a separate identification decision.

Measurement Diagnostics

Outcome Coding: Floor-Share Evidence and Item Correlations

Both substantive outcomes are floor indicators: the share of the three institutions that a respondent rates *nada*, and the share *nada satisfecho* with democracy. A binary at the bottom of a scale is only informative if the bottom is populated, so the coding needs its own evidence rather than a convention.

Table B10: Response distributions for the trust and satisfaction items, pooled across waves 2016-2019.


Item	None	A little	Some	Quite a bit	A lot	Modal	N
Trust: executive	0.448	0.271	0.206	0.060	0.015	1 none	10,541
Trust: judiciary	0.438	0.264	0.221	0.064	0.013	1 none	10,491
Trust: congress	0.582	0.241	0.148	0.025	0.005	1 none	10,463
Trust: police	0.195	0.199	0.285	0.245	0.076	3 some	10,569
Satisfaction w/ democracy	0.415	0.274	0.214	0.076	0.020	1 none	10,127

Cells are the share of non-missing responses in each category. ELSOC records both batteries on a five-point scale running none, a little, some, quite a bit, a lot; none is the floor. Modal gives the category with the largest share. The three items that make up the low-trust share – executive, judiciary, congress – have their mode at the floor, which is why the outcome is coded there: the indicator separates a plurality from the rest of the sample rather than isolating a rare response. Trust in police is the exception, with its mode at the scale midpoint, and is reported as its own outcome rather than folded into the index. Computed on the full analytical panel, since the coding decision is a property of the instrument rather than of the estimation sample.

In ELSOC the floor is where the mass is concentrated. *Nada* is the modal response for all three institutional trust items — 44.8 percent for the executive, 43.8 for the judiciary, 58.2 for congress — and for satisfaction with democracy, 41.5 percent. The indicator therefore separates a plurality from the rest of the sample rather than isolating a rare response. A respondent who moves to the floor joins a category already occupied by a plurality of their compatriots.

The floor is also the category whose share changes across waves. The congress floor share runs 58, 56 and 49 percent across the three pre-treatment waves and then 68 percent in 2019; satisfaction with democracy runs 44, 37, 30 and then 53 percent. A measure built on the upper categories would be tracking a thin and shrinking part of the distribution over exactly the period of interest.

Trust in police is the exception, so it is reported as its own outcome rather than folded into the index. Its modal response is the scale midpoint, *algo*, at 28.5 percent, with only 19.5 percent at the floor. The same coding applied to police measures something structurally different: a tail rather than a plurality. That difference is why police is reported separately from a three-item index whose components behave alike.



Pooled across waves 2016–2019, non-missing responses. The three items behind the low-trust share – executive, judiciary, congress – peak at none, the floor of the scale; trust in police peaks at the midpoint. This is why the trust outcome is coded at the floor and why police is reported separately.

Figure B3: Response distributions, pooled 2016–2019. The three index items peak at the floor; trust in police peaks at the midpoint.

The item correlations decide which items to combine and which to keep apart. The correlations below are Spearman rank correlations rather than product-moment, because every response is an ordered five-point category whose steps cannot be assumed equal.

Table B11: Rank correlations among the trust, satisfaction and police-attitude items.

Item	(1)	(2)	(3)	(4)	(5)	(6)	(7)	(8)
(1) Trust: executive								
(2) Trust: judiciary	0.41							
(3) Trust: congress	0.53	0.56						
(4) Satisfaction w/ democracy	0.43	0.22	0.33					
(5) Trust: police	0.40	0.32	0.28	0.19				
(6) Police treat class w/ respect	0.19	0.10	0.08	0.13	0.33			
(7) Justified: police repress marches	0.08	0.05	0.06	0.11	0.15	0.04		
(8) Justified: police evict schools	0.12	0.07	0.09	0.15	0.24	0.11	0.55	
(9) Justified: stones thrown at police	-0.10	-0.03	-0.01	0.02	-0.18	-0.11	0.32	0.15

Spearman rank correlations, estimation sample, waves pooled, pairwise complete. Rank rather than product-moment because every item is an ordered five-point response whose steps cannot be assumed equal in size. Column heads are the row numbers. Minimum pairwise n is 768 and applies to the c35 rows, which run on a rotating module. Items 1 to 3 are the components of the low-trust share.

The three components of the low-trust share correlate 0.41 to 0.56 with one another, the tightest block in the matrix; that block structure licenses combining them. Satisfaction with democracy correlates with the block, most strongly with the executive at 0.43.

Trust in police behaves as an institutional-confidence item rather than a police item. It correlates 0.40 with trust in the executive and 0.28 to 0.32 with the judiciary and congress, but only 0.15 and 0.24 with approval of police repressing marches or evicting occupied schools. A respondent rating confidence in Carabineros appears to be rating the state in part: the item falls nationally in 2019 rather than with local exposure, the same pattern the results show.

Approval of the police force is a separate dimension. The two justification items are correlated 0.55 with each other and at most 0.24 with police trust. Justifying stones thrown *at* the police correlates -0.18 with police trust and $+0.32$ with justifying repression. The common factor there is tolerance of political violence rather than a position on the police. Perceived respectful treatment is weakly related to the rest of the battery; its strongest association is 0.33, with police trust. It measures something the other items do not. That makes it attractive as a proxy, and it also leaves it unanchored to the constructs the design estimates. The item-by-item decomposition below (Trust Index, by Component) extends this: it estimates each component of the trust index both as a level and as a floor share.

Trust Index, by Component

The composite result could reflect one component alone. Table [B12](#) and Figure [B4](#) repeat the focused no-surge specification on each item separately, both as its own 1–5 level and as its own floor share.

Table B12: Effect of own-zone exposure on each trust-index component separately, level and floor share.

Item	Level			Low-share			Obs.
	effect (SE)	p	pre-trend p	effect (SE)	p	pre-trend p	
<i>A. Municipalities where protest did not surge</i>							
Trust: executive	-0.025 (0.103)	0.812	0.698	0.133 (0.077)	0.083	0.846	10,140
Trust: judiciary	-0.238 (0.161)	0.139	0.265	0.249 (0.089)	0.005	0.784	10,078
Trust: congress	0.015 (0.139)	0.913	0.618	0.120 (0.085)	0.156	0.496	10,045
<i>B. High-minus-low protest: DDD contrast</i>							
Trust: executive	-0.011 (0.121)	0.929	0.698	-0.127 (0.084)	0.130	0.846	10,140
Trust: judiciary	0.230 (0.175)	0.188	0.265	-0.205 (0.096)	0.032	0.784	10,078
Trust: congress	-0.050 (0.150)	0.739	0.618	-0.116 (0.092)	0.206	0.496	10,045
<i>C. Municipalities that surged: implied</i>							
Trust: executive	-0.035 (0.065)	0.584	0.698	0.007 (0.033)	0.839	0.846	10,140
Trust: judiciary	-0.007 (0.070)	0.916	0.265	0.044 (0.036)	0.222	0.784	10,078
Trust: congress	-0.035 (0.056)	0.537	0.618	0.005 (0.035)	0.897	0.496	10,045

Each row swaps the outcome for one trust-index component and keeps the main results table's specification and sample: the full-sample triple difference, own-zone treatment interacted with wave, respondent and municipality-by-wave fixed effects, standard errors clustered on the zone. Panels A, B and C are the same three quantities the main table reports, so a row here is directly comparable to its counterpart there. Panel C is a linear combination of A and B, and its standard error comes from the covariance matrix rather than from adding the other two. The pre-trend column is a joint test over all four pre-period terms, as in the main table, so it is a property of the fit and repeats across that fit's three panels. These rows used a narrower no-surge sample until 2026-08-14; those numbers are not comparable to these. Level: the raw 1-5 item. Low-share: the share answering nada, the item's own floor – the same construction that, averaged across the three rows, gives the political trust: low share outcome in the main results table. Effect (SE): the 2019 quantity named by the panel. p: two-sided test that it is zero. Pre-trend p: a joint Wald test that the four 2016 and 2017 terms are zero, computed separately for each column since the level and low-share models differ in scale. Obs.: respondent-wave observations, identical for the level and low-share fits of the same item since both share the item's own missingness pattern.




Figure B4: Own-zone effect on each trust-index component, level and floor share. 95% CI. The two panels are on different scales and do not share an axis.

95% CI (design df). Level and low-share panels are on different scales – a 1-5 item and a 0/1 floor

indicator – so axes are not shared. Same specification as the main results table, one trust-index component at a time; see the accompanying table for exact values.

On the floor-share metric, all three items move in the same direction as the composite, with point estimates of 0.133 (executive), 0.249 (judiciary), and 0.120 (congress). Only the judiciary reaches conventional significance on its own ($p = 0.007$); the executive falls just short ($p = 0.089$) and congress further short ($p = 0.163$). None of the three pre-trend tests reject. For the raw 1–5 level the estimates are small and inconsistent in sign, and none approaches significance. This is consistent with the level being the noisier of the two codings for these items.

The composite’s no-surge effect, 0.187 (the main estimates table, Panel A), does not come from a single outlying item: it lies within the range the three component estimates span. Their shared direction is the expected pattern when an index responds but each item alone is under-powered to detect the response.

Specification and Pre-Trend Protocol Sensitivity

Pretreatment Balance

Treated and control respondents are not identical. Table B13 reports the differences separately within both protest strata. In the no-surge stratum, zones that experienced a registered event have higher mean schooling, lower unemployment, and a smaller share of low-education households. Their residents also began the panel more critical, 0.51 against 0.35 on democratic dissatisfaction in 2016. The surge stratum is much larger and has smaller baseline outcome gaps, but it also differs in neighborhood composition.

Table B13: Treated and control respondents in both protest strata: neighbourhood covariates, outcome levels, pre-period change, and panel composition.

Measure	S < 1		S ≥ 1	
	Treated	Control	Treated	Control
<i>A. Zone covariates (z-scores)</i>				
Population density	-0.28	-0.52	0.15	0.12
Mean schooling	0.07	-0.28	0.68	0.17
Unemployment rate	0.00	0.23	-0.17	-0.05
Immigrant share	-0.39	-0.31	0.55	0.14
Low-education household heads	-0.05	0.25	-0.67	-0.19
Overcrowding	-0.38	-0.17	-0.21	-0.05
<i>B. Outcome in 2016</i>				
Political trust: low share	0.58	0.48	0.52	0.51
Satisfaction with democracy: not satisfied share	0.51	0.35	0.50	0.45
Police trust: level	3.38	3.17	3.11	3.04
Police trust: low share	0.07	0.12	0.09	0.14
<i>C. Change in outcome, 2016–2018</i>				
Political trust: low share	-0.21	-0.15	-0.13	-0.11
Satisfaction with democracy: not satisfied share	-0.32	-0.08	-0.19	-0.16
Police trust: level	-0.32	0.03	-0.05	-0.09
Police trust: low share	0.05	-0.04	0.01	-0.01
<i>D. Panel composition</i>				
Share observed all four waves	0.54	0.67	0.62	0.53
Chi-squared p, completion balance	0.05		0.00	
Mean waves observed	3.13	3.40	3.29	3.12
Respondents	69	355	377	2417

S < 1 is the no-surge DDD reference stratum; S ≥ 1 is the surge stratum. Treated denotes an INDH-registered event in the respondent's own zone by interview. Panel A gives zone-level means of the six standardised CIT/Censo controls; Panel B gives 2016 outcome means; Panel C gives raw 2016–2018 changes among respondents observed in both waves; Panel D reports panel completion. Levels are absorbed by respondent fixed effects. The raw changes pool municipalities and are descriptive, not the within-municipality pre-trend tests reported in Table 4.

None of this is what the design requires. The respondent fixed effect absorbs any difference

that does not change over time, so neither the covariates nor the 2016 levels can bias β_{2019} or δ_{2019} . They are reported so that the effect sizes in the main estimates table can be read against each group's starting level and its distance from the scale bounds. What the design does require is that the treated and control zones within each protest stratum were not already diverging. The main estimates table reports the component-specific tests from the selected models, while Table B15 below reports the model-level trigger and probe.

Panel C of the main estimates table shows that, in the no-surge stratum, the raw change pooled across municipalities differs between the groups: dissatisfaction falls by 0.32 among the eventually treated against 0.08 among controls. Yet the corresponding within-municipality pre-trend test does not reject. The pooled change compares treated respondents in one municipality with control respondents in another, while the estimator and the test compare them only within the same municipality and wave. The municipality-by-wave fixed effects remove the gap between the two numbers.

Identifying Condition: Relative Pre-Trends Across Protest Strata

Gruber (1994) introduced the triple-difference design. Olden and Møen (2022) show that it rests on a single identifying condition, weaker than parallel trends holding in both strata separately. Absent treatment, the treated–control differential must evolve similarly across the third-difference groups — here, the $S_c \geq 1$ and $S_c < 1$ protest strata. That relative parallel-trends condition supports δ_{2019} , the DDD contrast. Interpreting β_{2019} , the $S_c < 1$ treatment DiD, on its own additionally requires ordinary parallel trends within that stratum.

Table B14 reports the testable pre-period implications of both conditions, taken from the same canonical fits as the main estimates. Panel A reports the 2016 and 2017 treatment-by-year-by-stratum interactions (the Olden–Møen condition); Panel B reports the 2016 and 2017 treatment-by-year interactions within $S_c < 1$. Each panel carries a joint Wald test. No outcome rejects either condition: Panel A joint p ranges from .45 to .95, Panel B from .45 to .62. These are non-rejections built on two pre-period coefficients, so power against slow-moving violations is limited — the police-trust-level pre-period interactions are sizeable though imprecise — but the observable implication of the identifying condition holds for every outcome.

Table B14: Identifying condition for the triple difference: pre-period placebo tests (Olden and Moen 2022).

Outcome	2016 (SE)	2017 (SE)	Joint F	p
<i>A. Relative pre-trend across protest strata: DDD interactions</i>				
Political trust: low share	-0.038 (0.091)	-0.106 (0.086)	0.79	0.455
Democratic dissatisfaction	-0.036 (0.111)	-0.011 (0.079)	0.06	0.946
Police trust: level	-0.139 (0.310)	-0.236 (0.227)	0.61	0.544
Police trust: low share	0.040 (0.062)	0.038 (0.073)	0.23	0.795
<i>B. Pre-trend where protest did not surge: treatment interactions</i>				
Political trust: low share	0.027 (0.085)	0.079 (0.078)	0.54	0.584
Democratic dissatisfaction	0.064 (0.105)	0.059 (0.067)	0.48	0.619
Police trust: level	0.179 (0.296)	0.227 (0.210)	0.63	0.535
Police trust: low share	-0.068 (0.054)	-0.044 (0.067)	0.80	0.450

Gruber (1994) introduced the triple-difference design; Olden and Moen (2022) derive its identifying condition: absent treatment, the treated–control differential must evolve similarly across the third-difference groups, here the high- and no-surge strata. Panel A is that condition’s testable pre-period implication: the 2016 and 2017 treatment-by-year-by-stratum interactions from the same canonical fit as the main estimates, with 2018 omitted, jointly zero under the condition. This is the assumption supporting the DDD contrast. Interpreting the no-surge stratum’s treatment DiD on its own additionally requires ordinary parallel trends within that stratum; Panel B tests its 2016 and 2017 treatment-by-year interactions. Cells report the coefficient with its zone-clustered standard error in parentheses; Joint F and p test the panel’s two pre-period terms together, with denominator degrees of freedom equal to clusters minus one. Democratic dissatisfaction uses the protocol-selected baseline-outcome-trends specification; the other outcomes the unadjusted DDD. Non-rejection is consistent with the condition but cannot establish it; with two pre-period coefficients the tests have limited power against slow-moving violations.

Protocol Trigger and Probe Decision

Table B15: Programmatic pre-trend protocol and selected specification by outcome.

Outcome	Initial p	Triggered	Probe p	Pre L2 ratio	Clears 0.25	Canonical specification
Political trust: low share	0.799	No	–	–	–	Unadjusted DDD
Satisfaction with democracy: not satisfied share	0.179	Yes	0.659	0.34	Yes	Baseline-outcome trends
Police trust: level	0.817	No	–	–	–	Unadjusted DDD
Police trust: low share	0.661	No	–	–	–	Unadjusted DDD

Initial p is the joint Wald p-value for the four 2016 and 2017 treatment interactions in the unadjusted full-sample DDD. A value below 0.25 triggers one probe: the respondent's latest observed dependent variable in 2016–2018 interacted with the 2016, 2017 and 2019 indicators, allowing outcome-baseline-specific differential trends. Respondents without an observed pre-treatment dependent variable are retained by median filling that baseline before estimation; no missingness-indicator term enters the regression. The probe is canonical only when its joint p-value increases and the L2 norm of the four pre-period treatment coefficients decreases; Pre L2 ratio is adjusted over initial. The estimation sample is asserted identical across the initial and probe fits. The complete machine-readable audit, formulas, missing-baseline counts and decision reasons are in 07_pretrend_protocol_log.csv. This is a declared, outcome-specific specification rule adopted after examining the pre-period diagnostics; it improves the displayed pre-period fit but does not prove conditional parallel trends, and the conventional post-treatment p-values do not adjust for specification selection.

Only democratic dissatisfaction triggers the protocol because its model-level joint pre-trend p rises from .179 to .659 and its pre-period coefficient L2 norm falls to 0.34 of the unadjusted value. Its baseline-outcome differential-trends specification is therefore canonical; the other three outcomes retain the unadjusted DDD.

Table B16 shows the consequence of the pre-trend protocol rule by presenting both candidate fits regardless of which the protocol selected. The right-hand column for the three non-triggering outcomes is the specification the protocol declined, shown so that its effect can be read rather than assumed.

Table B16: Every estimate under both specifications, whether or not the protocol selected the adjustment.

Outcome	Unadjusted DDD (SE)	Baseline trends (SE)	Shift	Shift in SE	Reported
<i>A. Municipalities where protest did not surge: treatment DiD</i>					
Political trust: low share	0.187 (0.062)	0.169 (0.060)	-0.017	-0.28	same verdict
Democratic dissatisfaction	0.313 (0.088)	0.201 (0.066)	-0.112	-1.27	same verdict
Police trust: level	-0.130 (0.191)	-0.170 (0.171)	-0.039	-0.21	same verdict
Police trust: low share	0.041 (0.087)	0.069 (0.067)	0.028	0.32	same verdict
<i>B. High-minus-low protest: DDD contrast</i>					
Political trust: low share	-0.171 (0.067)	-0.168 (0.065)	0.003	0.04	same verdict
Democratic dissatisfaction	-0.319 (0.095)	-0.230 (0.073)	0.089	0.94	same verdict
Police trust: level	0.260 (0.206)	0.312 (0.185)	0.052	0.25	same verdict
Police trust: low share	-0.079 (0.091)	-0.108 (0.072)	-0.028	-0.31	same verdict
<i>C. Municipalities that surged: implied treatment DiD</i>					
Political trust: low share	0.016 (0.025)	0.002 (0.023)	-0.015	-0.60	same verdict
Democratic dissatisfaction	-0.006 (0.035)	-0.029 (0.031)	-0.023	-0.67	same verdict
Police trust: level	0.130 (0.077)	0.142 (0.071)	0.012	0.16	DIFFERS
Police trust: low share	-0.039 (0.027)	-0.039 (0.027)	0.000	0.00	same verdict

The protocol chooses one specification per outcome. This table sets that choice aside and reports both fits for all four outcomes and all three components, on identical rows. Unadjusted DDD is Equation 1; Baseline trends is Equation 2, the single pre-declared probe. Only democratic dissatisfaction adopts it, so for the other three outcomes the right-hand column is the specification the protocol declined. Shift is baseline trends minus unadjusted in the outcome's own units, and Shift in SE the same quantity in unadjusted standard errors. Reported flags cells where the two disagree at 5 percent. Democratic dissatisfaction is the not-satisfied share; see the main estimates table for full outcome definitions.

The estimates are close to insensitive to which model is used. Across twelve outcome-by-component cells there is no sign change, and eleven agree at the 5 percent level. The single disagreement is the implied surge effect on the police level, $p = .094$ against $p = .047$ — a quantity about which this paper makes no claim, and in favor of the model that was not selected. Democratic dissatisfaction's no-surge coefficient moves the most, by 1.27 unadjusted standard errors, and remains significant with the same sign. Nine of the twelve cells move by less than a third of a standard error, and three by less than a tenth. The protocol tightens the pre-period fit for one outcome; it does not choose the paper's conclusions.

The conclusions of this paper are the same under either model, for every outcome and every component. The adjustment reduces the standard error in all twelve cells, as a genuine pre-trend correction should.

Missing Outcomes: Multivariate-Imputation Sensitivity

The main estimates use the observed outcome rows available for each model. As a sensitivity analysis, Table B17 instead completes the four-wave outcome panel using multivariate imputation by chained equations (MICE). The imputation is respondent-wide: the five source items underlying democratic dissatisfaction, institutional trust, and police trust borrow information from their other waves and from repeated measures in the panel. Treatment, the protest stratum, geography, and missing values of S_c are never imputed. The four analysis outcomes are reconstructed only after the source items have been imputed. The pre-trend protocol is then applied anew to the pooled MICE models rather than carrying over the observed-outcome decision. For any triggered outcome, the latest completed pre-treatment outcome is centred within each completed dataset and interacted with wave before the adjusted models are pooled; no missingness-indicator term enters estimation.

Table B17: Protocol-selected full-sample triple-difference estimates after multivariate imputation by chained equations.

Outcome	Effect in 2019 (SE)	t	p	FMI	Pre-trend F	Pre-trend p	Obs.	Respondents	Zones (treated)
<i>A. Municipalities where protest did not surge: treatment DiD</i>									
Political trust: low share	0.143 (0.064)	2.24	0.025	0.02	0.19	0.824	12,868	3,217	749 (123)
Satisfaction with democracy: not satisfied share	0.182 (0.074)	2.45	0.015	0.03	0.84	0.436	12,868	3,217	749 (123)
Police trust: level	-0.112 (0.190)	-0.59	0.556	0.02	0.58	0.563	12,868	3,217	749 (123)
Police trust: low share	0.036 (0.085)	0.42	0.675	0.01	0.94	0.393	12,868	3,217	749 (123)
<i>B. High-minus-low protest: DDD contrast</i>									
Political trust: low share	-0.119 (0.069)	-1.73	0.083	0.03	0.14	0.874	12,868	3,217	749 (123)
Satisfaction with democracy: not satisfied share	-0.202 (0.080)	-2.53	0.012	0.03	0.21	0.811	12,868	3,217	749 (123)
Police trust: level	0.245 (0.205)	1.20	0.232	0.02	0.37	0.688	12,868	3,217	749 (123)
Police trust: low share	-0.077 (0.090)	-0.86	0.390	0.01	0.36	0.702	12,868	3,217	749 (123)
<i>C. Municipalities that surged: implied treatment DiD</i>									
Political trust: low share	0.025 (0.025)	1.00	0.319	0.04	0.08	0.919	12,868	3,217	749 (123)
Satisfaction with democracy: not satisfied share	-0.020 (0.030)	-0.68	0.496	0.02	1.99	0.139	12,868	3,217	749 (123)
Police trust: level	0.133 (0.076)	1.75	0.080	0.02	0.11	0.900	12,868	3,217	749 (123)
Police trust: low share	-0.041 (0.027)	-1.53	0.126	0.01	0.51	0.599	12,868	3,217	749 (123)

Respondent-wide MICE jointly imputes the five underlying satisfaction and trust items across 2016–2019, using their other waves and repeated auxiliary attitudes, health, neighbourhood, and economic measures. The four outcomes are derived after imputation; auxiliary variables enter the imputation model only, not the DDD regression. The observed panel contributes 10,204 rows; completing all four waves for respondents with observed S_c yields 12,872 rows. MICE uses $m = 20$ completed datasets and 30 iterations. Within the imputed analysis, the four-restriction model-level pre-trend trigger and baseline-outcome differential-trends probe are pooled across completed datasets and apply the same p-value/L2 selection rule as Table 5. When triggered, the latest completed pre-treatment outcome is centred within each completed dataset and interacted with wave; no missingness-indicator term enters the regression. The adjacent protocol table reports the selected specification by outcome. Coefficients and zone-clustered covariance matrices are combined with Rubin’s rules. FMI is the fraction of total scalar variance attributable to missing data; the component pre-trend p-values displayed here use a pooled two-restriction Wald approximation. Treatment, S_c group, geography,

Table B18: Pooled pre-trend protocol and selected MICE specification by outcome.

Outcome	Initial p	Triggered	Probe p	Pre L2 ratio	Clears 0.25	Canonical specification
Political trust: low share	0.971	No	–	–	–	Unadjusted DDD
Satisfaction with democracy: not satisfied share	0.104	Yes	0.255	0.51	Yes	Baseline-outcome trends
Police trust: level	0.841	No	–	–	–	Unadjusted DDD
Police trust: low share	0.534	No	–	–	–	Unadjusted DDD

Initial p jointly tests the four 2016 and 2017 treatment interactions after pooling the unadjusted coefficient vectors and zone-clustered covariance matrices across completed datasets. A value below 0.25 triggers the baseline-outcome differential-trends probe in every completed dataset. The latest completed pre-treatment outcome is centred within its dataset and interacted with 2016, 2017 and 2019. The probe is selected only when its pooled joint p-value rises and the L2 norm of the pooled four-term pre-period coefficient vector falls; Pre L2 ratio is adjusted over initial. The protocol uses the same rule as Table 5, but applies it anew after MICE rather than carrying the observed-outcome decision into the imputed analysis.

Only democratic dissatisfaction triggers the pooled MICE protocol. Its joint four-restriction pre-trend p -value rises from .078 to .199, and the pooled pre-period coefficient L2 norm falls to 0.53 of its unadjusted value. The baseline-outcome differential-trends probe is therefore selected even though the adjusted joint test remains below .25. Under that selected model, the no-surge effect is 0.183 (SE 0.074, $p = .014$), the surge-minus-no-surge DDD contrast is -0.203 (SE 0.080, $p = .011$), and the implied surge effect is -0.021 (SE 0.030, $p = .490$). The other three outcomes do not trigger and retain their unadjusted pooled DDD specifications.

The initial regression is otherwise the full-sample DDD used in the main estimates table, with respondent and municipality-by-wave fixed effects, 2018 omitted, and standard errors clustered by census zone. Rubin's rules pool both the coefficient vectors and their clustered covariance matrices. The pooled four-restriction test supplies the trigger, and the pooled coefficient vector supplies the L2 comparison. This diagnostic therefore replaces outcome-wise row deletion while leaving the treatment assignment and estimand unchanged. It relies on the additional missing-at-random assumption stated in the table note.

Predicted and Latent Trait Controls

The respondent fixed-effect specification absorbs all stable respondent attributes without measuring them. As a diagnostic, we replace that fixed effect with either a prediction-derived respondent score or a latent trait score estimated separately for each dependent variable k . Thus, θ_{ik} is specific to outcome Y_{kit} : no democracy-derived score is reused in a trust or police model. Table B19 reports the coefficient on the resulting control in every democracy, trust, police, and institutional-item model. Both constructions use only the 2016–2018 waves.

Table B19: Coefficients on the prediction-derived and latent-SEM trait controls.

Outcome	Predicted-signal theta			DV trait-state SEM		
	Coefficient (SE)	p	N	Coefficient (SE)	p	N
<i>Panel A. Democracy, composite, and police outcomes</i>						
Satisfaction with democracy (1–5)	0.895 (0.064)	<.001	1,332	1.715 (0.045)	<.001	1,332
Satisfaction: lowest category	0.386 (0.032)	<.001	1,332	1.974 (0.065)	<.001	1,332
Institutional low-trust share	0.449 (0.023)	<.001	1,380	1.504 (0.034)	<.001	1,380
Police trust (1–5)	0.995 (0.088)	<.001	1,411	1.395 (0.037)	<.001	1,411
Police trust: lowest category	0.248 (0.029)	<.001	1,411	1.811 (0.075)	<.001	1,411
<i>Panel B. Constituent institutional items</i>						
Executive-government trust (1–5)	1.410 (0.068)	<.001	1,414	1.784 (0.056)	<.001	1,414
Executive-government trust: lowest category	0.628 (0.031)	<.001	1,414	1.913 (0.045)	<.001	1,414
Judiciary trust (1–5)	0.788 (0.066)	<.001	1,395	1.455 (0.040)	<.001	1,395
Judiciary trust: lowest category	0.345 (0.021)	<.001	1,395	1.688 (0.047)	<.001	1,395
Congress trust (1–5)	1.024 (0.057)	<.001	1,393	1.588 (0.041)	<.001	1,393
Congress trust: lowest category	0.520 (0.027)	<.001	1,393	1.706 (0.046)	<.001	1,393

Entries are the unstandardized coefficient on the respondent trait control, with zone-clustered standard error in parentheses, from the corresponding event-study specification. Every outcome k has its own predicted-signal θ_{ik} , obtained from a lavaan trait-state decomposition of that outcome’s full-sample 2016–2018 elastic-net fitted values and oriented in the outcome’s natural direction. DV trait-state SEM is the lavaan factor score from $Y_{it} = \text{Trait}_i + \text{State}_{it}$, estimated separately for each DV using 2016–2018 only. The two coefficients are not directly comparable because the trait scores have different scales. The SEM trait is generated from the same outcome and therefore measures decomposition, not independent predictive validity. These generated scores are Mundlak-like stable-heterogeneity controls, but shrinkage and measurement error mean they are not algebraically equivalent to respondent fixed effects. Binary outcomes and their SEMs are on the observed linear-probability scale. N is the regression observation count.

The prediction-derived score is the lavaan Trait factor score extracted from the fitted values of an outcome-specific elastic net refitted on all 2016–2018 observations. Respondent-grouped cross-validation is used only to diagnose the prediction stage, not to construct the control.

Its direction follows the dependent-variable coding, so a higher score means a higher ex-

pected level for level outcomes and a higher expected probability or share for floor-coded outcomes. In that sense it is a Mundlak-like correlated-random-effects control for stable person-specific heterogeneity. It is not algebraically identical to a respondent fixed effect: the generated score is shrunk and measured with error, and it summarizes only three pre-period observations. Its mixed significance reflects that limitation. The fitted-value measurement models identify one common Trait and three orthogonal wave-specific States with unit loadings and zero observed residual variance; consequently, each State combines transitory variation and measurement error. All 11 models converged (robust CFI = .988–1.000, RMSEA = .000–.040, and SRMR = .007–.019). The direct-DV SEM coefficients are positive and precisely estimated because each latent trait score is constructed from repeated observations of the same dependent variable. They therefore quantify how strongly the extracted stable component predicts the outcome it was built from; they are not evidence that an independently measured covariate predicts it. The two columns also use different score scales, so their raw coefficient magnitudes should not be compared. The strict unit-loading SEM fits all outcomes adequately except executive-government trust at the level scale (robust CFI = .753; RMSEA = .173), whose latent-trait coefficient should be read with particular caution.

Zone Fixed Effect

The main design identifies the treatment interactions from within-respondent change over time. An alternative structure identifies the same interactions from within-zone variation pooled across respondents instead, replacing the respondent fixed effect with a zone fixed effect and changing nothing else. The two structures answer different questions. The zone fixed effect is not protected against change in who is observed across waves, and the pre-period waves have roughly a quarter fewer observations than 2018 and 2019. Agreement between the structures therefore shows how much of the estimate depends on panel composition, and divergence would show where it does.

Table B20: The main estimates under a zone fixed effect in place of the respondent fixed effect, protest partition ($S_c \leq 1$).

Outcome	Effect in 2019 (SE)	p	2016 (SE)	2017 (SE)	Pre-trend p	Obs.	Zones (treated)
<i>A: respondent FE (main design)</i>							
Political trust: low share	0.187 (0.063)	0.004	0.027 (0.085)	0.079 (0.079)	0.591	1,379	91 (13)
Satisfaction with democracy: not satisfied share	0.313 (0.089)	0.001	0.168 (0.110)	0.153 (0.082)	0.125	1,325	91 (13)
Police trust: level	-0.130 (0.192)	0.500	0.179 (0.299)	0.227 (0.211)	0.543	1,410	91 (13)
Police trust: low share	0.041 (0.088)	0.644	-0.068 (0.054)	-0.044 (0.068)	0.459	1,410	91 (13)
<i>B: zone FE</i>							
Political trust: low share	0.177 (0.068)	0.011	0.038 (0.094)	0.061 (0.087)	0.761	1,387	91 (13)
Satisfaction with democracy: not satisfied share	0.294 (0.097)	0.003	0.196 (0.111)	0.176 (0.083)	0.075	1,344	91 (13)
Police trust: level	-0.117 (0.191)	0.544	0.256 (0.266)	0.335 (0.200)	0.249	1,411	91 (13)
Police trust: low share	0.036 (0.086)	0.678	-0.080 (0.049)	-0.075 (0.068)	0.252	1,411	91 (13)

Effect in 2019: the treatment-by-2019 interaction in the outcome's own units, standard error in parentheses, with its two-sided p-value. 2016 (SE) and 2017 (SE): the pre-period interactions. Pre-trend p: a joint Wald test that both pre-period interactions are zero. Zones (treated): clusters in the estimation sample, with the number containing a registered event in parentheses. Panel A reproduces the main estimates table's specification; Panel B replaces the respondent fixed effect with a zone fixed effect and changes nothing else. Year-by-municipality fixed effects, 2018 the reference year by omission, standard errors clustered on the zone, no covariates, under both panels.

Table B21: Effect in 2019 under each fixed-effect structure, and the difference between them.

Outcome	Respondent FE		Zone FE		Difference	Ratio
	Effect (SE)	Pre-trend p	Effect (SE)	Pre-trend p		
<i>A: protest partition ($S_c \leq 1$)</i>						
Political trust: low share	0.187 (0.063)	0.591	0.177 (0.068)	0.761	-0.010	0.95
Satisfaction with democracy: not satisfied share	0.313 (0.089)	0.125	0.294 (0.097)	0.075	-0.019	0.94
Police trust: level	-0.130 (0.192)	0.543	-0.117 (0.191)	0.249	0.014	0.89
Police trust: low share	0.041 (0.088)	0.459	0.036 (0.086)	0.252	-0.005	0.88

COLUMNS. Respondent FE and Zone FE: the treatment-by-2019 interaction under each structure, standard error in parentheses. Difference: zone minus respondent, in the outcome's own units. Ratio: zone over respondent. The two estimates are not independent – they are fit on overlapping samples – so the difference is reported as a magnitude and no test of it is attached. Pre-trend p columns are the joint Wald test on the 2016 and 2017 interactions under each structure. Everything other than the fixed effect is held fixed across the two columns; see the accompanying table for the full specification.

The two structures agree closely on every outcome. Every difference in Table B21 is small relative to the coefficients' own standard errors, and the zone-FE to respondent-FE ratio sits close to one throughout (0.88–0.95). Little of the main estimate depends on which respondents are observed in which wave.

Zone Fixed Effect with Individual Trends

Table B22: Individual-level predictors interacted with the year dummies under a zone fixed effect: protest partition ($S_c \leq 1$).

Outcome	Effect in 2019 (SE)	2016	2017	Pre-trend p
<i>A: respondent FE, no controls</i>				
Political trust: low share	0.197 (0.061)	0.033	0.086	0.589
Satisfaction with democracy: not satisfied share	0.291 (0.091)	0.155	0.140	0.175
Police trust: level	-0.170 (0.183)	0.159	0.207	0.605
Police trust: low share	0.064 (0.089)	-0.056	-0.032	0.599
<i>B: zone FE, no controls</i>				
Political trust: low share	0.196 (0.067)	0.057	0.081	0.618
Satisfaction with democracy: not satisfied share	0.274 (0.100)	0.197	0.177	0.084
Police trust: level	-0.155 (0.182)	0.215	0.294	0.351
Police trust: low share	0.059 (0.087)	-0.068	-0.063	0.364
<i>C: zone FE, individual predictors x year</i>				
Political trust: low share	0.191 (0.068)	0.018	0.037	0.907
Satisfaction with democracy: not satisfied share	0.267 (0.102)	0.176	0.134	0.206
Police trust: level	-0.043 (0.171)	0.275	0.370	0.214
Police trust: low share	0.047 (0.085)	-0.084	-0.059	0.246

PANELS. A: the specification in the main estimates table, respondent fixed effect, no covariates. B: the respondent fixed effect replaced by a zone fixed effect, still no covariates. C: adds eight baseline individual predictors, each as a main effect and interacted with the 2016, 2017 and 2019 dummies. All three panels are fit on the same sample – respondents complete on the outcome and on every predictor – so the differences between panels are the specification and not a change in who is estimated on. PREDICTORS. Left-right self-placement (c15, with 'ninguna' recoded to the scale centre), educational level (m01), social class self-placement (c33), agreement with that placement (c34), an aggression item (f04_01), birth cohort, interest in politics (c13) and news consumption (c14_02). All are held at their 2018 baseline value with a 2017 then 2016 fallback, so none is measured after the estallido. WHY PANEL C ONLY FOLLOWS PANEL B. Under a respondent fixed effect anything fixed about a person is already absorbed, so baseline predictors have no compositional channel to act through. Under a zone fixed effect the estimate uses differences between people in the same zone, and the pre-period waves are roughly a quarter smaller than 2018 and 2019, so the composition of who is compared changes across years. Panel C asks how much of the pre-period treatment interaction that change accounts for. Columns 2016 and 2017 are the pre-period treatment interactions, point estimates only; their standard errors, the 2019 p-value and the observation counts are in the accompanying CSV. Pre-trend p is the joint Wald test that both pre-period interactions are zero. Eight predictors times four terms is 32 parameters against 13 treated clusters, so panel C is heavily parameterised relative to the variation available. Standard errors are clustered on the zone throughout.

Table B23: Pre-period treatment interactions with and without individual predictor by year terms, under a zone fixed effect.

Outcome	PT p: resp. FE	PT p: zone FE	PT p: zone + controls	2016		2017		2019		Controls joint p
				bare	controls	bare	controls	bare	controls	
<i>A: protest partition ($S_c \leq 1$)</i>										
Political trust: low share	0.589	0.618	0.907	0.057	0.018	0.081	0.037	0.196	0.191	0.0003
Satisfaction with democracy: not satisfied share	0.175	0.084	0.206	0.197	0.176	0.177	0.134	0.274	0.267	0.0000
Police trust: level	0.605	0.351	0.214	0.215	0.275	0.294	0.370	-0.155	-0.043	0.4027
Police trust: low share	0.599	0.364	0.246	-0.068	-0.084	-0.063	-0.059	0.059	0.047	0.0026

The three Pre-trend p columns are the joint Wald test on the 2016 and 2017 treatment interactions under each of the three specifications in the accompanying table. The 2016, 2017 and 2019 columns are the treatment interactions themselves, bare and with the individual predictor by year terms added, both under the zone fixed effect and both on the same sample. Controls joint p: a Wald test that all 24 predictor-by-year interactions are jointly zero, i.e. whether the added terms explain anything about the outcome at all. A small value there alongside unchanged treatment interactions means the predictors matter for the outcome but are not correlated with treatment, which is the case in which controlling for them cannot move the estimate.

Adding the individual trends moves the 2019 treatment interactions and the pre-trend tests only modestly. The joint test that the 24 added predictor-by-year terms are zero rejects for three of the four outcomes: low trust, dissatisfaction, and police's own floor share, though not its level. The treatment interactions themselves are not appreciably different with or without the added terms. That is the pattern expected when the predictors matter for an outcome without being correlated with treatment.

What Pooling the Protest Strata Costs

The year-by-municipality fixed effect absorbs anything common to a municipality in a wave, including a municipality-level protest surge. This subsection tests whether separating municipalities by their own 2009–2018 protest norm adds anything that fixed effect does not already remove.

The main specification addresses the question directly: it does not separate municipalities by filtering. It keeps every municipality and interacts treatment-by-wave with $\mathbf{1}\{S_c \geq 1\}$, so the difference between the strata is a coefficient with a standard error rather than a comparison between two fits. That coefficient is Panel B of the main estimates table: -0.171 (SE 0.067 , $p = .011$) for the low-trust share and -0.230 (SE 0.073 , $p = .002$) for democratic dissatisfaction. The strata differ, and the difference is roughly the size of the no-surge effect itself; Panel C's implied surge effects are therefore near zero.

This section keeps the older, untested version of the same comparison, because it reports a quantity the tested version does not: what happens to a single treatment coefficient when the two strata are pooled into one.

Table B24: What pooling the two protest strata does to a single treatment coefficient.

Outcome	Effect in 2019 (SE)	t	p	2016 (SE)	2017 (SE)	Pre-trend p	Obs.	Municipalities	Zones
<i>A: No-surge group only ($S < 1$)</i>									
Political trust: low share	0.187 (0.063)	2.98	0.004	0.027 (0.085)	0.079 (0.079)	0.591	1,379	14	91
Satisfaction with democracy: not satisfied share	0.313 (0.089)	3.51	0.001	0.168 (0.110)	0.153 (0.082)	0.125	1,325	14	91
Police trust: level	-0.130 (0.192)	-0.68	0.500	0.179 (0.299)	0.227 (0.211)	0.543	1,410	14	91
Police trust: low share	0.041 (0.088)	0.46	0.644	-0.068 (0.054)	-0.044 (0.068)	0.459	1,410	14	91
<i>B: Pooled (every municipality)</i>									
Political trust: low share	0.033 (0.023)	1.41	0.160	-0.005 (0.032)	-0.014 (0.032)	0.911	10,329	93	760
Satisfaction with democracy: not satisfied share	0.026 (0.033)	0.78	0.438	0.060 (0.039)	0.071 (0.041)	0.157	10,021	93	759
Police trust: level	0.107 (0.072)	1.49	0.136	0.057 (0.086)	0.014 (0.081)	0.790	10,554	93	761
Police trust: low share	-0.033 (0.026)	-1.27	0.205	-0.036 (0.029)	-0.009 (0.026)	0.436	10,554	93	761

COLUMNS. Effect in 2019: the treatment-by-2019 interaction in the outcome's own units, with its standard error in parentheses. 2016 (SE) and 2017 (SE): the corresponding pre-period interactions. Pre-trend p: a joint Wald test that both pre-period interactions are zero. Obs.: respondent-wave observations. Municipalities and Zones: clusters in the estimation sample. PANELS. A restricts estimation to municipalities below their own 2009-2018 protest norm. B pools every municipality in the panel, changing nothing else about the specification. NEITHER PANEL IS THE MAIN SPECIFICATION. Both are unadjusted difference-in-differences without the stratum interaction and without the pre-trend protocol's baseline-outcome trend terms, held identical across the two so that the sample is the only thing that differs. The main table estimates the two strata jointly and reports their difference as the DDD contrast in its Panel B, with a standard error attached; that coefficient, not the gap between the two panels here, is the study's test of whether the strata differ. SPECIFICATION. Treatment is an INDH-registered event in the respondent's own INE census zone by their interview date. Respondent and year-by-municipality fixed effects, 2018 the reference year by omission, standard errors clustered on the zone, under both panels.

Table B25: Effect in 2019 estimated on the no-surge group alone and on every municipality pooled, and the difference between them.

Outcome	No-surge only		Pooled		Difference	Ratio
	No-surge Effect (SE)	No-surge Pre-trend p	Effect (SE)	Pre-trend p		
Political trust: low share	0.187 (0.063)	0.591	0.033 (0.023)	0.911	-0.154	0.17
Satisfaction with democracy: not satisfied share	0.313 (0.089)	0.125	0.026 (0.033)	0.157	-0.287	0.08
Police trust: level	-0.130 (0.192)	0.543	0.107 (0.072)	0.790	0.238	-0.82
Police trust: low share	0.041 (0.088)	0.459	-0.033 (0.026)	0.436	-0.074	-0.82

COLUMNS. No-surge only and Pooled: the treatment-by-2019 interaction under each sample, standard error in parentheses. Difference: pooled minus no-surge, in the outcome's own units. Ratio: pooled over no-surge. The two estimates are not independent – the no-surge sample is a subset of the pooled one – so the difference is reported as a magnitude and no test of it is attached. The tested version of this comparison is the DDD contrast in the main table, which estimates both strata in one model and attaches a standard error to their difference. Pre-trend p columns are the joint Wald test on the 2016 and 2017 interactions under each sample. Everything other than the sample is held fixed across the two columns; see the accompanying table for the full specification and cluster counts.

Pooling collapses the effect. Political trust's low share falls from 0.187 (SE 0.063, 14 municipalities) to 0.033 (SE 0.023, 93 municipalities), and democratic dissatisfaction from 0.313 (SE 0.089) to 0.026 (SE 0.033). The pooled standard errors are smaller, not larger, so this is not power lost to noise; the point estimate itself moves toward zero on substantially more data. Both panels fit the same unadjusted difference-in-differences, with no stratum interaction and none of the pre-trend protocol's baseline-outcome trend terms, so only the sample differs. The no-surge figures here are therefore the unadjusted counterparts of Panel A, and for democratic dissatisfaction that unadjusted 0.313 is the number the protocol replaced with 0.201.

The year-by-municipality fixed effect removes a level confound: a shock to trust shared by every zone in a municipality-wave. It does not require the treatment effect itself to be the same across municipalities. Suppose own-zone violence moves trust where protest activity is at its own historical norm, but less so where a municipality is already surging. That could happen because a police event is a smaller relative shock against an already- mobilised backdrop, or because which zone records an event stops being close to haphazard once a municipality is surging. Pooling every municipality then estimates a composition-weighted average dominated by the 79 additional, higher-surge municipalities. Stratum separation is therefore not redundant with the fixed effect: it selects which population the average treatment effect describes, rather than removing a confound the fixed effect leaves behind. The full DDD separates the strata without discarding anyone.

Sub-group Moderation and Political Heterogeneity

Ideology: Left vs. Non-Left

Table B26 and Figure B5 test whether the estimates differ by respondent ideology. Table B26 and Figure B5 split the focused non-shocked 2019 effect by whether the respondent placed themselves on the left (0–4 on ELSOC’s 0–10 left-right scale, baseline wave) or not, where “not left” includes respondents who did not place themselves at all. The point estimates run in the more adverse direction among left respondents on all four outcomes, most visibly for police trust’s level, where the not-left effect is 0.050 against -0.512 among the left. But none of the four differences is distinguishable from zero at conventional levels ($p = 0.201, 0.935, 0.091, 0.548$).

Table B26: Effect in 2019 by respondent ideology: left (baseline left-right self-placement 0-4 on ELSOC’s 0-10 scale) versus everyone else, including respondents who did not place themselves.

Outcome	Not-left		Left		Difference (SE)	p (difference)
	effect (SE)	N	effect (SE)	N		
Political trust: low share	0.130 (0.078)	325	0.297 (0.104)	89	0.167 (0.130)	0.201
Satisfaction with democracy: not satisfied share	0.317 (0.118)	313	0.333 (0.154)	89	0.017 (0.205)	0.935
Police trust: level	0.050 (0.230)	334	-0.512 (0.251)	88	-0.562 (0.329)	0.091
Police trust: low share	0.009 (0.100)	334	0.102 (0.128)	88	0.093 (0.154)	0.548

Pooled triple-interaction specification – $tr:t_year$, $mod_left_d:t_year$, and $tr:t_year:mod_left_d$ together, with respondent and municipality-by-wave fixed effects and standard errors clustered on the zone – the design’s own specification with ideology and its interaction with treatment added, not a split-sample re-estimation. Not-left effect: the treatment-by-2019 effect for respondents not coded left ($mod_left_d = 0$), i.e. everyone placing themselves at 5-10 on the scale plus those who did not place themselves at all. Left effect: the same effect for $mod_left_d = 1$, computed as the not-left effect plus the difference below, with its standard error from the joint covariance of the two coefficients. Difference: the $tr:t_2019:mod_left_d$ coefficient itself, the moderation estimate, with its own p-value. The $mod_left_d:t_year$ terms (not shown) absorb any ideology-specific trend common to both arms, so the difference column isolates treatment-specific heterogeneity rather than a pre-existing ideological gap. This specification does not itself test parallel trends separately by ideology.




Figure B5: Treatment effect in 2019 by ideology: left (0–4) vs. everyone else, including respondents who did not place themselves.

Left vs. Right, Centre Dropped


A stricter version of the same comparison codes left as 0–4 and right as 6–10 on the same base-line scale. It drops the centre (5), independiente, ninguna, and true missing from the estimation sample entirely, rather than pooling them into either arm. Table B27 and Figure B6 repeat the focused non-shocked specification on this narrower sample. The pattern is broadly the same as the left-vs-everyone comparison above: point estimates run larger among left respondents on three of four outcomes, most visibly for police trust's level, where the left effect is -0.702 against -0.099 among the right. But none of the four differences clears conventional significance ($p = 0.794, 0.608, 0.160, 0.855$). With fewer than 90 respondents in either arm, the design has limited power to detect one.

Table B27: Effect in 2019 by respondent ideology: left (baseline left-right self-placement 0-4 on ELSOC's 0-10 scale) versus right (6-10), the centre and non-responders dropped from the estimation sample.

Outcome	Left		Right		Difference (SE)	p (difference)
	effect (SE)	N	effect (SE)	N		
Political trust: low share	0.364 (0.124)	88	0.407 (0.180)	83	0.043 (0.163)	0.794
Satisfaction with democracy: not satisfied share	0.373 (0.150)	88	0.537 (0.266)	84	0.165 (0.320)	0.608
Police trust: level	-0.702 (0.242)	87	-0.099 (0.394)	84	0.603 (0.425)	0.160
Police trust: low share	0.183 (0.140)	87	0.214 (0.157)	84	0.031 (0.168)	0.855

Pooled triple-interaction specification – $tr:t_year$, $mod_right_d:t_year$, and $tr:t_year:mod_right_d$ together, with respondent and municipality-by-wave fixed effects and standard errors clustered on the zone – the design's own specification with ideology and its interaction with treatment added, not a split-sample re-estimation. Unlike the left-vs-everyone comparison in the main ideology table, here the estimation sample itself is restricted to respondents who placed themselves at 0-4 or 6-10: placement at the centre (5), independiente (11), ninguna (12), and true missing are all dropped, not folded into either arm. Left effect: the treatment-by-2019 effect for the left arm ($mod_right_d = 0$). Right effect: the same effect for the right arm ($mod_right_d = 1$), computed as the left effect plus the difference below, with its standard error from the joint covariance of the two coefficients. Difference: the $tr:t_2019:mod_right_d$ coefficient itself, the moderation estimate, with its own p-value. The $mod_right_d:t_year$ terms (not shown) absorb any ideology-specific trend common to both arms, so the difference column isolates treatment-specific heterogeneity rather than a pre-existing ideological gap. This specification does not itself test parallel trends separately by ideology.

Treatment effect by ideology: left vs. right



Left (0-4) vs. right (6-10) on ELSOC's baseline 0-10 self-placement; centre and non-responders dropped from the estimation sample rather than pooled with either arm. From the pooled triple-interaction specification in the accompanying table. Not the same model as the main-text figure or the left-vs-everyone ideology figure. The three share outcomes share a common axis; police trust's level does not, since it is a different unit.

Figure B6: Treatment effect in 2019, left (0–4) vs. right (6–10), centre and non-responders dropped from the estimation sample.

2017 Presidential Vote

Table B28 and Figure B7 split the design's own 2019 effect by whether the respondent voted for right-wing Sebastián Piñera in the first round of the November 2017 presidential election, among the subsample with a recorded first-round vote. ELSOC has no item for the December runoff, and a missing vote recall is dropped rather than coded not- Piñera, so this comparison is narrower than the ideology moderator's. Police trust's level is the one heterogeneity result in this section that reaches conventional significance: the effect is -0.544 among respondents who voted for another candidate against $+0.526$ among Piñera voters, a difference of 1.070 ($p = 0.016$). The other three differences are not distinguishable from zero ($p = 0.239, 0.305, 0.194$).

Table B28: Effect in 2019 by 2017 vote: Sebastian Pinera versus any other first-round candidate.

Outcome	Not-Pinera		Pinera		Difference (SE)	p (difference)
	effect (SE)	N	effect (SE)	N		
Political trust: low share	0.278 (0.088)	128	0.026 (0.178)	123	-0.252 (0.212)	0.239
Satisfaction with democracy: not satisfied share	0.348 (0.125)	127	0.086 (0.261)	117	-0.262 (0.254)	0.305
Police trust: level	-0.544 (0.327)	127	0.526 (0.354)	124	1.070 (0.433)	0.016
Police trust: low share	0.149 (0.142)	127	-0.066 (0.164)	124	-0.215 (0.164)	0.194

Pooled triple-interaction specification – `tr:t_year`, `mod_pinera_d:t_year`, and `tr:t_year:mod_pinera_d` together, with respondent and municipality-by-wave fixed effects and standard errors clustered on the zone – the design’s own specification with 2017 vote choice and its interaction with treatment added, not a split-sample re-estimation. The moderator is ELSOC c39, a single-wave (ola 3/2018) recall of the respondent’s FIRST-ROUND vote in the November 2017 presidential election; ELSOC carries no item for the December runoff. Unlike the ideology moderator, which codes missing self-placements as not-left, a missing vote recall here is dropped rather than coded not-Pinera, so this is Pinera voters versus other first-round candidates’ voters, on the subsample with a recorded vote – a narrower and smaller comparison than the ideology table’s. Not-Pinera effect: the treatment-by-2019 effect for respondents who recalled voting for another candidate. Pinera effect: the same effect for Pinera voters, computed as the not-Pinera effect plus the difference below, with its standard error from the joint covariance of the two coefficients. Difference: the `tr:t_2019:mod_pinera_d` coefficient itself, the moderation estimate, with its own p-value. The `mod_pinera_d:t_year` terms (not shown) absorb any vote-choice-specific trend common to both arms, so the difference column isolates treatment-specific heterogeneity rather than a pre-existing gap between Pinera and non-Pinera voters. This specification does not itself test parallel trends separately by 2017 vote.




Figure B7: Treatment effect in 2019 by 2017 vote: Piñera vs. every other first-round candidate.

Pinera vs. every other first-round candidate in the November 2017 presidential election (respondents with no recorded vote are dropped, not coded not-Pinera), from the pooled triple-interaction specification in the accompanying table. Not the same model as the main-text figure, which carries no vote interaction. The three share outcomes share a common axis; police trust's level does not, since it is a different unit.

Police Dynamics, Cutpoints, and Conduct

Wave-to-Wave Police Trust Transitions

the main estimates table reports no police-trust effect distinguishable from zero at conventional levels in any of its three DDD components, and that result is easy to misread. It does not say that nothing happened to police trust in 2019. In the focused no-surge stratum, the share rating Carabineros *nada* is 11.7, 15.2 and 10.3 percent across the three pre-treatment waves and then 24.5 percent in 2019, while the mean falls from 3.13 to 2.61 on the five-point scale. Trust in the police collapsed. Within that stratum it collapsed in treated and untreated zones alike: the floor share rises from 0.119 to 0.275 among treated respondents and from 0.100 to 0.239 among control respondents. The municipality-by-wave fixed effects remove a shock common to every zone in a municipality. The design is built to detect differences between neighbors, and on this outcome there were none.

A mean can conceal how such a fall is produced. The same drop follows from everyone sliding one category, or from part of the sample collapsing to the floor while another part recovers, and those are different phenomena. Table B29 reports the wave-to-wave movement directly.

Table B29: Movement in trust in police between consecutive waves, by exposure.

Group	Stays at floor	Top to floor	Stays at top
<i>2016 to 2017</i>			
treated	0.67 (2/3)	0.10 (2/20)	0.40 (8/20)
control	0.23 (7/30)	0.10 (12/115)	0.55 (63/115)
<i>2017 to 2018</i>			
treated	0.43 (3/7)	0.00 (0/12)	0.58 (7/12)
control	0.26 (9/34)	0.01 (1/86)	0.70 (60/86)
<i>2018 to 2019</i>			
treated	0.38 (3/8)	0.18 (4/22)	0.32 (7/22)
control	0.48 (16/33)	0.13 (19/145)	0.41 (60/145)

Cells are the share of respondents in the origin category who end in the destination category, with counts in parentheses as destination over origin. Categories collapse the five-point scale to three: none; a little or some; quite a lot. Top to floor is the share of respondents rating the police quite a lot in the first wave who rate them none in the second. Only respondents observed in both waves of a pair contribute, so the counts differ across rows. Estimation sample; treated means an INDH-registered event in the respondent's own zone by their interview date.

The collapse is a genuine movement to the floor, and it is specific to 2019. Among respondents who rated the police *bastante* or *mucha* in one wave, the share rating them *nada* in the next is 0.10 (treated) and 0.10 (control) across 2016–2017, then 0.00 (treated) and 0.01 (control) across 2017–2018, and 0.18 (treated) and 0.13 (control) across 2018–2019. Retention at the top falls in step, from roughly 0.58 (treated) and 0.70 (control) in the last pre-treatment pair to 0.32 (treated) and 0.41 (control). The gap between treated and control in that flow is based on 4 respondents against 19 and should not be read as a difference.




Figure B8: Marginal distribution of trust in police by wave.

Marginal distribution of trust in police by wave, estimation sample. Bars sum to one within each wave and group.




Figure B9: Movement in trust in police from 2018 to 2019, respondents observed in both waves. Colour is the 2018 category.

Movement in trust in police from 2018 to 2019, respondents observed in both waves. Colour is the 2018 category; width of each stratum is the percentage of respondents in that category, within group – the two panels share one y-scale so treated and control are directly comparable despite their different Ns.




Figure B10: Transition probabilities for trust in police across the treatment wave. Rows sum to one.

Row-normalised transition probabilities for trust in police across the treatment wave. Each row sums to one. Only respondents observed in both 2018 and 2019 contribute: 67 treated and 337 control.

Threshold Cutpoint Regressions

That pattern also settles how the outcome should be specified. Trust in police is ordinal, and Panel A of the main estimates table estimates a no-surge effect on the mean of the coded 1–5 variable. That mean is a summary of the distribution only if its steps are equal in size — if moving from *bastante* to *algo* is the same amount of change as moving from *poca* to *nada*. Respondents jumped from the top of the scale to its floor. The Panel A estimate is then a number in units whose equal spacing is not established by the instrument.

The threshold specification drops the assumption without changing anything else. For each cutpoint k , the focused no-surge equation is estimated on the binary indicator that the rating falls at or below k . The fixed effects, the clustering and the pre-trend test are untouched; only the coding of the outcome differs. Four coefficients then describe where in the distribution treated and control zones diverge, if anywhere.

Table B30: Trust in police at every threshold of the response scale, rather than as a mean.

Outcome	Share in category	Effect in 2019 (SE)	p	Pre-trend p	Obs.
Trust at none or below	0.158	0.041 (0.088)	0.644	0.459	1,410
Trust at none or a little or below	0.344	0.101 (0.110)	0.359	0.378	1,410
Trust at none, a little or some or below	0.637	0.027 (0.092)	0.765	0.935	1,410
Trust at anything but a lot or below	0.905	-0.039 (0.027)	0.144	0.190	1,410

Each row is the design's own specification – respondent and year-by-municipality fixed effects, 2018 the reference year, standard errors clustered on the zone – applied to a binary outcome that equals one if the respondent's rating falls at or below that threshold. Rows are not independent: each nests the one above it, and the fourth is the complement of rating the police a lot. Share in category is the unconditional share of observations at or below the threshold, given so each effect can be read against the size of the group it concerns. Reporting thresholds rather than a mean avoids treating the five-point scale as though its steps were equal in size, which the transition patterns in the transitions table give no reason to assume.

They do not diverge anywhere. The estimates run +0.041, +0.102, +0.028 and -0.039 across the four thresholds, none distinguishable from zero, and pre-trends hold at each ($p = 0.19$ to 0.94). The null on police in the main estimates table is not an artifact of collapsing an ordinal scale to a mean. No part of the distribution moves differently in treated zones. The same four rows also confirm the earlier point: 90.5 percent of observations fall at or below *bastante*, so the linear model's implicit weighting is dominated by distinctions among categories that almost everyone occupies.

Common Shift vs. Local Differential Decomposition

Every outcome in this study moves in 2019, and the design estimates only the local part of that movement. Setting the two components side by side shows where each outcome's change is located.

Table B31: Where each outcome's 2019 movement sits: the part common to all zones and the part specific to exposed ones.

Outcome	Common shift	Local differential (SE)	p	Specification	Local as share of common
Political trust: low share	0.132	0.187 (0.062)	0.003	Unadjusted DDD	1.42
Satisfaction with democracy: not satisfied share	0.220	0.201 (0.066)	0.002	Baseline-outcome trends	0.91
Police trust: low share	0.139	0.041 (0.087)	0.641	Unadjusted DDD	0.29
Police trust: level	-0.501	-0.130 (0.191)	0.495	Unadjusted DDD	0.26
Justified: repress marches	-0.315	-0.063 (0.172)	0.717	Unadjusted DiD (S < 1 only)	0.20
Justified: evict schools	-0.552	-0.186 (0.222)	0.403	Unadjusted DiD (S < 1 only)	0.34

Common shift is the raw change from 2018 to 2019 among respondents in untreated zones within the $S < 1$ group. It is not a causal estimate: it contains the estallido, the national media environment, and everything else that moved the whole country. Local differential is the design's own 2019 coefficient for that group, the only part the estimator identifies, since the municipality-by-wave fixed effects absorb anything common by construction. The last column is the local differential over the common shift in absolute value. THE TWO COLUMNS ARE NOT THE SAME KIND OF QUANTITY. The common shift is always a raw difference in means. The local differential is whichever specification that outcome reports, named in the Specification column: the four headline outcomes are read from the protocol-selected full-sample DDD, and for democratic dissatisfaction that selected model carries baseline-outcome differential trends, which reallocate part of the wave movement the raw common shift leaves in place. The two police-force justification outcomes are not in the main table at all and retain their own unadjusted DiD fit on this group. Read the ratio as an order of magnitude, not as a partition of a single number. A zero local differential is produced equally by an outcome the treatment does not move and by an outcome on which every respondent was effectively treated. This decomposition does not separate those, and no comparison available in these data does.

Within the no-surge group, the two substantive outcomes divide their movement roughly evenly. The low-political-trust share rises 0.132 among respondents in untreated zones and a further 0.187 in exposed ones, and dissatisfaction with democracy 0.220 and 0.201. Every police measure divides it the other way. Trust in police falls 0.501 scale points across untreated zones and a further 0.130 — not distinguishable from zero — in exposed ones. Approval of police repressing marches falls 0.315 and 0.063, and approval of evicting occupied schools falls 0.552 and 0.186.

The two columns are not the same kind of quantity, and the table's Specification column says which model each local differential comes from. The common shift is always a raw difference in means. The local differential comes from whichever specification that outcome reports. For democratic dissatisfaction that is the baseline-outcome differential-trends DDD, whose trend terms reallocate part of the wave movement the raw common shift leaves in place. The ratio should therefore be read as an order of magnitude, not a point value. Democratic dissatisfaction's ratio is 0.91 under the selected model and would be 1.42 under the unadjusted one. Both readings say the local and common components are of comparable size, and the column is used only for that claim.

The police measures are not inert: they show the largest aggregate movements of any out-

come here, larger than the two on which the design finds effects. They lack variation between neighbors. Police use of force during the estallido was nationally visible, and a respondent did not need a registered event in their own zone to learn about it and revise. On this account every respondent was in some measure exposed, and a design built on the contrast between exposed and unexposed neighbours has little left to detect. Democratic satisfaction and institutional trust, by contrast, appear to respond to something proximate that national coverage did not supply.

The reading is consistent with the evidence and is not established by it. A local differential of zero has two possible sources: an outcome that does not respond to treatment, and an outcome on which everyone was effectively treated. The municipality-by-wave fixed effects remove the common component by construction, so the two cannot be separated here. Nor can redrawing the exposure boundary help: saturation and absence both leave the estimate insensitive to where the line is drawn. Distinguishing them would need variation in national exposure — in media reach, or in timing relative to the most publicised incidents — that these data do not contain.

Police Conduct vs. Police Confidence

Trust in an institution and approval of what it does are not the same attitude, and a respondent may hold the government rather than Carabineros responsible for how Carabineros behave. If so, an item about police conduct would register what a confidence item misses. ELSOC has one: c35_03 asks how often the respondent believes Carabineros treat people of their own social class with respect.

The item does not supply the missing evidence for the police channel, and it cannot be read as evidence itself.

Table B32: Perceived respectful treatment by Carabineros: an attitude about police conduct rather than confidence. Reported with its assumption failures, not as a result.

Specification	Effect in 2019 (SE)	p	2016 placebo (SE)	p (2016)	Within R2	Obs.
(1) bare	-0.580 (0.322)	0.075	0.561 (0.282)	0.050	0.011	776
(2) + individual	-0.532 (0.301)	0.080	0.525 (0.284)	0.068	0.022	770
(3) + neighbourhood x year	-0.623 (0.482)	0.199	0.394 (0.281)	0.165	0.063	770

Outcome is c35_03, how often the respondent believes Carabineros treat people of their social class with respect, on a five-point scale from never to always. Waves 2016, 2018 and 2019; the item is not fielded in 2017 and 2018 is the reference. Column 2 adds individual time-varying controls: educational attainment, subjective class placement and agreement with it, and march participation at the previous observed wave. Contemporaneous political interest, news consumption, left-right placement and march participation are excluded as plausibly downstream of exposure. Column 3 adds the six neighbourhood covariates interacted with year. Within R2 is the share of residual variation the specification explains after the fixed effects. TWO FAILURES. First, parallel trends does not hold: the 2016 placebo is of comparable size and opposite sign to the 2019 estimate in every column, the signature of a series drifting across the window rather than stepping at treatment. Second, the item runs on a rotating module and its 2019 coverage depends on treatment: 40.6 percent of treated respondents have it against 26.2 percent of controls, chi-squared $p = 0.023$. These estimates are therefore computed on a treatment-selected subsample and would need a specification that models the pre-period divergence, and a coverage correction, before they could be read as evidence.

The bare estimate is -0.580 (SE 0.322) — treated respondents report less respectful treatment in 2019 — but the 2016 placebo is $+0.561$ (SE 0.282), almost the same magnitude with the opposite sign. Column 2 adds individual time-varying controls: educational attainment, subjective class placement and agreement with it, and march participation at the previous observed

wave. Contemporaneous political interest, news consumption, left–right placement and march participation are excluded, on the same reasoning that leads the design to lag march participation everywhere else. Exposure to repression could have moved all four in 2019, and conditioning on a descendant of treatment would bias the treatment-effect estimate that remains. Column 3 adds the neighborhood covariates interacted with year.

The controls do explain residual variation: within R^2 rises from 0.011 to 0.022 and then 0.063. The estimates move across the columns: -0.580 , -0.532 , -0.623 for 2019, against $+0.561$, $+0.525$, $+0.394$ for the 2016 placebo, with p -values on the placebo of 0.05, 0.07 and 0.16.

A second problem is unrelated to the first. The c35 battery runs on a rotating module: it is absent in 2017 and reaches 37 percent of respondents in 2019. That coverage is not independent of treatment — 40.6 percent of treated respondents were administered the item against 26.2 percent of controls ($\chi^2 p = 0.023$). The estimates are computed on a treatment-selected subsample.

The item is reported because it asks a genuine question about police conduct, and because its failures show what a better measure would need. It is not part of the evidence for anything in this report.

Alternative Definitions of Exposure

An effect attributed to living in the zone where an event occurred could be an artifact of where exactly the exposure line falls. In that case the boundary is arbitrary and the treatment measures something more diffuse. The data layer records four alternative lines built from the same geocoded event file and the same interview-date timing rule as the main definition: an event within 500, 1000, or 1500 m of the respondent, and an event in the respondent's own or any adjacent census zone. The 2000 m radius was retired: at that distance 2354 of 3417 placed respondents counted as exposed, so the comparison group was largely respondents who lived only a little further from an event. Each is substituted for the treatment indicator in the main full-sample triple difference with everything else held fixed — the estimation sample rule, the stratum interaction, the fixed effects, zone clustering, and each outcome's protocol-selected specification. The specification is deliberately not re-selected per definition, so only the exposure line varies.

Table B33: The main triple difference under alternative definitions of own-exposure.

Definition	Treated	S < 1 effect (SE)	p	DDD contrast (SE)	p	Pre-trend p
<i>A: Political trust: low share</i>						
Own census zone (main definition)	437	0.187 (0.062)	0.003	-0.171 (0.067)	0.011	0.799
Event within 500 m	649	0.129 (0.065)	0.047	-0.135 (0.068)	0.047	0.342
Event within 1000 m	1,400	0.118 (0.058)	0.044	-0.108 (0.061)	0.077	0.500
Own or adjacent census zone	1,527	0.168 (0.060)	0.005	-0.160 (0.063)	0.011	0.462
Event within 1500 m	1,970	0.118 (0.061)	0.051	-0.115 (0.064)	0.073	0.668
<i>B: Satisfaction with democracy: not satisfied share</i>						
Own census zone (main definition)	431	0.201 (0.066)	0.002	-0.230 (0.073)	0.002	0.659
Event within 500 m	644	0.181 (0.074)	0.015	-0.193 (0.078)	0.014	0.505
Event within 1000 m	1,373	0.180 (0.065)	0.006	-0.173 (0.069)	0.012	0.455
Own or adjacent census zone	1,504	0.115 (0.081)	0.159	-0.100 (0.084)	0.238	0.241
Event within 1500 m	1,940	0.136 (0.071)	0.056	-0.143 (0.076)	0.059	0.771
<i>C: Police trust: level</i>						
Own census zone (main definition)	444	-0.130 (0.191)	0.495	0.260 (0.206)	0.207	0.817
Event within 500 m	659	-0.067 (0.185)	0.716	0.150 (0.198)	0.447	0.026
Event within 1000 m	1,418	0.030 (0.179)	0.866	-0.079 (0.189)	0.677	0.283
Own or adjacent census zone	1,550	-0.055 (0.201)	0.785	0.035 (0.209)	0.867	0.404
Event within 1500 m	2,001	-0.047 (0.164)	0.775	0.037 (0.177)	0.833	0.736
<i>D: Police trust: low share</i>						

Continued on next page.

Table B33 continued from previous page.

Definition	Treated	S < 1 effect (SE)	p	DDD contrast (SE)	p	Pre-trend p
Own census zone (main definition)	444	0.041 (0.087)	0.641	-0.079 (0.091)	0.384	0.661
Event within 500 m	659	0.100 (0.079)	0.203	-0.143 (0.083)	0.083	0.023
Event within 1000 m	1,418	-0.036 (0.076)	0.639	0.045 (0.079)	0.572	0.015
Own or adjacent census zone	1,550	0.044 (0.073)	0.549	-0.045 (0.076)	0.557	0.647
Event within 1500 m	2,001	0.002 (0.093)	0.982	0.005 (0.096)	0.961	0.809

DEFINITIONS. Own census zone is the main definition: an INDH-registered event in the respondent's own INE census zone by their interview date. The buffer rows redraw that boundary at 500, 1000, or 2000 m of the respondent; own or adjacent adds the neighbouring zones. All five share the same event file and timing rule, and rows are ordered by treated share (the Treated column) within each panel. COLUMNS. S < 1 effect: the treatment-by-2019 interaction among municipalities below their own 2009–2018 protest norm, standard error in parentheses. DDD contrast: the high-minus-low differential on that effect. Pre-trend p: joint Wald test on all four pre-period interactions. SPECIFICATION. Each row refits the main full-sample triple difference changing only the treatment indicator: respondent and year-by-municipality fixed effects, 2018 the reference year, standard errors clustered on the zone, and each outcome keeping the specification the pre-trend protocol selected in the main table – the protocol is not re-run per definition. As the line widens the untreated group increasingly consists of respondents far from any event, so wide definitions average own-zone exposure with mere proximity; attenuation along that ordering is the hyperlocal-mechanism prediction, not instability.

Both outcomes keep the sign and rough size of their effect under every redefinition. The no-surge effect on the low-trust share runs 0.187 (own zone), 0.129 (500 m), 0.118 (1000 m), 0.168 (own or adjacent), 0.118 (1500 m); on democratic dissatisfaction, 0.201, 0.181, 0.180, 0.115, 0.136. Every DDD contrast on these two outcomes stays negative.

The wider definitions lose precision, and two of them stop separating from zero at conventional levels: dissatisfaction under the own-or-adjacent line ($p = .159$), and both outcomes at 1500 m ($p = .051$ and $p = .056$). The point estimates also shrink as the line widens, and the ordering of that shrinkage is informative. The widest definition treats 2,018 of the panel's 3,343 respondents, so its untreated comparison increasingly consists of respondents further from any event, and its coefficient averages own-zone exposure together with proximity. A hyperlocal mechanism predicts attenuation along that ordering. The attenuation is not monotone in radius, though — dissatisfaction is larger at 1500 m (0.136) than under the own-or-adjacent line (0.115) — so the ladder should be read as a pattern rather than a dose-response curve.

The police outcomes stay null under the redefinitions as they are in the main table, with two caveats, both reported in the table. The widest line produces one nominally significant police-trust coefficient (+0.300, $p = 0.046$, with a matching negative DDD contrast). Under that definition 2,018 of the panel's 3,343 respondents are treated, so the comparison is closer to event-distant versus everyone else than to the design's own contrast. The joint pre-trend test also rejects at the 5 percent level for the police outcomes under the 500 m and 1000 m lines ($p = 0.015$ to 0.026), so those cells do not support a clean reading in either direction. Neither caveat applies to the two outcomes on which the study's results rest, whose pre-trend tests hold

under all five definitions ($p = 0.17$ to 0.80).

Statistical Power and Minimum Detectable Effects

All three reported components rest on the same full-sample fit, but they are not identified off the same clusters. The no-surge DiD and, through it, the DDD contrast depend on the 13 treated zones in the 14 municipalities below their own protest norm; the implied surge effect rests on 110 treated zones in 75 municipalities. The usual $2.8 \times \text{SE}$ rule for a minimum detectable effect assumes a normal sampling distribution for the cluster-robust t statistic. Whether that holds is a question rather than a premise at the first of those cluster counts. It is checked by simulation, for every component rather than only the first.

The observed outcomes are held fixed — every marginal, every within-respondent and within-zone dependence is the data’s own — and we simulate assignment. Each replication permutes treatment across zones within municipality, preserving each municipality’s count of treated zones, then injects a known effect into treated 2019 observations. Because the protest shock is measured at the municipality, no permutation moves a zone between strata, so both groups keep their observed treated-cluster counts.

The placement of the injection determines which component’s null is false, and it separates the three panels. Injecting into no-surge treated observations moves the no-surge DiD alone. Injecting into surge treated observations instead moves the DDD contrast and the implied surge effect by the same amount while leaving the no-surge effect at zero, so one pass calibrates both. At zero injected effect, nothing is injected under either scheme, every component’s null holds, and the rejection rate is that component’s empirical size.

Each replication estimates the outcome’s own reported specification, including the baseline-outcome differential-trend terms wherever the pre-trend protocol selected them. The implied surge p -value is computed by the same delta method the main table uses. The simulation therefore calibrates the numbers the table reports. Before running, the script asserts two things: that a replication matches the real fit in observations, fixed-effect cells, cluster count and outcome marginals, and that the simulated estimator reproduces the reported no-surge coefficient exactly when handed the observed assignment.

Table B34: How large an effect each reported component can detect, and how often it reports one when none exists, by simulation.

Outcome	Specification	False positives	Detectable effect	Rule of thumb	Estimate
<i>A. No-surge treatment DiD</i>					
Political trust: low share	Unadjusted DDD	0.100	0.175	0.174	0.187
Satisfaction with democracy: not satisfied share	Baseline-outcome trends	0.087	0.227	0.185	0.201
Police trust: level	Unadjusted DDD	0.095	–	0.535	-0.130
Police trust: low share	Unadjusted DDD	0.101	0.236	0.244	0.041
<i>B. High-minus-low DDD contrast</i>					
Political trust: low share	Unadjusted DDD	0.095	0.192	0.187	-0.171
Satisfaction with democracy: not satisfied share	Baseline-outcome trends	0.085	0.234	0.204	-0.230
Police trust: level	Unadjusted DDD	0.086	–	0.577	0.260
Police trust: low share	Unadjusted DDD	0.098	0.247	0.255	-0.079
<i>C. Implied surge DiD</i>					
Political trust: low share	Unadjusted DDD	0.067	0.071	0.069	0.016
Satisfaction with democracy: not satisfied share	Baseline-outcome trends	0.067	0.081	0.086	-0.029
Police trust: level	Unadjusted DDD	0.063	0.220	0.217	0.130
Police trust: low share	Unadjusted DDD	0.069	0.081	0.077	-0.039

OUTCOMES. Political trust: low share: share of three institutional trust items answered none. Satisfaction with democracy: not satisfied share: not at all satisfied with democracy. Police trust: level is the 1-5 item; Police trust: low share is the share answering none, the item's floor. PANELS. One per component of the main estimates table. All three are calibrated here; earlier versions covered the no-surge component only. COLUMNS. Specification: the protocol-selected model that outcome reports, and the model simulated for it. False positives: how often that component rejects at 5 percent with no effect injected; above 0.05 means the corresponding p-values in the main table are optimistic. Detectable effect: the smallest injected effect found in 80 percent of replications, a dash meaning 80 percent is never reached. Rule of thumb: 2.8 times the standard error. METHOD. Each replication estimates that outcome's own reported specification on 9962 respondent-waves in 89 municipalities, holding the observed outcomes fixed, and re-randomizes treatment at the zone within municipality. The protest surge is measured at the municipality, so no permutation moves a zone between strata. Before running, the script asserts that a replication matches the real fit and that the simulated estimator reproduces the reported Panel A coefficient exactly. INJECTION. Which treated 2019 observations receive the effect is what makes a component's null false. Panel A injects into no-surge treated observations, 13 treated zones in 14 municipalities. Panels B and C inject into surge treated observations instead – 110 treated zones in 75 municipalities – which moves the contrast and the implied surge effect together and leaves the no-surge effect at zero. Panel C's p-value is the delta-method test on the sum of the two coefficients, referred to the normal, as in the main table. 10000 replications per point, seed 20260804.

The test over-rejects on the two components identified off few treated clusters, and is close to nominal on the third. For the no-surge DiD the cluster-robust t rejects at 9.5 to 11 percent of replications against a nominal 5 for three outcomes. Democratic dissatisfaction is the exception at 6.8 percent: it is the one outcome with baseline-outcome trend terms, which absorb pre-

period variation that the other three specifications leave in the residual. The DDD contrast behaves almost identically, 5.8 to 10.8 percent: it is identified from the same 13 no-surge treated zones as the no-surge difference it subtracts. The implied surge effect is the one component that is roughly correctly sized, 4.0 to 7.8 percent, because it rests on 110 treated zones in 75 municipalities rather than 13 in 14.

A treated-cluster count is not, however, a count of the information in the design, and the two are easy to conflate. The cluster-robust correction responds to the correlation remaining among within-zone residuals of the estimating equation once the fixed effects are applied. That quantity is near zero here, so the design effect is essentially one and the effective sample is the full set of respondent-waves rather than anything approaching the cluster count. Clustering on the zone yields a *smaller* standard error on the reported coefficient than the heteroskedasticity-robust alternative, 0.062 against 0.073 for the low-trust share, and clustering more coarsely on the municipality gives 0.070. If clustering were discarding information, the ordering would be reversed.


One caveat is that respondents do not move between zones in this panel, so respondent fixed effects absorb every time-invariant zone characteristic by construction. A low residual intra-zone correlation is therefore partly structural rather than a discovery about Chilean neighborhoods. The residual correlation remains the quantity to which the correction is applied; the standard-error comparison is therefore reported beside it.

The over-rejection above is therefore a calibration problem. The number of ways treatment could have been assigned at the zone level is small, and the analytic t poorly approximates the sampling distribution that assignment induces. The permutation exercise is therefore the binding check rather than a supplement to one.

The same asymmetry appears in the minimum detectable effects. For the two components on which the argument rests, the simulated MDE is at or above the analytic rule: 0.175 against 0.174 and 0.216 against 0.185 for the no-surge DiD, 0.189 against 0.187 and 0.240 against 0.204 for the DDD contrast. The estimates lie at the boundary. Political trust's no-surge effect of 0.187 is just above its 0.175. Democratic dissatisfaction's 0.201 is just *below* its 0.216, and both DDD contrasts, -0.171 and -0.230 , are likewise just below their 0.189 and 0.240. Effects of that size are found in just under 80 percent of replications, so detecting them here was not guaranteed by the design. Their conventional p-values are optimistic to the extent of the over-rejection reported above. The point estimates stand; the intervals around them should not be read as tighter than the simulation indicates. Table B16 shows the two candidate models give the same verdict on both of these estimates, and on nine of the other ten cells.

The implied surge effects are the opposite case, and the place in this study where a null is informative. Their minimum detectable effects are 0.069, 0.079 and 0.083 on the three share outcomes, roughly a third the size of anything the no-surge comparison can resolve. The estimates are $+0.016$, -0.029 and -0.039 . An effect of even a tenth of a share point in municipalities that surged would have been detected in 80 percent of replications and was not. The near-zero Panel C coefficients are therefore precisely estimated nulls, not underpowered ones. This justifies reading the DDD contrast as attenuation in municipalities that surged rather than as noise.

Trust in police as a level is the exception throughout. For the no-surge DiD and the DDD contrast, 80 percent power is never reached anywhere on the grid; the rejection rate is still 53 percent at an injected effect of 0.40, half a point on a five-point scale. Only its implied surge effect is resolvable, at 0.217. The two no-surge nulls on that outcome constrain nothing. The police *nada* share is informative only against effects of roughly 0.23 or larger in the no-surge group, but against 0.083 or larger in the surge group.



Rejection rate against injected effect, from 10000 replications per point, under each outcome's own reported specification. Treatment is re-randomized at the zone within municipality each replication. The three curves are three different alternatives, not three views of one: the no-surge DiD is probed by injecting into no-surge treated 2019 observations, while the DDD contrast and the implied surge effect are probed by injecting into surge treated observations, which moves both by the same amount. At zero injected effect nothing is injected, so each curve's leftmost point is that component's empirical size. Dashed line: 80 percent power. Grey line: nominal 5 percent.

Figure B11: Rejection rate against injected effect, under each outcome's own reported specification, for all three components of the main table. The three curves are three different alternatives, not three views of one: each is driven by the injection scheme that breaks its own null. Each curve's leftmost point is that component's empirical size.

Methodological Standards Comparison

The closest published design is Ang and Tebes (2024), which estimates the civic response to police killings in California by comparing Census blocks where a killing occurred to other blocks in the same block group. The structure is the one used here: a hyperlocal treatment, a neighborhood-level comparison set, and neighborhood-by-period fixed effects. Table [B35](#) sets their battery of tests against this study, item by item, together with two additional items produced here.

Table B35: Tests and inference standards in Ang and Tebes (2024) set against this study.

Item	Ang and Tebes (2024)	This study	Status
<i>A. Design</i>			
Comparison set	Treated Census block vs other blocks in the same block group; block and block-group-by-election fixed effects (Eq. 2)	Treated zone vs other zones in the same municipality; respondent and municipality-by-wave fixed effects	matched
Treatment timing	First killing in the block between the 2002 and 2010 elections; one event date per block	INDH event in the respondent's own zone on or before their wave-4 interview date	matched
Reference period	Last election before the killing, omitted	2018 wave, omitted	matched
Identifying units	285 blocks with a killing across 68,326 blocks; 341,420 block-elections	123 treated zones across 749 zones in 89 municipalities (53 with both treated and untreated zones); 10,170 respondent-waves	differs (scale)
<i>B. Inference</i>			
Cluster level	Census block group – the fixed-effect unit, one level above treatment; Bertrand, Duflo and Mullainathan (2004)	Zone – the unit treatment is assigned on; Abadie, Athey, Imbens and Wooldridge (2023)	differs (unit)
Number of clusters	6,400	749	differs (scale)
Multiplicity adjustment	None; conventional stars reported	None; four outcomes reported side by side	matched
<i>C. Assumption tests</i>			
Event study	Leads and lags plotted with 95% intervals (Fig. 2)	2016, 2017 and 2019 no-surge interactions plotted with 95% intervals in the main event-study figure	matched
Joint pre-trend test	F = 0.33, p = 0.806 (registrations); F = 1.09, p = 0.350 (votes)	Joint Wald across four outcomes, p = 0.45–0.62	matched

Continued on next page.

Table B35 continued from previous page.

Item	Ang and Tebes (2024)	This study	Status
Treated vs control balance	Demographics, party, age, prior turnout and registration for killing vs non-killing blocks in the same block groups (Table 1, Panel B)	Six neighbourhood covariates, 2016 outcome baselines, pre-period changes, and wave completion for treated vs control zones within both protest strata (balance table)	matched
Placebo treatment	Blocks treated after the outcome window, estimated on referenda (Table 3, cols. 3-4)	Not implemented; pre-period interactions serve the same role and are estimable here	not implemented
<i>D. Robustness</i>			
Alternative specifications	Nine columns: homicide controls, minority-share, adolescent-population and population-change interactions, no population controls, dropping multi-event blocks, full sample, registration threshold (Table 2)	Full-sample DDD: $S < 1$ treatment-by-year effects, surge DDD contrasts, and implied surge effects	differs (DDD)
Proximity check	Coefficients by 0.1-mile distance bins to 2 miles, showing decay away from the incident block (Eq. 1, Fig. 1)	Exposure-redefinition battery: the main DDD refit with the treatment line redrawn at 500/1000/2000 m of the respondent and at own-or-adjacent zone. Effects persist and attenuate as the line widens; a ring design was tried and retired – under municipality-by-wave fixed effects the ring was identified in only 5 municipalities	differs (redefinition, not bins)
Local-crime confound	Equation 2 re-estimated with criminal homicides as the outcome	Not implemented; municipality-by-wave fixed effects absorb municipality-level crime and policing trends, but not within-municipality ones	not implemented

Continued on next page.

Table B35 continued from previous page.

Item	Ang and Tebes (2024)	This study	Status
Selective migration	Student transfer records, ACS same-residence shares (86.6 vs 86.8 percent), and a population difference-in-differences (b = 0.006, p = 0.78)	Movers are dropped at panel construction; differential completion by exposure is reported in the balance table but not otherwise tested	differs
<i>E. Extensions</i>			
Heterogeneity	Voter race, deceased race, age, registration length, party, incident location (Fig. 3)	Exploratory moderation by left-right placement and 2017 Piñera vote, reported for the focused S < 1 comparison	differs (moderators)
Mechanism	Armed vs unarmed victims (Fig. 5); referenda on criminal-justice reform (Table 3); media coverage (Fig. 4)	Police-trust subitems and a limited respectful-treatment item are examined; neither identifies the incident mechanism used by Ang and Tebes	differs (measurement)
<i>B. Inference</i>			
Size of the test	Not reported	For all three DDD components, simulated by re-randomizing treatment within municipality under each outcome's own reported specification: the cluster-robust test rejects at 6.8-11.0 percent for the S < 1 DiD, 5.8-10.8 for the DDD contrast and 4.0-7.8 for the implied surge effect, against a nominal 5 (power table)	beyond
<i>D. Robustness</i>			

Continued on next page.

Table B35 continued from previous page.

Item	Ang and Tebes (2024)	This study	Status
Minimum detectable effect	Not reported	Simulated power curves per outcome, holding the observed outcomes fixed and fitting the specification that outcome reports (power figure). All three reported components are calibrated: the S < 1 DiD, the DDD contrast, and the implied surge effect	beyond

E. Extensions

Replication data	APSR Dataverse	This package: run_all.R reproduces every table and figure from raw inputs	matched
------------------	----------------	---	---------

Ang and Tebes, Civic Responses to Police Violence, American Political Science Review 118(2):972-987, with the figure, table or equation carrying each item. Status is descriptive: matched, the same class of evidence is produced here; differs, a deliberate difference in implementation or a difference in scale; not implemented, the item is absent, with a note on whether the data layer supports it; beyond, produced here with no counterpart in the article. Counts and pre-trend p-values in the right-hand column are read from 07_main_estimates.csv.

The design logic, the event study with a joint pre-period test, and the treated/control comparability table are produced here in the same form. Inference differs in unit and in scale: Ang and Tebes (2024) cluster at the block group, the level of their fixed effect, with 6,400 clusters; this study clusters at the zone, the level at which treatment is assigned, with the outcome-specific cluster counts reported in the main estimates table. Their robustness battery is broader, while this report includes an exposure-redefinition battery and exploratory ideology and vote heterogeneity. It does not implement their placebo-treatment or local-crime-outcome designs. The article reports neither the size of the test nor a minimum detectable effect; both appear in Table B35.

Supplementary Information C: Study 3: Cross-National Evidence

Ignacio Urbina and Vitória Sgorlon

Stony Brook University

September 2026

Contents

What this appendix documents	1
Measurement and sample	3
Robustness and sensitivity	8
Full model estimates	31
Moderation: ideology and race	44
Country-level patterns and scale validation	47
How each result was estimated	55
References	60

What this appendix documents

The question. This study asks whether trust in institutions tracks satisfaction with democracy across Latin America, and whether people who report that the police mistreated them trust institutions less. The study is observational and descriptive. It does not identify a causal effect. The checks reported below ask whether the descriptive pattern is stable, not whether it supports a causal claim.

Two analyses. The first analysis relates institutional trust and trust in the police to satisfaction with democracy. We estimate it separately in two 2023 surveys, the Latinobarómetro and the LAPOP AmericasBarometer. We keep the surveys apart, so agreement between them counts as a replication rather than as one result reported twice. The second analysis relates what respondents say about their own treatment by the police to those same trust measures, and to satisfaction with democracy. It uses the 2008 LAPOP wave. That wave is the only cross-national round in the region that ever asked respondents directly whether the police mistreated them.

Countries. All three analyses use the same sixteen countries: Argentina, Bolivia, Brazil, Chile, Colombia, Costa Rica, the Dominican Republic, Ecuador, El Salvador, Guatemala, Honduras, Mexico, Panama, Paraguay, Peru and Uruguay. Both 2023 surveys fielded these sixteen, and we restrict the 2008 wave to the same set. Fixing the country set means that any difference between the three sets of estimates reflects the survey or the period, never which countries were counted. The 2008 wave also fielded Belize, Haiti, Jamaica, Nicaragua and Venezuela, which we drop. The Latinobarómetro also fields Venezuela, and LAPOP fields Nicaragua, Haiti and several Anglophone Caribbean countries. All of these fall outside the matched set.

Samples. We drop respondents who are missing any variable a model uses. That leaves 13,275 respondents in the Latinobarómetro 2023 sample, out of 18,004 fielded across the sixteen countries; 19,279 in LAPOP 2023, out of 25,059; and 17,703 in LAPOP 2008, out of 26,894 in the sixteen matched countries and 38,053 fielded across twenty-three. All sixteen countries survive in every analysis. Individual checks also drop respondents missing the extra variables those checks require, so each check runs on its own smaller sample and reports that sample.

Two companion studies. Two tables and one figure below compare this study's trust scale with two companion studies in the same project: Study 1, a survey experiment in Brazil ($N = 1,248$), and Study 2, a panel survey in Chile (ELSOC, $N = 6,270$). They appear only where we compare scale reliability across the three studies. They play no other part in the analyses documented here.

How this appendix is organised. **Measurement and sample** documents the survey items, how we coded them, and what the samples look like. **Robustness and sensitivity** lists the checks, states what each would have had to show to overturn the result, and reports what it did show. **Full model estimates** gives every coefficient in every main equation. **Moderation** compares subgroups by ideology and by race. **Country-level patterns and scale validation** reports what differs across countries and tests the trust scale against a published alternative.

What the coefficients mean on the original scales. Every outcome is standardised, so coefficients are in standard deviations rather than in the units respondents answered in. Trust in police is a seven-point item whose standard deviation is 1.86 in the 2008 wave, so the victimization coefficient of -0.402 amounts to about three-quarters of a scale point — roughly the gap between neighbouring response categories. Satisfaction with democracy is a four-point item whose standard deviation is 0.71, so the victimization coefficient of -0.122 amounts to under a tenth of a scale point. The political-trust coefficient of 0.283 per standard deviation of trust amounts to about a fifth of one. In the 2008 estimation sample 5.8% of respondents report police abuse, weighted (5.9% unweighted), so these coefficients compare a small minority with everyone else.

Scales and terms used throughout. The institutional-trust measure combines three items: confidence in the legislature, in the executive and in the courts. We rescale each item to run from 0 to 1, add them, and standardise the sum. All three studies in this project use this scale. The reliability comparison below therefore covers all three. Trust in police is a single item. We standardise every outcome and index within its own dataset, so a coefficient gives the change in outcome standard deviations per standard deviation of the predictor, or per unit for a predictor that is either 0 or 1.

Several statistics recur in the tables below. Marginal R^2 is the share of the outcome explained by the predictors alone, as against conditional R^2 , which also credits the differences between countries. ΔR^2 is how much marginal R^2 rises when the predictors of interest are added. The ICC is the share of the outcome that lies between countries rather than within them. Where we pool country-specific estimates, τ is how much the true effect differs from country to country. I^2 is the share of the observed variation between countries that reflects real differences rather than sampling error, and Q tests whether any real difference is present at all. An E-value states how strong an association an unmeasured variable would need with both the predictor and the outcome, on the risk-ratio scale, to account for an estimate entirely (VanderWeele and Ding 2017). CR2 is a standard-error correction built for the small number of clusters that sixteen countries provide (Pustejovsky and Tipton 2018). The Shapley share divides explained variance among predictors that correlate with one another, by averaging over every order in which they could enter (Grömping 2006). ω measures reliability without assuming, as α does, that every item counts equally.

Measurement and sample

One kind of model produces every number reported here. In all three datasets it is a survey-weighted linear multilevel regression: each country gets its own intercept and its own slope on the predictors of interest, and the remaining controls enter with one slope shared across countries. Country enters as a random effect rather than as a set of dummy variables. That is a choice, and Table C14 reports its consequences by refitting every one of these models with country dummies instead. We standardize every outcome within its own dataset, so coefficients are in outcome standard deviations unless we say otherwise.

Tables C1 and C2 document the survey items, their coding, and the samples.

Table C1: Measurement: source items and coding by construct and dataset.

Construct	Latinobarómetro 2023	LAPOP 2023	LAPOP 2008
Satisfaction with democracy	P11STGBS.A, 4-pt, reverse-coded so higher = more satisfied	pn4, 4-pt, reverse-coded so higher = more satisfied	pn4, 4-pt, reverse-coded
Trust in police	P13STGBS.B, 4-pt, reverse-coded so higher = more trust	b18, 7-pt, higher = more trust (no recoding)	b18, 7-pt, higher = more trust
Political trust index	P13ST.D/I/F (Congress, president, judiciary), each reverse-coded so higher = more trust, then scaled 0–1 and summed	b13, b21a, b31 (Congress, president, supreme court), each scaled 0–1 and summed	b13, b21a, b10a (Congress, president, justice system), each scaled 0–1 and summed
Economic perceptions	4 items (P5STGBS, P6STGBS, P7ST, P8STGBS), reverse-coded, scaled 0–1, summed, then standardised	2 items (soct2, idio2), reverse-coded, scaled 0–1, summed, then standardised; soct1/idio1 not fielded in this wave	4 items (soct1, soct2, idio1, idio2), reverse-coded, scaled 0–1, summed, then standardised
Political interest	2 items (P40STGBS, P44ST.A), reverse-coded then standardised	pol1 only, reverse-coded then standardised; pol2 not fielded	2 items (pol1, pol2), reverse-coded, scaled 0–1, summed, then standardised
Interpersonal trust	P9STGBS, single binary item, reverse-coded then standardised	it1, single 4-point item, reverse-coded then standardised	it1a, it1b, reverse-coded, scaled 0–1, summed, then standardised
Police abuse victimization	not asked	not asked	vic27, count of incidents in the past 12 months in which police mistreated the respondent verbally or physically, or assaulted them; coded 1 if one or more
Crime victimization	P58ST, coded 1 if the respondent or a relative was a crime victim	vic1ext, coded 1 if a crime victim in the past 12 months	vic1, coded 1 if a crime victim in the past 12 months
Age	EDAD, years	q2, years	q2, banded into five cohorts (18–23 to 60+)
Sex	SEXO, 1 if female	q1tc_r, 1 if female	q1, 1 if female

Continued on next page.

Table C1 continued from previous page.

Construct	Latinobarómetro 2023	LAPOP 2023	LAPOP 2008
Education	S11	edre	ed, banded into five levels
Household income	–	q10inc, last two digits, five brackets	q10, banded into five groups
Interviewer-rated status	S24, reverse-coded	–	–
Household assets	S20.A–S20.K (11 items), count owned, requires 8 answered, standardised	r3, r6, r7, r12, r15, r16, r18 (7 items), count owned, requires 5 answered, standardised	–
Ideology	P16ST, 0–10 left to right	11n, 1–10 left to right	11, 1–10 left to right
Urban residence	–	ur, 1 if urban	ur, 1 if urban
National capital	TAMCIUD = 8	–	tamano = 1
Large city	TAMCIUD in 6, 7, 8	–	tamano = 2
Ethnicity	–	–	etid, six categories
Religion	–	–	q3, five categories
Occupation	–	–	ocup4a with ocup1a, five categories
Household type	–	–	banded from marital status and co-resident children (q11, q12, q12a)
Police bribe solicitation	–	–	exc2, 1 if a police officer asked for a bribe
Corruption perception	P60ST, reverse-coded, standardised	exc7new, standardised	–
Winner–loser status	P43STGBS.B: voted government, opposition, or neither	–	–
Presidential approval	P15STGBS, 1 if approves	m1, reverse-coded, standardised	–

We standardise every outcome and trust index within its dataset before estimation, so coefficients are in outcome standard deviations. That makes them comparable across datasets in scale, though not in question wording. Which items each wave asked differs. Where a wave did not ask an item, we build the measure from the remaining items and say so. Only the 2008 wave asked respondents directly whether the police mistreated them. Rows below the first eight are controls, and a dash means the variable does not enter that dataset's models. The 2008 codes are built from the raw wave and checked against the analysis file. One variable name is misleading: the 2008 income variable is called a quintile but is a fixed recode of the income bracket, holding roughly 24, 28, 19, 8 and 4 per cent of respondents.

Table C2: Descriptive statistics for the variables entering the reported models, by dataset and role.

Role in model	Variable	N	Mean	SD	Min	Max
<i>Latinobarómetro 2023</i>						
Outcome	Satisfaction with democracy	17535	0.020	1.000	-1.196	2.013
Predictor of interest	Political trust	17159	0.009	0.993	-1.280	2.472
	Trust in police	17876	0.030	0.998	-1.296	1.755
Control for a competing channel	Economic perceptions	16450	-0.006	0.992	-2.200	2.516
	Political interest	17825	0.002	1.000	-1.274	2.479
	Interpersonal trust	17524	0.010	1.009	-0.447	2.239
Demographic / socioeconomic control	Age	18004	41.222	16.514	16.000	93.000
	Female	18004	0.521	0.500	0.000	1.000
	Education	18004	10.716	4.187	1.000	17.000
	SES (interviewer-rated)	18004	3.560	0.857	1.000	5.000
	Ideology	15396	5.365	3.053	0.000	10.000
	Crime victimization	17822	0.321	0.467	0.000	1.000
	National capital	18004	0.152	0.359	0.000	1.000
	Large city	18004	0.654	0.476	0.000	1.000
Competing predictor (extended specifications only)	Perceived progress reducing corruption	17522	0.001	1.000	-1.041	1.879
	Approves of president	17043	0.429	0.495	0.000	1.000
	Winner/loser status: government	18004	0.206	–	0.000	1.000
	Winner/loser status: none	18004	0.558	–	0.000	1.000
	Winner/loser status: opposition	18004	0.236	–	0.000	1.000
<i>LAPOP 2023</i>						
Outcome	Satisfaction with democracy	24344	-0.000	1.000	-1.756	2.091
Predictor of interest	Political trust	24019	-0.000	1.000	-1.488	2.100
	Trust in police	24882	-0.000	1.000	-1.434	1.609
Control for a competing channel	Economic perceptions	24685	0.000	1.000	-0.942	2.216
	Political interest	24988	0.000	1.000	-1.026	1.935
	Interpersonal trust	24506	0.000	1.000	-1.778	1.368
Demographic / socioeconomic control	Age	25059	40.554	16.641	16.000	95.000
	Female	24666	0.502	0.500	0.000	1.000

Continued on next page.

Table C2 continued from previous page.

Role in model	Variable	N	Mean	SD	Min	Max
Competing predictor (extended specifications only)	Education	24961	3.535	1.615	0.000	6.000
	Income bracket	22415	2.331	1.324	1.000	5.000
	Urban	25059	0.723	0.448	0.000	1.000
	Crime victimization	25034	0.232	0.422	0.000	1.000
	Perceived corruption among politicians	24294	0.000	1.000	-2.655	1.148
<i>LAPOP 2008</i>						
Outcome	Political trust	24892	0.013	0.987	-1.818	2.191
	Trust in police	26474	-0.044	0.982	-1.494	1.730
Predictor of interest	Police abuse victimization	26803	0.053	0.224	0.000	1.000
Demographic / socioeconomic control	Age	26829	39.173	15.988	18.000	101.000
	Female	26894	0.514	0.500	0.000	1.000
	Education (6 categories)	26894	—	—	—	—
	Income quintile (5 categories)	23885	—	—	—	—
	Ideology	21567	5.715	2.413	1.000	10.000
	Crime victimization	26767	0.182	0.386	0.000	1.000
	Occupation (6 categories)	26894	—	—	—	—
	Urban	26894	0.682	0.466	0.000	1.000
	National capital	26894	0.218	0.413	0.000	1.000
	Large city	26894	0.185	0.389	0.000	1.000
	Police bribe solicitation	25759	0.108	0.311	0.000	1.000
Competing predictor (extended specifications only)	Economic perceptions	26008	0.122	0.970	-1.797	2.696
	Political interest	26504	-0.001	0.999	-1.236	2.437
	Interpersonal trust	26236	-0.052	0.946	-0.922	3.453

Statistics are unweighted and computed before any respondent is dropped, so N varies by item and exceeds the estimation samples in the following table. Variables in standardised or scaled form are standardised within dataset, which is why their means sit near zero and their standard deviations near one. Competing predictors that are categorical are broken out by level, where Mean is the share of answering respondents in that level. Categorical demographic controls are summarised by their number of categories instead, because no reported comparison uses their internal splits. The competing predictors enter only the extended specifications. Economic perceptions, political interest and interpersonal trust are controls in the 2023 surveys but serve as the competing predictors in the 2008 analysis, which is why they appear under different roles across datasets.

Robustness and sensitivity

We apply the same set of checks to all datasets. Table C3 states, for each check, the worry it addresses, the result that would have counted against the claim, and the result we observed.

Trust and democratic satisfaction in 2023. Trust explains variance the controls do not. Adding political trust and trust in police raises marginal R^2 from 0.190 to 0.269 in Latinobarómetro and from 0.084 to 0.208 in LAPOP. When each country in turn is held out and predicted from the other fifteen, out- of- sample R^2 rises from 0.195 to 0.287 in Latinobarómetro and from 0.099 to 0.227 in LAPOP. Fit improves in every held-out country in both surveys.

The trust coefficients also remain after we adjust for the strongest competing predictors each questionnaire offers. Table C10 adds them in sequence. Political trust falls from 0.268 to 0.219 in Latinobarómetro once winner–loser status, perceptions of corruption and presidential approval enter, and from 0.337 to 0.314 in LAPOP once perceptions of corruption enter. The political-trust coefficient remains larger than that of every competing measure of how people rate the regime, including presidential approval, which enters at 0.147 in Latinobarómetro (Panel B). It is not the largest coefficient in the model: economic perceptions enters at 0.236 and so exceeds political trust’s adjusted 0.219 in Latinobarómetro, but we treat economic perceptions as a control rather than as a competitor. The police-trust coefficient holds steady or grows slightly as these predictors are added.

Four different ways of computing the estimates give the same answer: survey-weighted maximum likelihood, design-based estimation using each survey’s own sampling units and strata, a bootstrap that resamples whole countries, and fits with and without survey weights. In the 2008 analysis, for example, the victimization coefficient on political trust is -0.211 (standard error 0.033) with weights and -0.205 (standard error 0.037) without. Table C8 tabulates the design-based comparison. The weighted and unweighted figures are not tabulated; both are given here.

Trust and satisfaction with democracy are related, but they are not the same thing. The correlation between them, corrected for measurement error, is 0.55 in Latinobarómetro and 0.52 in LAPOP. Political trust predicts satisfaction 3.3 times more strongly in Latinobarómetro, and 5.1 times more strongly in LAPOP, than it predicts an outcome with which it has no conceptual link. One test is less reassuring. A single unrotated component accounts for 52.2% of the variance in both surveys, just past the 50% mark conventionally treated as a warning, so the two constructs should not be called independent.

E-values for political trust are 2.04 in Latinobarómetro and 2.22 in LAPOP. Both exceed 1.35, which is the strongest association any measured covariate has on the same scale, so an unmeasured confounder would have to be stronger than anything we did measure. The corresponding values for police trust are 1.27 and 1.25. The first falls *below* the 1.35 benchmark: for police trust in Latinobarómetro, a confounder no stronger than a variable already in the model could account for the estimate. This is the weakest result in the study. It is why the substantive claims rest on political trust, and on the indirect role of trust in police, rather than on the direct police-trust coefficient.

A limit that no check addresses. The two analyses are fifteen years apart and come from different surveys. The link between police abuse and trust comes from the 2008 wave. The link between trust and satisfaction with democracy comes from the 2023 surveys. It is tempting to read them as two halves of one pathway — abuse erodes trust, and lost trust depresses satisfaction — but the data cannot support that reading. No wave measures both links at once, and the years

between them cover the end of the commodity boom and a broad regional fall in satisfaction with democracy. We therefore report the two associations separately, and never multiply them into a single chained estimate. The 2008 results show that a pathway of this shape is plausible. They do not measure a link of it.

Reported police abuse in 2008. Self-reported physical mistreatment by the police is negatively associated with political trust and with trust in police. The associations hold under design-based inference and under controls for the strongest competing predictors. They hold when police trust enters as ordered categories rather than as a continuous scale: an ordered model on the seven-point item gives -0.809 on the log-odds scale (standard error 0.063), matching the sign and significance of the linear estimate. They hold under three ways of handling missing data: dropping incomplete cases, weighting respondents by how likely they were to answer, and multiple imputation.

The main analysis pools the countries, so we ask separately here how far the association differs between them. Figure C7 estimates each country's effect first and then combines the sixteen, treating the true country effects as draws from a common distribution, $\hat{\beta}_c \sim \mathcal{N}(\mu, \tau^2)$, fitted by restricted maximum likelihood. This is the two-stage counterpart of letting the slope vary by country inside a single model, and it agrees with that approach to within $.003$ on every outcome. Countries differ hardly at all on political trust ($\tau = .034$; $I^2 = 8\%$; $Q = 16.5$ on 15 degrees of freedom, $p = .35$), moderately on trust in police ($\tau = .115$; $I^2 = 48\%$; $p = .016$), and somewhere in between on satisfaction with democracy ($\tau = .072$; $I^2 = 26\%$; $p = .098$).

A further country's true effect should fall between $-.30$ and $-.12$ for political trust, and between $-.64$ and $-.16$ for trust in police, both entirely negative, but between $-.28$ and $+.04$ for satisfaction with democracy, which includes zero.

Two further checks qualify this. When every control coefficient varies by country, not just the coefficient of interest, τ falls to roughly zero on all three outcomes ($.000$ for political trust, $.000$ for trust in police, $.000$ for satisfaction with democracy), so the detected differences disappear under that freer specification. Multiple imputation has the opposite effect. It keeps the third of respondents lost when incomplete cases are dropped — a loss concentrated in ideology and income. It *raises* the differences between countries on all three outcomes ($\tau = .076$, $.145$ and $.080$), moving political trust to $p = .068$ (Figure C8). Dropping incomplete cases was therefore hiding differences between countries rather than creating them, and the case for one common political-trust effect is weaker than the complete-case estimate alone suggests.

Reported police abuse also differs from ordinary crime victimization. The two measures correlate at 0.15 , and the police-abuse coefficient is two to three times larger on both trust outcomes. Police abuse predicts trust in other people three-and-a-half to nearly six times more weakly than it predicts trust in institutions. An unmeasured confounder would need associations with both police abuse and trust of at least 1.79 for political trust and 2.39 for trust in police, on the risk-ratio scale. Both exceed the strongest benchmark we can measure, ordinary crime victimization, whose own association is 1.14 on that scale. We report benchmarks as $\max(RR, 1/RR)$, so that an association running in the protective direction is compared with the E-value on the same side of one. The association with trust in police differs appreciably between countries. The association with political trust differs little, and is negative in all sixteen when incomplete cases are dropped, subject to the qualifications above.

The tables in this section. Table C3 is described at the top of this section. Table C4 gives the principal coefficient for every association in all three analyses, and is the reference point for every number that follows. Table C5 reports how much of each outcome the multilevel

models attribute to differences between countries, including for the 2008 models whose main-text estimates pool the countries instead. Table C6 gives the decomposition described above. Each remaining table takes one concern. Table C7 asks whether the predictors of interest add anything beyond the controls. Table C8 recomputes the standard errors three ways, to check that they do not depend on the model's own assumptions about variance. Table C9 asks whether the constructs are genuinely distinct rather than restatements of one another. Table C10 adds the strongest competing predictors each questionnaire offers. Table C11 varies both the measure and the treatment of missing data. Table C12 holds out each country in turn and tests prediction on the country left out. Table C13 asks how strong an unmeasured confounder would have to be to account for the estimates.

Table C3: Roadmap: what each check addresses, what would have overturned it, and the result.

Check	Concern addressed	What would have falsified it	Result
Incremental validity	The coefficient is real but adds nothing beyond the controls already in the model.	Marginal R-squared rising by less than about .01 when the terms of interest enter, or a likelihood-ratio test that does not reject.	Mixed by analysis. Substantial in the 2023 trust-satisfaction models (marginal R-squared rises by .079 and .124). Small in the 2008 police-abuse models (.002 and .009), significant but close to the threshold set above.
Design-based inference	Standard errors ignore the survey design (clustering, stratification, weights), or rest on too few country clusters.	A coefficient of interest moving by more than about a fifth of itself, or crossing conventional significance, under design-based estimation or a bootstrap that resamples countries.	Passed: three methods agree; weighted and unweighted coincide.
Discriminant validity	Trust and satisfaction with democracy are the same thing measured twice, so the association is true by definition or an artifact of asking both in one questionnaire.	A correlation above about .9 once measurement error is corrected, or a comparison outcome responding at least half as strongly as the outcome of interest.	Passed, with one qualification: the constructs are related but distinct, and the first-component share sits near its threshold.
Competing predictors	Something we left out, already known to shape how people rate the regime, drives the association.	The coefficient of interest collapsing once competing predictors enter.	Passed: the coefficients shrink modestly but hold up.
Measurement & functional form	The result comes from treating an ordered item as a continuous scale, or from a noisy multi-item index.	An ordered model, or an index built differently, reversing the effect or wiping it out.	Passed: the ordered model reproduces the result, and the indices are reliable.

Continued on next page.

Table C3 continued from previous page.

Check	Concern addressed	What would have falsified it	Result
Differential attrition	Respondents who are missing a model variable differ from those who answer, which biases the estimate once they are dropped.	Estimates shifting once non-response is modelled and corrected.	Passed: dropping incomplete cases, reweighting, and imputing all give the same answer.
Out-of-sample prediction	The association is overfitting and would not hold in a country the model never saw.	No overall gain out of sample, or gains in fewer than half the countries held out.	Mixed by analysis. Pooled gains and improvement in every held-out country in the 2023 analyses. In the 2008 analysis the overall gain is real but small, and prediction improves in 9 of 16 countries for political trust and 13 of 16 for trust in police. At the political-trust margin that is close to a coin flip.
Unmeasured confounding	An unobserved third variable explains the whole association.	An E-value below the association of the strongest covariate we did measure.	Passed for political trust. Trust in police does not clear its benchmark.
Mediation structure	Trust in police and political trust are not side-by-side predictors: one comes before the other, so adjusting for one understates the role of the other.	Both orderings implying that trust in police plays no part.	Trust in police plays a part under either ordering: mostly indirectly under Ordering A (67-79% mediated), and with a small direct component under Ordering B (5.8% mediated in both 2023 surveys). In LAPOP 2008, where victimization is observed and the chain runs from its origin, it plays a part under either ordering there too. See the mediation table in this section.

Continued on next page.

Table C3 continued from previous page.

Check	Concern addressed	What would have falsified it	Result
Between-country heterogeneity	A pooled coefficient averages countries that in fact differ, so one number describes none of them well and the shared sign is an artifact of pooling.	Countries differing so much that the pooled estimate describes none of them well, or country effects of mixed sign.	Passed on sign, qualified on size. Country effects are negative in 16 of 16 on both trust outcomes, and how far countries differ is modest and depends on the specification.
Missing-data composition	Dropping incomplete cases discards a third of the 2008 sample. Because non-response differs by country, it changes each country's sample by a different amount, which is a problem specific to a cross-national design.	Imputed estimates parting company with complete-case estimates, or the country pattern changing once the discarded third is recovered.	Passed on the coefficients, which barely move. Imputation does raise the differences between countries, so dropping incomplete cases was hiding them rather than creating them.
Temporal comparability	The police-abuse association is measured in 2008 and the trust-to-satisfaction association in 2023, in different surveys fifteen years apart. Chaining them would assume the second link already held in the first period, across an interval covering the end of the commodity boom and a regional fall in satisfaction with democracy.	The chain estimated inside a single wave, where every link is measured on the same respondents, disagreeing with the account assembled from the two links separately.	Passed for the chain itself. LAPOP 2008 measures victimization, both trust measures and satisfaction on the same respondents, so it is estimated inside one wave rather than assembled across two (see the mediation table's 2008 panel). No 2008 link is multiplied by a 2023 link. What no survey here can establish is whether the 2008 structure still held in 2023.

Continued on next page.

Table C3 continued from previous page.

Check	Concern addressed	What would have falsified it	Result
Estimator choice	Every estimate treats country as a random effect. If that choice is doing the work, treating country as a fixed effect instead would give a different answer.	A coefficient moving enough under country fixed effects to change its sign or its standing against zero.	Passed: refitting all 17 coefficients with country dummies and CR2 clustered standard errors moves none of them by more than 0.008 standard deviations, and none changes sign or its standing against zero. The clustered errors come out slightly smaller than the multilevel ones, so the reported intervals are not the narrower pair.

We run each check independently on all three analyses, and the tables that follow give the details. These checks test whether the descriptive findings hold up. None of them, alone or together, turns an observational design into a causal one.

Table C4: Principal coefficients across all three analyses.

Predictor (X)	Outcome (Y)	Estimate	SE	95% CI	N	Countries
<i>Latinobarómetro 2023</i>						
Political trust	Satisfaction with democracy	0.269	0.017	[0.236, 0.302]	13275	16
Police trust	Satisfaction with democracy	0.040	0.009	[0.024, 0.057]	13275	16
<i>LAPOP 2023</i>						
Political trust	Satisfaction with democracy	0.334	0.014	[0.306, 0.362]	19279	16
Police trust	Satisfaction with democracy	0.039	0.015	[0.010, 0.068]	19279	16
<i>LAPOP 2008</i>						
Police abuse victimization	Political trust	-0.205	0.032	[-0.267, -0.143]	17703	16
Police abuse victimization	Police trust	-0.402	0.042	[-0.485, -0.319]	17703	16

The 2023 rows are linear multilevel models that give each country its own intercept and its own slope on the predictors of interest. The 2008 rows are pooled regressions with a separate intercept for each country and standard errors clustered by country, using the CR2 correction and Satterthwaite degrees of freedom; their multilevel counterparts appear later in this appendix. We apply survey weights throughout and standardize every outcome within its dataset, so estimates are in outcome standard deviations. Intervals use the critical value each model's own degrees of freedom warrant. In the trust-satisfaction rows, political trust or trust in police is the predictor and satisfaction with democracy is the outcome. In the 2008 rows, reported police abuse is the predictor and trust is the outcome. "Countries" counts the countries in the estimation sample; all three analyses use the same sixteen. These are the reference values for the study. The same associations appear in later tables at slightly different numbers, for three reasons. Some later tables fit a multilevel model where this one pools the countries. Each check drops respondents missing the variables it needs, so it runs on its own sample, which its own note states. And some tables add predictors that are absent here. For reported police abuse on trust in police, for example, this table gives -0.402, the specification table gives -0.412 on a smaller shared sample under the multilevel fit, and the discriminant table gives -0.386 with ordinary crime victimization entered alongside. The spread across those variants is smaller than any of their standard errors.

Table C5: How much the multilevel models attribute to differences between countries.

Dataset	Model	N	Countries	SD(intercept)	SD(slopes)	SD(residual)	ICC	Form	Singular
Latinobarómetro 2023	Trust to satisfaction	13,275	16	0.206	0.008; 0.054	0.815	0.060	correlated	no
LAPOP 2023	Trust to satisfaction	19,279	16	0.150	0.050; 0.046	0.848	0.030	correlated	no
LAPOP 2008	Trust to satisfaction	17,337	16	0.246	0.071; 0.069	0.832	0.080	correlated	no
LAPOP 2008	Victimization to political trust	17,703	16	0.381	0.033	0.865	0.162	correlated	no
LAPOP 2008	Victimization to trust in police	17,703	16	0.303	0.113	0.873	0.107	correlated	no
LAPOP 2008	Victimization to satisfaction	17,337	16	0.320	0.069	0.877	0.117	correlated	no

Country-level standard deviations from the multilevel fits, on the same samples as the corresponding coefficient tables. SD(slopes) lists one value per slope that varies by country, in model order: the trust-to-satisfaction models have two (police trust, political trust) and the police-abuse models have one. Form records whether we estimated the covariance between the country-level terms freely or held it to zero, which this package does only when the free form breaks down. Singular marks a fit whose country-level variances sit at the edge of what can be estimated; none of the models does so on the sixteen matched countries reported here. For 2008 these models are the counterpart of the pooled estimates in the main text, and the two-stage figures in the country-level section below examine the differences between countries directly.

The mediation decomposition. Several tables and figures below split an association into a part that runs through an intermediate measure and a part that does not. Write the association of interest as running from a predictor X to satisfaction with democracy Y through an intermediate trust measure M . The first link is the association of X with M . The second link is the association of M with Y , holding X fixed. Multiplying the two gives the indirect part. Whatever is left of the X – Y association is the direct part. The proportion mediated is the indirect part divided by the total.

Both 2023 surveys measure the two trust constructs at the same moment, and cross-sectional data cannot say which comes first. We fit both orderings rather than assuming one. **Ordering A** treats trust in police as the predictor and institutional trust as the intermediate: police, then institutional, then satisfaction. **Ordering B** reverses them: institutional, then police, then satisfaction. A conclusion that holds under both orderings does not depend on that assumption. A conclusion that flips between them does, and we say so where it happens.

Table C6 reports the decomposition, and the pattern repeats across the two surveys. Under Ordering A, most of the association between trust in police and satisfaction runs through institutional trust: 67.2% in Latinobarómetro and 79.2% in LAPOP. Under Ordering B, almost none of institutional trust’s association runs through trust in police: 5.8% in both. The direct association belongs to institutional trust, and trust in police works mostly through it.

LAPOP 2008 supplies what neither 2023 survey can. It records the respondent’s own experience of police abuse alongside both trust measures and satisfaction, so the chain is estimated from its origin within one wave rather than assembled from links measured fifteen years apart. Police abuse is fixed as the predictor there, since neither trust nor satisfaction can produce a past assault, and the two orderings concern only the trust measures between it and satisfaction. Each ordering is then a chain of two intermediates rather than one. The indirect part is everything that reaches satisfaction through whichever trust measure the ordering places first, and the direct part is the remainder. Because the predictor is the same under both, the two orderings decompose a single total: 63.6% of the association between police abuse and satisfaction runs through trust in police under Ordering A, and 51.4% through institutional trust under Ordering B. Trust in police is implicated under both, as in the 2023 surveys, but here the two orderings disagree

about the path rather than about the magnitude. One consequence needs careful reading: that panel's proportion-mediated interval reaches past 100%. Its denominator is about a tenth of the trust associations above it, and the direct part's own interval crosses zero. A resample that moves the direct part across zero therefore leaves a total smaller than the part running through the mediators. We report the interval as estimated rather than truncating it.

Table C6: The mediation decomposition, both orderings, all three datasets.

Ordering	Direct	Indirect	Total	% mediated	E-value	Benchmark RR
<i>LAPOP 2008</i>						
A: abuse → police → institutional → satisfaction	-0.044	-0.077**	-0.121**	63.6%	1.292	1.061
B: abuse → institutional → police → satisfaction	-0.059	-0.062**	-0.121**	51.4%	1.292	1.061
<i>LAPOP 2023</i>						
A: police → institutional → satisfaction	0.039**	0.148**	0.187**	79.2%	2.216	1.134
B: institutional → police → satisfaction	0.334**	0.020**	0.354**	5.8%	1.254	1.134
<i>Latinobarómetro 2023</i>						
A: police → institutional → satisfaction	0.040**	0.083**	0.123**	67.2%	2.037	1.346
B: institutional → police → satisfaction	0.269**	0.017**	0.285**	5.8%	1.266	1.346

Note: 95% CIs from a 1,000-resample bootstrap over countries. Stars: *** $p < .001$, ** $p < .01$, * $p < .05$. Each country's own indirect effect is negative – the direction the pathway predicts, since LAPOP 2008's exposure is police abuse rather than trust – in 16/16 (A) and 16/16 (B) in LAPOP 2008; positive in 16/16 (A) and 14/16 (B) in LAPOP 2023, and 16/16 (A) and 15/16 (B) in Latinobarómetro. E-value is the smaller of the two links', since a mediated effect resists confounding only as well as its weaker link. Benchmark RR is the strongest association any measured covariate has with the outcome, as $\max(RR, 1/RR)$; an E-value above it means the confounder would have to beat every covariate we measured.

Table C7: What the trust and police-abuse terms add beyond the controls.

Dataset / outcome	ΔR^2 (marginal)	R^2 (marginal)	LRT χ^2	df	p	Top predictor (Shapley share)
Latinobarómetro 2023 (joint model)	0.079	0.269	985.365	2	<.001	Economic perceptions (0.101)
LAPOP 2023 (joint model)	0.124	0.208	2165.720	2	<.001	Political trust (0.103)
LAPOP 2008 → political trust	0.002	0.053	45.373	1	<.001	Ideology (0.021)
LAPOP 2008 → police trust	0.009	0.063	170.928	1	<.001	Ideology (0.015)

The likelihood-ratio test compares models fitted by maximum likelihood rather than REML, because a test on these terms requires it; the full-estimate tables elsewhere report REML fits of the same specifications. Rows marked joint test both trust terms at once, so one increment and one test statistic cover both trust terms of that dataset. The Shapley share divides the model's explained variance among all the predictors; the column names only the largest. Marginal R-squared is the share of the outcome the predictors explain on their own, and it is the denominator the increment should be read against: the same increment means one thing against a model that explains a fifth of the variance and another against one that explains a twentieth. We do not use the conditional version, which also credits the differences between country means, because those differences would inflate it without saying anything about the predictors.

Table C8: Three alternatives to the model's own standard errors.

Predictor (X)	Outcome (Y)	WeMix		svyglm		Wild-cluster p
		est.	SE	est.	SE	
<i>Latinobarómetro</i>						
Political trust	Satisf. w/ democracy	0.271	0.015	0.269	0.011	<.001
Police trust	Satisf. w/ democracy	0.042	0.008	0.042	0.010	<.001
<i>LAPOP 2023</i>						
Political trust	Satisf. w/ democracy	0.337	0.016	0.336	0.009	<.001
Police trust	Satisf. w/ democracy	0.037	0.014	0.037	0.008	0.009
<i>LAPOP 2008</i>						
Police abuse victimization	Political trust	-0.211	0.033	-0.211	0.034	<.001
Police abuse victimization	Police trust	-0.383	0.035	-0.384	0.033	<.001

Each column re-estimates the same coefficient under a different treatment of the survey design. WeMix applies survey weights by pseudo-maximum-likelihood inside the multilevel model. The svyglm column uses each survey's own sampling units nested within country strata (LAPOP 2023 and 2008 use their own sampling-unit and stratum variables; Latinobarómetro uses its city variable). The wild cluster bootstrap resamples whole countries, which is the recommended approach when there are few clusters, as here. What matters is whether all three agree, and whether the weighted and unweighted fits agree.

Table C9: Are the constructs distinct? Tests for institutional trust and reported police abuse

Dataset / context	Diagnostic	Estimate
<i>Panel A: Trust versus satisfaction</i>		
Latinobarómetro 2023	Correlation, corrected for measurement error: satisfaction × political trust	0.551
Latinobarómetro 2023	Ratio: political trust → satisfaction, against → trust in other people	3.252×
Latinobarómetro 2023	Share of variance in the first component	52.2%
LAPOP 2023	Correlation, corrected for measurement error: satisfaction × political trust	0.517
LAPOP 2023	Ratio: political trust → satisfaction, against → trust in other people	5.103×
LAPOP 2023	Share of variance in the first component	52.2%
<i>Panel B: Police abuse versus ordinary crime</i>		
LAPOP 2008	Correlation: police abuse with ordinary crime victimization	0.155
LAPOP 2008	Police abuse → political trust	-0.204
LAPOP 2008	Ordinary crime victimization → political trust	-0.097
LAPOP 2008	Police abuse → trust in police	-0.386
LAPOP 2008	Ordinary crime victimization → trust in police	-0.117
LAPOP 2008	Police abuse → trust in other people (comparison outcome)	-0.055
LAPOP 2008	Correlation, corrected for measurement error: political trust with trust in police	0.607

Panel A asks whether institutional trust is distinct from satisfaction with democracy. The ratio compares political trust's coefficient on satisfaction against its coefficient on trust in other people. A first component above 50% signals that answers may be running together because one questionnaire collected them all, so we treat these constructs as related rather than independent. Panel B asks whether reported police abuse is distinct from ordinary crime victimization, and whether its associations are specific to trust in institutions. The police-abuse and crime coefficients come from a single model containing both, and the last row reports the correlation between political trust and trust in police once measurement error is corrected. Trust in other people serves as the comparison outcome because the argument under test concerns confidence in institutions and gives no reason for confidence in other people to move with it. The two are not unrelated: they correlate, and trust in other people enters the main models as a control. The test is therefore not that the comparison association is zero, but that it is much weaker than the association of interest, which is what the ratios report.

Table C10: Extended specifications: the coefficients of interest and the predictors that compete with them

Dataset	Outcome (Y)	Specification / competing predictor	Estimate (SE)	95% CI
<i>Panel A: Coefficients of interest across every specification</i>				
Latinobarómetro	Satisf. w/ democracy	Political trust — (1) Base controls; ΔR^2 trust block = 0.078	0.268*** (0.017)	[0.236, 0.301]
Latinobarómetro	Satisf. w/ democracy	Political trust — (2) + winner/loser, corruption; ΔR^2 trust block = 0.072	0.243*** (0.016)	[0.210, 0.275]
Latinobarómetro	Satisf. w/ democracy	Political trust — (3) + presidential approval; ΔR^2 trust block = 0.058	0.219*** (0.017)	[0.186, 0.252]
Latinobarómetro	Satisf. w/ democracy	Trust in police — (1) Base controls; ΔR^2 trust block = 0.078	0.039*** (0.009)	[0.021, 0.056]
Latinobarómetro	Satisf. w/ democracy	Trust in police — (2) + winner/loser, corruption; ΔR^2 trust block = 0.072	0.039*** (0.009)	[0.023, 0.056]
Latinobarómetro	Satisf. w/ democracy	Trust in police — (3) + presidential approval; ΔR^2 trust block = 0.058	0.043*** (0.009)	[0.026, 0.060]
LAPOP 2023	Satisf. w/ democracy	Political trust — (1) Base controls; ΔR^2 trust block = 0.128	0.337*** (0.014)	[0.310, 0.363]

Continued on next page.

Table C10 continued from previous page.

Dataset	Outcome (Y)	Specification / competing predictor	Estimate (SE)	95% CI
LAPOP 2023	Satisf. w/ democracy	Political trust — (2) + corruption perception; ΔR^2 trust block = 0.103	0.314*** (0.013)	[0.288, 0.340]
LAPOP 2023	Satisf. w/ democracy	Trust in police — (1) Base controls; ΔR^2 trust block = 0.128	0.040** (0.015)	[0.011, 0.069]
LAPOP 2023	Satisf. w/ democracy	Trust in police — (2) + corruption perception; ΔR^2 trust block = 0.103	0.036* (0.015)	[0.008, 0.065]
LAPOP 2008	Political trust	Police abuse victimization — (1) Base controls; ΔR^2 trust block = 0.002	-0.204*** (0.031)	[-0.265, -0.144]
LAPOP 2008	Political trust	Police abuse victimization — (2) + econ., pol. interest, social trust; ΔR^2 trust block = 0.001	-0.178*** (0.030)	[-0.236, -0.120]
LAPOP 2008	Trust in police	Police abuse victimization — (1) Base controls; ΔR^2 trust block = 0.010	-0.411*** (0.031)	[-0.472, -0.350]
LAPOP 2008	Trust in police	Police abuse victimization — (2) + econ., pol. interest, social trust; ΔR^2 trust block = 0.008	-0.390*** (0.031)	[-0.450, -0.330]

Continued on next page.

Table C10 continued from previous page.

Dataset	Outcome (Y)	Specification / competing predictor	Estimate (SE)	95% CI
<i>Panel B: How large the competing predictors are in the fullest specification</i>				
Latinobarómetro	Satisf. w/ democracy	Perceived progress reducing corruption (higher = more progress)	0.052*** (0.008)	[0.036, 0.068]
Latinobarómetro	Satisf. w/ democracy	Voted for government (vs. neither)	0.130*** (0.021)	[0.088, 0.171]
Latinobarómetro	Satisf. w/ democracy	Voted for opposition (vs. neither)	0.021 (0.020)	[-0.018, 0.061]
Latinobarómetro	Satisf. w/ democracy	Approves of president	0.147*** (0.019)	[0.109, 0.185]
Latinobarómetro	Satisf. w/ democracy	Economic perceptions (control)	0.236*** (0.009)	[0.218, 0.254]
LAPOP 2023	Satisf. w/ democracy	Perceived corruption among politicians (higher = more corruption)	-0.079*** (0.007)	[-0.093, -0.065]
LAPOP 2023	Satisf. w/ democracy	Economic perceptions (control)	0.112*** (0.007)	[0.098, 0.126]
LAPOP 2008	Political trust	Economic perceptions	0.269*** (0.007)	[0.255, 0.284]
LAPOP 2008	Political trust	Political interest	0.078*** (0.007)	[0.064, 0.092]
LAPOP 2008	Political trust	Interpersonal trust	0.085*** (0.007)	[0.071, 0.100]
LAPOP 2008	Trust in police	Economic perceptions	0.153*** (0.008)	[0.137, 0.168]
LAPOP 2008	Trust in police	Political interest	0.022** (0.007)	[0.007, 0.036]

LAPOP 2008	Trust in police	Interpersonal trust	0.079*** (0.008)	[0.064, 0.094]
------------	-----------------	---------------------	---------------------	----------------

Panel A keeps every specification. Each row adds the named predictors to the row above it, on one shared sample within each dataset. The trust block's added marginal R-squared appears in the specification label because it is the same for both rows of a specification. Panel B gives the competing predictors' own coefficients in the fullest specification. Sample sizes hold constant within each dataset (Latinobarómetro N=12,750; LAPOP 2023 N=18,977; LAPOP 2008 N=16,928), so the body omits them. Each is smaller than that dataset's main estimation sample, because the competing predictors have missing answers of their own. Stars: *** p<.001, ** p<.01, * p<.05. The 2023 surveys score their corruption items in opposite directions; the labels make that explicit.

Table C11: Measurement validity and missing-data sensitivity

Panel A: Scale reliability and agreement with LAPOP's published scales

Source / dataset	Scale	Items	Std. α	N
Study 1 · Brazil experiment	Three-item scale	3	0.809	1,248
Study 2 · ELSOC Chile	Three-item scale	3	0.720	6,270
Latinobarómetro 2023	Three-item scale	3	0.756	17,159
Latinobarómetro 2023	LAPOP official (4 items)	4	0.797	16,861
Latinobarómetro 2023	LAPOP official (5 items)	5	0.851	16,828
LAPOP 2023	Three-item scale	3	0.802	24,019
LAPOP 2023	LAPOP official (4 items)	4	0.832	23,919
LAPOP 2023	LAPOP official (5 items)	5	0.839	23,823
LAPOP 2008	Three-item scale	3	0.715	24,892
LAPOP 2008	LAPOP official (4 items)	4	0.770	24,534
LAPOP 2008	LAPOP official (5 items)	5	0.832	24,429

Panel B: Coefficient stability across missing-data treatments

Dataset	Predictor → outcome	Std. α	ω	Complete cases β	Reweighting β	Imputed β
Latinobarómetro 2023	Political trust → Satisfaction with democracy	0.756	0.761	0.263	0.271	0.263
Latinobarómetro 2023	Police trust → Satisfaction with democracy	—	—	0.048	0.041	0.039
LAPOP 2023	Political trust → Satisfaction with democracy	0.802	0.809	0.333	0.332	0.334
LAPOP 2023	Police trust → Satisfaction with democracy	—	—	0.038	0.035	0.032
LAPOP 2008	Police abuse → Political trust	0.715	0.723	-0.199	-0.207	-0.177
LAPOP 2008	Police abuse → Police trust	—	—	-0.393	-0.396	-0.361

Panel A reports standardised Cronbach's α for the three-item institutional-trust scale across all studies, and for LAPOP's published four- and five-item scales where they were fielded. Within-country correlations between the three-item scale and the published LAPOP scales exceed .90 in all 16 matched countries of both 2023 surveys (lowest $r = .922$ against the four-item scale, .905 against the five-item). Panel B reports how far the coefficient of interest moves across three treatments of missing data: dropping incomplete cases, reweighting respondents by how likely they were to answer, and multiple imputation. α and ω apply to the multi-item political trust index; trust in police is a single item, so measures of internal consistency do not apply.

Table C12: Holding out one country at a time: how much trust and reported police abuse improve prediction.

Dataset / outcome	Out-of-sample R ² without	Out-of-sample R ² with	Gain	Countries improved
Latinobarómetro	0.195	0.287	0.092	16/16
LAPOP 2023	0.099	0.227	0.128	16/16
LAPOP 2008 → political trust	0.015	0.016	0.001	9/16
LAPOP 2008 → police trust	0.031	0.038	0.007	13/16

“Countries improved” counts the held-out countries where adding the predictors of interest lowers prediction error. We hold out each country in turn, refit the model on the other fifteen, and predict the held-out country without using anything specific to it, since nothing country-specific would be known for a genuinely new country. The country-level section below plots the results one country at a time.

Table C13: How strong an unmeasured confounder would have to be (main specification).

Predictor (X)	Outcome (Y)	β	E-value (point)	E-value (CI limit)	Benchmark RR (strongest measured covariate)
<i>Latinobarómetro</i>					
Political trust	Satisf. w/ democracy	0.269	2.037	1.928	1.346
Police trust	Satisf. w/ democracy	0.040	1.266	1.192	1.346
<i>LAPOP 2023</i>					
Political trust	Satisf. w/ democracy	0.334	2.216	2.123	1.134
Police trust	Satisf. w/ democracy	0.039	1.254	1.115	1.134
<i>LAPOP 2008</i>					
Police abuse victimization	Political trust	-0.205	1.787	1.593	1.142
Police abuse victimization	Police trust	-0.397	2.388	2.122	1.144

The E-value is the weakest association, on the risk-ratio scale, that an unmeasured confounder would need with both the predictor and the outcome to account for the estimate entirely. The benchmark is the strongest association any measured covariate has on that scale – economic perceptions in the 2023 analyses, ordinary crime victimization in the 2008 analysis – written as $\max(\text{RR}, 1/\text{RR})$ so that it sits on the same side of 1 as the E-value and the two can be compared. An E-value above its benchmark means the confounder would have to be stronger than anything we measured. An E-value below it means a confounder no stronger than a variable already in the model could suffice. One row falls below: police trust in Latinobarómetro, where the E-value of 1.27 sits under its benchmark of 1.35. That is the weakest result in the study. The political-trust rows, which carry the substantive claims, clear their benchmarks in every dataset. Adding the competing predictors barely moves these numbers: the political-trust E-value falls from 2.04 to 1.95 in Latinobarómetro and from 2.22 to 2.15 in LAPOP 2023.

Whether treating country as random matters. Every estimate above gives each country its own intercept and its own slope, drawn from a common distribution. The alternative is to give each country a dummy variable and cluster the standard errors on country. That asks the same question while assuming nothing about how countries are distributed. Table C14 refits every

model in the study the second way, on the same respondents. Nothing depends on the choice. Across all fitted coefficients the two estimators agree to within 0.008, and no coefficient changes sign or its standing against zero. The two mediation decompositions agree to within about one percentage point on the share of the association that runs through institutional trust. The clustered standard errors run slightly smaller than the multilevel ones, so the reported intervals are the wider of the two.

Table C14: Every estimate refit with country fixed effects instead of random effects.

Quantity	Multilevel (SE)	Country fixed effects (SE)	Difference	N
<i>Latinobarometro 2023</i>				
Trust in police → satisfaction	0.040 (0.009)	0.042 (0.008)	+0.001	13,275
Political trust → satisfaction	0.269 (0.017)	0.269 (0.015)	+0.001	13,275
<i>LAPOP 2023</i>				
Trust in police → satisfaction	0.039 (0.015)	0.037 (0.014)	-0.002	19,279
Political trust → satisfaction	0.334 (0.014)	0.336 (0.016)	+0.002	19,279
<i>LAPOP 2008</i>				
Police abuse → political trust	-0.205 (0.032)	-0.205 (0.028)	-0.000	17,703
Police abuse → trust in police	-0.402 (0.042)	-0.402 (0.036)	-0.001	17,703
Police abuse → satisfaction	-0.121 (0.036)	-0.122 (0.033)	-0.001	17,337
Trust in police → satisfaction	0.049 (0.019)	0.051 (0.019)	+0.003	17,337
Political trust → satisfaction	0.282 (0.020)	0.283 (0.020)	+0.001	17,337
<i>Latinobarometro 2023</i>				
Mediation, b-path: trust in police	0.040 (0.009)	0.042 (0.008)	+0.001	13,275
Mediation, b-path: political trust	0.269 (0.017)	0.269 (0.015)	+0.001	13,275
Mediation, a-path: police → political trust	0.309 (0.030)	0.301 (0.030)	-0.008	13,275
Mediation, a-path: political → police trust	0.408 (0.038)	0.407 (0.030)	-0.001	13,275
Proportion mediated, ordering A	0.672	0.661	-0.011	13,275
Proportion mediated, ordering B	0.058	0.059	+0.001	13,275
<i>LAPOP 2023</i>				
Mediation, b-path: trust in police	0.039 (0.014)	0.037 (0.014)	-0.002	19,279
Mediation, b-path: political trust	0.334 (0.013)	0.336 (0.016)	+0.002	19,279
Mediation, a-path: police → political trust	0.444 (0.025)	0.446 (0.022)	+0.002	19,279
Mediation, a-path: political → police trust	0.525 (0.025)	0.526 (0.026)	+0.000	19,279
Proportion mediated, ordering A	0.793	0.802	+0.009	19,279

Proportion mediated, ordering	0.057	0.055	-0.002	19,279
B				

Each row is one quantity estimated twice on the identical sample: the multilevel form the study reports, which gives each country its own intercept and its own slope, and a pooled regression with a dummy for each country and standard errors clustered on country using the CR2 correction. Sixteen countries is few enough that the ordinary cluster-robust standard error can come out too small, which is why CR2 rather than CR1. Across all 17 fitted coefficients the two estimators differ by at most 0.008 standard deviations, and no coefficient changes sign or its standing against zero. The two proportions mediated are products of coefficients rather than fitted terms, so they are reported as point estimates: their confidence intervals come from a bootstrap over countries that we do not re-run under the second estimator.

Whether the controls’ functional form matters. The 2008 models adjust for the controls by entering them in the regression, linearly and additively. The specification imposes that restriction rather than testing it. The alternative is to adjust by reweighting: make the covariate distributions the same for respondents who do and do not report police abuse, and only then compare their outcomes. Table C15 does that seven ways: inverse-probability weighting on a logistic and on a post-double-selection lasso propensity, covariate balancing propensity scores, entropy balancing, a gradient-boosted propensity, and augmented and double-machine-learning estimators. The target is the average among the victimized over the region of common propensity support, and overlap weights are reported beside it. Weighting is a stronger check than a further regression specification because the adjustment can be inspected before any outcome is examined. Figure C1 reports the covariate differences before and after weighting, and no estimate enters the table until its specification clears overlap, positivity, balance, effective sample size and weight concentration. Balance on the higher-order terms is judged only on those that predict the outcome. That restriction makes the gate a test of confounding rather than of fit: a term that shifts who is victimized but bears no relation to trust cannot bias the comparison. Forcing the weights to balance such a term would lose precision and remove no bias. Roughly four in five of the interaction terms fall into that category here, and the screen that identifies them is the same one the post-double-selection propensity uses. Sixty-seven of seventy-five specifications clear the gate. The eight that fail are all one scheme under one target — the post-double-selection propensity weighted to the victimized, which balances worse than the simplest scheme rather than better. The same propensity weighted for overlap clears every time. They are listed rather than quietly dropped. Every scheme that clears the gate returns the same answer as the model reported in the main text, so those results do not depend on how the control set is specified. A second check asks the same question from the other side, and tests it rather than assuming it. Some of the expanded terms predict who reports police abuse while bearing no relation to trust — the terms the outcome screen discards. Such a term is a non-confounder by construction, so adding it to the adjustment should change the estimate by nothing beyond sampling noise. If it does change the estimate, the simpler specification was using the term as a proxy for something it had failed to capture. We add those terms and bootstrap the difference by country: no estimate moves detectably, jointly or for any single outcome ($p = 0.65$ across the three). The reported models may miss other things; the way they enter the controls is not one of them.

This check establishes one thing only: weighting balances what was measured. It addresses neither reverse causality — treatment and outcome are recorded in the same cross-section — nor confounding by variables the 2008 wave did not measure. The E-values in Table C13 bound the second of those and are not superseded here.

Table C15: Adjusted differences under eight weighting schemes, LAPOP 2008.

Outcome	Estimator	Difference	SE	N
Political trust	IPW logistic	-0.198	(0.035)	17,749
Political trust	IPW logistic, trim=crump	-0.198	(0.035)	17,749
Political trust	IPW logistic, trim=pct1	-0.217	(0.035)	17,393
Political trust	IPW post-DS, trim=pct1	-0.199	(0.036)	17,394
Political trust	CBPS	-0.206	(0.035)	17,748
Political trust	CBPS, trim=crump	-0.206	(0.035)	17,748
Political trust	CBPS, trim=pct1	-0.222	(0.035)	17,394
Political trust	Entropy bal.	-0.206	(0.035)	17,750
Political trust	Boosted PS	-0.201	(0.031)	17,733
Political trust	Boosted PS, trim=crump	-0.201	(0.031)	17,733
Political trust	Boosted PS, trim=pct1	-0.202	(0.031)	17,394
Political trust	IPW logistic, ATO	-0.207	(0.035)	17,749
Political trust	IPW logistic, ATO, trim=crump	-0.207	(0.035)	17,749
Political trust	IPW logistic, ATO, trim=pct1	-0.220	(0.036)	17,393
Political trust	IPW post-DS, ATO	-0.201	(0.033)	17,750
Political trust	IPW post-DS, ATO, trim=crump	-0.201	(0.033)	17,750
Political trust	IPW post-DS, ATO, trim=pct1	-0.208	(0.036)	17,394
Political trust	CBPS, ATO	-0.207	(0.035)	17,749
Political trust	CBPS, ATO, trim=crump	-0.207	(0.035)	17,749
Political trust	CBPS, ATO, trim=pct1	-0.220	(0.036)	17,393
Political trust	Boosted PS, ATO	-0.199	(0.028)	17,732
Political trust	Boosted PS, ATO, trim=crump	-0.199	(0.028)	17,732
Political trust	Boosted PS, ATO, trim=pct1	-0.197	(0.030)	17,394
Political trust	AIPW	-0.205	(0.030)	17,750
Political trust	DML cross-fitted	-0.215	(0.029)	17,750
Trust in police	IPW logistic	-0.385	(0.041)	18,456
Trust in police	IPW logistic, trim=crump	-0.385	(0.041)	18,456
Trust in police	IPW logistic, trim=pct1	-0.409	(0.042)	18,088
Trust in police	CBPS	-0.388	(0.041)	18,455
Trust in police	CBPS, trim=crump	-0.388	(0.041)	18,455
Trust in police	CBPS, trim=pct1	-0.413	(0.043)	18,087
Trust in police	Entropy bal.	-0.389	(0.042)	18,457
Trust in police	Boosted PS	-0.392	(0.038)	18,443
Trust in police	Boosted PS, trim=crump	-0.392	(0.038)	18,443
Trust in police	Boosted PS, trim=pct1	-0.390	(0.041)	18,083
Trust in police	IPW logistic, ATO	-0.399	(0.041)	18,456
Trust in police	IPW logistic, ATO, trim=crump	-0.399	(0.041)	18,456
Trust in police	IPW logistic, ATO, trim=pct1	-0.414	(0.042)	18,088
Trust in police	IPW post-DS, ATO	-0.400	(0.041)	18,457
Trust in police	IPW post-DS, ATO, trim=crump	-0.400	(0.041)	18,457
Trust in police	IPW post-DS, ATO, trim=pct1	-0.414	(0.043)	18,085
Trust in police	CBPS, ATO	-0.399	(0.041)	18,456
Trust in police	CBPS, ATO, trim=crump	-0.399	(0.041)	18,456
Trust in police	CBPS, ATO, trim=pct1	-0.414	(0.042)	18,088

Continued on next page.

Table C15 continued from previous page.

Outcome	Estimator	Difference	SE	N
Trust in police	Boosted PS, ATO	-0.385 (0.037)	18,436	
Trust in police	Boosted PS, ATO, trim=crump	-0.385 (0.037)	18,436	
Trust in police	Boosted PS, ATO, trim=pct1	-0.383 (0.038)	18,087	
Trust in police	AIPW	-0.378 (0.035)	18,457	
Trust in police	DML cross-fitted	-0.400 (0.035)	18,457	
Satisfaction w/ democracy	IPW logistic	-0.120 (0.040)	18,109	
Satisfaction w/ democracy	IPW logistic, trim=crump	-0.120 (0.040)	18,109	
Satisfaction w/ democracy	IPW logistic, trim=pct1	-0.114 (0.038)	17,745	
Satisfaction w/ democracy	CBPS	-0.124 (0.040)	18,108	
Satisfaction w/ democracy	CBPS, trim=crump	-0.124 (0.040)	18,108	
Satisfaction w/ democracy	CBPS, trim=pct1	-0.117 (0.038)	17,739	
Satisfaction w/ democracy	Entropy bal.	-0.123 (0.040)	18,109	
Satisfaction w/ democracy	Boosted PS	-0.119 (0.036)	18,093	
Satisfaction w/ democracy	Boosted PS, trim=crump	-0.119 (0.036)	18,093	
Satisfaction w/ democracy	Boosted PS, trim=pct1	-0.104 (0.040)	17,742	
Satisfaction w/ democracy	IPW logistic, ATO	-0.112 (0.039)	18,109	
Satisfaction w/ democracy	IPW logistic, ATO, trim=crump	-0.112 (0.039)	18,109	
Satisfaction w/ democracy	IPW logistic, ATO, trim=pct1	-0.110 (0.039)	17,745	
Satisfaction w/ democracy	IPW post-DS, ATO	-0.098 (0.041)	18,109	
Satisfaction w/ democracy	IPW post-DS, ATO, trim=crump	-0.098 (0.041)	18,109	
Satisfaction w/ democracy	IPW post-DS, ATO, trim=pct1	-0.107 (0.042)	17,745	
Satisfaction w/ democracy	CBPS, ATO	-0.112 (0.039)	18,109	
Satisfaction w/ democracy	CBPS, ATO, trim=crump	-0.112 (0.039)	18,109	
Satisfaction w/ democracy	CBPS, ATO, trim=pct1	-0.110 (0.039)	17,745	
Satisfaction w/ democracy	Boosted PS, ATO	-0.107 (0.036)	18,089	
Satisfaction w/ democracy	Boosted PS, ATO, trim=crump	-0.107 (0.036)	18,089	
Satisfaction w/ democracy	Boosted PS, ATO, trim=pct1	-0.097 (0.039)	17,744	
Satisfaction w/ democracy	AIPW	-0.123 (0.036)	18,109	
Satisfaction w/ democracy	DML cross-fitted	-0.112 (0.036)	18,109	
Political trust	Multilevel, reported	-0.205 (0.032)	17,703	
Trust in police	Multilevel, reported	-0.402 (0.042)	17,703	
Satisfaction w/ democracy	Multilevel, reported	-0.121 (0.036)	17,337	

Weighted differences in the police-abuse victimization contrast, LAPOP 2008, matched 16 countries. Each row re-estimates the same association under a different adjustment scheme. Standard errors are CR2 clustered on country with Satterthwaite degrees of freedom. 67 of 75 propensity specifications passed the diagnostic gate (overlap, positivity, main-effect balance below 0.10, effective sample size, weight concentration); specifications that failed are not estimated and are listed in r36_skipped.csv. These are adjusted differences, not causal effects: every scheme here assumes selection on observables, and neither reverse causality nor confounding by variables the 2008 wave does not measure is addressed by weighting.




Figure C1: Standardised covariate differences before and after weighting, LAPOP 2008. Main effects balance easily. The interactions are split by whether they predict the outcome: only those that do are held to the balance criterion (see text). The right-hand panel is subsampled for legibility.

Full model estimates

Tables C16–C21 report every coefficient in the equations behind the main-text results, one column per equation. The first two cover the 2023 surveys, one dataset each. The next three cover the 2008 police-abuse analysis, one outcome each. The last repeats the 2023 surveys’ trust-to-satisfaction specification on the 2008 wave, using the controls that wave itself offers rather than the ones the 2023 tables share. Each cell gives a coefficient with its standard error in parentheses, and an asterisk marks $p < .05$. We calculate significance using the Satterthwaite approximation rather than a normal approximation, because the predictors of interest have slopes that vary by country. That leaves them with far fewer effective degrees of freedom — closer to the number of countries than to the number of respondents.

Table C16: Full estimates, Latinobarómetro 2023: every outcome equation.

Term	Satisfaction w/ democracy	First link of the mediation		Satisfaction, extended controls
		Predicting institutional trust	Predicting trust in police	
Trust in police (z)	0.040 (0.009)*	0.309 (0.029)*		0.042 (0.008)*
Political trust (z)	0.269 (0.017)*		0.408 (0.038)*	0.243 (0.017)*
Economic perceptions	0.266 (0.009)*	0.357 (0.007)*	0.030 (0.009)*	0.248 (0.009)*
Political interest	0.042 (0.008)*	0.105 (0.007)*	0.005 (0.008)	0.030 (0.008)*
Interpersonal trust	0.058 (0.007)*	0.047 (0.006)*	0.032 (0.007)*	0.059 (0.007)*
Age (years)	0.001 (0.000)	-0.000 (0.000)	-0.001 (0.000)*	0.000 (0.000)
Female	-0.022 (0.014)	0.069 (0.013)*	0.042 (0.015)*	-0.026 (0.014)
Education level	-0.003 (0.002)	-0.004 (0.002)*	0.007 (0.002)*	-0.001 (0.002)
SES (interviewer rating)	0.024 (0.010)*	-0.008 (0.009)	0.018 (0.010)	0.023 (0.010)*
Asset index (z)	-0.002 (0.010)	-0.027 (0.009)*	0.039 (0.010)*	-0.001 (0.010)
Ideology (left-right)	0.009 (0.002)*	0.007 (0.002)*	0.024 (0.003)*	0.008 (0.002)*
Crime victimization	-0.055 (0.016)*	-0.051 (0.014)*	-0.053 (0.016)*	-0.053 (0.016)*
National capital	0.052 (0.023)*	0.003 (0.020)	-0.054 (0.024)*	0.046 (0.023)*
Large city	-0.018 (0.017)	-0.055 (0.015)*	-0.010 (0.017)	-0.010 (0.017)
Corruption perception (z)				0.054 (0.008)*
Electoral status: Non-voter / None				-0.160 (0.021)*
Electoral status: Opposition voter				-0.151 (0.024)*
Intercept	-0.091 (0.071)	0.097 (0.078)	-0.170 (0.080)*	0.035 (0.072)
N	13,275	13,396	13,396	13,146
Countries	16	16	16	16
SD(country intercept)	0.206	0.262	0.247	0.202

Continued on next page.

Table C16 continued from previous page.

Term	Satisfaction w/ democracy	First link of the mediation		Satisfaction, extended controls
		Predicting institutional trust	Predicting trust in police	
SD(slope: Trust in police (z))	0.008	0.114		0.000
SD(slope: Political trust (z))	0.054		0.145	0.054
SD(residual)	0.815	0.722	0.840	0.812
ICC (country)	0.060	0.116	0.080	0.058
RE form	correlated	correlated	correlated	correlated

Each cell gives a coefficient with its standard error in parentheses, and an asterisk marks $p < .05$, two-tailed. Only that one level is marked. We calculate p-values using the Satterthwaite approximation because the predictors of interest have slopes that vary by country, which leaves them with about as many effective degrees of freedom as there are countries rather than respondents; a normal approximation would overstate their significance appreciably. Models are weighted by the survey weight and fitted by REML. Each column drops respondents missing any variable it uses, so N differs between the base and extended specifications. Categorical controls enter as sets of indicator variables, and in each case the omitted reference category is the largest group: age 30–44, less than high school for education, Mestizo for ethnicity, Catholic for religion, small family for household type, and employed-paid for occupation. Every dummy therefore reads against the most common group rather than an arbitrary one. Respondents who did not answer a categorical item form their own Not reported level rather than being dropped or folded into the reference category. Columns 2 and 3 are the first link of the mediation decomposition, fitted in both directions because cross-sectional data cannot say which trust measure comes first. Column 4 adds the strongest competing predictors: perceptions of corruption and winner-loser status.

Table C17: Full estimates, LAPOP AmericasBarometer 2023: every outcome equation.

Term	Satisfaction w/ democracy	First link of the mediation		Satisfaction, extended controls
		Predicting institutional trust	Predicting trust in police	
Trust in police (z)	0.039 (0.015)*	0.445 (0.025)*		0.036 (0.015)*
Political trust (z)	0.334 (0.014)*		0.526 (0.025)*	0.314 (0.013)*
Economic perceptions	0.117 (0.007)*	0.202 (0.006)*	-0.041 (0.006)*	0.112 (0.007)*
Political interest	0.033 (0.007)*	0.098 (0.005)*	-0.009 (0.006)	0.027 (0.007)*
Interpersonal trust	0.021 (0.007)*	0.028 (0.006)*	0.065 (0.006)*	0.018 (0.007)*
Age (years)	-0.001 (0.000)	-0.001 (0.000)*	0.001 (0.000)*	-0.001 (0.000)
Female	-0.069 (0.013)*	0.032 (0.011)*	0.041 (0.012)*	-0.062 (0.013)*
Education level	-0.026 (0.005)*	-0.032 (0.004)*	-0.010 (0.005)*	-0.027 (0.005)*
Income bracket	0.004 (0.006)	-0.028 (0.005)*	-0.016 (0.005)*	0.005 (0.006)
Asset index (z)	-0.023 (0.008)*	-0.035 (0.007)*	-0.002 (0.007)	-0.022 (0.008)*
Urban	-0.016 (0.015)	-0.050 (0.013)*	-0.033 (0.014)*	-0.013 (0.015)
Crime victimization	-0.100 (0.015)*	-0.062 (0.013)*	-0.127 (0.014)*	-0.090 (0.015)*
Ideology (left-right)	0.001 (0.002)	0.025 (0.002)*	0.022 (0.002)*	0.001 (0.002)
Corruption perception (z)				-0.078 (0.007)*
Intercept	0.160 (0.050)*	0.132 (0.081)	-0.073 (0.079)	0.157 (0.047)*
N	19,279	19,485	19,485	19,029
Countries	16	16	16	16
SD(country intercept)	0.150	0.305	0.294	0.135
SD(slope: Trust in police (z))	0.050	0.096		0.051
SD(slope: Political trust (z))	0.046		0.097	0.040
SD(residual)	0.848	0.711	0.772	0.846
ICC (country)	0.030	0.156	0.127	0.025

RE form	correlated	correlated	correlated	correlated
---------	------------	------------	------------	------------

Each cell gives a coefficient with its standard error in parentheses, and an asterisk marks $p < .05$, two-tailed. Only that one level is marked. We calculate p-values using the Satterthwaite approximation because the predictors of interest have slopes that vary by country, which leaves them with about as many effective degrees of freedom as there are countries rather than respondents; a normal approximation would overstate their significance appreciably. Models are weighted by the survey weight and fitted by REML. Each column drops respondents missing any variable it uses, so N differs between the base and extended specifications. Categorical controls enter as sets of indicator variables, and in each case the omitted reference category is the largest group: age 30–44, less than high school for education, Mestizo for ethnicity, Catholic for religion, small family for household type, and employed-paid for occupation. Every dummy therefore reads against the most common group rather than an arbitrary one. Respondents who did not answer a categorical item form their own Not reported level rather than being dropped or folded into the reference category. Columns 2 and 3 are the first link of the mediation decomposition, fitted in both directions. Column 4 adds perceptions of corruption, the only competing item this wave asked with usable coverage.

Table C18: Full estimates, LAPOP 2008 police-abuse analysis: institutional trust.

Term	Base controls	Extended controls
Police-abuse victimization	-0.205 (0.032)*	-0.178 (0.030)*
Age: 18–23	0.059 (0.024)*	0.018 (0.024)
Age: 24–29	-0.013 (0.021)	-0.040 (0.021)
Age: 45–59	0.026 (0.020)	0.046 (0.019)*
Age: 60+	0.069 (0.026)*	0.096 (0.025)*
age_cohortNot reported	0.243 (0.135)	0.246 (0.135)
Female	-0.045 (0.016)*	0.013 (0.015)
Education: College graduate	-0.025 (0.037)	-0.083 (0.036)*
Education: High school graduate	-0.045 (0.021)*	-0.071 (0.020)*
educ_levelNot reported	-0.012 (0.106)	-0.113 (0.103)
Education: Postgraduate	-0.056 (0.031)	-0.121 (0.031)*
Education: Some college	-0.008 (0.025)	-0.060 (0.025)*
income_quintile1	0.038 (0.020)	0.074 (0.019)*
income_quintile3	-0.012 (0.019)	-0.061 (0.018)*
income_quintile4	0.013 (0.024)	-0.096 (0.024)*
income_quintile5	0.113 (0.032)*	-0.059 (0.031)
householdCouple	-0.022 (0.021)	-0.016 (0.020)
Household: Large family	0.034 (0.040)	0.036 (0.039)
Household: Medium family	0.021 (0.021)	0.016 (0.020)
Household: Single person	0.007 (0.020)	-0.010 (0.019)
Ethnicity: Black	0.003 (0.036)	0.039 (0.036)
Ethnicity: Indigenous	0.052 (0.031)	0.028 (0.031)
Ethnicity: Mulatto	0.003 (0.049)	0.019 (0.047)
ethnicidNot reported	-0.003 (0.047)	0.010 (0.047)
Ethnicity: Other	0.086 (0.065)	0.122 (0.064)
Ethnicity: White	0.031 (0.017)	0.033 (0.017)*
Religion: Evangelical	0.014 (0.024)	-0.002 (0.023)
Religion: Not religious	-0.112 (0.025)*	-0.118 (0.024)*
religionidNot reported	-0.211 (0.103)*	-0.295 (0.103)*
Religion: Other	-0.146 (0.045)*	-0.117 (0.044)*
Religion: Protestant	-0.027 (0.032)	-0.038 (0.031)
Ideology (left-right)	0.056 (0.003)*	0.051 (0.003)*
Crime victimization	-0.128 (0.018)*	-0.102 (0.018)*
Occupation: Employed (unpaid)	0.188 (0.088)*	0.215 (0.087)*
occupNot reported	0.073 (0.075)	0.107 (0.076)
Occupation: Non-active / other	0.023 (0.019)	0.039 (0.018)*
Occupation: Student	-0.033 (0.032)	-0.023 (0.031)
Occupation: Unemployed	-0.083 (0.031)*	-0.027 (0.030)
Urban	-0.106 (0.019)*	-0.097 (0.019)*
National capital	-0.084 (0.019)*	-0.072 (0.019)*
Large city	-0.106 (0.021)*	-0.073 (0.021)*
Asked for bribe by police	-0.158 (0.023)*	-0.134 (0.023)*
Economic perceptions		0.269 (0.008)*

Continued on next page.

Table C18 continued from previous page.

Term	Base controls	Extended controls
Political interest		0.078 (0.007)*
Interpersonal trust		0.085 (0.007)*
Intercept	-0.123 (0.099)	-0.150 (0.085)
N	17,703	16,928
Countries	16	16
SD(country intercept)	0.378	0.321
SD(slope: Police-abuse victimization)	0.035	0.000
SD(residual)	0.865	0.826
ICC (country)	0.161	0.132
RE form	correlated	correlated

Each cell gives a coefficient with its standard error in parentheses, and an asterisk marks $p < .05$, two-tailed. Only that one level is marked. We calculate p-values using the Satterthwaite approximation because the predictors of interest have slopes that vary by country, which leaves them with about as many effective degrees of freedom as there are countries rather than respondents; a normal approximation would overstate their significance appreciably. Models are weighted by the survey weight and fitted by REML. Each column drops respondents missing any variable it uses, so N differs between the base and extended specifications. Categorical controls enter as sets of indicator variables, and in each case the omitted reference category is the largest group: age 30–44, less than high school for education, Mestizo for ethnicity, Catholic for religion, small family for household type, and employed-paid for occupation. Every dummy therefore reads against the most common group rather than an arbitrary one. Respondents who did not answer a categorical item form their own Not reported level rather than being dropped or folded into the reference category. These are descriptive estimates, not causal ones. The predictor is the respondent's own report of being mistreated by the police, recorded in a single cross-section. The extended column adds the attitude measures that the 2023 control sets carry.

Table C19: Full estimates, LAPOP 2008 police-abuse analysis: trust in the police.

Term	Base controls	Extended controls
Police-abuse victimization	-0.401 (0.042)*	-0.387 (0.041)*
Age: 18–23	0.031 (0.025)	-0.008 (0.025)
Age: 24–29	-0.036 (0.022)	-0.054 (0.022)*
Age: 45–59	0.057 (0.020)*	0.065 (0.020)*
Age: 60+	0.065 (0.026)*	0.072 (0.026)*
age_cohortNot reported	0.108 (0.137)	0.050 (0.141)
Female	-0.041 (0.016)*	-0.016 (0.016)
Education: College graduate	-0.004 (0.038)	-0.022 (0.038)
Education: High school graduate	-0.052 (0.021)*	-0.053 (0.021)*
educ_levelNot reported	-0.137 (0.107)	-0.182 (0.107)
Education: Postgraduate	-0.028 (0.031)	-0.051 (0.032)
Education: Some college	-0.006 (0.026)	-0.024 (0.026)
income_quintile1	0.042 (0.020)*	0.070 (0.020)*
income_quintile3	0.010 (0.019)	-0.019 (0.019)
income_quintile4	0.010 (0.024)	-0.047 (0.025)
income_quintile5	0.101 (0.032)*	0.006 (0.033)
householdCouple	-0.025 (0.021)	-0.022 (0.021)
Household: Large family	0.033 (0.041)	0.023 (0.041)
Household: Medium family	0.029 (0.021)	0.026 (0.021)
Household: Single person	0.011 (0.020)	-0.002 (0.020)
Ethnicity: Black	-0.028 (0.037)	-0.003 (0.037)
Ethnicity: Indigenous	-0.019 (0.031)	-0.039 (0.032)
Ethnicity: Mulatto	-0.024 (0.049)	-0.027 (0.049)
ethnicidNot reported	-0.065 (0.048)	-0.056 (0.048)
Ethnicity: Other	-0.018 (0.066)	0.020 (0.067)
Ethnicity: White	0.025 (0.017)	0.025 (0.017)
Religion: Evangelical	-0.034 (0.024)	-0.044 (0.024)
Religion: Not religious	-0.145 (0.025)*	-0.153 (0.025)*
religionidNot reported	-0.288 (0.104)*	-0.317 (0.107)*
Religion: Other	-0.207 (0.046)*	-0.195 (0.046)*
Religion: Protestant	-0.066 (0.032)*	-0.074 (0.032)*
Ideology (left-right)	0.042 (0.003)*	0.038 (0.003)*
Crime victimization	-0.129 (0.018)*	-0.109 (0.018)*
Occupation: Employed (unpaid)	0.121 (0.089)	0.139 (0.090)
occupNot reported	-0.032 (0.076)	0.013 (0.079)
Occupation: Non-active / other	0.046 (0.019)*	0.054 (0.019)*
Occupation: Student	-0.031 (0.032)	-0.012 (0.033)
Occupation: Unemployed	-0.099 (0.031)*	-0.060 (0.032)
Urban	-0.073 (0.020)*	-0.063 (0.020)*
National capital	-0.113 (0.020)*	-0.113 (0.020)*
Large city	-0.119 (0.022)*	-0.099 (0.022)*
Asked for bribe by police	-0.218 (0.024)*	-0.206 (0.024)*
Economic perceptions		0.153 (0.008)*

Continued on next page.

Table C19 continued from previous page.

Term	Base controls	Extended controls
Political interest		0.022 (0.007)*
Interpersonal trust		0.079 (0.008)*
Intercept	-0.094 (0.081)	-0.094 (0.077)
N	17,703	16,928
Countries	16	16
SD(country intercept)	0.300	0.285
SD(slope: Police-abuse victimization)	0.113	0.104
SD(residual)	0.873	0.859
ICC (country)	0.106	0.099
RE form	correlated	correlated

Each cell gives a coefficient with its standard error in parentheses, and an asterisk marks $p < .05$, two-tailed. Only that one level is marked. We calculate p-values using the Satterthwaite approximation because the predictors of interest have slopes that vary by country, which leaves them with about as many effective degrees of freedom as there are countries rather than respondents; a normal approximation would overstate their significance appreciably. Models are weighted by the survey weight and fitted by REML. Each column drops respondents missing any variable it uses, so N differs between the base and extended specifications. Categorical controls enter as sets of indicator variables, and in each case the omitted reference category is the largest group: age 30–44, less than high school for education, Mestizo for ethnicity, Catholic for religion, small family for household type, and employed-paid for occupation. Every dummy therefore reads against the most common group rather than an arbitrary one. Respondents who did not answer a categorical item form their own Not reported level rather than being dropped or folded into the reference category. These are descriptive estimates, not causal ones. The predictor is the respondent's own report of being mistreated by the police, recorded in a single cross-section. The extended column adds the attitude measures that the 2023 control sets carry.

Table C20: Full estimates, LAPOP 2008 police-abuse analysis: satisfaction with democracy.

Term	Base controls	Extended controls
Police-abuse victimization	-0.121 (0.036)*	-0.096 (0.035)*
Age: 18–23	0.043 (0.025)	0.002 (0.025)
Age: 24–29	0.003 (0.022)	-0.030 (0.022)
Age: 45–59	0.013 (0.020)	0.028 (0.020)
Age: 60+	0.039 (0.026)	0.066 (0.026)*
age_cohortNot reported	0.119 (0.139)	0.092 (0.141)
Female	-0.063 (0.016)*	-0.026 (0.016)
Education: College graduate	-0.057 (0.038)	-0.079 (0.038)*
Education: High school graduate	-0.036 (0.021)	-0.048 (0.021)*
educ_levelNot reported	0.089 (0.109)	-0.039 (0.108)
Education: Postgraduate	-0.061 (0.032)	-0.086 (0.032)*
Education: Some college	0.003 (0.026)	-0.028 (0.026)
income_quintile1	0.002 (0.020)	0.032 (0.020)
income_quintile3	-0.006 (0.019)	-0.050 (0.019)*
income_quintile4	0.051 (0.025)*	-0.042 (0.025)
income_quintile5	0.128 (0.033)*	-0.010 (0.033)
householdCouple	0.018 (0.021)	0.020 (0.021)
Household: Large family	0.091 (0.042)*	0.105 (0.041)*
Household: Medium family	0.044 (0.021)*	0.051 (0.021)*
Household: Single person	0.037 (0.020)	0.026 (0.020)
Ethnicity: Black	-0.020 (0.037)	0.007 (0.037)
Ethnicity: Indigenous	0.019 (0.032)	-0.010 (0.032)
Ethnicity: Mulatto	-0.061 (0.050)	-0.064 (0.049)
ethnicaidNot reported	-0.130 (0.049)*	-0.115 (0.049)*
Ethnicity: Other	-0.140 (0.066)*	-0.106 (0.067)
Ethnicity: White	-0.031 (0.017)	-0.031 (0.017)
Religion: Evangelical	0.038 (0.024)	0.030 (0.024)
Religion: Not religious	-0.098 (0.025)*	-0.106 (0.025)*
religionidNot reported	0.228 (0.108)*	0.243 (0.110)*
Religion: Other	-0.090 (0.046)	-0.075 (0.046)
Religion: Protestant	-0.026 (0.032)	-0.024 (0.032)
Ideology (left-right)	0.039 (0.003)*	0.035 (0.003)*
Crime victimization	-0.095 (0.018)*	-0.071 (0.018)*
Occupation: Employed (unpaid)	-0.038 (0.092)	-0.008 (0.091)
occupNot reported	-0.156 (0.078)*	-0.129 (0.079)
Occupation: Non-active / other	-0.026 (0.019)	-0.017 (0.019)
Occupation: Student	-0.084 (0.033)*	-0.085 (0.033)*
Occupation: Unemployed	-0.116 (0.032)*	-0.068 (0.032)*
Urban	-0.016 (0.020)	-0.013 (0.020)
National capital	-0.101 (0.020)*	-0.084 (0.020)*
Large city	-0.098 (0.022)*	-0.063 (0.022)*
Asked for bribe by police	-0.108 (0.024)*	-0.091 (0.024)*
Economic perceptions		0.242 (0.008)*

Continued on next page.

Table C20 continued from previous page.

Term	Base controls	Extended controls
Political interest		0.023 (0.007)*
Interpersonal trust		0.053 (0.008)*
Intercept	-0.137 (0.085)	-0.149 (0.074)
N	17,337	16,611
Countries	16	16
SD(country intercept)	0.318	0.270
SD(slope: Police-abuse victimization)	0.066	0.070
SD(residual)	0.877	0.851
ICC (country)	0.116	0.092
RE form	correlated	correlated

Each cell gives a coefficient with its standard error in parentheses, and an asterisk marks $p < .05$, two-tailed. Only that one level is marked. We calculate p-values using the Satterthwaite approximation because the predictors of interest have slopes that vary by country, which leaves them with about as many effective degrees of freedom as there are countries rather than respondents; a normal approximation would overstate their significance appreciably. Models are weighted by the survey weight and fitted by REML. Each column drops respondents missing any variable it uses, so N differs between the base and extended specifications. Categorical controls enter as sets of indicator variables, and in each case the omitted reference category is the largest group: age 30–44, less than high school for education, Mestizo for ethnicity, Catholic for religion, small family for household type, and employed-paid for occupation. Every dummy therefore reads against the most common group rather than an arbitrary one. Respondents who did not answer a categorical item form their own Not reported level rather than being dropped or folded into the reference category. These are descriptive estimates, not causal ones. The predictor is the respondent's own report of being mistreated by the police, recorded in a single cross-section. The extended column adds the attitude measures that the 2023 control sets carry.

Table C21: Full estimates, LAPOP 2008: institutional trust and police trust predicting satisfaction with democracy.

Term	Base controls	Extended controls
Political trust (z)	0.281 (0.020)*	0.233 (0.018)*
Trust in police (z)	0.049 (0.019)*	0.046 (0.017)*
Age: 18–23	0.022 (0.024)	-0.004 (0.024)
Age: 24–29	0.006 (0.021)	-0.020 (0.021)
Age: 45–59	0.003 (0.019)	0.015 (0.019)
Age: 60+	0.018 (0.025)	0.041 (0.025)
age_cohortNot reported	0.044 (0.132)	0.027 (0.136)
Female	-0.046 (0.015)*	-0.026 (0.015)
Education: College graduate	-0.051 (0.036)	-0.059 (0.036)
Education: High school graduate	-0.017 (0.020)	-0.026 (0.020)
educ_levelNot reported	0.098 (0.104)	-0.002 (0.104)
Education: Postgraduate	-0.041 (0.030)	-0.051 (0.031)
Education: Some college	0.009 (0.025)	-0.010 (0.025)
income_quintile1	-0.012 (0.019)	0.010 (0.019)
income_quintile3	-0.003 (0.018)	-0.035 (0.018)
income_quintile4	0.043 (0.024)	-0.021 (0.024)
income_quintile5	0.090 (0.031)*	0.004 (0.031)
householdCouple	0.024 (0.020)	0.024 (0.020)
Household: Large family	0.079 (0.040)*	0.094 (0.040)*
Household: Medium family	0.037 (0.020)	0.046 (0.021)*
Household: Single person	0.032 (0.019)	0.028 (0.019)
Ethnicity: Black	-0.025 (0.036)	-0.007 (0.036)
Ethnicity: Indigenous	0.008 (0.030)	-0.011 (0.031)
Ethnicity: Mulatto	-0.063 (0.047)	-0.069 (0.048)
ethnicaidNot reported	-0.123 (0.047)*	-0.114 (0.047)*
Ethnicity: Other	-0.171 (0.063)*	-0.137 (0.064)*
Ethnicity: White	-0.039 (0.017)*	-0.038 (0.017)*
Religion: Evangelical	0.033 (0.023)	0.031 (0.023)
Religion: Not religious	-0.061 (0.024)*	-0.072 (0.024)*
religionidNot reported	0.300 (0.102)*	0.325 (0.107)*
Religion: Other	-0.039 (0.044)	-0.039 (0.044)
Religion: Protestant	-0.014 (0.031)	-0.012 (0.031)
Ideology (left-right)	0.021 (0.003)*	0.020 (0.003)*
Crime victimization	-0.057 (0.017)*	-0.044 (0.018)*
Occupation: Employed (unpaid)	-0.107 (0.087)	-0.068 (0.088)
occupNot reported	-0.180 (0.074)*	-0.158 (0.076)*
Occupation: Non-active / other	-0.039 (0.018)*	-0.031 (0.018)
Occupation: Student	-0.068 (0.031)*	-0.076 (0.031)*
Occupation: Unemployed	-0.089 (0.030)*	-0.061 (0.030)*
Urban	0.019 (0.019)	0.016 (0.019)
National capital	-0.073 (0.019)*	-0.063 (0.019)*
Large city	-0.059 (0.021)*	-0.040 (0.021)

Continued on next page.

Table C21 continued from previous page.

Term	Base controls	Extended controls
Asked for bribe by police	-0.058 (0.022)*	-0.056 (0.023)*
Economic perceptions		0.170 (0.008)*
Political interest		0.005 (0.007)
Interpersonal trust		0.030 (0.007)*
Intercept	-0.100 (0.068)	-0.112 (0.064)
N	17,337	16,611
Countries	16	16
SD(country intercept)	0.245	0.227
SD(slope: Political trust (z))	0.071	0.062
SD(slope: Trust in police (z))	0.069	0.060
SD(residual)	0.832	0.821
ICC (country)	0.080	0.071
RE form	correlated	correlated

Each cell gives a coefficient with its standard error in parentheses, and an asterisk marks $p < .05$, two-tailed. Only that one level is marked. We calculate p-values using the Satterthwaite approximation because the predictors of interest have slopes that vary by country, which leaves them with about as many effective degrees of freedom as there are countries rather than respondents; a normal approximation would overstate their significance appreciably. Models are weighted by the survey weight and fitted by REML. Each column drops respondents missing any variable it uses, so N differs between the base and extended specifications. Categorical controls enter as sets of indicator variables, and in each case the omitted reference category is the largest group: age 30–44, less than high school for education, Mestizo for ethnicity, Catholic for religion, small family for household type, and employed-paid for occupation. Every dummy therefore reads against the most common group rather than an arbitrary one. Respondents who did not answer a categorical item form their own Not reported level rather than being dropped or folded into the reference category. This repeats, on the 2008 wave, the specification the two 2023 surveys report. It uses the controls available in LAPOP 2008 rather than the ones the two 2023 tables share, so it asks the same question with a different set of controls. These are descriptive estimates; they support no causal claim.

Moderation: ideology and race

Both comparisons below are descriptive. Neither identifies a mechanism by which ideology or race changes the effect of police violence. Each says only which group the association is larger for.

Table C22 and panel A of Figure C2 split every association by where respondents place themselves ideologically. We fit one model per association and let the coefficient differ between the two groups inside it, rather than estimating each group separately. Left means 0–4 on Latinobarómetro’s 0–10 scale and 1–4 on LAPOP’s 1–10 scale; right means 6–10 on both. We exclude respondents at the centre (5) and those who decline to place themselves, so the comparison is between the two wings rather than between one wing and everyone else.

Table C22: Study 3 associations by ideological self-placement: left compared with right.

Outcome	Association	Left effect (SE)	Right effect (SE)	Right - left (SE)	Difference p	Left N	Right N	Countries	RE structure
<i>Latinobarometro 2023: trust -> satisfaction</i>									
Satisfaction with democracy	Trust in police	0.015 (0.015)	0.050 (0.013)	0.034 (0.018)	0.056	4215	5255	16	decorrelated (correlated form singular)
Satisfaction with democracy	Political trust index	0.283 (0.020)	0.286 (0.018)	0.003 (0.018)	0.889	4215	5255	16	decorrelated (correlated form singular)
<i>LAPOP 2023: trust -> satisfaction</i>									
Satisfaction with democracy	Trust in police	0.014 (0.018)	0.048 (0.018)	0.034 (0.015)	0.023	6901	7587	16	correlated
Satisfaction with democracy	Political trust index	0.334 (0.017)	0.353 (0.016)	0.019 (0.015)	0.210	6901	7587	16	correlated
<i>LAPOP 2008: police abuse -> trust</i>									
Political trust index	Police-abuse victimization	-0.304 (0.054)	-0.135 (0.049)	0.170 (0.068)	0.013	4599	8755	16	correlated
Trust in police	Police-abuse victimization	-0.423 (0.054)	-0.379 (0.049)	0.044 (0.068)	0.517	4768	9118	16	decorrelated (correlated form singular)
Satisfaction with democracy	Police-abuse victimization	-0.243 (0.052)	-0.047 (0.047)	0.196 (0.068)	0.004	4664	8964	16	decorrelated (correlated form singular)

We fit one model per dataset and outcome, and let the association differ between the two groups inside it rather than estimating each group separately. Left means 0–4 on Latinobarómetro’s 0–10 scale and 1–4 on LAPOP’s 1–10 scale; right means 6–10 on both. Respondents at the centre (5) and those who decline to place themselves are excluded. The Difference column is how much the association changes as one moves from left to right, and its standard error comes from the fitted model. Each country keeps its own intercept and its own slope on the association of interest. These are descriptive comparisons. They are not causal, and they do not test any psychological mechanism.

Table C23 and panel B of Figure C2 split the same associations by race, comparing non-white with white respondents. This matches the comparison Study 1 makes. All three datasets support it. The Latinobarómetro asks which ethnicity or race the respondent identifies with (item S7), and LAPOP asks the same in both waves (item etid). Each is present in all sixteen countries; the item is missing for 11% of Latinobarómetro respondents and 6% of LAPOP 2023 respondents. Non-white pools every category other than White. The 2023 models are the main models with the race interaction added and nothing removed; the 2008 model drops ethnicity from its controls, because the comparison is built from it.

Race makes no difference to any association in any dataset. All seven differences are statistically indistinguishable from zero; the largest is 0.052 standard deviations, on police abuse and political trust in 2008, against a white-respondent estimate of -0.240 for that same association. It says that the measured associations do not differ by race. It does not say that police violence falls equally on racial groups. Who is exposed and what exposure does are different questions, and only the second is answered here.

Table C23: Study 3 associations by race: non-white compared with white.

Outcome	Association	White effect (SE)	Non-white effect (SE)	Non-white - white (SE)	Difference p	White N	Non-white N	Countries	RE structure
<i>Latinobarometro 2023: trust -> satisfaction</i>									
Satisfaction with democracy	Trust in police	0.028 (0.017)	0.045 (0.012)	0.017 (0.018)	0.344	3162	9089	16	decorrelated (correlated form singular)
Satisfaction with democracy	Political trust index	0.290 (0.022)	0.268 (0.018)	-0.022 (0.018)	0.223	3162	9089	16	decorrelated (correlated form singular)
<i>LAPOP 2023: trust -> satisfaction</i>									
Satisfaction with democracy	Trust in police	0.040 (0.018)	0.037 (0.015)	-0.003 (0.015)	0.834	4852	13659	16	correlated
Satisfaction with democracy	Political trust index	0.327 (0.018)	0.341 (0.015)	0.013 (0.015)	0.379	4852	13659	16	correlated
<i>LAPOP 2008: police abuse -> trust</i>									
Political trust index	Police-abuse victimization	-0.240 (0.056)	-0.187 (0.036)	0.052 (0.065)	0.422	5600	11776	16	decorrelated (correlated form singular)
Trust in police	Police-abuse victimization	-0.429 (0.066)	-0.385 (0.048)	0.044 (0.069)	0.524	5832	12215	16	decorrelated (correlated form singular)
Satisfaction with democracy	Police-abuse victimization	-0.107 (0.061)	-0.121 (0.044)	-0.014 (0.067)	0.834	5751	11975	16	correlated

We fit one model per dataset and outcome, and let the association differ between the two groups inside it rather than estimating each group separately. Non-white pools every category other than White – Mestizo, Black, Indigenous, Mulatto and Other, plus Latinobarómetro’s Asian and the nationally specific categories a few LAPOP countries add. Respondents with no recorded ethnicity are excluded rather than assigned to either group. Ethnicity is dropped from the 2008 control set, because the comparison is built from it; the 2023 control sets never included it, so they are the main models’ unchanged. Each country keeps its own intercept and its own slope on the association of interest. These are descriptive comparisons, not causal ones.




Figure C2: Every Study 3 association under each moderator. Panel A splits by ideological self-placement, panel B by race. Each panel is drawn in two blocks: the upper shows the 2023 trust-to-satisfaction associations and the lower the LAPOP 2008 police-abuse ones; the outcome differs between the two blocks. Predictors run down the left, outcomes across the top, and the dataset is given at the right. Each cell shows both groups of that panel's comparison as a filled and an open mark with 95% confidence intervals — circles for the ideological comparison, squares for the racial one, so the two are told apart by the mark and not only by the legend. Δ is the difference between the two groups — right minus left in panel A, non-white minus white in panel B — and * marks $p < .05$ on that difference, not on the individual associations. Both panels share one horizontal scale.

Country-level patterns and scale validation

This section asks what differs across the sixteen countries, and whether the trust scale behaves the same way in each. We present these results as figures because what matters is whether the direction holds across countries, not how large any one country's estimate is. Figure C3 shows, country by country, whether leaving that country out of estimation makes prediction worse. Figure C4 shows each country's indirect effect under both orderings of the mediation decomposition. Figures C5 and C6 plot each country's slope in the two 2023 models as a departure from the average slope, so a country above the line is one where trust tracks satisfaction more strongly than average. Figures C7 and C8 take the 2008 country estimates and ask how far they genuinely differ, first after dropping incomplete cases and then under multiple imputation. Figure C9 closes the section by testing the three-item trust scale used throughout against LAPOP's own published index.




Figure C3: How much prediction improves in each country when that country is held out of estimation.

Note: we hold out each country in turn, refit the model on the remaining countries, and predict the held-out country without using anything specific to it, since nothing country-specific would be known for a genuinely new country. Positive values mean that adding the predictors of interest lowers prediction error for that country, and open circles mark the countries where they do not. Each panel has its own horizontal scale, so panels cannot be compared with one another.




Figure C4: Each country's indirect effect, under both orderings of the mediation decomposition.

Note: each country's own indirect effect from the mediation decomposition, under each of the two possible orderings between trust in police and institutional trust. The indirect effect is positive in 16 of 16 (Ordering A) and 15 of 16 (Ordering B) Latinobarometro countries, and in 16 of 16 (Ordering A) or 14 of 16 (Ordering B) LAPOP countries; open circles mark the countries where it is negative. Each panel has its own horizontal scale.




Figure C5: How each country's slope on police trust and on political trust departs from the average slope, Latinobarómetro 2023.




Figure C6: How each country's slope on police trust and on political trust departs from the average slope, LAPOP 2023.




Figure C7: The effect of reported police abuse in each country, and how far the sixteen countries differ.

Note: the effect of reported police abuse in each country, LAPOP 2008, estimated in two stages. The first stage estimates one effect per country from a weighted linear model in which the intercept and the police-abuse slope vary by country while the control coefficients are shared, matching how the single-stage multilevel model treats the controls. The second stage combines those sixteen effects, treating them as draws from one common distribution, fitted by restricted maximum likelihood, and giving each country more weight the more precisely its own effect is estimated. Squares are country effects, sized by that weight, and thin horizontal lines are 95% confidence intervals. Countries run in order of their effect on trust in police, and the three panels keep that same order. The diamond is the pooled mean, the thick bar its 95% confidence interval, and the light grey bar the 95% prediction interval – the range in which a further country’s true effect would be expected to fall, which unlike the confidence interval does not narrow as more countries are added. The two trust outcomes share one sample; satisfaction with democracy has its own.




Figure C8: The same estimates under multiple imputation, which keeps the third of respondents that dropping incomplete cases discards.

Note: the same figure as the one before it, differing only in how missing data are handled. Dropping incomplete cases discards 34% of the rows in the sixteen matched countries (26,894 down to 17,703), a loss concentrated in ideology (19.8% missing) and household income (11.2%). Who fails to answer those two depends on education and political engagement, and depends on them differently in each country, so dropping incomplete cases changes each country's sample by a different amount. Here the same quantities are estimated on all 26,894 rows after imputing the missing answers 20 times, country by country, using predictive mean matching for the continuous and ordered items and logistic imputation for the yes-or-no ones. The country effects are combined across imputations at the first stage, so each country's standard error carries the extra uncertainty imputation introduces; the second stage then proceeds exactly as before. On average across countries, 0.03 of the information is missing. Squares are country effects sized by their weight, and thin lines are 95% confidence intervals. The diamond is the pooled mean, the thick bar its 95% confidence interval, and the light grey bar the 95% prediction interval for a further country. Countries run in order of their effect on trust in police.




Figure C9: The three-item institutional-trust scale tested against LAPOP's published index: reliability, agreement, and whether the choice of scale changes the results.

Note: the three-item scale used throughout this project combines trust in the legislature, the executive and the courts. LAPOP's own AmericasBarometer report instead uses four seven-point trust items plus a four-point item on the national government (Lupu et al. 2023, fn. 10). Panel A gives standardised Cronbach's alpha, with the dashed line at the conventional floor of .70. The wider LAPOP scales are blank for Studies 1 and 2 because those questionnaires do not carry the extra items. Panel B gives the correlation within each country between the three-item scale and each published LAPOP scale, across all 16 matched countries. The lowest of the 64 comparisons is $r = 0.905$, against the .90 rule marked by the dashed line. Panel C refits the multilevel models on one shared sample, changing nothing but the scale. Bars are 95% confidence intervals, and stars mark * $p < .05$, ** $p < .01$, *** $p < .001$.

How each result was estimated

Every number in this appendix comes from a specific model fitted by a specific package, and a reader checking the work should not have to guess which. The table below lists every model estimated: what it is, on which survey, with which controls, through which package, and which tables and figures it produces. It is traced from the code rather than written by hand: a tool walks backwards from each exhibit to the script that renders it, to the files that script reads, to the script that wrote those, and records the estimation calls it finds there. Only estimation is recorded — model fits, standard errors, internal- consistency coefficients, imputations, decompositions of explained variance. Data preparation and plotting are left out, because they say nothing about how a quantity was produced. The last column is derived from the code, and the tool checks its own reading of the source against the R parser and fails if the two disagree. The model name, dataset and control roster in the other columns are authored rather than derived, because no parser recovers them; they are checked against the code by hand.

Table C24: Every model estimated, and the exhibits it produces.

Model	Dataset	Controls	Package (call)	Appears in
Trust to satisfaction with democracy, multilevel with country random slopes	Latinobarometro 2023	Latinobarometro roster	lmerTest (lmer)	Table A4, Table A5, Figure A5, Figure A6
Trust to satisfaction with democracy, multilevel with country random slopes	LAPOP 2023	LAPOP 2023 roster	lmerTest (lmer)	Table A4, Table A5, Figure A5, Figure A6
Reported police abuse to trust, pooled country fixed effects with cluster-robust errors	LAPOP 2008	LAPOP 2008 roster	stats + clubSandwich (lm, coef_test (CR2))	Table A4, Table A5, Figure A5, Figure A6
Reported police abuse to trust, multilevel counterpart of the pooled fit	LAPOP 2008	LAPOP 2008 roster	lmerTest (lmer)	Table A4, Table A5, Figure A5, Figure A6
Association by ideological self-placement, one pooled interaction model per association	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster, less ideology	lmerTest (lmer)	Table A18, Table A19, Figure A1, Figure A2
Association by race, one pooled interaction model per outcome	LAPOP 2008	LAPOP 2008 roster, less ethnicity	lmerTest (lmer)	Table A18, Table A19, Figure A1, Figure A2
Nested models for what trust adds beyond the controls – controls only	Latinobarometro 2023	Latinobarometro roster	lme4 (lmer, anova)	Table A7
Nested models for what trust adds beyond the controls – plus political trust	Latinobarometro 2023	Latinobarometro roster	lme4 (lmer, anova)	Table A7
Nested models for what trust adds beyond the controls – plus both trust measures	Latinobarometro 2023	Latinobarometro roster	lme4 (lmer, anova)	Table A7
Nested models for what trust adds beyond the controls – controls only	LAPOP 2023	LAPOP 2023 roster	lme4 (lmer, anova)	Table A7
Nested models for what trust adds beyond the controls – plus political trust	LAPOP 2023	LAPOP 2023 roster	lme4 (lmer, anova)	Table A7
Nested models for what trust adds beyond the controls – plus both trust measures	LAPOP 2023	LAPOP 2023 roster	lme4 (lmer, anova)	Table A7
Nested models for what reported police abuse adds beyond the controls – controls only	LAPOP 2008	LAPOP 2008 roster	lme4 (lmer, anova)	Table A7
Nested models for what reported police abuse adds beyond the controls – plus reported police abuse	LAPOP 2008	LAPOP 2008 roster	lme4 (lmer, anova)	Table A7
Shapley decomposition of explained variance across correlated predictors	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	relaimpo (calc.relimp)	Table A7

Continued on next page.

Table C24 continued from previous page.

Model	Dataset	Controls	Package (call)	Appears in
Design-based re-estimation using the survey's own sampling units and strata	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	survey (svydesign, svyglm)	Table A8
Survey-weighted multilevel model by pseudo-maximum-likelihood	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	WeMix (mix)	Table A8
Latent correlation between trust and satisfaction, corrected for measurement error	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	none	lavaan (cfa)	Table A9
Comparison outcome: the same predictor against trust in other people	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster, reduced	stats (lm)	Table A9
Reported police abuse entered alongside ordinary crime victimization	LAPOP 2008	LAPOP 2008 roster, less ideology	stats (lm)	Table A9
Specifications adding the strongest competing predictors, in sequence – base controls	Latinobarometro 2023	Latinobarometro roster, extended	lme4 (lmer)	Table A10
Specifications adding the strongest competing predictors, in sequence – plus corruption and winner-loser status	Latinobarometro 2023	Latinobarometro roster, extended	lme4 (lmer)	Table A10
Specifications adding the strongest competing predictors, in sequence – plus presidential approval	Latinobarometro 2023	Latinobarometro roster, extended	lme4 (lmer)	Table A10
Specifications adding the strongest competing predictors, in sequence – base controls	LAPOP 2023	LAPOP 2023 roster, extended	lme4 (lmer)	Table A10
Specifications adding the strongest competing predictors, in sequence – plus perceived corruption	LAPOP 2023	LAPOP 2023 roster, extended	lme4 (lmer)	Table A10
Specifications adding the strongest competing predictors, in sequence – plus executive approval	LAPOP 2023	LAPOP 2023 roster, extended	lme4 (lmer)	Table A10
Specifications adding the strongest competing predictors, in sequence – base controls	LAPOP 2008	LAPOP 2008 roster, extended	lme4 (lmer)	Table A10
Specifications adding the strongest competing predictors, in sequence – plus economic perceptions, political interest and interpersonal trust	LAPOP 2008	LAPOP 2008 roster, extended	lme4 (lmer)	Table A10

Continued on next page.

Table C24 continued from previous page.

Model	Dataset	Controls	Package (call)	Appears in
Ordered model treating the outcome as categories rather than a scale	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	ordinal (clmm)	Table A11
Response model for who answers, used to reweight by probability of response	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	stats (glm)	Table A11
Internal consistency of the trust scales	Latinobarometro 2023, LAPOP 2023, LAPOP 2008, Study 1 Brazil, Study 2 Chile	none	psych (alpha, omega, pool)	Table A10, Table A11, Figure A9
The same specification refitted with each alternative trust scale in turn	Latinobarometro 2023, LAPOP 2023	Latinobarometro roster	lmerTest (lmer)	Table A10, Figure A9
Leave-one-country-out cross-validation, predicting the held-out country	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	lme4 (lmer)	Table A12, Table A14, Table A15, Table A16, Table A17, Figure A3
Refits underlying the sensitivity of each estimate to an unmeasured confounder	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	lme4 (lmer)	Table A6, Table A13
Mediation decomposition, both orderings, with a bootstrap over countries	Latinobarometro 2023, LAPOP 2023	Latinobarometro roster	lme4 (lmer)	Table A6, Figure A4
Country-specific effects, first stage of the two-stage synthesis	LAPOP 2008	LAPOP 2008 roster	stats (lm)	Figure A7, Figure A8
Random-effects synthesis of the country-specific effects, second stage – restricted maximum likelihood	LAPOP 2008	none	metafor (rma)	Figure A7, Figure A8
Random-effects synthesis of the country-specific effects, second stage – Knapp-Hartung	LAPOP 2008	none	metafor (rma)	Figure A7, Figure A8
Random-effects synthesis of the country-specific effects, second stage – DerSimonian-Laird	LAPOP 2008	none	metafor (rma)	Figure A7, Figure A8
Random-effects synthesis of the country-specific effects, second stage – fixed effect	LAPOP 2008	none	metafor (rma)	Figure A7, Figure A8
The same two-stage synthesis after multiple imputation	LAPOP 2008	LAPOP 2008 roster	mice + stats + metafor (mice, lm, rma)	Figure A8
Every coefficient of every reported equation – base controls	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	lmerTest (lmer)	Table A14, Table A15, Table A16, Table A17

Continued on next page.

Table C24 continued from previous page.

Model	Dataset	Controls	Package (call)	Appears in
Every coefficient of every reported equation – extended controls	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	lmerTest (lmer)	Table A14, Table A15, Table A16, Table A17
Trust to satisfaction with democracy, pooled country fixed effects	LAPOP 2008	LAPOP 2008 roster	stats + clubSandwich (lm, coef_test (CR2))	Table A4, Table A5, Table A14, Table A15, Table A16, Table A17, Figure A5, Figure A6
One-stage comparator: the slope allowed to vary by country within a single model	LAPOP 2008	LAPOP 2008 roster	lme4 (lmer)	Figure A7, Figure A8
Wild cluster bootstrap-t over countries	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	lme4 (lmer)	Table A8
Weighted and unweighted fits of the same specification, compared	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	lme4 (lmer)	Table A8
The same specification under three treatments of missing data – complete cases	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	stats + mice (lm, mice, pool)	Table A11
The same specification under three treatments of missing data – reweighted by probability of response	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	stats + mice (lm, mice, pool)	Table A11

The same specification under three treatments of missing data – multiply imputed	Latinobarometro 2023, LAPOP 2023, LAPOP 2008	Latinobarometro roster	stats + mice (lm, mice, pool)	Table A11
--	--	------------------------	-------------------------------	-----------

One row per estimated equation. Where a specification was fitted to several outcomes or on several surveys, the dataset column says so and the row covers each fit. Control rosters are named rather than spelled out on every row: the Latinobarometro roster is economic perceptions, political interest, interpersonal trust, ideology, crime victimization, demographics (age, sex, education, interviewer-rated status, assets), capital and large-city residence; the LAPOP 2023 roster is economic perceptions, political interest, interpersonal trust, ideology, crime victimization, demographics (age, sex, education, income bracket, assets), urban residence; the LAPOP 2008 roster is ideology, crime victimization, police bribe solicitation, demographics (age cohort, sex, education, income quintile, household, ethnicity, religion, occupation), urban, capital and large-city residence; the Latinobarometro roster, less ideology is the Latinobarometro roster without ideology, which is the moderator here; the LAPOP 2008 roster, less ethnicity is the LAPOP 2008 roster without ethnicity, from which the moderator is derived; the Latinobarometro roster, reduced is economic perceptions, political interest, ideology, crime victimization, age, sex, education, interviewer-rated status; the LAPOP 2008 roster, less ideology is the LAPOP 2008 roster without ideology; the Latinobarometro roster, extended is the Latinobarometro roster plus perceived corruption, winner-loser status and presidential approval; the LAPOP 2023 roster, extended is the LAPOP 2023 roster plus perceived corruption and executive approval; the LAPOP 2008 roster, extended is the 2008 roster plus economic perceptions, political interest and interpersonal trust. The exact variables appear in the full-estimate tables, which list every coefficient. The last column is derived from the code rather than compiled by hand: each exhibit is followed back through the script that renders it, the files that script reads, and the script that wrote those, to the estimation calls made there. The model name, dataset and control roster are authored rather than derived, and checked against the code by hand. lmerTest is named rather than lme4 where a script attaches it, because it replaces lme4's lmer() to supply the Satterthwaite degrees of freedom the text reports. Four further models were estimated but are not reported here – an ordering test between the two trust measures, a graded response model for the trust battery, democracy support as an alternative outcome, and a probe of whether the country-specific slopes differ under sub-national clustering. They are in the replication materials and listed in the study README.

References

- Grömping, Ulrike. 2006. “Relative Importance for Linear Regression in R: The Package relaimpo.” *Journal of Statistical Software* 17(1): 1–27.
- Lupu, Noam, Mariana Rodríguez, and Elizabeth J. Zechmeister, eds. 2023. *Pulse of Democracy*. Nashville: LAPOP.
- Pustejovsky, James E., and Elizabeth Tipton. 2018. “Small-Sample Methods for Cluster-Robust Variance Estimation and Hypothesis Testing in Fixed Effects Models.” *Journal of Business & Economic Statistics* 36(4): 672–683.
- Rubin, Donald B. 1987. *Multiple Imputation for Nonresponse in Surveys*. New York: Wiley.
- VanderWeele, Tyler J., and Peng Ding. 2017. “Sensitivity Analysis in Observational Research: Introducing the E-Value.” *Annals of Internal Medicine* 167(4): 268–274.
- Viechtbauer, Wolfgang. 2010. “Conducting Meta-Analyses in R with the metafor Package.” *Journal of Statistical Software* 36(3): 1–48.