

Repressive Democracies: State Repression, Protest, and Policy Change

Ignacio Urbina*

Ph.D. Candidate in Political Science, Stony Brook University

September 17, 2026

Abstract

This paper studies how the cost of preparing for mass protest affects a democratic state's choice of repression. Citizens coordinate under uncertainty about the participation threshold required for policy change. The paper uses a psychological-rewards global game as its participation stage. It adds a state that values the policy status quo, plans operational resources for expected protest volume, and bears a convex political cost from authorizing expected harm per participant. The central result is that a higher marginal cost of preparing for protesters induces greater expected participant harm when the resources saved through deterrence exceed the additional operational cost of harsher tactics. Two police technologies show why the effect is positive in deployment-intensive systems and can be negative when the same inputs are important for implementing force. Democratic accountability enters through the marginal political-cost schedule. Stronger constraints reduce the state's chosen policy. A concise extension allows repression to change protest motivation. Repression backfires precisely when the elasticity of success-contingent motivation to expected harm exceeds one. A state that anticipates a stronger emotional response chooses less expected harm in the deployment-separable benchmark.

*Email: ignacio.urbina@stonybrook.edu. Website: www.ignacio-urbina.com.

1 Introduction.

Studies disagree about how state repression affects mobilization (Davenport 2007a). Some find that threats of repression deter mobilization by raising the cost of protest (Chenoweth, Perkoski, and Kang 2017; della Porta 1995; Christopher M. Sullivan 2016). Others find that repression can backfire by radicalizing victims of state violence (Adam-Troian, Çelebi, and Mahfud 2020; Christopher M. Sullivan and Davenport 2017). Repression and dissent become strategically interdependent when states choose the severity of repression (Chenoweth, Perkoski, and Kang 2017; Ritter 2014).

Scholars call this phenomenon the “protest-repression” nexus (Davenport 2007a; Curtice and Behlendorf 2021), and they have studied it primarily in autocracies. An empirical strand examines protest in democracies (Carey 2006; Nam 2007; Machado, Scartascini, and Tommasi 2011; Su 2015; Aytaç, Schiumerini, and Stokes 2017). It asks mainly what makes protest more or less likely across democratic settings. The literature has paid less attention to how severely states choose to repress protest. Competitive elections, protected speech, and organized civil society make protest a viable complement to representative channels, while legal and electoral accountability make the costs of coercion especially visible. Democratic protest policing therefore places deterrence, operational capacity, and political accountability in the same strategic problem.

Mass protests are costly for incumbent elites in democracies for two reasons. A massive protest may force politicians to change policies they prefer to preserve (Bosi, Giugni, and Uba 2016; Klein and Regan 2018). Gridlocked constitutional provisions that benefit elites are one example, including the tax code or the electoral system. Massive protests also require the state to deploy personnel, equipment, transport, detention capacity, and other operational resources.

To study this strategic problem, I develop a sequential model of policy-change protest. The elected executive first sets the expected repression cost borne by a participant. A continuum of citizens then decides whether to protest. The protest changes policy only if participation reaches an uncertain critical threshold, and each citizen observes a noisy private signal of that threshold. Participants receive a success-contingent expressive reward and bear the expected repression cost. The state values preservation of the disputed status quo. It plans operational resources for the expected volume of protest and bears a convex political cost from expected participant harm. This

structure places mobilization, deterrence, and the state's two distinct costs in a single strategic system.

Democratic institutions determine how likely a participant is to be victimized and how much physical or psychological damage state agents may inflict. These institutions include legal rules, standing procedures, and political oversight. Within that framework, elected executives and police commanders set deployments, authorized tactics, and rules of engagement for anticipated protest episodes. The model represents these choices by the expected repression cost announced or credibly anticipated before mobilization and fixed while an episode unfolds. Political leaders weigh its deterrent value against legal, electoral, administrative, and reputational consequences (Nordås and Davenport 2013; Ritter 2014).

The paper isolates one mechanism. A costly crowd gives the state an incentive to deter participation, but harsher policy creates political costs and can require additional operational resources. I represent repression by the expected cost borne by a participant. This measure combines the reach and intensity of coercion. The state's optimal policy balances the status quo it protects and the planned resources it saves by reducing expected protest volume against the political and operational costs of harsher tactics. I identify the precise technological condition under which more expensive protest preparation produces harsher repression.

The participation benchmark is inherited from Bueno de Mesquita and Shadmehr's (2023) psychological-rewards setting. A continuum of nonpivotal citizens receives private signals about the critical participation threshold, and a success-contingent reward generates strategic complementarity. The present paper changes the government's problem. Its payoff distinguishes the political value of the disputed status quo, the planned operational cost of expected protest volume, and the political cost of expected participant harm.

Table 1 identifies the exact analytic elements retained from the benchmark and the changes that generate this paper's contribution.

Table 1: Inherited benchmark and this paper’s contribution

Analytic element	Inherited from Bueno de Mesquita and Shadmehr (2023)	Treatment in this paper
Citizen stage	A continuum receives noisy private signals about an unknown participation threshold; success is determined by a monotone cutoff equilibrium	<i>Retained.</i> The same coordination game is applied to democratic protest for policy change
Marginal incentive	A nonpivotal citizen compares a success-contingent psychological reward with the individual cost of repression	<i>Reparameterized.</i> d_p is participation-contingent motivation; $c_r = q_r s_r$ is expected harm per participant
State timing	The state chooses repression before citizens mobilize	<i>Institutionally specified.</i> The executive chooses after the policy dispute is known and holds the directive fixed during mobilization
State payoff	Repression deters participation and carries a direct convex cost	<i>New payoff.</i> It separates the status quo stake B_s , planned operational cost $H(\bar{N}, c_r; \theta_0)$, and political cost $K(c_r; \theta_1)$
Central result	Optimal repression balances deterrence against the direct cost of repressive capacity	<i>New result.</i> A planning-cost shifter raises c_r^* exactly when the resources saved through deterrence exceed its direct effect on the cost of harsher tactics
Motivation extension	Success-contingent motivation is fixed	<i>New extension.</i> Repression may change motivation; backlash occurs above unit elasticity, and anticipated emotional responsiveness changes the state’s policy

The rest of the paper reviews the relevant protest-repression and coordination literatures, presents the participation stage, develops the state’s crowd-cost mechanism and its comparative statics, and discusses the paper’s theoretical and empirical contribution.

2 Literature review.

2.1 State repression, political attitudes, and social movements.

Research on state repression examines why governments coerce citizens and how coercion affects political behavior, especially protest. Political repression is the use of institutional or physical force, most commonly violent, to subdue or coerce individuals and groups that challenge political actors or the status quo (deMeritt 2016). Much of this research focuses on authoritarian regimes

or flawed democracies, where rulers face weak accountability and citizens lack protected speech or competitive elections.

Comparative research has persistently found that democratic countries exhibit lower levels of state repression (Davenport 2007a, 2007b; Davenport and Armstrong 2004). Yet violent coercion by state agents also occurs in developed democracies (Belur 2010; Whitehead, Fair, and Payne 2010; Worden 1995). Citizen backlash is not confined to force used against protesters. George Floyd was killed during an arrest, not while participating in a demonstration. The ensuing protests reduced favorability toward the police among low-prejudice and politically liberal Americans (Reny and Newman 2021). Police violence outside a protest can therefore mobilize citizens and raise the political costs that surround later police action. Victims of police repression may also experience anxiety, depression, anger, and fear (Bor et al. 2018; Çelebi, Adam-Troian, and Mahfud 2022; Cooper et al. 2004; J. E. DeVlyder et al. 2017; J. DeVlyder, Fedina, and Link 2020).

Yet not all citizens perceive the use of force by repressive bureaus the same way. In the United States, for example, studies consistently find that conservatives support the police at higher rates than liberals (Ekins 2016; Fingerhut 2017; Hindelang 1974; Zamble and Annesley 1987). Voters may perceive and evaluate state repression differently, to the extent that they have directional motivations and are affective partisans. Attitudes about the police are also racially divided (Ekins 2016; Fingerhut 2017).

Davenport (2007a) calls the conflicting evidence about how repression affects social movements the “punishment puzzle.” Some studies find that repression discourages mass and peaceful protest but encourages radicalization (della Porta 1995). Others find that it increases both protest escalation and mobilization (Adam-Troian, Çelebi, and Mahfud 2020; Christopher M. Sullivan and Davenport 2017).

2.2 Empirical regularities and theoretical foundations.

The literature has identified several determinants that systematically explain state repression, among them regime type, economic development, population size, intrastate civil conflicts, and the degree of grievances and dissent (Davenport 2007b; Nordås and Davenport 2013; Christopher M. Sullivan 2016; Christopher M. Sullivan and Davenport 2017). For example, Nordås and

Davenport (2013) argue that the demographic structure of the population is a key determinant of state repression. States respond with more repression where youth bulges are larger, meaning cohorts aged 15 to 24 relative to the overall adult population. States that anticipate more dissent raise repression preemptively, to deter it.

Chenoweth, Perkoski, and Kang (2017) summarize several consistent findings in the literature. First, dissent precedes state repression in most settings, a regularity Davenport (2007a) calls the “law of coercive responsiveness.” Second, democratic countries tend to exhibit lower levels of repression. The authors also discuss the “punishment puzzle” mentioned above. They stress that repression and dissent, whether protest or other tactics, are two complex interlinked processes.

Formal models explain the relationship between state repression, protests, and regime or policy change in several ways. No definite consensus has emerged about the best modeling approach. Ritter (2014) notes that models choose different objects of investigation, and that each one, taken alone, misses important strategic features of the relation between repression and dissent. For example, some studies model the determinants of regime change or dissent in general as a process of an ongoing or expected conflict. In one approach, citizens protest to alter the status quo, and their success depends on the state’s concessions. A revolutionary vanguard’s level of violence may signal the level of dissent in society, which provokes citizens to cooperate (Bueno de Mesquita 2010). Other formal models turn to uncertainty and asymmetric information about government types (Pierskalla 2010). A “strong” state may repress citizens at high cost to signal its resolve. Citizens protest over repeated play and learn what type of government they face. Closest to the present paper in timing, De Jaegher and Hoyer (2019) model *preemptive* repression. A government sets repressive capacity before protest occurs, and dissidents respond by investing in countermeasures. De Jaegher and Hoyer derive both sides of the punishment puzzle. Repression deters when the benefit of protest is large relative to the cost of countermeasures, and backfires when it is small. Their complete-information coordination game has two dissidents, each of whom can determine the outcome. The relative size of the equilibria’s basins of attraction resolves the resulting multiplicity.

Thus, as Ritter (2014) notes, some models study the conditions that ignite dissent, and others study how the state chooses the severity of repression. But studies miss an important part of the puzzle when they treat these two processes separately. She therefore proposes a model that

considers the conditions sparking dissent together with the state's choice of severity. She models the strategic setting as a bargaining problem over policy provision and its effect on incumbent survival. Ritter finds that a higher likelihood of leader survival decreases the level of dissent and repression. Yet the same increase raises the severity of both.

2.3 Global games of protest and regime change.

This paper borrows its analytic setup from the global games literature on coordination under incomplete information (Carlsson and van Damme 1993; Morris and Shin 1998, 2003). The canonical regime change game has a continuum of citizens who simultaneously decide whether to attack a regime. The regime collapses if and only if the mass of attackers exceeds its strength. That strength is an uncertain fundamental, and citizens receive noisy private signals about it. Under complete information about the fundamental, the game exhibits multiple equilibria in an intermediate range of regime strength. The incomplete information selects a unique equilibrium in threshold strategies.

Bueno de Mesquita and Shadmehr (2023) provide the closest application of this framework to repression. They develop “a standard global game of regime change.” A continuum of citizens receive noisy private signals $x_i = \theta + \sigma \varepsilon_i$ about the regime's strength θ , and the rebellion succeeds if and only if the fraction of rebels exceeds θ . Their psychological-rewards setting yields the threshold $1 - c$, where c is the participation cost, and their government chooses costly repressive capacity before citizens move. The present paper takes that participation stage as its benchmark. Its contribution begins with the government's payoff: the state values a particular policy status quo and separately accounts for the operational resources planned for expected protest volume and the political cost of expected participant harm. This distinction produces the crowd-management mechanism studied below.

The success-contingent reward also connects the model to work on selective incentives and democratic participation. Aytaç and Stokes (2019) explain participation through costs that citizens associate with abstention, including emotional responses and social obligations. Wood's (2003) pleasure-in-agency similarly locates a reward in contributing rather than in passively receiving a public good. These approaches distinguish participation-contingent motivation from the policy benefit shared by participants and nonparticipants. The distinction supplies the strate-

gic complementarity in the benchmark and clarifies the conditional role of the citizens' policy prize.

Little (2016) distinguishes two coordination problems that potential protesters face. The first is "political coordination," deciding whether to participate without knowing how many others are dissatisfied enough to join. The second is "tactical coordination," deciding when, where, and how to protest. His political coordination model also has a continuum of citizens with heterogeneous assessments and equilibria in cutoff strategies. He shows that better tactical information "almost always lowers the effective cost of protest, leading to larger (and more cohesive) antiregime action." This distinction matters for the present model because communication technology informs citizens about the critical threshold they face here. I return to the point in the discussion of measurement and research design.

3 Modeling protests, repression, and policy change.

In this section, I develop a model of citizens who decide to protest to change a disliked policy. The citizens face a coordination and collective action problem. The protest changes the policy only if citizens' participation reaches a given threshold. The threshold is an uncertain fundamental, and citizens receive noisy private signals of it. On the other side, the state would prefer not to change the policy. To influence participation, it can set the expected repression cost borne by a participant. A realized mass protest consumes operational resources, while a policy that imposes greater expected harm invites legal, electoral, and reputational backlash. I begin with the citizens' participation problem and its complete-information benchmark. Next, I solve the global game. Then, I propose the two-stage sequential game in which the state sets expected participant cost first, and I inspect the strategic implications of the model.

3.1 The citizens' participation problem.

There is a continuum of citizens of mass one, indexed by $i \in [0, 1]$. Each citizen chooses between participating in the protest (P) or not (NP). Let $n_i \in \{0, 1\}$ indicate participation, and let $N^* \in [0, 1]$ denote the realized mass of participants. Protesters want to change a policy so gridlocked that ordinary representative channels struggle to move it. The protest is successful if participation

Table 2: Notation and role in the model

Object	Meaning and units	Role
τ	Critical participating share required for policy change	Realized fundamental
N^*	Share of citizens who participate	Equilibrium outcome
B_p	Public policy benefit from successful protest	Citizen payoff primitive
d_p	Success-contingent participation reward	Citizen payoff primitive
q_r	Probability or share of participants targeted	Measured component of the directive
s_r	Harm conditional on being targeted	Measured component of the directive
$c_r = q_r s_r$	Expected repression cost per participant	State choice
B_s	Government value of preserving the policy status quo	State payoff primitive
$\tilde{N}(c_r)$	Expected share of citizens who participate	Anticipated outcome
$H(\tilde{N}, c_r; \theta_0)$	Planned operational cost for expected protest volume and the chosen repression policy	State cost technology
$K(c_r; \theta_1)$	Political cost of expected participant harm	State cost technology
σ	Noise in citizens' private signals	Information parameter

reaches a critical threshold τ , so the policy changes if and only if $N^* \geq \tau$.

The payoff primitives are as follows. A change in the policy provides a collective and indivisible benefit B_p to all individuals in the society (i.e., a public good). Individuals who participate may get an expressive or emotional benefit d_p . They may feel, for example, that fighting for a just cause is good in and of itself. Participants also pay c_r , the expected repression cost borne by a participant. This object combines the probability of victimization with the conditional harm suffered if victimized. If q_r denotes reach and s_r conditional severity, then $c_r = q_r s_r$.

Assumption 1 (One-dimensional directive). The executive chooses from an institution-specific, one-dimensional menu of policing directives indexed by $c_r \in [0, \infty)$. The menu assigns a single policy package to each index value, with reach $q_r : [0, \infty) \rightarrow [0, 1]$, conditional severity $s_r : [0, \infty) \rightarrow [0, \infty)$, and $q_r(c_r)s_r(c_r) = c_r$.

The directive is the policy unit; q_r and s_r are its measurable components. This restriction fits institutions in which rules of engagement bundle deployment, targeting, and authorized force into stable policy packages. The menu retains the cost-minimizing package whenever an institution makes several packages with the same product feasible, and it absorbs the minimized cost

into H . The functions H and K below evaluate costs along that feasible policy menu. Expected participant harm is therefore sufficient for the citizen's deterrence decision and indexes the packages compared by the executive. Conditional severity refers only to s_r .

Two features of the continuum setting must be made explicit, because they discipline how the primitives enter the participation margin.

First, each individual citizen is nonpivotal. With finitely many players, Δp_i would place the public good benefit B_p in the individual calculus. That term is the change in the probability of meeting the threshold induced by one's own participation. In the continuum, a single citizen has measure zero, so $\Delta p_i = 0$. The citizen receives B_p when the protest succeeds whether she participated or abstained, and B_p cancels from that comparison. The benchmark therefore isolates the participation-contingent reward as the motive of the marginal participant.

Second, the expressive benefit must be success-contingent for participation to respond to the prospects of success at all. A participant enjoys the expressive reward d_p only if the movement succeeds. That satisfaction arrives only when the change actually happens. This success-contingency is a substantive assumption, and it generates the coordination problem.¹ Were d_p unconditional, participation would reduce to the comparison of d_p against c_r , independent of any belief about others, and the strategic structure would disappear.

With these two features, the net payoff of participating relative to abstaining, for citizen i , is

$$u_i(P) - u_i(NP) = \mathbb{1}\{N^* \geq \tau\} \cdot d_p - c_r,$$

since $B_p \cdot \mathbb{1}\{N^* \geq \tau\}$ enters both sides and cancels. A citizen therefore participates when her belief that the protest will succeed, times the expressive reward, exceeds the expected repression cost.

3.2 The complete-information benchmark.

Suppose first, as a benchmark, that τ is common knowledge. Assume $0 < \tau \leq 1$, the interesting range, and $c_r < d_p$, so that participation is worthwhile when success is assured. Then two pure strategy equilibria coexist, and the small-noise limit selects a unique interior prediction (Carlsson and van Damme 1993; Morris and Shin 1998, 2003). If every citizen participates, $N^* = 1 \geq \tau$,

¹See Bueno de Mesquita and Shadmehr (2023) and Wood (2003).

the protest succeeds, and each participant receives $d_p - c_r > 0$. No citizen gains by deviating, since abstaining forfeits $d_p - c_r$ while leaving success (and hence B_p) unchanged. If no citizen participates, $N^* = 0 < \tau$, the protest fails, and a lone deviator would pay c_r for a failed protest. Universal abstention is therefore also an equilibrium. If instead $c_r > d_p$, universal abstention is the unique equilibrium, because participation is dominated even under assured success.

The complete-information benchmark presents two problems.² The multiplicity makes prediction difficult, and the equilibrium set is insensitive to the parameters, including the repression policy itself. One could respond by selecting an equilibrium, for instance the mixed strategy equilibrium, and reading its participation probabilities as frequencies. I instead take seriously that the critical threshold is a variable citizens cannot observe. The size a protest must reach to force a policy change is written nowhere, unlike the threshold to win an election under plurality rule, which is common knowledge. How citizens form beliefs about τ is therefore central to any behavioral prediction.

3.3 The global game.

Let the critical threshold τ be drawn from the common prior $U[0, 1]$ and remain unobserved by citizens. Citizens and the state use this same proper prior. Each citizen receives a noisy private signal

$$x_i = \tau + \sigma \varepsilon_i,$$

where the ε_i are i.i.d. with a log-concave density f and c.d.f. F with full support on $\mathbb{R}$, independent of τ , and $\sigma > 0$ scales the noise. A citizen with a low signal believes the required crowd is small and success likely. A citizen with a high signal believes success is out of reach.

A strategy is a mapping from signals to participation. The participation payoff is decreasing in the (believed) threshold, so cutoff strategies are the natural place to look for equilibria. A citizen participates if and only if $x_i < \hat{x}$. Given such a cutoff, the mass of participants at true threshold τ is

$$N^*(\tau) = \Pr(x_i < \hat{x} \mid \tau) = F\left(\frac{\hat{x} - \tau}{\sigma}\right),$$

which is continuous and strictly decreasing in τ . The protest therefore succeeds if and only if τ is

²The result mirrors Bueno de Mesquita and Shadmehr (2023, Proposition 1).

below a unique critical value $\hat{\tau}$. The equilibrium is characterized by two conditions:³

$$\text{(belief consistency / critical mass):} \quad F\left(\frac{\hat{x} - \hat{\tau}}{\sigma}\right) = \hat{\tau},$$

$$\text{(indifference of the marginal citizen):} \quad \Pr(\tau \leq \hat{\tau} \mid x_i = \hat{x}) \cdot d_p = c_r.$$

The first condition says that at the critical threshold the crowd is exactly critical. The second says that a citizen whose signal equals the cutoff is indifferent between participating and abstaining. Under the proper prior, the marginal citizen's exact posterior success probability is

$$\pi_\sigma(\hat{x}, \hat{\tau}) = \frac{F(\hat{x}/\sigma) - F((\hat{x} - \hat{\tau})/\sigma)}{F(\hat{x}/\sigma) - F((\hat{x} - 1)/\sigma)}.$$

Thus a positive-noise monotone cutoff equilibrium solves

$$F\left(\frac{\hat{x}_\sigma - \hat{\tau}_\sigma}{\sigma}\right) = \hat{\tau}_\sigma, \quad d_p \pi_\sigma(\hat{x}_\sigma, \hat{\tau}_\sigma) = c_r. \quad (1)$$

These equations retain the boundary terms generated by $U[0, 1]$. At an interior threshold they converge, as $\sigma \rightarrow 0$, to $(1 - \hat{\tau})d_p = c_r$. The limiting success threshold is therefore

$$\hat{\tau}(c_r) = \max\left\{1 - \frac{c_r}{d_p}, 0\right\}. \quad (2)$$

Normalized expected participant harm c_r/d_p determines the limiting threshold.⁴ Proposition 1 states the proper-prior convergence result.

Proposition 1 (Proper-prior small-noise selection). *Fix a policy with $\rho \equiv c_r/d_p \in (0, 1)$ and let $\hat{\tau}_0 = 1 - \rho$. For every $\delta \in (0, \min\{\hat{\tau}_0, 1 - \hat{\tau}_0\})$, there is a $\bar{\sigma} > 0$ such that equation (1) has an interior monotone cutoff equilibrium with $\hat{\tau}_\sigma \in [\hat{\tau}_0 - \delta, \hat{\tau}_0 + \delta]$ for every $\sigma \in (0, \bar{\sigma})$. Every selection from these local equilibria satisfies*

$$\hat{x}_\sigma \longrightarrow \hat{\tau}, \quad \hat{\tau}_\sigma \longrightarrow \hat{\tau} = 1 - \frac{c_r}{d_p}.$$

The interior monotone limiting selection is therefore unique. If $c_r > d_p$, participation is strictly dom-

³These conditions correspond to the canonical regime-change game in Bueno de Mesquita and Shadmehr (2023, equation 1).

⁴Equation (2) is the analog of $\theta^p = 1 - c$ in Bueno de Mesquita and Shadmehr (2023, equation 2 and Proposition 2), with c_r/d_p in the role of c .

inated. At $c_r = d_p$, full deterrence is the continuous extension of the interior small-noise selection, although other weak equilibria can coexist at equality. At $c_r = 0$, full participation is the corresponding limiting selection.

The full proof is in the appendix. Three remarks.

First, the selected interior branch exists for all sufficiently small positive noise and remains in a fixed compact neighborhood of $1 - c_r/d_p$. The bounded proper prior can also generate boundary-adjacent roots at positive noise. The theory follows the local branch that continues the interior global-games prediction. In the limiting game, protest succeeds if and only if $\tau \leq \hat{\tau}(c_r)$, so the state's comparative statics use its unique small-noise limit.

Second, the limiting outcome changes smoothly with the repression policy. The deterrence derivative is⁵

$$\frac{\partial \hat{\tau}}{\partial c_r} = -\frac{1}{d_p}.$$

Each unit of expected participant cost shrinks the set of fundamentals at which the protest succeeds at a constant rate, inversely proportional to the expressive reward. Repression is blunt against a strongly motivated citizenry (high d_p), and sharp against a weakly motivated one. This is the global games rendering of a familiar intuition: intrinsic motivation erodes the deterrent effect of expected participant harm.

Third, the direct deterrence margin is governed by d_p alone. The zero partial effect of B_p is a benchmark result conditional on three assumptions: citizens form a continuum, the policy prize is public and available equally to participants and abstainers, and policy stakes are separable from success-contingent motivation. A larger policy stake may still raise d_p by intensifying grievance, identity, efficacy, or pleasure-in-agency, and it may raise the government's B_s . The model isolates the direct pivotality channel by fixing those pathways.

The state's problem is evaluated in the small-noise limit $\sigma \rightarrow 0$.⁶ The cutoff $\hat{x}$ then converges to $\hat{\tau}$ and participation becomes sharp: $N^*(\tau) \rightarrow 1$ for $\tau < \hat{\tau}$ and $N^*(\tau) \rightarrow 0$ for $\tau > \hat{\tau}$. Expected protest

⁵The notation $\partial \hat{\tau}/\partial c_r$ denotes the change in the limiting equilibrium success threshold induced by a marginal change in expected participant cost, with d_p fixed. The expression applies on the interior region $0 < c_r < d_p$; at the clamped boundary the relevant one-sided derivative is zero.

⁶The executive's optimization problem below is defined over the limiting policy mapping $c_r \mapsto \hat{\tau}(c_r)$. The positive-noise calculations in the appendix illustrate the robustness of the selected citizen branch; they do not supply a separate convergence theorem for a finite-noise state optimizer.

mass is therefore $\bar{N}(c_r) \equiv \mathbb{E}_\tau[N^*(\tau) \mid c_r] = \hat{\tau}(c_r)$. The limiting environment captures the extensive margin between a mass protest and no mobilization and keeps the state's planning problem transparent.

3.4 The state's problem.

The elected executive knows the policy dispute and its stakes, then fixes an operational directive before citizens receive their private information and mobilize. Police implement that directive during the episode.

Assumption 2 (Within-episode commitment). The state chooses c_r after the policy dispute is known and before citizen mobilization, and c_r remains fixed during the protest episode.

The timing captures advance deployment, rules of engagement, authorized tactics, and expectations formed from directives and prior practice. The directive is publicly announced, legally recorded, or credible under standing practice before citizens mobilize. Legal review, administrative compliance, press scrutiny, and anticipatory electoral response can therefore impose a political cost when the directive is authorized, even if deterrence prevents its use.

Let $H(\bar{N}, c_r; \theta_0)$ denote the operational cost of preparing for expected protest mass $\bar{N}$ under policy c_r . It includes personnel, equipment, transport, command capacity, and other resources committed in anticipation of protest, including preparedness expenditures incurred before turnout is realized.

Assumption 3 (Planned operational cost). The function $H(\bar{N}, c_r; \theta_0)$ is twice continuously differentiable on $\bar{N} \in [0, 1]$ and $c_r \geq 0$, satisfies $H(0, 0; \theta_0) = 0$, $H_N > 0$, $H_c \geq 0$, and $H_{N\theta_0} > 0$.

The shifter θ_0 raises the marginal cost of preparing for additional protesters. The cross-partial $H_{c\theta_0}$ permits the same shifter to change the direct marginal operational cost of harsher tactics.⁷

The same subscript convention applies to the political-cost function.⁸

⁷Subscripts on H denote partial derivatives evaluated at $(\bar{N}, c_r; \theta_0)$. Thus $H_N \equiv \partial H / \partial \bar{N}$, $H_c \equiv \partial H / \partial c_r$, $H_{NN} \equiv \partial^2 H / \partial \bar{N}^2$, $H_{Nc} \equiv \partial^2 H / (\partial \bar{N} \partial c_r)$, $H_{cc} \equiv \partial^2 H / \partial c_r^2$, $H_{N\theta_0} \equiv \partial^2 H / (\partial \bar{N} \partial \theta_0)$, and $H_{c\theta_0} \equiv \partial^2 H / (\partial c_r \partial \theta_0)$. Every partial derivative fixes the function's remaining arguments.

⁸Thus $K_c \equiv \partial K(c_r; \theta_1) / \partial c_r$, $K_{cc} \equiv \partial^2 K(c_r; \theta_1) / \partial c_r^2$, and $K_{c\theta_1} \equiv \partial^2 K(c_r; \theta_1) / (\partial c_r \partial \theta_1)$. The function's other argument remains fixed in each derivative.

Assumption 4 (Political cost of the authorized directive). The function $K(c_r; \theta_1)$ is twice continuously differentiable on $c_r \geq 0$, satisfies $K(0; \theta_1) = 0$, $K_c(0; \theta_1) = 0$, $K_c(c_r; \theta_1) > 0$ for $c_r > 0$, $K_{cc}(c_r; \theta_1) > 0$, $K_{c\theta_1}(c_r; \theta_1) > 0$, and $K_c(c_r; \theta_1) \rightarrow \infty$ as $c_r \rightarrow \infty$.

Strict convexity captures rising marginal legal, electoral, administrative, and reputational costs of authorizing greater expected harm per participant. The cross-partial gives θ_1 a transparent institutional interpretation among democracies: a larger θ_1 raises the marginal political price of every positive c_r .

The state keeps $B_s > 0$ when the status quo survives; its policy payoff after protest-induced change is normalized to zero. With the proper planning distribution $\tau \sim U[0, 1]$, the success probability and expected protest mass in the small-noise limit both equal $\hat{\tau}(c_r)$. Expected utility is

$$u_S(c_r) = (1 - \hat{\tau}(c_r))B_s - H(\hat{\tau}(c_r), c_r; \theta_0) - K(c_r; \theta_1). \quad (3)$$

For $c_r \in (0, d_p)$, the marginal payoff is⁹

$$u_c(c_r) = \frac{B_s + H_N}{d_p} - H_c - K_c. \quad (4)$$

The term $(B_s + H_N)/d_p$ values the status quo and planned operational resources saved when deterrence reduces expected protest mass. The term H_c is the direct marginal operational cost of implementing harsher tactics. The term K_c is the marginal political cost of expected participant harm.

Assumption 5 (Interior curvature). The objective is twice continuously differentiable and strictly concave on $(0, d_p)$, with $u_{cc}(c_r) < 0$ and right derivative $u_c(0^+) > 0$.

Using equation (4), strict concavity is equivalent to

$$\frac{H_{NN}}{d_p^2} - \frac{2H_{Nc}}{d_p} + H_{cc} + K_{cc} > 0.$$

This condition allows planned operational cost to be nonlinear in expected protest mass and expected participant harm while keeping a unique policy response.

⁹Here $u_c \equiv \partial u_S / \partial c_r$ and $u_{cc} \equiv \partial^2 u_S / \partial c_r^2$. These derivatives vary the state's policy c_r and allow the induced expected mass $\tilde{N}(c_r) = \hat{\tau}(c_r)$ to adjust, while fixing $B_s, d_p, \theta_0, \theta_1$, and the functional forms of H and K .

Proposition 2 (The optimal repression policy). *Under Assumptions 2–5, the state has a unique optimal policy on $[0, d_p]$. The solution is interior and satisfies $u_c(c_r^*) = 0$ when the left derivative $u_c(d_p^-) < 0$; otherwise the state chooses the full-deterrence corner $c_r^* = d_p$. Expected harm above d_p is dominated because the success threshold and expected protest mass are already clamped at zero while H and K are weakly increasing in c_r .*

3.5 Comparative statics.

The general technology separates effects that follow directly from the model from the condition governing the crowd-cost mechanism.

Proposition 3 (Comparative statics). *At the interior solution,¹⁰*

$$\begin{aligned} \frac{\partial c_r^*}{\partial B_s} &= -\frac{1/d_p}{u_{cc}} > 0, & \frac{\partial c_r^*}{\partial \theta_1} &= \frac{K_{c\theta_1}}{u_{cc}} < 0, \\ \frac{\partial c_r^*}{\partial d_p} &= \frac{B_s + H_N + c_r^* H_{Nc} - (c_r^*/d_p) H_{NN}}{d_p^2 u_{cc}}, & \frac{\partial c_r^*}{\partial B_p} &= 0, \end{aligned}$$

Thus $\partial c_r^*/\partial d_p < 0$ when $B_s + H_N + c_r^* H_{Nc} > (c_r^*/d_p) H_{NN}$. For the crowd-cost shifter,

$$\frac{\partial c_r^*}{\partial \theta_0} = -\frac{H_{N\theta_0}/d_p - H_{c\theta_0}}{u_{cc}}. \quad (5)$$

Hence a higher marginal cost of preparing for protesters increases expected participant harm if and only if

$$\frac{H_{N\theta_0}}{d_p} > H_{c\theta_0}. \quad (\text{Crowd-cost condition})$$

A rise in B_s lowers the equilibrium success probability, whereas a rise in θ_1 raises it. When the crowd-cost condition holds, a rise in θ_0 lowers the success probability. A rise in d_p raises it whenever the stated condition for $\partial c_r^*/\partial d_p < 0$ holds.

The crowd-cost condition is the paper's central technological result. The term $H_{N\theta_0}/d_p$ is the additional planning cost saved when deterrence reduces expected protest mass. The term $H_{c\theta_0}$

¹⁰Each expression $\partial c_r^*/\partial z$ or $\partial \hat{c}_r^*/\partial z$ is an equilibrium comparative-static derivative. It allows the optimizing policy, expected protest mass, and success threshold to adjust to primitive z while the other listed primitives and both cost-function families remain fixed. All derivatives of H on the right-hand side are evaluated at $(\bar{N}, c_r; \theta_0) = (\hat{c}_r^*, c_r^*; \theta_0)$; derivatives of K are evaluated at $(c_r^*; \theta_1)$.

captures how the same shifter changes the direct marginal operational cost of harsher tactics. The first force raises the state's optimal expected harm per participant; the second tempers that increase and can reverse it.

Two organizational technologies give the condition substantive content. Let w denote the price of personnel and logistical inputs, and let $A(c_r)$ collect operational costs that do not vary with w . In a *deployment-separable* system,

$$H^D(\tilde{N}, c_r; w) = w\tilde{N} + A(c_r). \quad (6)$$

The input price raises the marginal cost of preparing for an additional protester, $H_{Nw}^D = 1$, without changing the direct marginal cost of the directive, $H_{cw}^D = 0$. Thus $\partial c_r^*/\partial w > 0$: higher personnel, transport, or crowd-control deployment costs make deterrence more valuable.¹¹

In a *tactic-intensive* system, the same inputs also implement the authorized level of expected harm:

$$H^T(\tilde{N}, c_r; w) = w(\tilde{N} + \lambda c_r) + A(c_r), \quad \lambda > 0. \quad (7)$$

Now $H_{Nw}^T = 1$ and $H_{cw}^T = \lambda$. The price increase raises expected participant harm when $\lambda < 1/d_p$ and lowers it when $\lambda > 1/d_p$.¹² The second case describes organizations in which specialized officers, less-lethal munitions, medical support, detention processing, or documentation requirements make stronger directives especially input-intensive. Equations (6) and (7) yield a discriminating implication: comparable input-price increases lead the state to authorize more expected participant harm where they load mainly on personnel deployment and less expected participant harm where they load heavily on implementing force.

The institutional comparison follows directly from $K_{c\theta_1} > 0$. Environments with stronger legal liability, judicial oversight, electoral accountability, or media scrutiny raise the marginal political price of expected harm and therefore reduce c_r^* . This gives the democratic setting a formal role through the cost schedule faced by the elected executive.

¹¹The derivative $\partial c_r^*/\partial w$ allows the optimal directive and anticipated protest mass to adjust to the input price w ; B_s , d_p , θ_1 , and the functional forms of A and K remain fixed. The other arguments of H remain fixed in the cross-partial H_{Nw} and H_{cw} .

¹²The loading λ converts one unit of expected participant harm into the same operational-input units used to measure $\tilde{N}$; its units are therefore units of anticipated participation per unit of harm. Because d_p and c_r share payoff units, $1/d_p$ has the same units as λ . The numerical cutoff is tied to those normalizations and is not directly portable across differently scaled data.

For the transparent benchmark, use

$$H(\bar{N}, c_r; \theta_0) = \theta_0 \bar{N}, \quad K(c_r; \theta_1) = \theta_1 c_r^2.$$

This is the deployment-separable technology with $A(c_r) = 0$. Then $H_{N\theta_0} = 1$, $H_{c\theta_0} = 0$, the crowd-cost condition holds strictly, and the solution reduces to

$$c_r^* = \min \left\{ \frac{B_s + \theta_0}{2\theta_1 d_p^2}, d_p \right\}, \quad \hat{t}^* = \max \left\{ 1 - \frac{B_s + \theta_0}{2\theta_1 d_p^2}, 0 \right\}.$$

Figure 1 uses this additive-quadratic specialization to show the mechanism.




Figure 1: The crowd-management mechanism in the additive-quadratic benchmark, $H(\bar{N}, c_r; \theta_0) = \theta_0 \bar{N}$ and $K(c_r; \theta_1) = \theta_1 c_r^2$, with $\theta_1 = 1$ and $d_p = 3$. Solid lines show the effective policy for three values of the government's status quo stake; dotted lines show the unconstrained interior solution. A higher marginal cost of preparing for protesters raises expected participant harm until the full-deterrence cap binds.

In the benchmark, a higher θ_0 raises expected participant harm. It increases the planned resources saved by reducing expected protest mass without changing the direct marginal operational cost of harsher tactics. A higher θ_1 lowers c_r by steepening the marginal political-cost schedule. Greater intrinsic motivation d_p lowers the interior policy because motivated citizens are

harder to deter. At the full-deterrence corner, c_r tracks d_p one for one.

Two features of these results deserve emphasis.

First, the status quo prize B_s has an unambiguously positive effect at every interior optimum. The monotone cutoff equilibrium makes the deterrence technology the linear relation (2), so the marginal value of expected participant harm rises with the government's stake in the disputed policy.

Second, the direct partial effect of B_p equals zero under the benchmark's nonpivotality and separability assumptions. Policy stakes affect repression through the participation reward d_p and the government's status quo value B_s . The relevant distinction is between the prize's direct public-good component and the motivational and governmental channels through which politically important disputes alter behavior.

3.6 Directive costs and exposure-dependent political costs.

The baseline $K(c_r; \theta_1)$ is paid when the executive authorizes an observable directive. Democratic institutions may also impose consequences after citizens are exposed to it. In the small-noise limit, total expected harm in the population is

$$T(c_r) = c_r \hat{\tau}(c_r).$$

This aggregation uses the benchmark's homogeneous expected harm per participant and sharp small-noise participation.¹³ Let $L(T; \theta_2)$ be an increasing, convex political cost of exposure, in addition to the directive cost K . The state then maximizes

$$u_S^E(c_r) = (1 - \hat{\tau})B_s - H(\hat{\tau}, c_r; \theta_0) - K(c_r; \theta_1) - L(c_r \hat{\tau}; \theta_2).$$

¹³At positive noise with homogeneous exposure, the corresponding object is $c_r \mathbb{E}_\tau[N^*(\tau)]$. With heterogeneous exposure c_{ri} , it is $\mathbb{E}_\tau[\int_0^1 n_i(\tau) c_{ri} di]$. Casualty, injury, or detention totals therefore map to T only after accounting for turnout and the distribution of exposure among participants.

At an interior policy its marginal payoff is¹⁴

$$u_c^E(c_r) = \frac{B_s + H_N}{d_p} - H_c - K_c - L_T \left(\hat{t} - \frac{c_r}{d_p} \right). \quad (8)$$

Exposure-dependent accountability changes the optimal level because total expected harm first rises and then falls as deterrence reduces the crowd. Differentiating equation (8) in the planning-cost shifter again gives $H_{N\theta_0}/d_p - H_{c\theta_0}$ when the extended objective remains strictly concave and L does not depend on θ_0 . The crowd-cost condition therefore continues to determine the sign of $\partial c_r^*/\partial \theta_0$. The directive cost remains positive at full deterrence, while the exposure cost reaches zero with expected turnout.

3.7 An extension: Repression, protests, and emotions.

The benchmark holds success-contingent motivation fixed. To isolate one feedback, let citizens observe the directive c_r before choosing whether to participate and let $d_p(c_r) = d_0 + e(c_r)$. The observable object can be a public rule, announced tactic, or credible expectation formed from prior police practice. The threshold becomes

$$\hat{t}(c_r) = \max \left\{ 1 - \frac{c_r}{d_0 + e(c_r)}, 0 \right\}.$$

At an interior threshold, differentiation gives¹⁵

$$\frac{\partial \hat{t}}{\partial c_r} = - \frac{d_0 + e(c_r) - c_r e'(c_r)}{(d_0 + e(c_r))^2}.$$

Define the elasticity of success-contingent motivation to expected participant harm as $\varepsilon(c_r) = c_r e'(c_r)/d_p(c_r)$.

Proposition 4 (Deterrence and backlash). *Expected participant harm deters, $\partial \hat{t}/\partial c_r < 0$, if and only if $\varepsilon(c_r) < 1$; it backfires, $\partial \hat{t}/\partial c_r > 0$, if and only if $\varepsilon(c_r) > 1$.*

The extension converts the punishment puzzle into a scale-free condition. Anticipated repres-

¹⁴Here $L_T \equiv \partial L(T; \theta_2)/\partial T$, evaluated at $T = c_r \hat{t}(c_r)$. The derivative varies c_r , allows $\hat{t}(c_r)$ and total expected harm to adjust, and fixes $B_s, d_p, \theta_0, \theta_1, \theta_2$, and the cost-function families. Since $\partial \hat{t}/\partial c_r = -1/d_p$, $\partial(c_r \hat{t})/\partial c_r = \hat{t} - c_r/d_p$.

¹⁵The derivative $\partial \hat{t}/\partial c_r$ varies the observable policy c_r and fixes d_0 and the response function $e(\cdot)$; $e'(c_r)$ is the ordinary derivative of that function evaluated at c_r .

sion deters when motivation rises less than proportionally with expected harm and recruits when it rises more than proportionally.

The emotional response also changes the directive the state is willing to choose. Index the strength of that response by $\alpha \geq 0$ and write $d_p(c_r; \alpha) = d_0 + \alpha e(c_r)$, so $\alpha = 1$ recovers the preceding specification. At an interior threshold, define the marginal deterrent return

$$D(c_r, \alpha) \equiv -\frac{\partial \hat{t}(c_r; \alpha)}{\partial c_r} = \frac{d_p(c_r; \alpha) - c_r \alpha e'(c_r)}{d_p(c_r; \alpha)^2} = \frac{1 - \varepsilon(c_r; \alpha)}{d_p(c_r; \alpha)}, \quad (9)$$

where $\varepsilon(c_r; \alpha) = c_r \alpha e'(c_r) / d_p(c_r; \alpha)$.¹⁶

Proposition 5 (Emotions and the state's policy). *Suppose $e(c_r) \geq 0$ and $e'(c_r) \geq 0$, and let $c_r^*(\alpha) > 0$ be an interior optimum with $\hat{t}(c_r^*; \alpha) \in (0, 1)$ at which the state's objective is strictly concave. First, any such optimum lies in the deterrence region: $\varepsilon(c_r^*; \alpha) < 1$. Second, under the deployment-separable technology $H^D(\bar{N}, c_r; w) = w\bar{N} + A(c_r)$, a stronger nonzero emotional response reduces the state's chosen expected harm:*

$$\frac{\partial c_r^*}{\partial \alpha} < 0 \quad \text{if} \quad e(c_r^*) + c_r^* e'(c_r^*) > 0.$$

A state does not choose a locally backfiring directive at an interior optimum: such a directive increases expected protest while also carrying operational and political costs. Even before backfire occurs, stronger emotional responsiveness weakens the marginal deterrent return. In a deployment-separable system, the state responds by selecting less expected harm. Evidence that mobilization follows a coercive tactic is informative about this mechanism when the tactic causally changes both expected participant harm and success-contingent motivation. The timing concerns an observable ex ante policy. Across repeated episodes, prior realized violence can determine citizens' current expectations of that policy.

¹⁶The partial derivative $\partial \hat{t}(c_r; \alpha) / \partial c_r$ varies the directive while fixing d_0 , α , and the response function $e(\cdot)$. In the proof, $D_\alpha \equiv \partial D / \partial \alpha$, $G_\alpha \equiv \partial G / \partial \alpha$, and $G_{c_r} \equiv \partial G / \partial c_r$ hold the other displayed argument fixed. The derivative $\partial c_r^* / \partial \alpha$ varies the strength of the emotional response while allowing the state to reoptimize; B_s , w , θ_1 , $A(\cdot)$, $K(\cdot)$, and $e(\cdot)$ remain fixed.

4 Discussion: contribution, observables, and research design.

The theory distinguishes the planned operational cost of expected protest volume from the political cost of authorizing expected harm per participant. These costs affect the state's choice differently. A higher marginal cost of preparing for protesters makes deterrence more valuable, whereas the marginal legal, electoral, administrative, and reputational costs of the directive restrain the policy used to deter. The derived crowd-cost condition determines the direction of the crowd-cost effect when the same cost environment also changes the direct marginal operational cost of harsher tactics. Expected participant harm rises only when the planning resources saved through deterrence exceed that additional marginal operational cost. Equation (8) adds political consequences from total expected harm while preserving this sign condition.

The model assigns separate roles to three political quantities. The government's B_s is the value of preserving the disputed status quo. Citizens' d_p is the participation-contingent reward from contributing to a successful movement. The public policy benefit B_p has no direct pivotality effect in the continuum benchmark, but politically important disputes can affect both d_p and B_s . Klein and Regan's (2018) cross-national protest-event analysis provides suggestive evidence about status quo stakes. They find that high concession costs, meaning demands that are expensive for the ruling coalition to accept, are associated with coercive state responses. This pattern is consistent with the positive effect of B_s . They also find that high disruption costs encourage accommodation. That result points in the opposite direction from the additive crowd-cost benchmark, but it concerns economic and public-order losses tied to a protest's location, duration, and size. The model's θ_0 instead changes the marginal operational cost of preparing for an additional participant. The contrast shows why the two costs should not be combined into a single measure of how expensive a protest is for the government.

The observable counterpart of c_r combines reach and intensity. Let q_r be the probability that a participant is targeted and s_r the harm conditional on targeting. An event-level measure of expected participant harm joins participant exposure to evidence about deaths, injuries, detention, sexual violence, psychological trauma, and other coercive tactics, so that $c_r = q_r s_r$. Under Assumption 1, observed directives trace an institution-specific policy menu. The executive does not choose q_r and s_r independently. Measuring both components constructs the deterrence index

and reveals how police organizations package reach and conditional severity. Castro (2025) shows the value of tactic-specific information by coding civilian complaints during Chile’s 2019 protest cycle by tactic, municipality, and day and validating them against media reports and police ammunition records. She finds that selective high-severity tactics generated short-run mobilization and that low-severity tactics increased protest regardless of selectivity. Those results are consistent with the backfire region of Proposition 4 for the observed tactic combinations, although they do not estimate the model’s motivational elasticity.

Estimating the operational cost $H(\bar{N}, c_r; \theta_0)$ requires the government’s anticipated protest volume and the resources committed before an episode. The relevant inputs include scheduled personnel, overtime capacity, transport, equipment and consumables, temporary command infrastructure, detention capacity, and medical support. Benchmarking their prices against comparable labor and procurement costs measures opportunity cost across jurisdictions. Budget documents, deployment plans, procurement records, and after-action accounts can distinguish standing capacity from resources added in anticipation of protest. They can also estimate the loading λ in equation (7): input-price changes that mainly raise deployment cost predict greater expected participant harm, while price changes that strongly raise the cost of implementing force predict restraint. Within-jurisdiction changes are especially informative because ordinary crime also affects resource demand and police capacity.

Institutions shift $K(c_r; \theta_1)$ by changing the consequences of authorizing a directive for elected officials. Published rules of engagement, judicial preclearance, legislative oversight, internal compliance review, and press access expose the authorization itself to scrutiny. Civil liability, casualty investigations, and electoral reactions to documented injuries instead load on $L(T; \theta_2)$. Aytaç, Schiumerini, and Stokes (2017) study backlash movements in Turkey, Brazil, and Ukraine and find that governments with stable electoral support were more willing to use harsh measures, whereas electorally insecure governments had stronger incentives to reduce violence. This pattern is consistent with the model’s claim that a lower marginal electoral cost permits greater expected participant harm. At a broader level, Davenport and Armstrong (2004) find that democracy reduces human-rights violations only above a threshold of democratic development. Their country-level outcome is not specific to protest directives, but the nonlinear result is consistent with interpreting θ_1 as meaningful institutional variation among democracies. Research designs can use the

timing of these institutions to distinguish changes in the marginal directive cost K_c from changes in the marginal exposure cost L_T .

Existing evidence also points to the motivational channel, while showing that participants need not respond uniformly. Adam-Troian, Çelebi, and Mahfud's (2020) cross-sectional study of 523 Yellow Vest protesters finds that exposure to police violence is positively associated with intentions to attend future demonstrations and to sacrifice for the movement. Perceived loss of significance and identification with the Yellow Vests mediate those associations. Sullivan and Davenport's (2017) panel study of the Republic of New Africa finds organizational backlash but different individual responses: members who personally experienced repression became more likely to participate, while members who were not directly repressed withdrew. These findings are consistent with a positive motivational response for some citizens, but they also caution against treating $e(c_r)$ as homogeneous. Survey and experimental data can measure d_p and the elasticity in Proposition 4 more directly. Items can measure political efficacy, anticipated satisfaction from contributing to collective success, and willingness to participate under stated probabilities of success and victimization. A factorial design can vary victimization reach, conditional harm, and the probability of policy success separately. The resulting responses distinguish the direct deterrent effect of expected harm from any change in success-contingent motivation.

Event records can first describe the joint variation in anticipated turnout, planned operational inputs, coercive reach, conditional harm, and institutional constraints. Suitable designs can then use institutional reforms and predetermined changes in labor or procurement costs to isolate particular margins. The comparative statics specify the expected pattern. Status quo stakes raise expected participant harm, and stronger marginal accountability lowers it. A planning-cost shifter raises expected harm only where the resources saved through deterrence exceed its effect on the marginal operational cost identified by the crowd-cost condition.

5 Conclusion.

This paper explains why a government that finds large protests expensive to manage may threaten harsher repression before people mobilize. The logic has two steps. Expected harm discourages some citizens from joining. A smaller expected crowd lets the government plan fewer officers, ve-

hicles, detention spaces, and supplies. Harsher repression becomes attractive when the resources saved by preparing for a smaller crowd exceed the added operational and political costs of authorizing force. What matters is the cost of managing one more protester. The total price of maintaining a police force is a different quantity.

This tradeoff also explains why governments facing similar protests can respond differently. Consider an economic downturn that cuts tax revenue while raising demand for unemployment benefits, emergency transfers, and other social spending. The resulting fiscal pressure leaves less money for police overtime, recruitment, and equipment. If officials still expect a large protest, they may be unable to expand deployment as they otherwise would. Where revising the use-of-force directive is comparatively cheap, the executive may compensate by authorizing tactics that expose each participant to greater expected harm, hoping that stronger deterrence allows a smaller force to manage the episode. The result still depends on how the police organization produces force. If stronger tactics require specialized officers, equipment, documentation, medical support, or detention capacity, the budget pressure can instead produce restraint. Legal review, political oversight, and electoral accountability further change the price that an elected government pays when it authorizes expected harm. Repression depends on which resources have become scarce and on how the directive uses them.

Citizens complete the argument because the crowd is not fixed in advance. Each person weighs the expected harm of participation against the satisfaction of helping a protest succeed. A stronger directive usually discourages participation and makes policy change less likely. Anger can reverse that result, but only when the added motivation grows faster than the expected harm. Repression can therefore anger citizens without producing a larger protest; anger and backfire are separate outcomes. When the state anticipates a stronger emotional response, it chooses less expected harm in the deployment-separable benchmark.

The argument points to a dynamic question for future research. Governments announce directives and plan deployments before an episode, but commanders and citizens also learn after turnout and injuries become visible. Comparing the initial directive with later tactical revisions would show when governments use an early threat to avoid a costly crowd, when operational demands force restraint, and when observed violence creates the mobilization that the directive was meant to prevent.

References

- [1] Adam-Troian, Jais, Elif Çelebi, and Yara Mahfud. 2020. “‘Return of the Repressed’: Exposure to Police Violence Increases Protest and Self-Sacrifice Intentions for the Yellow Vests.” *Group Processes & Intergroup Relations* 23(8): 1171–86. <https://doi.org/10.1177/1368430220920707>.
- [2] Aytaç, S. Erdem, Luis Schiumerini, and Susan Stokes. 2017. “Protests and Repression in New Democracies.” *Perspectives on Politics* 15(1): 62–82. <https://doi.org/10.1017/S1537592716004138>.
- [3] Aytaç, S. Erdem, and Susan C. Stokes. 2019. *Why Bother? Rethinking Participation in Elections and Protests*. Cambridge: Cambridge University Press.
- [4] Belur, Jyoti. 2010. “Police and the Use of Deadly Force.” In *Permission to Shoot? Police Use of Deadly Force in Democracies*, ed. Jyoti Belur. New York: Springer, 1–20. https://doi.org/10.1007/978-1-4419-0975-6_1.
- [5] Bor, Jacob, Atheendar S. Venkataramani, David R. Williams, and Alexander C. Tsai. 2018. “Police Killings and Their Spillover Effects on the Mental Health of Black Americans: A Population-Based, Quasi-Experimental Study.” *The Lancet* 392(10144): 302–10. [https://doi.org/10.1016/S0140-6736\(18\)31130-9](https://doi.org/10.1016/S0140-6736(18)31130-9).
- [6] Bosi, Lorenzo, Marco Giugni, and Katrin Uba. 2016. “The Consequences of Social Movements: Taking Stock and Looking Forward.” In *The Consequences of Social Movements*, eds. Lorenzo Bosi, Marco Giugni, and Katrin Uba. Cambridge: Cambridge University Press, 3–38.
- [7] Bueno de Mesquita, Ethan. 2010. “Regime Change and Revolutionary Entrepreneurs.” *American Political Science Review* 104(3): 446–66. <https://doi.org/10.1017/S0003055410000274>.
- [8] Bueno de Mesquita, Ethan, and Mehdi Shadmehr. 2023. “Rebel Motivations and Repression.” *American Political Science Review* 117(2): 734–50. <https://doi.org/10.1017/S0003055422000600>.

- [9] Carey, Sabine C. 2006. "The Dynamic Relationship between Protest and Repression." *Political Research Quarterly* 59(1): 1–11. <https://doi.org/10.1177/106591290605900101>.
- [10] Carlsson, Hans, and Eric van Damme. 1993. "Global Games and Equilibrium Selection." *Econometrica* 61(5): 989–1018. <https://doi.org/10.2307/2951491>.
- [11] Castro, Francisca. 2025. "Deterrence or Defiance? How Severity and Selectivity Shape Protest Responses to Repression." *Mobilization: An International Quarterly* 30(4): 387–408. <https://doi.org/10.17813/1086-671X-30-4-387>.
- [12] Çelebi, Elif, Jais Adam-Troian, and Yara Mahfud. 2022. "Positive Links between Exposure to Police Violence, PTSD, and Depression Symptoms among Yellow Vests Protesters in France." *Journal of Interpersonal Violence* 37(5–6): 2267–88. <https://doi.org/10.1177/0886260520935863>.
- [13] Chenoweth, Erica, Evan Perkoski, and Sooyeon Kang. 2017. "State Repression and Nonviolent Resistance." *Journal of Conflict Resolution* 61(9): 1950–69. <https://doi.org/10.1177/0022002717721390>.
- [14] Cooper, Hannah, Lisa Moore, Sofia Gruskin, and Nancy Krieger. 2004. "Characterizing Perceived Police Violence: Implications for Public Health." *American Journal of Public Health* 94(7): 1109–18. <https://doi.org/10.2105/AJPH.94.7.1109>.
- [15] Curtice, Travis B., and Brandon Behlendorf. 2021. "Street-Level Repression: Protest, Policing, and Dissent in Uganda." *Journal of Conflict Resolution* 65(1): 166–94. <https://doi.org/10.1177/0022002720939304>.
- [16] Davenport, Christian. 2007a. "State Repression and Political Order." *Annual Review of Political Science* 10: 1–23. <https://doi.org/10.1146/annurev.polisci.10.101405.143216>.
- [17] ———. 2007b. *State Repression and the Domestic Democratic Peace*. Cambridge: Cambridge University Press.
- [18] Davenport, Christian, and David A. Armstrong. 2004. "Democracy and the Violation of Human Rights: A Statistical Analysis from 1976 to 1996." *American Journal of Political Science* 48(3): 538–54. <https://doi.org/10.1111/j.0092-5853.2004.00086.x>.

- [19] De Jaegher, Kris, and Britta Hoyer. 2019. "Preemptive Repression: Deterrence, Backfiring, Iron Fists, and Velvet Gloves." *Journal of Conflict Resolution* 63(2): 502–27. <https://doi.org/10.1177/0022002717750450>.
- [20] deMeritt, Jacqueline H. R. 2016. "The Strategic Use of State Repression and Political Violence." *Oxford Research Encyclopedia of Politics*. Oxford University Press. <https://doi.org/10.1093/acrefore/9780190228637.013.32>.
- [21] DeVyllder, Jordan E. et al. 2017. "Prevalence, Demographic Variation and Psychological Correlates of Exposure to Police Victimization in Four US Cities." *Epidemiology and Psychiatric Sciences* 26(5): 466–77. <https://doi.org/10.1017/S2045796016000810>.
- [22] DeVyllder, Jordan, Lisa Fedina, and Bruce Link. 2020. "Impact of Police Violence on Mental Health: A Theoretical Framework." *American Journal of Public Health* 110(11): 1704–10. <https://doi.org/10.2105/AJPH.2020.305874>.
- [23] Ekins, Emily E. 2016. *Policing in America: Understanding Public Attitudes toward the Police*. Washington, DC: Cato Institute.
- [24] Fingerhut, Hannah. 2017. *Deep Racial, Partisan Divisions in Americans' Views of Police Officers*. Washington, DC: Pew Research Center.
- [25] Gehlbach, Scott, Konstantin Sonin, and Milan W. Svobik. 2016. "Formal Models of Nondemocratic Politics." *Annual Review of Political Science* 19: 565–84. <https://doi.org/10.1146/annurev-polisci-042114-014927>.
- [26] Ginkel, John, and Alastair Smith. 1999. "So You Say You Want a Revolution: A Game Theoretic Explanation of Revolution in Repressive Regimes." *Journal of Conflict Resolution* 43(3): 291–316. <https://doi.org/10.1177/0022002799043003002>.
- [27] Hindelang, Michael J. 1974. "Public Opinion Regarding Crime, Criminal Justice, and Related Topics." *Journal of Research in Crime and Delinquency* 11(2): 101–16. <https://doi.org/10.1177/002242787401100202>.
- [28] Klein, Graig R., and Patrick M. Regan. 2018. "Dynamics of Political Protests." *International Organization* 72(2): 485–521. <https://doi.org/10.1017/S0020818318000061>.

- [29] Little, Andrew T. 2016. "Communication Technology and Protest." *The Journal of Politics* 78(1): 152–66. <https://doi.org/10.1086/683187>.
- [30] Machado, Fabiana, Carlos Scartascini, and Mariano Tommasi. 2011. "Political Institutions and Street Protests in Latin America." *Journal of Conflict Resolution* 55(3): 340–65. <https://doi.org/10.1177/0022002711400864>.
- [31] Morris, Stephen, and Hyun Song Shin. 1998. "Unique Equilibrium in a Model of Self-Fulfilling Currency Attacks." *American Economic Review* 88(3): 587–97.
- [32] Morris, Stephen, and Hyun Song Shin. 2003. "Global Games: Theory and Applications." In *Advances in Economics and Econometrics: Theory and Applications, Eighth World Congress*, eds. Mathias Dewatripont, Lars Peter Hansen, and Stephen J. Turnovsky. Cambridge: Cambridge University Press, 56–114.
- [33] Nam, Taehyun. 2007. "Rough Days in Democracies: Comparing Protests in Democracies." *European Journal of Political Research* 46(1): 97–120. <https://doi.org/10.1111/j.1475-6765.2006.00645.x>.
- [34] Nordås, Ragnhild, and Christian Davenport. 2013. "Fight the Youth: Youth Bulges and State Repression." *American Journal of Political Science* 57(4): 926–40. <https://doi.org/10.1111/ajps.12025>.
- [35] Pierskalla, Jan Henryk. 2010. "Protest, Deterrence, and Escalation: The Strategic Calculus of Government Repression." *Journal of Conflict Resolution* 54(1): 117–45. <https://doi.org/10.1177/0022002709352462>.
- [36] della Porta, Donatella. 1995. "Social Movements and the State: Thoughts on the Policing of Protest." EUI Working Paper RSC No. 95/13. Florence: European University Institute, Robert Schuman Centre.
- [37] Reny, Tyler T., and Benjamin J. Newman. 2021. "The Opinion-Mobilizing Effect of Social Protest against Police Violence: Evidence from the 2020 George Floyd Protests." *American Political Science Review* 115(4): 1499–1507. <https://doi.org/10.1017/S0003055421000460>.

- [38] Ritter, Emily Hencken. 2014. "Policy Disputes, Political Survival, and the Onset and Severity of State Repression." *Journal of Conflict Resolution* 58(1): 143–68. <https://doi.org/10.1177/0022002712468724>.
- [39] Su, Yen-Pin. 2015. "Anti-Government Protests in Democracies: A Test of Institutional Explanations." *Comparative Politics* 47(2): 149–67. <https://doi.org/10.5129/001041515814224444>.
- [40] Sullivan, Christopher M. 2016. "Undermining Resistance: Mobilization, Repression, and the Enforcement of Political Order." *Journal of Conflict Resolution* 60(7): 1163–90. <https://doi.org/10.1177/0022002714567951>.
- [41] Sullivan, Christopher M., and Christian Davenport. 2017. "The Rebel Alliance Strikes Back: Understanding the Politics of Backlash Mobilization." *Mobilization* 22(1): 39–56. <https://doi.org/10.17813/1086-671X-22-1-39>.
- [42] Whitehead, Neil L., Jo Ellen Fair, and Leigh A. Payne. 2010. *Violent Democracies in Latin America*. Durham, NC: Duke University Press.
- [43] Wood, Elisabeth Jean. 2003. *Insurgent Collective Action and Civil War in El Salvador*. Cambridge: Cambridge University Press.
- [44] Worden, Robert E. 1995. "The 'Causes' of Police Brutality: Theory and Evidence on Police Use of Force." In *And Justice for All: Understanding and Controlling Police Abuse of Force*, eds. William A. Geller and Hans Toch. Washington, DC: Police Executive Research Forum, 31–60.
- [45] Zamble, Edward, and Phyllis Annesley. 1987. "Some Determinants of Public Attitudes toward the Police." *Journal of Police Science & Administration*.




Figure 2: Interior and full-deterrence regions for $H(\bar{N}, c_r; \theta_0) = \theta_0 \bar{N}$ and $K(c_r; \theta_1) = \theta_1 c_r^2$. The boundary is $B_s + \theta_0 = 2\theta_1 d_p^2$.




Figure 3: Region-conditional comparative statics for the additive-quadratic benchmark. Panel A varies the government's status quo stake. Panel B varies success-contingent motivation. Panel C shows the corresponding probability of policy change. Dotted curves are unconstrained interior solutions and solid curves are effective policies. The vertical axes label c_r as expected participant harm, not harm conditional on targeting.

6 Appendix.

6.1 Policy regions and comparative statics for the additive-quadratic benchmark.

6.2 Posterior beliefs under the common proper prior.

Lemma 1 (Proper-prior posterior). *Suppose $\tau \sim U[0, 1]$ and $x_i = \tau + \sigma \varepsilon_i$, where ε_i has density f and c.d.f. F . For $t \in [0, 1]$,*

$$\Pr(\tau \leq t \mid x_i = x) = \frac{F(x/\sigma) - F((x-t)/\sigma)}{F(x/\sigma) - F((x-1)/\sigma)}.$$

Proof. The likelihood of x conditional on τ is $f((x-\tau)/\sigma)/\sigma$. The uniform prior restricts its support to $[0, 1]$. Substituting $u = (x-\tau)/\sigma$ gives

$$\begin{aligned} \int_0^t \frac{1}{\sigma} f\left(\frac{x-\tau}{\sigma}\right) d\tau &= F(x/\sigma) - F((x-t)/\sigma), \\ \int_0^1 \frac{1}{\sigma} f\left(\frac{x-\tau}{\sigma}\right) d\tau &= F(x/\sigma) - F((x-1)/\sigma). \end{aligned}$$

Their ratio is the posterior probability. □

Lemma 2 (Monotone posterior). *Under the log-concavity assumption on f , the posterior probability $\Pr(\tau \leq t \mid x_i = x)$ is nonincreasing in x for every fixed $t \in (0, 1)$. A marginal-indifference solution to equation (1) therefore supports a monotone cutoff best response.*

Proof. For $x' > x$, the likelihood ratio of the two signals is

$$\frac{f((x' - \tau)/\sigma)}{f((x - \tau)/\sigma)}.$$

Log-concavity of f makes this ratio nondecreasing in τ : writing $u = (x - \tau)/\sigma$ and $a = (x' - x)/\sigma > 0$, the ratio $f(u + a)/f(u)$ is nonincreasing in u , while u decreases with τ . The posterior after x' therefore dominates the posterior after x in the monotone-likelihood-ratio order. Higher signals place more probability on higher thresholds, so $\Pr(\tau \leq t \mid x_i = x)$ is nonincreasing in x . At a solution to equation (1), citizens below the indifferent signal weakly prefer participation and citizens above it weakly prefer abstention. □

6.3 Proof of Proposition 1.

Proof of Proposition 1. Fix a cutoff strategy profile with cutoff $\hat{x}$: citizen i participates if and only if $x_i < \hat{x}$.

Step 1: participation and the induced success threshold. At true threshold τ , the mass of participants is

$$N^*(\tau) = \Pr(x_i < \hat{x} \mid \tau) = \Pr(\tau + \sigma \varepsilon_i < \hat{x}) = \Pr\left(\varepsilon_i < \frac{\hat{x} - \tau}{\sigma}\right) = F\left(\frac{\hat{x} - \tau}{\sigma}\right),$$

which is continuous and strictly decreasing in τ . At $\tau = 0$, $N^*(0) > 0$, while at $\tau = 1$, $N^*(1) < 1$. Hence $N^*(\tau) - \tau$ crosses zero exactly once on $(0, 1)$. Denote the crossing by $\hat{\tau}_\sigma$. It satisfies

$$F\left(\frac{\hat{x}_\sigma - \hat{\tau}_\sigma}{\sigma}\right) = \hat{\tau}_\sigma, \quad (10)$$

and protest succeeds exactly when $\tau \leq \hat{\tau}_\sigma$.

Step 2: reduce the system to one dimension. A citizen with signal x_i gains $\Pr(\tau \leq \hat{\tau}_\sigma \mid x_i)d_p - c_r$ from participating. Lemma 1, evaluated at $t = \hat{\tau}_\sigma$ and $x_i = \hat{x}_\sigma$, gives

$$\Pr(\tau \leq \hat{\tau}_\sigma \mid x_i = \hat{x}_\sigma) = \frac{F(\hat{x}_\sigma/\sigma) - F((\hat{x}_\sigma - \hat{\tau}_\sigma)/\sigma)}{F(\hat{x}_\sigma/\sigma) - F((\hat{x}_\sigma - 1)/\sigma)}.$$

Together with marginal indifference, this expression and equation (10) give the exact system in equation (1).

For any candidate threshold $t \in (0, 1)$, define $z(t) = F^{-1}(t)$ and

$$x_\sigma(t) = t + \sigma z(t).$$

This pair satisfies the crossing condition because $F((x_\sigma(t) - t)/\sigma) = F(z(t)) = t$. Substitution into the marginal citizen's posterior gives the continuous function

$$P_\sigma(t) = \frac{F(t/\sigma + z(t)) - t}{F(t/\sigma + z(t)) - F((t - 1)/\sigma + z(t))}. \quad (11)$$

An interior solution is a root of $P_\sigma(t) = \rho$, where $\rho = c_r/d_p$.

Step 3: local existence and compactness. Let $t_0 = 1 - \rho$ and fix $\delta \in (0, \min\{t_0, 1 - t_0\})$. On the

compact interval $I_\delta = [t_0 - \delta, t_0 + \delta]$, $z(t)$ is bounded. Uniformly for $t \in I_\delta$,

$$F(t/\sigma + z(t)) \longrightarrow 1, \quad F((t-1)/\sigma + z(t)) \longrightarrow 0$$

as $\sigma \rightarrow 0$. Equation (11) therefore implies the uniform convergence

$$\sup_{t \in I_\delta} |P_\sigma(t) - (1-t)| \longrightarrow 0.$$

At the left endpoint, the limiting value is $1-(t_0-\delta) = \rho+\delta$; at the right endpoint it is $1-(t_0+\delta) = \rho-\delta$.

For all sufficiently small positive σ , continuity then gives

$$P_\sigma(t_0 - \delta) > \rho > P_\sigma(t_0 + \delta).$$

The intermediate value theorem supplies a root $\hat{t}_\sigma \in I_\delta$ for every such σ . Lemma 2 makes the associated cutoff $\hat{x}_\sigma = x_\sigma(\hat{t}_\sigma)$ a best response, so the root is an interior monotone cutoff equilibrium. Its compactness follows from the construction rather than from a separate assumption.

Step 4: convergence of the selected branch. For any selection of roots inside I_δ ,

$$|\hat{t}_\sigma - t_0| = |(1 - \hat{t}_\sigma) - \rho| \leq \sup_{t \in I_\delta} |P_\sigma(t) - (1-t)| \longrightarrow 0.$$

Because $F^{-1}(\hat{t}_\sigma)$ is bounded on I_δ ,

$$\hat{x}_\sigma - \hat{t}_\sigma = \sigma F^{-1}(\hat{t}_\sigma) \longrightarrow 0.$$

Hence both equilibrium objects converge to $t_0 = 1 - c_r/d_p$.

Step 5: boundary policies. If $c_r > d_p$, the participation gain is strictly negative even when success is certain, so abstention is uniquely optimal. At $c_r = d_p$, the gain is weakly negative and equals zero when success is certain. Universal abstention and universal participation can therefore both be supported as weak equilibria. The selected full-deterrence outcome is the continuous extension $\hat{t}(c_r) \downarrow 0$ as $c_r \uparrow d_p$. At $c_r = 0$, participation is weakly dominant, and full participation is the corresponding continuation $\hat{t}(c_r) \uparrow 1$ as $c_r \downarrow 0$. $\square$

Remark 1. At positive noise, equation (1) can have additional roots near the boundaries of the common prior. Proposition 1 selects the root in a fixed neighborhood of $1 - c_r/d_p$ whose existence is guaranteed for all sufficiently small noise. Every selection from that neighborhood has the same limit. The state's problem is defined using this limiting branch.¹⁷

6.4 Positive-noise numerical check.

This subsection evaluates equation (1) under Gaussian noise. Let Φ denote the standard normal c.d.f. For any candidate threshold $t \in (0, 1)$, the first equation implies $x_\sigma(t) = t + \sigma\Phi^{-1}(t)$. Substituting this expression into the marginal citizen's posterior reduces the two-equation system to the one-dimensional residual in equation (11). I evaluate that residual on 1,201 equally spaced points in $(0, 1)$, bracket every sign change, and refine each root by bisection. When the bounded prior generates more than one root, I follow the interior branch by selecting the root closest to $1 - c_r/d_p$, as specified in Proposition 1.

For each selected root $\hat{t}_\sigma$, expected protest mass is calculated directly as

$$\mathbb{E}_\tau[N^*(\tau)] = \int_0^1 \Phi\left(\frac{x_\sigma(\hat{t}_\sigma) - \tau}{\sigma}\right) d\tau,$$

using Simpson's rule. The state payoff then uses this integrated expected mass rather than imposing its limiting equality with the success probability. I maximize that payoff over $c_r \in (0, d_p)$ by golden-section search and compare the resulting interior candidate with both endpoints. Repeated objective evaluations use 1,000 integration subintervals; the reported expected-mass comparisons use 4,000. For $B_s = 3$, $d_p = 3$, $\theta_1 = 1$, and $H(\bar{N}, c_r; \theta_0) = \theta_0 \bar{N}$, this procedure gives:

¹⁷The locally diffuse limiting argument follows the standard global-games logic in Morris and Shin (2003) and Bueno de Mesquita and Shadmehr (2023).

σ	c_r^* at $\theta_0 = 1.0$	c_r^* at $\theta_0 = 1.1$
0.05	0.6947	0.7134
0.03	0.6832	0.7013
0.02	0.6778	0.6954
0.01	0.6722	0.6894

The limiting policies are 0.6667 and 0.6833, respectively. The positive planning-cost effect is present at every reported noise level, and the numerical solutions move toward their analytic limits. The table is a robustness illustration, not a convergence theorem for the finite-noise optimizer. The expected-mass approximation also tightens with noise: at $c_r/d_p = 0.2$, the absolute gap between expected mass and the success probability falls from 0.0421 at $\sigma = 0.05$ to 0.0084 at $\sigma = 0.01$.

6.5 Proof of Proposition 2.

Proof of Proposition 2. For $c_r \in [0, d_p]$, substitute $\hat{\tau}(c_r) = 1 - c_r/d_p$ into equation (3):

$$u_S(c_r) = \frac{c_r}{d_p} B_s - H\left(1 - \frac{c_r}{d_p}, c_r; \theta_0\right) - K(c_r; \theta_1).$$

Differentiation gives

$$\begin{aligned} u_c(c_r) &= \frac{B_s + H_N}{d_p} - H_c - K_c, \\ u_{cc}(c_r) &= -\frac{H_{NN}}{d_p^2} + \frac{2H_{Nc}}{d_p} - H_{cc} - K_{cc}. \end{aligned}$$

Assumption 5 makes the objective strictly concave on $(0, d_p)$ and gives $u_c(0^+) > 0$. If $u_c(d_p^-) < 0$, continuity of the interior derivative and strict concavity imply a unique root of $u_c(c_r) = 0$ in $(0, d_p)$, which is the unique maximizer. If $u_c(d_p^-) \geq 0$, the objective is increasing up to the full-deterrence corner and is maximized at $c_r = d_p$ on this interval. At every $c_r > d_p$, the clamped threshold and expected protest mass are zero, so $u_S(c_r) = B_s - H(0, c_r; \theta_0) - K(c_r; \theta_1)$ and its derivative is

$-H_c(0, c_r; \theta_0) - K_c(c_r; \theta_1) < 0$. Thus the policy characterized in the proposition is globally optimal. This argument uses one-sided derivatives at the kink $c_r = d_p$. $\square$

6.6 Proof of Proposition 3.

Proof of Proposition 3. At an interior solution define¹⁸

$$G(c_r; B_s, \theta_0, \theta_1, d_p) \equiv u_c(c_r) = 0.$$

Because $G_{c_r} = u_{cc} < 0$, the implicit function theorem gives

$$\frac{\partial c_r^*}{\partial B_s} = -\frac{1/d_p}{u_{cc}} > 0, \quad \frac{\partial c_r^*}{\partial \theta_1} = \frac{K_{c\theta_1}}{u_{cc}} < 0,$$

where $K_{c\theta_1} > 0$. Fixing c_r and differentiating G in d_p gives

$$G_{d_p} = -\frac{B_s + H_N + c_r H_{Nc} - (c_r/d_p) H_{NN}}{d_p^2},$$

and therefore

$$\frac{\partial c_r^*}{\partial d_p} = \frac{B_s + H_N + c_r^* H_{Nc} - (c_r^*/d_p) H_{NN}}{d_p^2 u_{cc}}.$$

This derivative is negative under the condition stated in Proposition 3. For the crowd-cost shifter,

$$G_{\theta_0} = \frac{H_{N\theta_0}}{d_p} - H_{c\theta_0}, \quad \frac{\partial c_r^*}{\partial \theta_0} = -\frac{H_{N\theta_0}/d_p - H_{c\theta_0}}{u_{cc}}.$$

Since $u_{cc} < 0$, this derivative is positive exactly when the numerator is positive. This is the crowd-cost condition. The primitive B_p is absent from G , which gives its zero direct partial effect.

Finally, $\hat{\tau}^* = 1 - c_r^*/d_p$. Its response to B_s , θ_0 , and θ_1 has the opposite sign from the corresponding response of c_r^* . For motivation,

$$\frac{\partial \hat{\tau}^*}{\partial d_p} = \frac{c_r^*}{d_p^2} - \frac{1}{d_p} \frac{\partial c_r^*}{\partial d_p} > 0,$$

¹⁸For the first-order-condition function G , $G_z \equiv \partial G/\partial z$ fixes every displayed primitive other than z . In particular, $G_{c_r} = u_{cc}$, while G_{d_p} and G_{θ_0} differentiate the first-order condition at a fixed candidate policy c_r . In the limiting state problem, expected mass equals $\hat{\tau} = 1 - c_r/d_p$, so G_{d_p} includes the induced change $\partial \hat{\tau}/\partial d_p = c_r/d_p^2$ at fixed c_r . The implicit-function step then allows the equilibrium policy to adjust.

because $c_r^* > 0$ and $\partial c_r^*/\partial d_p < 0$. The direct partial effect of B_p on $\hat{t}^*$ is also zero. $\square$

6.7 Proof of Proposition 4.

Proof of Proposition 4. With $d_p(c_r) = d_0 + e(c_r)$, apply the quotient rule to $\hat{t}(c_r) = 1 - c_r/d_p(c_r)$:

$$\frac{\partial \hat{t}}{\partial c_r} = -\frac{d_p(c_r) \cdot 1 - c_r e'(c_r)}{d_p(c_r)^2} = -\frac{d_p(c_r) - c_r e'(c_r)}{d_p(c_r)^2}.$$

The denominator is positive, so the derivative has the sign of $-(d_p(c_r) - c_r e'(c_r))$. Dividing the parenthesis by $d_p(c_r) > 0$,

$$d_p(c_r) - c_r e'(c_r) \geq 0 \iff 1 - \frac{c_r e'(c_r)}{d_p(c_r)} = 1 - \varepsilon(c_r) \geq 0,$$

so $\partial \hat{t}/\partial c_r < 0$ exactly when $\varepsilon(c_r) < 1$ and $\partial \hat{t}/\partial c_r > 0$ exactly when $\varepsilon(c_r) > 1$. $\square$

6.8 Proof of Proposition 5.

Proof of Proposition 5. At an interior threshold, equation (9) gives $D = -\hat{t}_c$. Differentiating the state payoff at a fixed α gives the first-order condition

$$u_c(c_r; \alpha) = (B_s + H_N)D - H_c - K_c = 0.$$

At a positive policy, $B_s + H_N > 0$, $H_c \geq 0$, and $K_c > 0$. The equality is therefore impossible when $D \leq 0$. Every positive interior optimum must have $D > 0$, which is equivalent to $\varepsilon(c_r^*; \alpha) < 1$.

Under the deployment-separable technology, the first-order condition becomes

$$G(c_r, \alpha) \equiv (B_s + w)D(c_r, \alpha) - A'(c_r) - K_c(c_r; \theta_1) = 0.$$

Holding c_r fixed and differentiating the marginal deterrent return in α gives

$$D_\alpha(c_r, \alpha) = -\frac{e(c_r) \left[d_p(c_r; \alpha) - c_r \alpha e'(c_r) \right] + c_r e'(c_r) d_0}{d_p(c_r; \alpha)^3} < 0$$

whenever the emotional response is nonzero and $D > 0$. Hence $G_\alpha = (B_s + w)D_\alpha < 0$. Because

$G_{c_r} = u_{cc} < 0$ at the strictly concave interior optimum, the implicit function theorem yields

$$\frac{\partial c_r^*}{\partial \alpha} = -\frac{G_\alpha}{G_{c_r}} < 0.$$

□

6.9 Numerical verification protocol.

I evaluate the remaining numerical checks directly from the displayed equations. For an interior state optimum, I solve $u_c(c_r) = 0$ by bisection and compare each analytic comparative static in Proposition 3 with a symmetric finite difference obtained after re-solving the state's problem at the perturbed primitive. The perturbation is reduced until both solutions remain interior. I perform this comparison under planning-cost technologies for which $H_{N\theta_0}/d_p - H_{c\theta_0}$ is positive, negative, and zero. For Proposition 4, I differentiate the threshold numerically and compare its sign with $\varepsilon(c_r) - 1$, including values on both sides of one. These calculations are checks on the algebra; the propositions continue to follow from the analytic proofs above.